

# **Effective and Practical Strategies for Combatting Misinformation**

Catherine King

March 14, 2025

Software and Societal Systems Department  
School of Computer Science  
Carnegie Mellon University  
Pittsburgh, PA 15213

## **Thesis Committee:**

Prof. Kathleen M. Carley, Chair, Carnegie Mellon University

Prof. Hong Shen, Carnegie Mellon University

Prof. Chris Labash, Carnegie Mellon University

Prof. Pablo Barberá, University of Southern California and Meta

*Submitted in partial fulfillment of the requirements  
for the degree of Doctor of Philosophy in Societal Computing.*

Copyright © 2025 Catherine King

This work was supported in part by the Office of Naval Research (ONR) Award N000142112229 for “Scalable Tools for Social Media Assessment”, the Knight Foundation, the Center for Computational Analysis of Social and Organizational Systems (CASOS), and the Center for Informed Democracy and Social-cybersecurity (IDeAS). The views and conclusions contained in this document are those of the author and should not be interpreted as representing the official policies, either expressed or implied, of the ONR or the U.S. government.

March 14, 2025  
DRAFT

**Keywords:** misinformation interventions, countermeasures, social network analysis, social media

*To my family.*

March 14, 2025  
DRAFT



## Abstract

Social media platforms, which are becoming a primary news source for many individuals, can quickly spread mis/disinformation online faster than ever before. These information disorders contribute to increased polarization and extremism, threatening to undermine democracy and trust in public institutions worldwide. Because of this growing problem, researchers have begun investigating the effectiveness of possible interventions to counter this misinformation. This research is critical given the many societal challenges we face that are associated with the spread of false or misleading information.

Most research in the countermeasures space focuses on the effectiveness of some more easily studied interventions. Some interventions, like fact-checking, are studied more than others because they can be analyzed without complete access to comprehensive social media data. Most researchers also focus on determining the effectiveness of an intervention without considering if the public would support the countermeasure. Platforms and governments will likely only implement changes that have public support.

In this thesis, I develop a framework for designing and evaluating misinformation interventions that integrates current research on effectiveness with user acceptance to enable more effective implementation strategies. To accomplish this task, I present a detailed categorization of interventions. Then, I conduct a citation network analysis of the literature in this field to identify research gaps, areas of disagreement, and potential next steps. I administered a comprehensive survey asking the American public about their social media behavior and their opinions on various interventions. The survey also examines how certain factors may affect user acceptance. These factors include perceived effectiveness, fairness, and intrusiveness. Next, I conducted an effectiveness study on a training game to contribute to the currently contentious literature regarding the effectiveness, or potential lack thereof, of training games. Finally, I combine this research to evaluate countermeasures across a comprehensive list of features, aiming to identify the shared characteristics that make countermeasures effective and practical.



## **Acknowledgments**

...



# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Thesis Objective . . . . .	1
1.2	Definitions . . . . .	2
1.3	Misinformation Categorization . . . . .	2
1.3.1	Literature Review . . . . .	2
1.3.2	Proposed Typology . . . . .	4
1.3.3	Summary . . . . .	8
1.4	Countermeasures Categorization . . . . .	8
1.4.1	Literature Review . . . . .	8
1.4.2	Proposed Typology . . . . .	10
1.5	Current Platform Policies . . . . .	11
1.5.1	Community Guidelines . . . . .	11
1.5.2	Enforcement . . . . .	13
1.6	Operationalized Interventions . . . . .	14
1.7	Data . . . . .	14
1.8	Chapter Overview . . . . .	18
<b>2</b>	<b>Citation Network Analysis of the Interventions Literature</b>	<b>19</b>
2.1	Introduction . . . . .	19
2.2	Data and Methods . . . . .	20
2.2.1	Paper Labels . . . . .	20
2.2.2	Inclusion Criteria . . . . .	21
2.2.3	Literature Search . . . . .	21
2.2.4	Training of Labelers . . . . .	22
2.2.5	ChatGPT Prompts . . . . .	23
2.2.6	Label Agreement . . . . .	25
2.2.7	Data Availability . . . . .	25
2.3	Citation Network Analysis . . . . .	25
2.3.1	Data Attributes . . . . .	27
2.3.2	Nodesets . . . . .	27
2.3.3	Networks . . . . .	28
2.4	Results . . . . .	29
2.4.1	Descriptive Statistics . . . . .	29
2.4.2	Publication Venue Analysis . . . . .	31

2.4.3	Topic Analysis . . . . .	33
2.4.4	Author Analysis . . . . .	37
2.5	Discussion . . . . .	39
2.6	Conclusions . . . . .	40
2.6.1	Limitations . . . . .	40
2.6.2	Contributions . . . . .	40
<b>3</b>	<b>Characterizing User-based Countermeasures</b>	<b>43</b>
3.1	Introduction . . . . .	43
3.2	Data and Methods . . . . .	45
3.2.1	Ethics Information . . . . .	45
3.2.2	Pilot Data . . . . .	45
3.2.3	Survey Design and Sampling Plan . . . . .	45
3.2.4	Measures . . . . .	47
3.3	Closeness Analysis . . . . .	48
3.3.1	Related Work . . . . .	49
3.3.2	Bayesian Power Analysis . . . . .	50
3.3.3	Statistical Methods . . . . .	53
3.3.4	Results . . . . .	55
3.3.5	Closeness Discussion . . . . .	64
3.4	Platform Analysis . . . . .	65
3.4.1	Related Work . . . . .	65
3.4.2	Results . . . . .	70
3.4.3	Platform Discussion . . . . .	75
3.5	Overall Discussion . . . . .	76
3.6	Conclusions . . . . .	77
3.6.1	Limitations . . . . .	77
3.6.2	Contributions . . . . .	78
<b>4</b>	<b>Improving User-based Countermeasures</b>	<b>79</b>
4.1	Introduction . . . . .	79
4.2	Related Work . . . . .	80
4.2.1	Media Literacy Interventions . . . . .	80
4.2.2	User-Based Interventions . . . . .	80
4.3	Data and Methods . . . . .	81
4.3.1	OMEN Overview . . . . .	81
4.3.2	Participant Demographics . . . . .	81
4.3.3	Survey Design . . . . .	82
4.3.4	Training Sessions . . . . .	87
4.3.5	Analysis Methods . . . . .	90
4.3.6	Ethics Information . . . . .	91
4.4	Results: Misinformation Detection . . . . .	91
4.4.1	Accuracy . . . . .	91
4.4.2	Confidence . . . . .	93

4.5	Results: Countering Misinformation . . . . .	94
4.5.1	Selected Interventions . . . . .	94
4.5.2	Effort Level . . . . .	95
4.5.3	Platform and Poster Impact . . . . .	95
4.5.4	Factors Affecting Countering Actions . . . . .	96
4.6	Participant Feedback . . . . .	101
4.7	Discussion . . . . .	103
4.8	Conclusions . . . . .	104
4.8.1	Limitations . . . . .	104
4.8.2	Contributions . . . . .	104
<b>5</b>	<b>Characterizing Platform and Government Countermeasures</b>	<b>107</b>
5.1	Introduction . . . . .	107
5.2	Related Work . . . . .	108
5.2.1	Intervention Selection . . . . .	109
5.3	Data and Methods . . . . .	111
5.3.1	Platform and Government Survey Questions . . . . .	111
5.3.2	Measures . . . . .	112
5.3.3	Pre-Registered Analysis Plan . . . . .	112
5.4	Results . . . . .	113
5.4.1	Factors that Influence Support for Interventions . . . . .	113
5.4.2	Overall Support and Perceptions of Interventions . . . . .	115
5.4.3	Individual Differences in Support and Perceptions of Interventions . . . . .	116
5.5	Discussion . . . . .	116
5.6	Conclusion . . . . .	121
5.6.1	Limitations . . . . .	121
5.6.2	Contributions . . . . .	122
<b>6</b>	<b>Recommendations for Effective and Practical Countermeasures</b>	<b>123</b>
6.1	Introduction . . . . .	123
6.2	Related Work . . . . .	124
6.2.1	Descriptive Characteristics . . . . .	124
6.2.2	User Perceptions . . . . .	125
6.2.3	Main Criteria . . . . .	126
6.3	Data and Methods . . . . .	127
6.3.1	Survey Design . . . . .	127
6.3.2	Participant Demographics . . . . .	129
6.3.3	Analysis Methods . . . . .	130
6.3.4	Measures . . . . .	130
6.4	Results . . . . .	130
6.4.1	Comparison of General Categories . . . . .	131
6.4.2	Content Distribution . . . . .	133
6.4.3	Content Moderation . . . . .	135
6.4.4	Account Moderation . . . . .	136

6.4.5	Content Labeling . . . . .	136
6.4.6	User-based Measures . . . . .	137
6.4.7	Media Literacy and Education . . . . .	138
6.4.8	Institutional Measures . . . . .	138
6.4.9	Generative AI . . . . .	139
6.4.10	Summary . . . . .	140
<b>7</b>	<b>Conclusion</b>	<b>141</b>
7.1	Summary . . . . .	141
7.2	Contributions . . . . .	141
7.3	Limitations . . . . .	141
7.4	Future Work . . . . .	141
<b>A</b>	<b>Countermeasures Categorization</b>	<b>143</b>
A.1	Content Distribution . . . . .	143
A.2	Content Moderation . . . . .	144
A.3	Account Moderation . . . . .	144
A.4	Content Labeling . . . . .	145
A.5	Media Literacy and Education . . . . .	145
A.6	User-based Measures . . . . .	146
A.7	Institutional Measures . . . . .	146
A.8	Other . . . . .	147
<b>B</b>	<b>Chapter 2: ChatGPT Prompts</b>	<b>149</b>
B.1	Effectiveness Labeling Prompt . . . . .	149
B.2	Acceptance Labeling Prompt . . . . .	150
B.3	Interventions Labeling Prompt . . . . .	150
<b>C</b>	<b>Chapter 3: Pre-Registered Hypotheses and Analyses</b>	<b>153</b>
<b>D</b>	<b>Chapter 3: Supplemental Categorical Analysis</b>	<b>157</b>
<b>E</b>	<b>Chapter 3: Demographics Summary Table</b>	<b>161</b>
<b>F</b>	<b>Chapter 4: OMEN Survey Posts</b>	<b>165</b>
F.1	Misinformation Detection Posts . . . . .	165
F.2	Countering Posts . . . . .	181
<b>G</b>	<b>Chapter 5: Robustness Analysis</b>	<b>189</b>
<b>Bibliography</b>		<b>193</b>

# List of Figures

2.1	Literature review process. . . . .	22
2.2	Number of articles in this review published per year. . . . .	29
2.3	Publication Venue x Publication Venue (Co-Authorship) Network. Nodes are sized by Total Degree Centrality and colored by discipline (red for computer science, yellow for psychology, blue for social sciences, black for multidisciplinary, and grey for other) . . . . .	32
2.4	Number of articles in this review by venue discipline. . . . .	33
2.5	The Intervention Topic x Intervention Topic Network. Nodes are sized by total degree centrality and colored by paper count. The lighter the blue, the more papers that study that topic. . . . .	36
2.6	Number of articles in this review by general intervention category studied. . . . .	36
2.7	The largest component of the Co-Authorship Network. Nodes are sized by total number of papers in the data set and colored by betweenness centrality, with lighter blue colors indicating higher betweenness. . . . .	38
2.8	The largest component of the Co-Authorship Network excluding one paper with 30 authors. Nodes are sized by total number of papers in the data set and colored by betweenness centrality, with lighter blue colors indicating higher betweenness.	39
3.1	Total number of participants who selected each intervention type from Table 3.1 at least once in their behavioral and opinion responses. For comparison purposes, only participants who had seen misinformation at that closeness level have their opinion counts included. . . . .	57
3.2	Number of participants who reported expending a maximum of no, low, or high effort when seeing misinformation compared with the total number of those same participants who <i>believe</i> one should expend no, low, or high effort when seeing misinformation at each of the three closeness levels and against oneself. . . . .	58
3.3	Total number of participants who had posted misinformation and selected each intervention type from Table 3.2 at least once in their behavioral and opinion responses. For comparison purposes, only participants who admitted to posting misinformation had their opinion counts included. . . . .	59
3.4	Highest effort level participants said one should respond with when seeing misinformation posted by others or oneself broken up by party, age, gender, race, education, and income. Chi-sq tests: $p < 0.05^*$ , $p < 0.01^{**}$ , and $p < 0.001^{***}$ . .	63

3.5	The number of participants who use each platform at least weekly or daily compared with the number of participants who have seen misinformation on that platform. . . . .	71
3.6	Highest effort level participants said they responded with when seeing misinformation posted by others on the top 6 platforms for misinformation. Chi-sq test: $p < 0.001^{***}$ . . . . .	72
3.7	Highest effort level participants said they responded with when seeing misinformation posted by <b>not close contacts only</b> on the top 6 platforms for misinformation. Chi-sq test: $p < 0.001^{***}$ . . . . .	73
3.8	Number of participants who admitted to posting misinformation on each platform. . . . .	74
3.9	Highest effort level participants said they responded with after realizing they had posted misinformation accidentally among the top 4 platforms. Chi-sq test was not significant. . . . .	75
4.1	An example of a false post shown to participants. . . . .	86
4.2	Summary training slide on misinformation detection. . . . .	88
4.3	The number of times participants said their answer would change depending on either the platform or the poster over both the pre- and post-training surveys. . . . .	94
4.4	The number of times participants said their answer would change depending on either the platform or the poster in both the pre and post-training surveys. . . . .	96
4.5	The number of times participants said their answer would change depending on either the poster or the platform combined over both the pre- and post-training surveys. . . . .	97
4.6	This figure displays the total number of times over both surveys that participants mentioned a specific sub-theme when asked if and how their answer would change depending on the misinformation poster and the platform. . . . .	99
4.7	This figure summarizes the responses to the feedback questions at the end of the study, indicating whether participants agreed or disagreed that the training was helpful. . . . .	102
4.8	The number of participants who said they used each technique in the post-training survey. . . . .	103
5.1	Estimate and 95% CI of the effect of perceived fairness, effectiveness and intrusiveness on support depending on implementer. . . . .	113
5.2	Average and 95% CI support, perceived fairness, perceived effectiveness, and perceived intrusiveness for each intervention (1-10) and implementer (government and platform) on a 1-5 Likert scale. . . . .	115
5.3	Average ratings by political party, age, gender, race, education, and income. One-way ANOVA tests were run on each grouping, with stars indicating the level of significance: $p < 0.05^*$ , $p < 0.01^{**}$ , and $p < 0.001^{***}$ . . . . .	118
6.1	Stacked bar plot representing the percentage of responses on the effectiveness and acceptance of the general categories of interventions. . . . .	131

6.2	Average effectiveness and acceptance Likert scores per general intervention category.	132
F.1	Pre-Test Misinformation Post #1.	165
F.2	Pre-Test Misinformation Post #2.	166
F.3	Pre-Test Misinformation Post #3.	167
F.4	Pre-Test Misinformation Post #4.	168
F.5	Pre-Test Real News Post #1.	169
F.6	Pre-Test Real News Post #2.	170
F.7	Pre-Test Real News Post #3.	171
F.8	Pre-Test Real News Post #4.	172
F.9	Post-Test Misinformation Post #1.	173
F.10	Post-Test Misinformation Post #2.	174
F.11	Post-Test Misinformation Post #3.	175
F.12	Post-Test Misinformation Post #4.	176
F.13	Post-Test Real News Post #1.	177
F.14	Post-Test Real News Post #2.	178
F.15	Post-Test Real News Post #3.	179
F.16	Post-Test Real News Post #4.	180
F.17	Pre-Test Countering Post #1.	181
F.18	Pre-Test Countering Post #2.	182
F.19	Pre-Test Countering Post #3.	183
F.20	Pre-Test Countering Post #4.	184
F.21	Post-Test Countering Post #1.	185
F.22	Post-Test Countering Post #2.	186
F.23	Post-Test Countering Post #3.	187
F.24	Post-Test Countering Post #4.	188



# List of Tables

1.1	Misinformation categorizations in the literature. . . . .	3
1.2	Misinformation suppliers: their motivation, news context, and targeted audiences. . . . .	7
1.3	Proposed misinformation categorization. . . . .	8
1.4	Categorizations of misinformation interventions in the literature. . . . .	9
1.5	Proposed categorization of misinformation interventions. . . . .	11
1.6	Proposed categorization of specific misinformation interventions. . . . .	12
1.7	Operationalized platform interventions. . . . .	15
1.8	Operationalized user-based and institutional interventions. . . . .	16
1.9	Summary of datasets used in this thesis. . . . .	17
2.1	The intervention labels used the least after Stage 1, sorted from highest to lowest based on the number of papers with those labels. . . . .	23
2.2	Summary of ChatGPT interventions labeling by different prompt and model types. The values indicate Jaccard similarity scores that compare the labeled output produced by the model run against the final label. . . . .	26
2.3	Jaccard Similarity agreement with the final labels broken up by rater type. Rater types are sorted by total number of papers labeled. . . . .	26
2.4	Jaccard Similarity agreement with the final labels broken up by specific raters. Raters are sorted by total number of papers labeled. . . . .	27
2.5	The top 10 most cited papers in the data set. . . . .	30
2.6	The top 10 venues in the Co-Publication Venue network. . . . .	31
2.7	The top venues by total number of publications in the data set. . . . .	32
2.8	The number of papers and unique authors who have studied each intervention type. . . . .	34
2.9	Statistical summary of the number of papers and unique authors by intervention type. Outliers are identified from the inter-quartile range. . . . .	35
2.10	The number of studies examining the effectiveness or acceptance of one or more interventions. . . . .	35
2.11	The top 10 most cited authors with more than one paper in the data set. . . . .	37
2.12	The top 10 authors in the Co-Authorship network. . . . .	37
3.1	Actions social media users can take when they see misinformation online. . . . .	46
3.2	Actions social media users can take when they realize they have posted misinformation online. . . . .	47
3.3	Misinformation exposure . . . . .	55

3.4	Pre-registered Bayesian paired hypothesis test results and interpretation for RQ1. Values in italics are for the hypotheses being run with the effort level summed as a robustness check. The p-value for the generalized McNemar's Chi-square Test is also included. . . . .	56
3.5	Pre-registered Bayesian paired hypothesis test results and interpretation for RQ2. Values in italics are for the hypotheses being run with the effort level summed as a robustness check. The p-value for the generalized McNemar's Chi-square Test is also included. . . . .	59
3.6	Pre-registered Bayesian paired hypothesis test results and interpretation for RQ3. Values in italics are for the hypotheses being run with the effort level summed as a robustness check. The p-value for the generalized McNemar's Chi-square Test is also included. . . . .	60
3.7	Social media platform categorization as developed by Zhu and Chen (2015) [249]. . . . .	66
3.8	Modified social media platform categorization. . . . .	67
3.9	Social media engagement on various platforms per dimension [227]. The innovation and empowerment dimensions are excluded, as all platforms scored low. . . . .	69
3.10	Number and percent of people who saw misinformation from close, somewhat, and not close contacts per platform. The top three platforms by percentage per closeness level are bolded. Platforms are sorted by the number of people who have seen misinformation on the platform. . . . .	72
3.11	Summary of Interventions Selected after Posting Misinformation . . . . .	74
4.1	The number of participants that used various social media platforms for work and for personal reasons. . . . .	83
4.2	Actions one can take when engaging with misinformation on social media. . . . .	85
4.3	The measured variables from Part 1 of the surveys . . . . .	91
4.4	The measured variables from Part 2 of the surveys . . . . .	92
4.5	The average and median number of questions answered correctly by participants. . . . .	93
4.6	Two-sided paired t-test results ( $df = 22$ ) and the estimated effect size, Cohen's $d$ , for the average accuracy scores given by each participant as defined in Table 4.3. . . . .	93
4.7	Two-sided paired t-test results ( $df = 22$ ) and the estimated effect size, Cohen's $d$ , for the average self-reported confidence scores on a scale 1-10 given by each participant as defined in Table 4.3. . . . .	94
4.8	The percentage and number of participants whose maximum effort level was in each of the three effort level categories described in Table 4.2. . . . .	95
4.9	The contingency table showing the number of participants that were no, low, and high effort in the pre vs. post-test. . . . .	95
4.10	Two-sided paired t-test results ( $df = 22$ ) and the estimated effect size, Cohen's $d$ , for the average poster and platform effect scores given by each participant as defined in Table 4.4. . . . .	96
4.11	The number of unique participants and total mentions of each major theme when considering the misinformation poster or the platform. . . . .	98

4.12	The number of unique participants that mentioned each sub-theme and the total number of times each sub-theme was mentioned when participants were asked to consider the poster of the misinformation. . . . .	98
4.13	The number of unique participants that mentioned each sub-theme and the total number of times each sub-theme was mentioned when participants were asked to consider the misinformation platform. . . . .	99
4.14	The average, standard deviation, and median feedback score on a 1-5 Likert scale, where 1 represents “strongly disagree” and 5 represents “strongly agree.” .	102
5.1	Selected misinformation interventions. Text in [brackets] was included when the specified implementer was government. . . . .	110
5.2	OLS regression predicting support for a misinformation intervention as a function of perceived fairness, perceived effectiveness, perceived intrusiveness, and implementer (reference level: social media company) with robust standard errors clustered on participant and intervention. Model 1 is the main model. Model 2 is the robustness check that includes responses from all participants who responded to at least part of the survey. . . . .	114
5.3	OLS regressions predicting average support, perceived fairness, perceived effectiveness, and perceived intrusiveness of interventions as a function of age (18-24 to 65+ age brackets mapped to numeric values 0 to 5), gender (reference level: Female), partisanship (reference level: Democrat, combined “Other” party with “Independent”), political ideology (very liberal to very conservative categories mapped to numeric values 0 to 4), highest education level (less than high school diploma to Doctorate or Professional Degree categories mapped to numeric values 0 to 6), income (less than \$20,000 to over \$200,000 income brackets mapped to numeric values 0 to 7), frequency seeing misinformation (never to very often mapped to numeric values 0 to 4). . . . .	117
5.4	Adhoc analysis: OLS regression predicting support for a misinformation intervention as a function of perceived fairness, perceived effectiveness, perceived intrusiveness, implementer (reference level: social media company), gender (reference level: Female), and partisanship (reference level: Democrat, with Independents combined with “Other”) with robust standard errors clustered on participant and intervention. . . . .	119
6.1	The percentage of respondents who agree or disagree that each general category of interventions is effective or acceptable to the public. . . . .	132
6.2	Descriptive characteristics of Content Distribution interventions. . . . .	133
6.3	Expert scores for content distribution interventions. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	134
6.4	Descriptive characteristics of Content Moderation interventions. . . . .	135
6.5	Expert scores for content moderation interventions. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	135
6.6	Descriptive characteristics of Account Moderation interventions. . . . .	136

6.7	Expert scores for Account Moderation interventions. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	136
6.8	Descriptive characteristics of Content Labeling interventions. . . . .	136
6.9	Expert scores for Content Labeling interventions. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	137
6.10	Descriptive characteristics of User-based interventions. . . . .	137
6.11	Expert scores for User-based Measures. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	137
6.12	Descriptive characteristics of Media Literacy interventions. . . . .	138
6.13	Expert scores for Media Literacy and Education interventions. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	138
6.14	Descriptive characteristics of Institutional Measures interventions. . . . .	138
6.15	Expert scores for Institutional Measures. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	139
6.16	Descriptive characteristics of Generative AI interventions. . . . .	139
6.17	Expert scores for Generative AI interventions. Mean values are shown with their corresponding standard deviations in parentheses. . . . .	139
6.18	Ranking of the top 10 overall interventions by the five metrics. . . . .	140
D.1	Contingency table of effort level, response type, and closeness level. Each cell contains the frequency of occurrences for each combination of the categorical variables. The chi-sq test statistic is shown: $p < 0.05^*$ , $p < 0.01^{**}$ , and $p < 0.001^{***}$ . . . . .	157
D.2	The generalized estimating equation (GEE) model predicting the average frequency counts as a function of effort level (reference level: no effort), response type (reference level: behavior), and closeness (reference level: close contacts) when considering responses against others. . . . .	159
D.3	Contingency table of effort level and response type for actions against oneself. Each cell contains the frequency of occurrences for each combination of the categorical variables. The chi-sq test statistic is shown: $p < 0.05^*$ , $p < 0.01^{**}$ , and $p < 0.001^{***}$ . . . . .	160
D.4	The generalized estimating equation (GEE) model predicting the average frequency counts as a function of effort level (reference level: no effort) and response type (reference level: behavior) when considering responses to oneself. . . . .	160
E.1	Demographics summary table. . . . .	161
G.1	OLS regressions predicting average support, perceived fairness, perceived effectiveness, and perceived intrusiveness of interventions as a function of demographic variables. Same analysis as Table 5.3, except partisanship separated out the “Other” party from “Independent”. . . . .	190

G.2 OLS regressions predicting average support, perceived fairness, perceived effectiveness, and perceived intrusiveness of interventions as a function of demographic variables. Same analysis as Table 5.3, except partisanship was mapped to numeric values 0 to 4, with strong Democrats mapped to 0 and strong Republicans mapped to 4. . . . .	191
--	-----



# Chapter 1

## Introduction

### 1.1 Thesis Objective

In recent years, there has been an increased research focus on the spread, impact, and mitigation of misinformation online. Social media is aiding in the dissemination of misinformation [136, 228, 247], and researchers are growing more concerned about how social media may be contributing to political polarization and distrust in institutions and the media. Information disorders like misinformation and disinformation have been shown to have pressing societal impacts ranging from undermining democracy [221], increasing extremism [236], and lowering the uptake of various public health measures during a pandemic like COVID-19 [167].

Countering misinformation is a challenging problem, as there are many possible solutions and aspects to consider. Researchers also often only have limited access to social media data, especially data that could be used to evaluate the effectiveness of various countermeasures [59]. Even if data is available, in some cases there are ethical challenges associated with sharing social media data with other researchers [33]. This lack of access contributes to why some countermeasures, like fact-checking, are studied significantly more than others.

According to a review of 223 countermeasures studies since 1972 by Courchesne et al. (2021), there has been a disproportionate amount of research on the effect of fact-checking [180, 230], debunking [51, 73], and prebunking [142, 194]. However, many countermeasures, including those that could target creators of disinformation, have not been studied at all [59]. Finally, most intervention papers focus on the effectiveness of the intervention without considering crucial aspects like user acceptance, political feasibility, and cost.

The goal of this thesis is to better understand the efficacy and practicality of misinformation countermeasures in order to provide analysis-driven recommendations. I propose an approach to developing misinformation interventions that (1) integrates current social science theory about effectiveness with (2) user opinions and acceptability while considering other relevant factors, such as transparency, cost, and fairness. My main research questions are as follows:

1. **How can we assess how practical and effective countermeasures are?**
2. **What do successful countermeasures have in common?**
3. **Can we develop a framework for providing analysis-driven recommendations on what to implement and why?**

This thesis is limited to user-based countermeasures, social media platform countermeasures, and possible government regulation. It focuses primarily on the US-based context. Section 1.2 defines the general terms used throughout the dissertation. Sections 1.3 and 1.4 outline the various types of misinformation and interventions. Section 1.6 translates these intervention categories into specific interventions that will be compared throughout the dissertation. Finally, Section 1.7 discusses the data used throughout the dissertation, and Section 1.8 concludes with an outline of the subsequent chapters.

## 1.2 Definitions

Social media platforms have allowed people and organizations to access and spread information faster than ever before, and they have also facilitated the increased speed and reach of misinformation [136, 228, 247]. Researchers typically make a distinction between misinformation and disinformation:

- **Misinformation** - False, inaccurate, and/or misleading information [211, 235, 241].
- **Disinformation** - Misinformation that is created with the *intent* to deceive and spread intentionally, often in the form of propaganda [211, 235, 241].

While this distinction is important, intent can be challenging to determine accurately. This thesis intends not to detect mis/disinformation but rather to develop measures against false and misleading content. Any false, inaccurate, or misleading information, regardless of intent, will typically be referred to as “misinformation” as an umbrella term for the remainder of this document.

## 1.3 Misinformation Categorization

This section will first provide an overview and comparison of how different researchers classify misinformation or “fake news”. Then, using this prior work, I will describe a proposed categorization.

### 1.3.1 Literature Review

Four prominent review articles are reviewed in this section. Table 1.1 shows the high-level misinformation categories described in these documents.

First, Wardle et al. (2017) define seven main types of misinformation, loosely categorized along an axis that ranges from low risk and intent to deceive or harm to high risk and intent to deceive and harm. These categories, listed from low to high risk, are: *satire or parody*, *false connection*, *misleading content*, *false context*, *imposter content*, *manipulated content*, and *fabricated content*. [235]. This typology has been adapted and used by several university library research guides, including Temple University<sup>1</sup>, Northeastern University<sup>2</sup>, and University of Iowa<sup>3</sup>.

<sup>1</sup><https://guides.temple.edu/c.php?g=646455&p=9816775>

<sup>2</sup><https://subjectguides.lib.neu.edu/fakenews>

<sup>3</sup><https://guides.lib.uiowa.edu/c.php?g=849536&p=6077637>

Table 1.1: Misinformation categorizations in the literature.

Type	Wardle et al. (2017) [235]	Tandoc et al. (2017) [211]	Zannettou et al. (2019) [245]	Wang et al. (2020) [234]
Satire and Parody	×	×	×	×
Fabricated Content	×	×	×	×
Manipulated Content	×	×		×
Imposter Content	×			×
False Context	×			
Misleading Content	×			
False Connection	×			
Error (False Content)				×
Propaganda		×	×	×
Conspiracy Theories			×	
Hoaxes			×	
Rumors			×	×
Advertising		×		×
Clickbait			×	×
Biased or one-sided			×	×
Other				×

In Tandoc et al. (2017), 34 articles that used the term “fake news” were reviewed to develop a typology [211]. The authors primarily considered two dimensions: the level of facticity (the extent to which the misinformation uses facts) and the intention to deceive. They created six categories of fake news. Among the categories with low intentions to deceive were *news satire* and *news parody*. This article was the only one of the four that differentiated between satire and parody as separate categories. *News satire* typically employs exaggeration to deliver the news and relies on a high level of facticity (e.g., the Daily Show). In contrast, *news parody* is more likely to rely on fabricated content (e.g., the Onion) [211]. Among the categories with higher intentions to deceive, listed from high to low facticity, were misleading *advertising*, *propaganda*, *manipulated content*, and *fabricated content*. The authors note that advertising and propaganda may sometimes have overlapping motives [211].

Zannettou et al. (2019) define eight types of misinformation and detail the actors behind it and their potential motives [245]. Like the previous two categorizations, they include categories for satirical news and fabricated content. They also include propaganda, similar to Tandoc et al. (2017), but explicitly note that it is a type of fabricated content with a political agenda. Additionally, they define *conspiracy theories*, *hoaxes*, and *rumors* as types of misinformation. While all three have distinct technical definitions, these definitions overlap. Conspiracy theories involve unsubstantiated rumors of a conspiracy, making them a type of rumor. Hoaxes typically consist of stories that contain either fabricated content or half-truths. They define *clickbait* as “the deliberate use of misleading headlines and thumbnails of content on the Web” and similar to yellow journalism [245]. This clickbait category overlaps with both Tandoc’s advertising category [211] and Wardle’s misleading content or false connection categories [235]. Their final category is for *biased* news, which is one-sided or hyperpartisan [245]. This also qualifies as a

type of misleading content.

Finally, Wang et al. (2020) build upon Tandoc et al.'s typology by adding six additional categories: *clickbait*, alarmist talk (categorized in Table 1.1 as *other*), subjective assumption (categorized as *biased or one-sided*), user-generated news impersonating real news (categorized as *imposter content*, hearsay (categorized as *rumors*, and incorrect content (categorized as *error*) [234].

In general, many of the categorizations in the literature incorporate the purpose and context of the misinformation, which are related to the misinformation supplier's motives rather than just the content's features. For a more precise analysis, I propose separating the purpose and context from the type of misinformation.

### 1.3.2 Proposed Typology

Wardle et al. (2017) define three elements of information disorder: the **misinformation supplier** who creates and disseminates the misinformation, the type of the misinformation message and its content, and finally, the **audience perceptions** of the message and the actions they take, if any, as a result of it [235]. These three elements correspond to the three phases where a countermeasure could intervene: the creation, the spread, and the belief in misinformation. These three main elements will organize my proposed misinformation typology.

#### 1. Misinformation Suppliers

Various types of actors create misinformation, whether intentionally or not. These actors may have different motivations. According to the literature, there are three main dimensions of actor types: official versus unofficial actors, level of organizational structure, and the use of automated technology [143, 235, 245, 247].

**Actor Types:** The different types of misinformation creators.

- *Official Organizations* - Official actors, including political parties, governments, corporations, and news organizations.
- *Unofficial Organizations* - Unofficial yet well-organized actors could include lobbying groups, terrorist or criminal networks, state-sponsored trolls, or groups of concerned citizens.
- *Individuals* - Includes journalists, influencers, or regular people.
- *Bots* - Agents that use automated technology generally can create and disseminate misinformation for cheaper and faster.

Actors have their own purposes and motivations. While some of these purposes can coexist (for example, one could promote misinformation for both a political agenda and financial reasons), this categorization helps to define motives. I propose the following primary purposes related to the creation and spread of misinformation, adapted from Wardle et al. (2017) [235] and Zannettou et al. (2019) [245].

**Purpose:** The purpose of a misinformation message lies in its intention and primary goal.

- *Political* - Political misinformation is intended to influence political attitudes and opinions. This type of misinformation can include intentionally increasing polarization, inciting violence, and feeding into extremism. This category addresses the *propaganda* and *biased* misinformation types proposed by many of the reviewed papers.
- *Financial* - This type of misinformation is intended to generate clicks and increase revenue, addressing the *advertising* and *clickbait* categories.
- *Distraction* - A message meant to distract the public with a different story, confuse, sow discord, or cause panic. Often, these messages are conspiratorial and intended to promote conspiracy theories, hoaxes, and rumors.
- *Accidental* - False information or context that spreads with no malicious intent. This spreading can occur in a variety of ways:
  - Individuals share something they believe others will correctly determine is a joke rather than something serious [235].
  - Individuals share something they truly believe is accurate [245]
- *Other* - False information created or spread for another reason. This could include sharing the misinformation for fun (trolling) or attention. Other reasons could include connecting with others in one's group [235, 245]

## 2. Misinformation Messages

The stylistic characteristics and news content in misinformation messages created by these actors can vary. We first consider the physical content associated with a piece of misinformation, which includes any text, images, videos, or audio related to the message, such as the headline and body of the article [247]. Actors can use a variety of techniques to generate content or misrepresent it. I propose the following categorization of the types of misinformation messages.

**Misinformation Types:** The type refers to how the message presents misinformation.

1. *Satire and Parody* - Humorous content that typically does not intend to cause harm, though it can fool some (e.g., The Daily Show, The Onion).
2. *False Connection* - Content containing headlines, captions, or images not supported by the rest of the content (e.g., clickbait).
3. *Misleading Content* - Misleading information or opinions presented as facts (e.g., cherry-picking, hyper-partisan news).
4. *False Context* - Correct information shared with false context (e.g., real images with incorrect captions or dates).
5. *Imposter Content* - Information posted while impersonating a genuine source or brand to gain credibility (e.g., the unofficial usage of an official logo or reputable individual's name).
6. *Manipulated Content* - Text, image, or video distortion (e.g., Photoshop, AI-generated images).

7. *Fabricated Content* - A false story, completely made-up.
8. *Error (False Content)* - Generally a mistake, often by a reputable news organization or honest person.

In addition to the characteristics that describe how the content is misleading, another important aspect of misinformation messages is their news content [247]. The content may be related to the difficulty of debunking misinformation, with some previous studies showing that political misinformation is among the most difficult to debunk [230].

**News Topics:** The topic indicates the main topic discussed in the message. Five general news categories and a well-known example of misinformation in that domain are listed below.

- *Political* - Articles on political figures, government policies, elections, and other newsworthy events. Example: “Pope endorses Trump in 2016 election”<sup>4</sup>.
- *Health* - Articles related to health and wellness. Topics include pandemics, illnesses, vaccines, and smoking. Example: “Garlic cures COVID-19”<sup>5</sup>.
- *Science and Technology* - Articles related to science and research. Topics include climate change, space, and evolution. Example: “The Earth is flat”<sup>6</sup>
- *Business and Finance* - Articles related to businesses, the economy, or consumer products. Example: “Corona beer sales in the U.S. plummet during the COVID-19 pandemic”<sup>7</sup>.
- *Entertainment* - Articles related to the entertainment or sports industries, including celebrities, movies, music, and athletes. Example: “Avril Lavigne died in 2003 and was replaced by a clone.”<sup>8</sup>.
- *Other* - Any other news topics.

### 3. Audience Perceptions

Misinformation suppliers may target either a general or specific audience to influence. The targeted users perceive the misinformation, and they may or may not believe it or take action because of it.

First, consider the **targeted groups**. Audiences are often targeted based on various demographic characteristics, including nationality, age, gender, race, sexuality, religion, or income. People may also be targeted based on their membership in a group, like a consumer group, non-profit, or company [235]. If we reconsider the main types of news topics and the motivation of misinformation creators and spreaders, Table 1.2 summarizes their likely targeted audience.

Next, we consider how the audience receives the message. Individuals determine whether to accept a message, either entirely or partially, based on their evaluations of the content and

<sup>4</sup><https://www.factcheck.org/2016/10/did-the-pope-endorse-trump/>

<sup>5</sup><https://www.factcheck.org/2020/02/fake-coronavirus-cures-part-2-garlic-isnt-a-cure/>

<sup>6</sup><https://www.reuters.com/fact-check/earth-is-not-flat-or-surrounded-by-an-ice-wall-2023-09-29/>

<sup>7</sup><https://www.factcheck.org/2020/03/coronavirus-fears-havent-sunk-sales-of-corona-beer-in-u-s/>

<sup>8</sup><https://www.forbes.com/sites/conormurray/2024/05/16/avril-lavigne-clone-conspiracy-explained-singer-laughs-off-false-rumor-heres-how-it-all-began/>

Table 1.2: Misinformation suppliers: their motivation, news context, and targeted audiences.

Supplier Motivation	News Context	Targeted Audience
Political	Political Health Science and Tech Business and Finance	Voters, citizens, funders
Financial	Health Science and Tech Business and Finance	Potential customers
Distraction	Any	General public, specific groups or people
Accidental	Any	General public, specific groups or people
Other	Any	General public, specific groups or people

supplier. For example, is the supplier a trusted messenger for the targeted individual? Is the language in the message highly emotional? According to previous research, individuals take multiple factors into account when evaluating content, including familiarity, simplicity, perceived credibility, motivated reasoning, and the persuasiveness of the message [143, 235]. Individuals have varying levels of misinformation susceptibility, often depending on these factors, which will be explored in later chapters.

**Message Reception:** Regardless of the reasons behind it, seminal work on reception theory posits that individuals perceive messages in one of three ways, enumerated below [95, 235].

1. Accept the message as is
2. Accept parts of the message
3. Decline the message, despite understanding it

Finally, we consider whether the targeted group or individual takes an action, specifically on social media. An individual could choose to ignore the message. They could like or comment positively if they agree with the message in part or in full, boosting the post's engagement metrics. They could additionally reshare the message to a larger audience, furthering its spread. If an individual disagrees with the message, they could report or flag the post, comment critically, or reshare the message in a critical manner. These actions are further summarized in Chapter 3.

**Actions:** The actions an individual can take on social media networks in response to being exposed to misinformation.

- Ignore the message
- Engage with the message positively (like or comment)
- Engage with the message negatively (comment, add context, report)
- Reshare the message

### 1.3.3 Summary

In summary, there are three elements to misinformation: the supplier, the message, and the audience. Table 1.3 summarizes the various aspects described in this categorization.

Table 1.3: Proposed misinformation categorization.

	Features	Examples
<b>Supplier</b>	Actor Types	Official or unofficial organizations, individuals, bots
	Purpose	Political, financial, distraction, accidental, other
<b>Message</b>	Type	Satire/parody, false connection, misleading content, false context, imposter content, manipulated content, fabricated content, errors
	News Topic	Political, health, science and technology, business and finance, entertainment, other
<b>Audience</b>	Targeted Group	Voters, citizens, funders, customers, general public, specific groups or people
	Message Reception	Full acceptance, partial acceptance, rejection
	Actions Taken	Ignore, engage with or share the content

## 1.4 Countermeasures Categorization

This section will first provide an overview and compare how researchers classify misinformation interventions. Then, using this prior work, I will outline a proposed categorization.

### 1.4.1 Literature Review

Several review papers and meta-analyses have been published in the field of misinformation interventions. Some reviews, such as Helmus and Keppe (2021) from the Rand Corporation, focus on relevant policy papers [100]. Others examine specific intervention categories, such as content moderation [108] or media literacy [106].

Four comprehensive review articles will be analyzed: Courchesne et al. (2021), Aghajari et al. (2023), Blair et al. (2024), and Kozyreva et al. (2024) [4, 35, 59, 129]. These articles were selected for their recency, the breadth of interventions covered, and their diverse disciplines. The Courchesne article primarily discusses platform interventions [59] and was published in the Harvard Misinformation Review, an interdisciplinary journal mainly associated with social sciences, as defined by the SCImago journal rank database<sup>9</sup>. The Aghajari article was published in the proceedings of a computer science conference [4]. The Blair article highlights research from both the Global North and the Global South, and was published in a psychology venue [35]. Finally, the Kozyreva article was published most recently with numerous high-profile scholars in this research area and appeared in the prestigious Nature Human Behavior journal [129]. Table 1.4 summarizes the primary interventions discussed by each paper in alphabetical order.

<sup>9</sup><https://www.scimagojr.com>

Table 1.4: Categorizations of misinformation interventions in the literature.

Type	Courchesne et al. (2021) [59]	Aghajari et al. (2023) [4]	Blair et al. (2024) [35]	Kozyreva et al. (2024) [129]
Account Moderation	×			
Account Removal		×		
Accuracy Prompts		×	×	×
Advertising Policy	×			
Content Distribution	×			
Content Labeling	×			
Content Moderation	×			
Context Labels			×	
Crowdsourcing		×		
Debunking			×	×
Fact-Checking	×	×		×
Friction			×	×
Inoculation			×	×
Journalist Training			×	
Lateral Reading/Verification				×
Media Literacy	×	×	×	×
Platform Alterations		×	×	
Politician Messaging			×	
Redirection	×			
Reporting	×			
Security/Verification	×			
Social Norms		×	×	×
Source Credibility Labels		×	×	×
Warning Labels		×		×
Other	×			

The Courchesne article classified types of interventions based on those publicly announced as having been implemented by various social media platforms [59]. These 10 categories were advertisement policy, content labeling, content/account moderation, content reporting, content distribution/sharing, disinformation disclosure, disinformation literacy, redirection, security/verification, and others. The article notes that fact-checking was included in the disinformation disclosure category, so this is categorized as fact-checking in Table 1.4.

The Aghajari article categorized interventions primarily based on the aspect of misinformation each intervention targeted: Content, Source, Individual Users, or Community [4]. For content-based interventions, they discussed fact-checking, warning labels, and platform alterations (reducing the size of misinformation). For source-based interventions, they discussed source credibility labels and the crowdsourcing of those labels. For user-based strategies, they considered media literacy, accuracy prompts, and account removal. Lastly, for community-based interventions, they addressed social norms.

The Blair article categorized 11 types of interventions into four groups: Informational, Educational, Sociopsychological, and Institutional [35]. For informational interventions, they defined inoculation/prebunking, debunking, credibility labels/tags, and contextual labels/tags. Under educational interventions, they defined media literacy. In the sociopsychological category, they defined accuracy prompts, friction, and social norms. Lastly, they defined platform alterations, politician messaging, and journalist training under institutional interventions.

Finally, the Kozyreva article classified nine types of interventions into three categories: Nudges, Boosts and Educational Interventions, and Refutation Strategies [129]. Under nudges, they included accuracy prompts, friction, and social norms. Under boosts, they included inoculation, lateral reading and verification strategies, and media literacy tips. Lastly, refutation strategies included debunking and rebuttals, warning and fact-checking labels, and source credibility labels.

Many of these categories overlap (e.g., lateral reading vs. media literacy) or are sub-categories of other categories. For example, redirection is a form of content distribution. Similarly, context labels and source credibility labels are types of content labeling. Additionally, specific categories are absent, especially in user-led or institutional interventions like government regulation.

User-based measures are an often overlooked aspect in the fight against misinformation. In addition to user reporting and social norms, users can block others or engage in social corrections. When misinformation is successfully posted on social media, other users serve as the first line of defense since they can flag or debunk it. While social corrections can be considered a type of debunking or fact-checking, the social context may be important to distinguish. Individual-level debunking, especially from trusted messengers, has been found to be effective in various contexts [21, 37, 144, 213, 233]. User-based countermeasures are addressed in greater detail in Chapters 3 and 4.

### 1.4.2 Proposed Typology

Many of the review articles used a similar categorization of countermeasures; however, there was no common typology [4, 35, 59, 129]. After reviewing the literature and these previous categorizations, I developed eight general categories of countermeasures, as shown in Table 1.5, and enumerated specific sub-categories of interventions, 1.6. The first five general categories could

apply to both platforms or governments; platforms could implement these changes willingly, or governments could require these changes.

The proposed interventions typology introduced in this section defines the general categories of misinformation interventions along with specific sub-categories. This typology is helpful in analyzing the literature in this field and will be used in the citation network analysis conducted in Chapter 2. See Appendix A for more details on each intervention and its associated references.

Table 1.5: Proposed categorization of misinformation interventions.

General Category	Example Interventions
Content distribution	accuracy prompts, friction, redirection
Content moderation	algorithmic downranking, fact-checking, remove posts
Account moderation	account removal, shadow banning
Content labeling	warning labels, source credibility labels, context labels
Media literacy and education	media literacy tips, inoculation
User-based measures	reporting users or posts, social corrections
Institutional measures	media support, data sharing, government regulation
Other	combining interventions, new interventions, generative AI

## 1.5 Current Platform Policies

The top social media platforms define and address misinformation differently, varying in what types of content they prohibit or limit and how they counter it. Reviewing current social media policies finds that all eight proposed general intervention categories are used in one or more platforms.

### 1.5.1 Community Guidelines

Before comparing how platforms respond to misinformation, we first investigate their community guidelines to determine what constitutes prohibited content. Platforms vary in how much offline harm false or misleading content can cause before determining whether to take action. Most platforms clearly state that misinformation with a serious risk of direct harm is prohibited. However, the two most prominent platforms, YouTube<sup>10</sup> and Meta<sup>11</sup>, focus their misinformation policies primarily on two news domains. Specifically, they focus on **health misinformation** (such as fake cures and misleading claims about vaccines and public health) and **election or civic misinformation** (such as false election dates, suppressing certain voter groups, or any effort intended to undermine electoral integrity).

Under its current leadership, X has loosened its policies and no longer appears to maintain a misinformation-specific policy. Instead, they enforce a policy against inauthentic content that could deceive users or cause serious harm<sup>12</sup>. However, they do not specify particular news con-

<sup>10</sup><https://www.youtube.com/howyoutubeworks/our-commitments/fighting-misinformation/>

<sup>11</sup><https://transparency.meta.com/policies/community-standards/misinformation>

<sup>12</sup><https://help.x.com/en/rules-and-policies#platform-integrity-and-authenticity>

Table 1.6: Proposed categorization of specific misinformation interventions.

General Category	Specific Intervention	Definition
Content Distribution	Content Distribution	The distribution of content on social media
	Redirection	Redirecting users to other content when searching
	Accuracy prompts	Reminding people about accuracy
	Friction	Pause and reflect before engaging with content
	Advertising policy	What ads are shown to which users
Content Moderation	Content Moderation	How content is shown or removed on social media
	Fact-Checking	Verification of information, often by experts
	Debunking	Fact-checking with context, narrative coherence
	Algo. Content Moderation	Algorithmic content moderation, like downranking
	Misinformation Detection	Automated detection of misinformation
Account Moderation	Account Moderation	Moderating through suspensions, bans, demonetization
	Account removal	The removal of a user from one or more platforms
	Shadow banning	Limiting the spread of posts from certain accounts
Content Labeling	Content Labeling	A type of misinformation disclosure through labels
	Crowdsourcing	Using regular people to verify and label information
	Warning Labels	General warnings about misinformation
	Source Credibility Labels	Disclosing or labeling a post's source
	Context Labels	Adding context to a post, like via community notes
Media Literacy and Education	Media Literacy	Improve the public's civic or digital reasoning
	Fake News Games	Games designed to help people detect misinformation
	Inoculation	Pre-bunking misinformation before exposure
User-based Measures	User-based Measures	How people respond to seeing misinformation
	Reporting	Users can report users or their posts
	Blocking	Users can block users or specific topics
	Social Corrections	Users that fact-check/debunk other users
	Social Norms	Using peer or community influence to change behavior
	Retractions	When accounts retract misinformation they posted
Institutional Measures	Institutional Measures	Measures by civic society, governments, or institutions
	Media Support	Investing in local news, journalist training
	Data Sharing	Sharing high-quality data with researchers
	Government Regulation	Any relevant laws, rules, or regulations
Other	Other	New interventions or those not fitting any category
	Generative AI	The usage of gen AI to counter or detect misinfo
	Combining Interventions	Using multiple interventions at once

texts that would be targeted, such as health-related or election misinformation.

Other platforms, such as Pinterest<sup>13</sup> and TikTok<sup>14</sup>, take a broader approach. Pinterest states that it limits or moderates content that may harm someone's "well-being, safety or trust." Including the erosion of trust as a potential form of harm is significant, as this reflects a more comprehensive strategy than most platforms. Additionally, both Pinterest and TikTok moderate election and health misinformation like other platforms, but they also include **science misinformation**, such as climate change denial, as content that is not permitted.

YouTube, Meta, and TikTok also require that AI-manipulated content that is realistic in nature or has the potential to mislead is labeled as such. All platforms, including X, prohibit scams and inauthentic accounts. Given that the platforms primarily focus on health and political misinformation as the most potentially dangerous forms worth addressing, this dissertation will focus primarily on these two types of misinformation.

### 1.5.2 Enforcement

In addition to differing community guidelines surrounding misinformation, platforms implement a range of interventions, engaging in nearly all the intervention types described in Table 1.7. YouTube outlines four primary strategies they use for combatting misinformation<sup>15</sup>:

- **Reducing the spread** of borderline violating content through interventions like algorithmic downranking and automated misinformation detection.
- **Increasing the distribution of high-quality content** by elevating authoritative sources and using fact-checking, warning labels, and source credibility labels.
- **Rewarding trusted accounts** with monetization policies and demonetizing those who violate community guidelines.
- **Giving users control** through reporting tools, blocking and filtering abilities, media literacy initiatives, and institutional measures like media support.

According to Meta's Community Standards<sup>16</sup>, they also employ misinformation detection and content distribution strategies to slow the spread of hoaxes, provide resources to improve media literacy, and require disclosure of AI-generated content. However, they recently removed fact-checking in the U.S. and replaced it with context labels [114].

Like YouTube and Meta, TikTok uses multiple intervention strategies. TikTok implements algorithmic content moderation, uses professional fact-checkers, labels government accounts, and requires that realistic AI-generated content is labeled<sup>17</sup>. Meanwhile, Pinterest focuses on removing or limiting the distribution of violating content, and they regularly release transparency reports<sup>18</sup>.

<sup>13</sup><https://policy.pinterest.com/en/community-guidelines>

<sup>14</sup><https://www.tiktok.com/transparency/en-us/combatting-misinformation/>

<sup>15</sup><https://www.youtube.com/howyoutubeworks/our-commitments/fighting-misinformation/>

<sup>16</sup><https://transparency.meta.com/policies/community-standards/misinformation>

<sup>17</sup><https://www.tiktok.com/transparency/en-us/combatting-misinformation/>

<sup>18</sup><https://policy.pinterest.com/en/community-guidelines>

## 1.6 Operationalized Interventions

Building on the proposed interventions typology in Section 1.4.2, this section examines how governments, platforms, or institutions implement these interventions in practice. The categorized interventions were operationalized by reviewing the literature and platform policies, and finding examples for each intervention type that had been implemented or proposed before [3, 192].

Table 1.7 outlines interventions primarily implemented by platforms and includes those focusing on the **Content Distribution**, **Content Moderation**, **Account Moderation**, and **Content Labeling** categories. Table 1.8 outlines interventions that are primarily led by users and institutions, and includes those focusing on the **User-based Measures**, **Media Literacy and Education**, **Institutional Measures**, and **Other** categories. The operationalized interventions are compared throughout the remainder of the dissertation. References for each intervention can be found in Appendix A.

## 1.7 Data

This thesis creates and uses five primary datasets. These datasets will be used to characterize countermeasures while considering several relevant factors. All datasets are summarized in Table 1.9.

Table 1.7: Operationalized platform interventions.

Specific Intervention	Intervention Implementation
Redirection	<b>Redirection</b> - Redirect users to other content, such as official content or no content, when searching for something potentially problematic or harmful
Accuracy Prompts	<b>Accuracy Prompts</b> - Nudge or remind people to consider accuracy before posting or sharing content, such as by being asked to rate the accuracy of the headline
Friction	<b>Friction</b> - Temporarily delay users from posting content they did not open via a pop-up
Platform Alterations	<b>Platform Alterations</b> - Reduce the size or visibility of a post containing misinformation <b>Ban political ads</b> - on social media platforms <b>Fact-check ads</b> - on social media platforms <b>Limit forwarding</b> - Cap the number of users one can forward a given message <b>Limit resharing</b> - Remove share buttons on posts after several levels of sharing
Advertising Policy	
Advertising Policy	
Content Distribution	
Content Distribution	
Fact-Checking, Debunking	<b>Fact-check labels</b> - Link information from 3rd party fact-checkers on misinformation posts
Misinfo. Detection	<b>Content removal</b> - Remove posts verified to contain misinformation.
Alg. Content Moderation	<b>Downranking</b> - De-emphasize posts in news feeds that contain misinformation
Alg. Content Moderation	<b>Algorithmic changes</b> - uprank high-quality news sources and downrank low-quality ones
Alg. Content Moderation	<b>Virality circuit breakers</b> - Automatically flag certain fast-spreading and unverified content, triggering a brief halt on algorithmic amplification until the information is verified
Content Moderation	<b>User control</b> - Give users more control over the algorithms powering their news feeds
Account Removal	<b>Account suspensions</b> - Permanently or temporarily ban users who post misinformation or violate other platform policies a certain number of times
Account Removal	<b>Deplatforming</b> - Coordinated efforts among several social media companies to remove especially problematic or dangerous user accounts.
Shadow Banning	<b>Shadow banning</b> - Limit the spread of posts from certain policy-violating accounts without explicitly banning or suspending them.
Demonetization	<b>Demonetization</b> - The removal or restriction of monetization features for a user account found repeatedly violating a platform's policies.
Crowdsourcing, Context Labels	<b>Crowdsourcing</b> - Labels generated by the general public rather than by professional fact-checkers, like Community Notes programs
Warning Labels	<b>Click-through warning labels</b> - Place misinformation posts behind click-through labels containing facts and context
Source Credibility Labels	<b>Source labeling</b> - Labels indicating the reliability of news sources using, for example, NewsGuard labels
Source Credibility Labels	<b>Government labels</b> - Label the accounts of government officials or state-run media
Content Labeling	<b>Misinformation disclosure</b> - Notify users if they posted content verified to contain misinformation.

Table 1.8: Operationalized user-based and institutional interventions.

Specific Intervention	Intervention Implementation
Reporting	<b>Reporting</b> - Improved reporting functionality and transparency, to encourage users to report misinformation and harassment
Social Norms, Social Corrections, Retractions	<b>Social norms</b> - The usage of social or community norms to encourage social corrections and self-corrections via platform changes
Social Norms	<b>Alter platform metrics</b> - Reward accuracy rather than engagement to discourage misinformation sharing
Media Literacy	<b>Digital Media Literacy</b> - Invest in and promote educational content on how to critically evaluate online information
Fake News Games, Inoculation	<b>Inoculation</b> - Inoculating people against misinformation with games or videos (For example, the Bad News Game)
Media Support	<b>Local Media Support</b> - Promote and invest in local media, which is thought to be most in tune with local norms, culture, and context.
Media Support	<b>Journalism support</b> - Supporting and training the next generation of journalists to engage in high-quality and independent reporting
Data Sharing	<b>Data Sharing</b> - Have social media companies regularly release data and/or internal research reports to the public and outside researchers.
Government Regulation	<b>Government Regulation</b> - Hold companies accountable for the content shared on their platforms. This could involve modifying Section 230, or regulating social media companies
Government Regulation	<b>Privacy Legislation</b> - The development of comprehensive privacy legislation, similar to Europe's GDPR
Government Regulation	<b>Anti-Trust Action</b> - Breaking up of monopolistic big tech companies
Government Regulation	<b>Taxes/Fines</b> - Taxing or fining social media companies for their use of personal user data
Gov or platform regulation	<b>Targeted advertising</b> - Limiting or banning micro-targeted advertising
Generative AI	<b>AI chatbots</b> - The use of generative AI chatbots to reduce belief in conspiracy theories or misinformation
Generative AI	<b>Gen AI Content</b> - The use of generative AI to generate rebuttals to misinformation or create educational initiatives
Generative AI	<b>Deepfakes</b> - Prohibit usage of AI or manipulated content to misrepresent the speech or actions of public figures
Generative AI	<b>AI Ads</b> - Prohibit usage of AI or manipulated content in political or targeted ads
Generative AI	<b>AI Disclosure</b> - Require clear disclosures on any ad that uses AI-generated images, video, or audio

Table 1.9: Summary of datasets used in this thesis.

Data	Size	Dates Collected	Ch. 2	Ch. 3	Ch. 4	Ch. 5	Ch. 6
Literature Corpus	412 papers	2004-2025	✓				✓
Countermeasures Survey	1010 responses	12 Jul to 23 Aug 2024		✓		✓	✓
Training Posts	40 posts	Various			✓		
Training Quizzes	23 responses	5 Feb 2024			✓		
Expert Opinions Survey	38 responses	23 Jan to 4 Mar 2025					✓

## Literature Corpus

The *Literature Corpus* data is a collection of papers primarily pulled from four prominent review papers and additional keyword searches to analyze gaps in the literature. The citation network analysis of this dataset is conducted in Chapter 2. Additionally, this data is summarized alongside the other data sets in Chapter 6.

## Countermeasures Survey Data

The *Countermeasures Survey* data was collected during the summer of 2024. Behavioral and opinion-based questions were asked of 1010 American-based social media users. This data set includes standard demographic questions, as well as questions about behaviors and opinions related to user-based, platform-based, and government-level countermeasures. The questions related to user-based interventions are analyzed in Chapter 3, while those related to platform and government interventions are analyzed in Chapter 5. Additionally, this data is summarized alongside the other data sets in Chapter 6.

## Social Media Training Posts

For Chapter 4, we curated a set of *Social Media Training Posts* for use in the pre- and post-training quizzes. These posts include a variety of misinformation, conspiracy theories, pink slime, and accurate news items to assess participants' misinformation detection and countering skills. We used fact-checking websites, such as FactCheck.org, to identify relevant misinformation posts or searched social media platforms directly for content on specific topics. We also reviewed existing COVID-19 Twitter datasets previously collected by the CASOS research group. Most of these posts came from Twitter and Facebook.

## Pre and Post-Training Quizzes

The *Training Quiz* data includes responses from both a pre-training and post-training quiz collected from 23 government analysts. The respondents completed a pre-quiz to assess their knowledge of misinformation detection and countering ability. They then underwent relevant training. Finally, they took a post-training quiz to see if there was any improvement in their detection abilities. This data set will be used primarily for Chapter 4, but will also be summarized in Chapter 6.

## **Expert Opinions Survey**

The *Expert Opinions* data set is a survey of misinformation researchers designed to gather their opinions on all operationalized countermeasures defined and analyzed in this thesis, along with the characteristics of those countermeasures, including effectiveness, acceptance, effort level, political feasibility, and cost. This survey was designed for analysis in Chapter 6.

## **1.8 Chapter Overview**

This thesis aims to establish a framework for assessing and developing countermeasures to misinformation that are both practical and effective. A citation network analysis is conducted in Chapter 2 to better understand the current state of the literature and find research gaps. User-based misinformation interventions are examined in Chapters 3 and 4. Chapter 3 outlines current user opinions and behaviors on this subject and considers ways this information could be leveraged to improve user-based countermeasures. Chapter 4 builds directly on Chapter 3 and focuses specifically on improving media literacy interventions, a type of user-based intervention, to determine if they can be adapted to improve countering abilities (not just misinformation detection abilities) among a highly skilled audience. Chapter 5 discusses platform and government interventions as well as the factors most associated with public support for these interventions: perceived fairness, effectiveness, and intrusiveness of the interventions. Chapter 6 synthesizes information from the previous chapters and a survey conducted with misinformation experts to provide analysis-driven recommendations.

# Chapter 2

## Citation Network Analysis of the Interventions Literature

To supplement the literature review in the Introduction, I conducted a systematic scoping review and a citation network analysis of relevant papers to generate a more comprehensive contextual background on the current literature in this field.

In this chapter, I begin with four prominent review papers published in different research areas and gather all relevant papers they reviewed, resulting in a total of 412 papers. This review analyzes the citation network and other metrics to provide policymakers and academics with relevant information on the current state of the literature in this domain and identify any research gaps. It will also provide guidance for what should be focused on in the remaining chapters of this dissertation.

The main research question for this chapter is: How are interventions currently studied in the academic literature? More specifically:

1. Which countermeasures are under-studied or over-studied in the literature?
2. What types of impacts is the literature examining? Impacts include effectiveness and user acceptance.
3. Are researchers working independently or collaborating across various disciplines?

### 2.1 Introduction

While several review articles in this field were discussed in the Introduction, most review articles focus on testing countermeasures and analyzing their effectiveness. However, they overlook the equally important metrics of practicality and acceptability to users. The Courchesne and Blair review articles explicitly focused only on experimental studies and excluded nonexperimental papers [35, 59]. Similarly, the Kozyreva article included only empirical studies, providing evidence of the efficacy of the interventions studied [129]. Only Aghajari included papers that simply presented an intervention rather than requiring all studies to examine their effectiveness [4].

Although the four selected review articles were among the broadest in the literature, their fo-

cus on effectiveness likely limited the types of interventions covered to predominantly platform-based rather than institutional or user-led interventions. I aim to fill this gap by considering user-level, platform-level, and policy interventions. This review begins with these four prominent review articles as seed papers and then supplements them with additional work from underrepresented intervention areas. This review aims to gain insights into the misinformation intervention space through a bibliometric analysis of the citation network of relevant articles.

## 2.2 Data and Methods

To identify under-studied interventions and factors and answer these research questions, we conducted a scoping literature review, a type of systematic review that is broader in nature. We followed the methods used in two previous computer science review articles focusing on interventions [4, 244]. More specifically, we adhered to the modified PRISMA guidelines for scoping review [219], as well as the more specific guidelines for systematic reviews in the information systems field developed by Okoli [166].

Given our research questions and objectives, we developed a protocol for this review that outlined the steps and procedures we would carry out in advance. We developed the intervention labels and eligibility criteria and defined our literature search strategy. Next, we trained the labelers, labeled the papers, and synthesized the results.

### 2.2.1 Paper Labels

We developed a comprehensive list of 35 specific intervention categories defined earlier in Table 1.6. Refer to Appendix A for more detailed definitions and references for each specific intervention. Note that each general category was always included as a specific intervention label. This label is applied when a paper studies an intervention that falls within that general category but is not explicitly defined by the other specific intervention categories. In addition to the specific intervention labels, we defined four additional labels:

- *Review Article*: A paper that reviews other papers in a specific field.
- *Meta-Analysis*: A review paper that quantitatively analyzes previous results.
- *Effectiveness*: A paper that studies and analyzes the effectiveness of one or more interventions at reducing the creation, spread, or belief in misinformation.
- *Acceptance*: A paper that studies the user acceptance or incorporates user feedback in the design and analysis of misinformation interventions.

Each paper was labeled with the countermeasures it discussed or analyzed. Papers were also assigned the review article, meta-analysis, effectiveness, and acceptance labels where appropriate. It is important to note that these labels are not mutually exclusive, as some papers can cover multiple interventions. Furthermore, although all interventions are assigned to one general category, some may apply to multiple categories. For example, fact-checking is a type of content moderation that is sometimes implemented as content labels.

## 2.2.2 Inclusion Criteria

We intentionally aimed to be as inclusive as possible in order to capture a wide variety of interventions and factors studied. To be included in our analysis, a paper must meet the following criteria:

1. **Content:** One of the article's main focuses should be interventions or countermeasures to misinformation. The paper does not need to address social media specifically but must primarily concentrate on interventions. The included articles could directly test the efficacy of one or more interventions through experimental studies or could be a review or discussion-based article.
2. **Article Type:** The article is a research article. It is not an opinion piece or simply an abstract.
3. **Venue:** The paper should come from a reputable venue or institution, but inclusion is not restricted to peer-reviewed publications only. To ensure quality, we enumerated the types of venues that could be included:
  - *Peer-reviewed Articles* - Papers from peer-reviewed indexed journals and conference proceedings, including workshop papers if they are in a proceedings.
  - *Technical Reports* - Reports from reputable and high-quality institutions like think tanks, non-profits, research centers, or governmental organizations. To ensure quality, these institutions should either be non-profit or affiliated with a university, foundation, or government, rather than being private or highly partisan organizations.
  - *Preprints* - Preprints posted in 2020 or later were included. Preprints are especially common in the fast-moving field of computer science. Preprints from before 2020 were excluded unless they had later been published in a peer-reviewed venue.
4. **Publication Date:** The paper was published in 2004 or later, since we primarily want to focus on the last twenty years of research in the social media era.
5. **Language:** The paper is written in English or translated into English. This criteria was due to a limitation of our labelers.

## 2.2.3 Literature Search

Our literature search was conducted in two stages. Figure 2.1 outlines our literature selection and review process.

First, we gathered all the studies analyzed by our four “seed papers.” We removed duplicate papers, applied the exclusion criteria, and labeled the papers according to the definitions provided in Section 2.2.1 and Table 1.6. This process resulted in 365 labeled papers by the end of Stage 1.

After labeling the initial set of papers, we proceeded to the second stage of the literature search. Using Scopus, we searched for any intervention papers published in 2024 or 2025 through keyword searches. Paper titles needed to include either the words “misinformation” or “disinformation” and either “intervention” or “counter\*”. Additionally, we searched for any

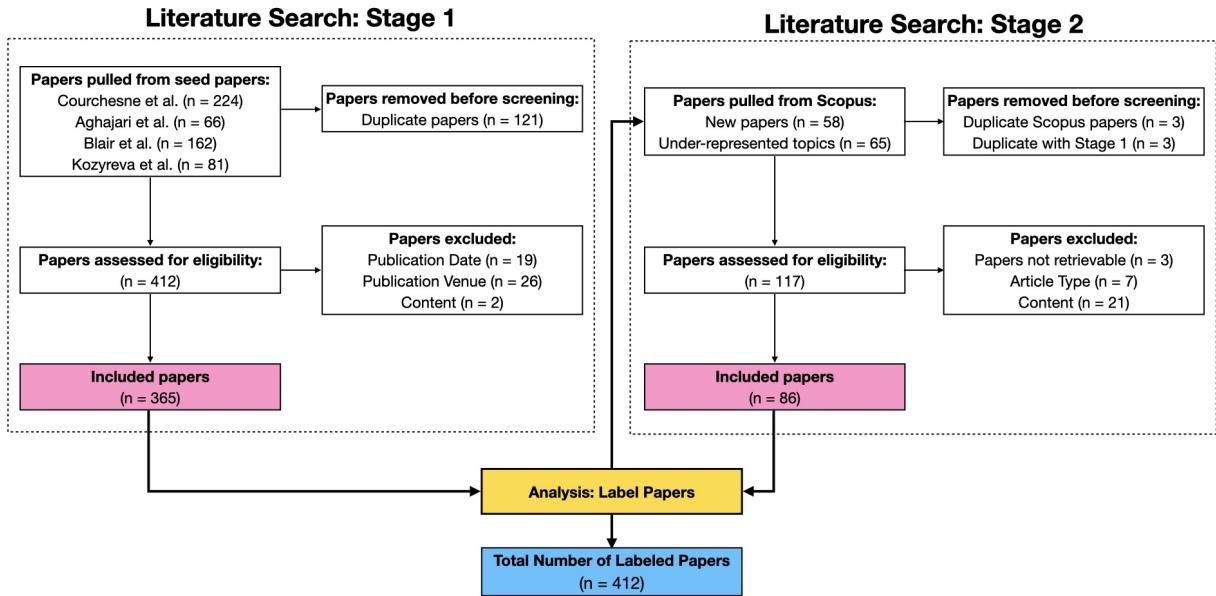


Figure 2.1: Literature review process.

specific intervention that had been labeled in fewer than 10 of the papers in the initial set. Table 2.1 displays the under-studied interventions and the associated keywords used in our Scopus keyword search. All these keywords were used alongside the words “misinformation” or “disinformation” in the title and “intervention” or “counter\*” in the title or abstract, limited to papers published after 2004. This search did not retrieve any papers for shadow banning and data sharing, prompting us to redo the same keyword search without requiring the words “intervention” or “counter\*” to be included. After applying the exclusion criteria and labeling the newly added papers, this process resulted in 86 labeled papers by the end of Stage 2, bringing the total to 412 labeled papers in our dataset.

## 2.2.4 Training of Labelers

Three advanced high school interns and ChatGPT assisted us with labeling for this project. The interns received individual training, as they began at different times in the summer of 2024. A training PowerPoint was created and shared with the interns for their reference. A summary of the training is provided below:

- **Background and Motivation:** Interns received an overview of misinformation and countermeasures. The literature review in Sections 1.3 and 1.4 was summarized.
- **Project Overview:** We discussed the overall project goals and research questions, and reviewed a summary of the 8-step systematic literature review guide as defined by Okoli [166].
- **Reading Academic Papers:** Interns received training on effectively reading and understanding academic papers using the CIMO framework and the SQ3R method for reading

Table 2.1: The intervention labels used the least after Stage 1, sorted from highest to lowest based on the number of papers with those labels.

Intervention Label	Scopus Keywords Used
Context Labels (9)	(“context” AND “label*”) OR (“community” AND “note”)
Algorithmic Content Moderation (8)	“downranking”
Advertising policy (5)	“advertising”
Redirection (5)	“redirect*”
Media Support (3)	“local news*” OR “media support”
Account Removal (2)	“deplatform*”
Reporting (2)	“user” AND “reporting”
Shadow Banning (0)	(“shadow*” AND “ban”) OR “monetiz*”
Blocking (0)	“blocking”
Data Sharing (0)	“data sharing”
Government Regulation (0)	(“government” AND “regulation”) OR (“government policy”)
Generative AI (0)	“gen* AI” OR “chatbot”

comprehension[166].

- *CIMO Framework* - The Context-Interventions-Mechanisms-Outcomes framework outlines the primary questions they should consider when reading the papers. This framework considers what is the context and systems that are involved, what is the event or intervention being studied, what is the relationship between the intervention and the outcomes, and finally, what the outcomes of the intervention are [38].
- *SQ3R Method* - The Survey, Question, Read, Recall, and Review technique is valuable for reading comprehension. During the survey step, interns were instructed to skim and scan relevant parts of each paper, including the title, abstract, keywords, and, if necessary, the introduction and conclusion, to determine whether the paper should be included. In the questioning step, they were asked to evaluate the paper’s relevance while keeping the CIMO framework in mind. Next, they were taught to read and make connections with previous papers. Finally, they were asked to recall the paper’s main points by reviewing their notes and writing a brief summary of those points [190].
- **Reference Manual:** Interns were given a checklist describing the inclusion and exclusion criteria, along with detailed definitions for all the labels defined in Section 2.2.1.
- **Tabular Documentation:** Interns were provided with a Google Sheets document to summarize the main findings, specify the interventions studied, indicate whether each paper studied effectiveness and acceptance, and determine if each was a meta-analysis or review article, among other items.

## 2.2.5 ChatGPT Prompts

ChatGPT prompts were designed to assist with labeling the specific interventions studied and whether the papers studied effectiveness or acceptance. These prompts were developed and

refined after all human labelers completed their labeling tasks. We applied standard prompt-engineering best practices to create the prompts, including clearly articulating the task, experimenting with different phrasing and techniques, and testing various model versions<sup>1</sup>. Refer to Appendix B for the final prompts.

Twenty Stage 1 papers were chosen as a test set while refining the prompt. These papers were selected to span as many intervention categories as possible. For the effectiveness and acceptance labeling tasks, the test prompts were executed using the 4o-mini and 4o models. Both models achieved the same accuracy for the effectiveness task, agreeing on 19 of the 20 papers. For the acceptance task, the 4o-mini model agreed with the final label on 17 papers, while the 4o model agreed on 18 papers. Cohen's kappa values could not be calculated because the categories were overwhelmingly unbalanced (all papers studied effectiveness but not acceptance.)

While the 4o model had a slightly higher agreement for acceptance, it is significantly more computationally intensive and financially costly. In February 2025, when these models were being run via the API, OpenAI charged \$0.15 per one million input tokens and \$0.60 per one million output tokens for the 4o-mini model. However, they charged \$2.50 per one million input tokens and \$10 per one million output tokens for the 4o model<sup>2</sup>. A token roughly equals four English characters or three-quarters of one word<sup>3</sup>. Considering we were using ChatGPT to process academic papers that are typically thousands of words long, the 4o model was avoided whenever possible. Therefore, the effectiveness and acceptance prompts were run on the rest of the papers using ChatGPT's API 4o-mini model.

For the intervention labeling task, there were 35 potential labels, which created significant room for disagreement. We improved the intervention prompt by making the definitions concise and explicitly stating at the beginning of the prompt that interventions should be labeled only as defined in the prompt. We tested three primary techniques to improve agreement with the final labels.

1. **Separate Intervention Prompts** - Creating eight prompts, one for each general category. These results were the worst, regardless of how the prompts were phrased. Using this technique, ChatGPT tended to assign too many labels to the papers. This over-labeling may have occurred because many of the general categories overlap (for example, fact-checking could be considered a content moderation technique or a content labeling technique).
2. **Alphabetical Combined Prompt** - Developing a single prompt and arranging the interventions in alphabetical order. These results represented an improvement over the separate intervention category prompts.
3. **Ordered Combined Prompt** - Creating a single prompt and ordering the interventions by category, much like how they are organized in Table 1.6. These results were the best.

The intervention labeling was run using both 4o-mini and 4o for each of the three techniques. Table 2.2 summarizes the twenty selected papers, their final labels, and the Jaccard similarity comparing the final labels with each model run. The ordered combined prompt demonstrated the highest average performance, with an overall average Jaccard similarity of 0.443 for the 4o-mini model and 0.579 for the 4o model, respectively. Given the substantial improvement provided by

<sup>1</sup><https://www.coursera.org/articles/what-is-prompt-engineering>

<sup>2</sup><https://platform.openai.com/docs/pricing>

<sup>3</sup><https://help.openai.com/en/articles/4936856-what-are-tokens-and-how-to-count-them>

the 4o model over the 4o-mini model, the 4o model was employed to label interventions on the rest of the papers.

## 2.2.6 Label Agreement

Approximately 25% of the included Stage 1 papers were assigned to two human labelers (101 papers) based on a randomly generated number. The remaining papers (265) were assigned to one human labeler and ChatGPT as the second labeler. There were five unique raters: three interns, myself, and ChatGPT. A collaborator, Peter Carragher, and I resolved any disagreements to determine the final assigned labels.

Given the varying number of labelers and potential labels per paper, the Jaccard similarity was calculated for both inter-rater reliability metrics and agreement with the final labels. This calculation corresponds to a straightforward percentage agreement for effectiveness and acceptance labeling.

- **Effectiveness - Average Inter-Rater Jaccard Similarity:** 0.956
- **Acceptance - Average Inter-Rater Jaccard Similarity:** 0.855
- **Interventions - Average Inter-Rater Jaccard Similarity:** 0.545

Agreement with the final labels varied by rater but was consistently high. Table 2.3 displays the Jaccard similarity agreement with the final label categorized by labeler type (ChatGPT, Expert, Intern). On average, ChatGPT matched or surpassed the agreement level of the interns on the intervention labeling task.

However, there was variability among the three interns. Table 2.4 shows the Jaccard similarity agreement with the final labels for each individual rater. ChatGPT had the second highest level of agreement with the final intervention labels, though it was slightly behind the interns on both effectiveness and acceptance. Interns are numbered in order of start date and overall internship duration, with Intern 1 starting and being trained first. Intern 3, who stayed on the project the longest, was by far the most accurate of the interns, indicating that experience likely improved their labeling accuracy.

Considering ChatGPT's relatively high level of agreement with the final intervention labels, the ChatGPT prompts were used to assist with labeling the 86 papers added in Stage 2 of the literature review. I reviewed ChatGPT's labels prior to assigning the final labels to these papers.

## 2.2.7 Data Availability

Here is the list of all reviewed papers and their labels on a public Zotero repository: **TO DO**.

## 2.3 Citation Network Analysis

As Figure 2.1 summarizes, a total of 412 papers were labeled following both the Stage 1 and Stage 2 literature searches. All papers were imported into Zotero, a citation management tool.

Table 2.2: Summary of ChatGPT interventions labeling by different prompt and model types. The values indicate Jaccard similarity scores that compare the labeled output produced by the model run against the final label.

<b>Paper Ref.</b>	<b>Final Labels</b>	<b>Separate Prompts</b>		<b>Alphabetized Prompt</b>		<b>Ordered Prompt</b>	
		<b>4o-mini</b>	<b>4o</b>	<b>4o-mini</b>	<b>4o</b>	<b>4o-mini</b>	<b>4o</b>
[5]	advertising policy, debunking	0	0.25	1	0.5	1	0.5
[11]	accuracy prompts, crowdsourcing, media literacy	0.176	0.3	1	1	0.667	1
[21]	social corrections, source credibility labels	0.111	0.143	0.333	0.5	0.5	0.5
[22]	friction	0	0	0	0	0	0
[36]	redirection	0	0	0	0.333	0.333	1
[36]	algorithmic content moderation, social corrections	0.111	0.222	0.333	1	0.333	0.333
[72]	retractions, source credibility labels	0.25	0.5	0.25	0.25	0.5	0.333
[75]	combining interventions, debunking, social norms	0.333	0.5	0.25	0.25	0.25	0.25
[77]	accuracy prompts	0.143	0.167	1	0.5	0.5	1
[79]	context labels, source credibility labels	0	0	0	0.333	0	0
[80]	fact-checking, inoculation	0.111	0.4	0.667	0.667	0.667	1
[111]	media literacy	0.25	0.5	0.5	1	1	1
[120]	platform alterations, social norms	0.111	0.5	0.333	0.5	0.333	0.33
[146]	misinformation detection, warning labels	0.067	0.1	0.25	0.25	0.333	0
[147]	fake news games, inoculation	0.167	0.25	0.25	0.333	0.25	0.333
[161]	source credibility labels, warning labels	0.077	0.125	0	0	0	0.333
[205]	friction, platform alterations, warning labels	0.077	0.125	0.2	0.667	0.2	0.667
[216]	account removal	0.125	0.143	1	1	0.5	1
[222]	media literacy	0.091	0.143	0.5	1	1	1
[231]	fact-checking	0.125	0	0.5	1	0.5	1
<b>Avg. Jaccard Similarity:</b>		0.116	0.218	0.418	0.554	0.443	<b>0.579</b>

Table 2.3: Jaccard Similarity agreement with the final labels broken up by rater type. Rater types are sorted by total number of papers labeled.

<b>Rater</b>	<b>Number Labeled</b>	<b>Final Effectiveness</b>	<b>Final Acceptance</b>	<b>Final Interventions</b>
Interns (3)	390	0.971	0.920	0.667
ChatGPT	264	0.962	0.864	0.746
Expert	76	0.987	0.974	0.938

Rater	Number Labeled	Final Effectiveness	Final Acceptance	Final Interventions
ChatGPT	264	0.962	0.864	0.746
Intern 3	201	0.970	0.930	0.737
Intern 2	96	0.979	0.948	0.668
Intern 1	93	0.968	0.871	0.515
Expert	76	0.987	0.974	0.938

Table 2.4: Jaccard Similarity agreement with the final labels broken up by specific raters. Raters are sorted by total number of papers labeled.

The final effectiveness, acceptance, and intervention labels were assigned as tags to each paper. The included papers were exported as a .bib file and loaded into ORA, where further data cleaning occurred. Authors who may have published under different names or sometimes omitted middle initials on their work were consolidated into a single author entry. Conference proceedings from different years were aggregated into one venue with the conference title for a more straightforward analysis. This process resulted in 1,147 unique authors and 209 unique publication venues from these 412 papers. The ORA software was used to conduct the citation network analysis and produce network visualizations from this set of papers [50]. The additional data attributes, nodesets, and networks are summarized below.

### 2.3.1 Data Attributes

Each paper’s meta-data was augmented with citation counts obtained from Semantic Scholar on February 25th, 2025<sup>4</sup>. These citation counts were retrieved using the Zotero Citation Counts Manager plug-in<sup>5</sup>. Discipline information for all 209 venues was pulled from Scimago’s journal rank data from 2023<sup>6</sup>. By far, the most popular areas were the social sciences ( $n = 117$ ), computer science ( $n = 57$ ), and psychology ( $n = 50$ ). Only 34 venues were exclusively from a discipline that was not one of the three main disciplines or a multidisciplinary journal. To reduce the number of discipline tags, any discipline not among the primary three or multidisciplinary was labeled as “other.”

### 2.3.2 Nodesets

There were four nodesets in this dataset.

- *Article* ( $n = 451$ ) - The Article nodeset contains the article title, publication year, and DOI extracted directly from the paper’s metadata. Additional attributes pulled into ORA are citation counts (sourced from Semantic Scholar) and whether the paper studies effectiveness (Yes/No) or acceptance (Yes/No).
- *Author* ( $n = 1147$ ) - The Author nodeset contains the unique author names. It additionally contains the following attributes: citation counts (summed over the Article citation counts

<sup>4</sup><https://www.semanticscholar.org>

<sup>5</sup><https://github.com/eschnett/zotero-citationcounts>

<sup>6</sup><https://www.scimagojr.com/journalrank.php>

for each author), the number of articles per author, and the average number of citations per article for each author.

- *Publication Venue* (n = 209) - The Publication Venue nodeset includes the venue title, the ISSN, and the venue type (journal, conference proceedings, or book). Discipline is an externally sourced attribute from SCImago, representing the primary research areas that the venue publishes. Venues can be classified as belonging to more than one discipline area.
- *Topic* (n = 35) - The Topic nodeset includes 35 unique intervention labels. Additional attributes are the General Intervention Category and Count (the total number of papers associated with each intervention label).

### 2.3.3 Networks

ORA creates multiple networks from this data. The networks that are analyzed further are described below:

- *Publication Venue x Publication Venue (Co-Authorship)* - This symmetric network is auto-generated by ORA. A link exists between two venues if there is an author who has published in both venues. Link values are weighted and represent the number of authors who have published in both venues.
  - Network Density: 0.0349
  - Link Statistics: 1516 total links, values range from 1-12, mean value is 1.89
  - Component Statistics: 73 isolates, 6 dyads, 2 triads, 2 larger components
  - Largest Component: 114 authors
- *Topic x Topic (Co-Topic)* - This symmetric network is auto-generated by ORA. A link exists between two topics if there is an article that covers both topics. Link values are weighted and represent the number of papers that study both intervention topics.
  - Network Density: 0.396
  - Link Statistics: 472 total links, values range from 1-26, mean value is 2.71
  - Component Statistics: 1 isolate (data sharing), 1 large component
  - Largest Component: 34 topics
- *Author x Author (Co-Authorship)* - This symmetric network is auto-generated by ORA. A link exists between two authors if they have published together. Link values are weighted and indicate the number of papers the authors have co-authored together.
  - Network Density: 0.005
  - Link Statistics: 6654 total links, values range from 1-19, mean value is 1.12
  - Component Statistics: 24 isolates, 53 dyads, 48 triads, 93 larger components
  - Largest Component: 326 authors
- *Publication Venue x Topic* - This network was created by folding the *Publication Venue x Article* and *Article x Topic* auto-generated by ORA. The links represent the number of

papers published in that venue discussing each intervention topic.

- Network Density: 0.082
- Link Statistics: 603 total links, values range from 1-7, mean value is 1.29
- Component Statistics: 1 large component
- Largest Component: 244

## 2.4 Results

In this section, we use descriptive statistics and network analysis to report on the growing number of papers over time, as well as to examine leading venues, topics, and authors.

### 2.4.1 Descriptive Statistics

The ORA Report ‘‘Bibliography and Citation’’ was run on this dataset. This report provides network metrics and analyzes leading authors, venues, and topics. Figure 2.2 shows the number of papers in this review published each year. There is an exponential growth of articles in this field. The spike in 2021 and the subsequent slight drop-off can likely be attributed to the fact that our seed paper that reviewed the most articles originated from a 2021 review article [59].

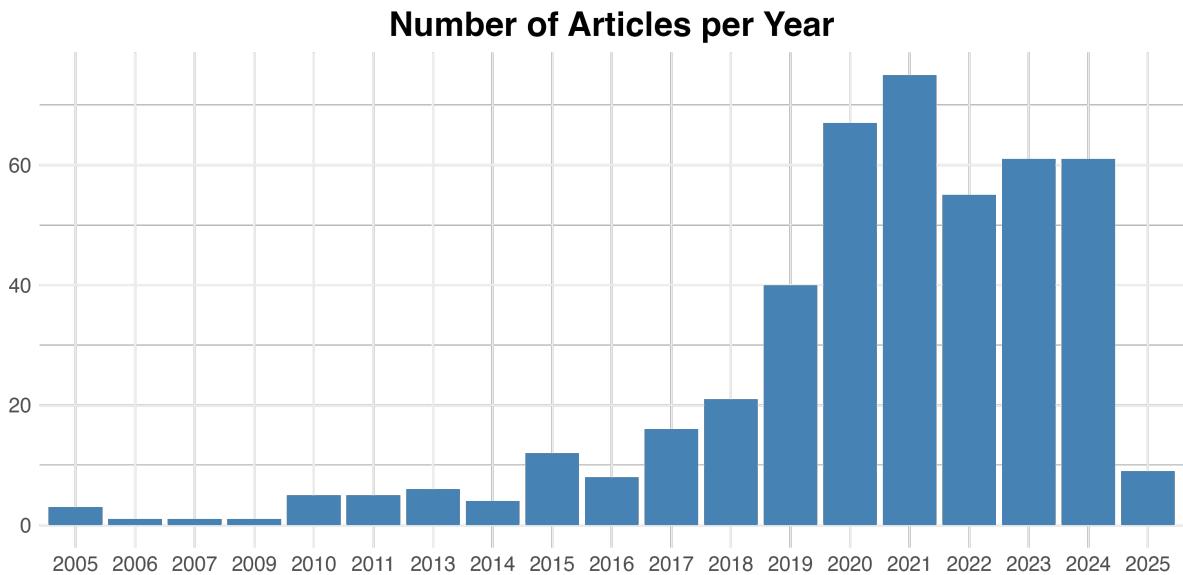


Figure 2.2: Number of articles in this review published per year.

Next, we analyze the most cited papers overall. Table 2.5 shows the top 10 most cited papers in the dataset. These papers are predominantly in multidisciplinary journals, with only one conference paper making the list.

Table 2.5: The top 10 most cited papers in the data set.

Short Paper Title	DOI	Year	Venue	Discipline	Citations
1. Information credibility on twitter	10.1145/1963405.1963500	2011	<i>Web Conference</i>	Computer Science	2370
2. When Corrections Fail: The Persistence of Political Misperceptions	10.1007/s11109-010-9112-2	2010	<i>Political Behavior</i>	Social Sciences	2317
3. Fighting COVID-19 Misinformation on Social Media	10.1177/0956797620939054	2020	<i>Psychological Science</i>	Psychology	1280
4. Effective Messages in Vaccine Promotion: A Randomized Trial	10.1542/peds.2013-2365	2014	<i>Pediatrics</i>	Medicine	1089
5. Prior exposure increases perceived accuracy of fake news	10.1037/xge0000465	2018	<i>Journal of Experimental Psychology: General</i>	Psychology	831
6. Inoculating the Public against Misinformation about Climate Change	10.1002/gch2.201600008	2017	<i>Global Challenges</i>	Multidisciplinary	674
7. Shifting attention to accuracy can reduce misinformation online	10.1038/s41586-021-03344-2	2021	<i>Nature</i>	Multidisciplinary	622
8. Analysing How People Orient to and Spread Rumours in Social Media	10.1371/journal.pone.0150989	2016	<i>PLoS ONE</i>	Multidisciplinary	622
9. Fighting misinformation on social media using crowdsourced judgments of news source quality	10.1073/pnas.1806781116	2019	Multidisciplinary	PNAS	594
10. Neutralizing misinformation through inoculation	10.1371/journal.pone.0175799	2017	<i>PLoS ONE</i>	Multidisciplinary	592

## 2.4.2 Publication Venue Analysis

In this section, we examine the leading publication venues in the dataset and visualize *Publication Venue x Publication Venue (Co-Authorship)* network.

Table 2.6 shows the top 10 publication venues in the *Publication Venue x Publication Venue* network by two centrality metrics. The total degree centrality for each venue indicates how many other venues it is connected to in the co-publication venue network. It represents venues that can be considered “central locations,” or venues that have the most connections to other venues. The second metric is eigenvector centrality, which is another way to measure how influential a node is in a network. A venue has high eigenvector centrality if it is connected to other high-scoring venues.

Table 2.6: The top 10 venues in the Co-Publication Venue network.

Rank	Total Degree Centrality	Eigenvector Centrality
1.	Nature Human Behavior	Nature Human Behaviour
2.	Harvard Misinformation Review	Science Advances
3.	Science Advances	Harvard Misinformation Review
4.	Royal Society Open Science	Psychological Science
5.	PNAS	PNAS
6.	Psychological Science	Royal Society Open Science
7.	J. Experimental Psychology: General	J. of Experimental Psychology: General
8.	Journal of Communication	Nature
9.	CSCW, J. of Applied Social Psych., Political Psych.	CSCW
10.	Cognitive Research: Principles and Implications	J. of Applied Social Psych.

Table 2.7 ranks the venues according to the number of publications they have in this data set. This table demonstrates that having a high number of publications does not necessarily ensure a top ranking for that venue based on centrality metrics. Several of these venues, including *Scientific Reports* and the *CHI Conference*, do not rank highly in terms of centrality within the Co-Publication Venue network (Table 2.6).

Next, we analyze the Co-Publication Venue network. Figure 2.3 illustrates the largest component. Nodes are sized by total degree centrality and colored by discipline. A selection of venues is highlighted. The density of this network is relatively low, at 0.0349. Of the 209 publication venues in our dataset, only about half (114) belong to the largest component. This finding indicates a degree of disjointedness in the literature in this area. *Nature Human Behavior*, *Science Advances*, and *Harvard Misinformation Review* appear highly central in the network, underscoring their high rankings in both total degree centrality and eigenvector centrality. The centrality of these venues suggests that they are relatively interdisciplinary journals connecting various fields and authors typically publishing in other journal disciplines.

Additionally, the top-left side of the network mainly consists of psychology journals, highlighted in yellow. We find the social science journals on the right side of the network. The top-right is primarily made up of communication and journalism venues, while the bottom-right features many political science venues. Finally, on the bottom-left, we have the computer science venues. Although these fields are connected through several interdisciplinary journals, the venues within each discipline are clustered together and primarily linked to one another.

Table 2.7: The top venues by total number of publications in the data set.

Publication Venue	Number of Publications	Discipline
1. CSCW Conference	18	Computer Science
2. Harvard Misinfo Review	15	Social Sciences
3. PLOS ONE	11	Multidisciplinary
4. CHI Conference	9	Computer Science
4. Science Communication	9	Social Sciences
5. Scientific Reports	8	Multidisciplinary
6. Journal of Communication	7	Social Sciences
6. Nature Human Behavior	7	Psychology, Other
7. Cognitive Research: Principles and Implications	6	Psychology, Other
7. Health Communication	6	Social Sciences
7. J. of Applied Research in Memory & Cognition	6	Psychology
7. Memory & Cognition	6	Psychology, Other
7. Political Behavior	6	Social Sciences
7. PNAS	6	Multidisciplinary

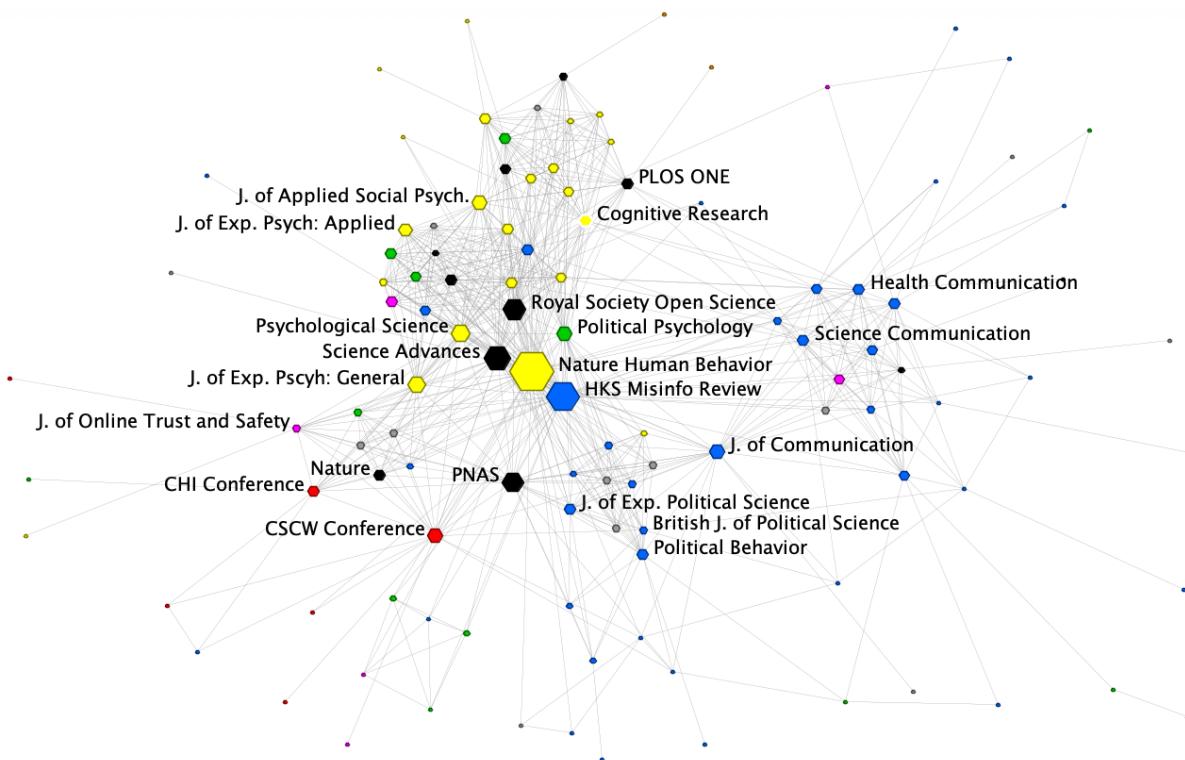


Figure 2.3: Publication Venue x Publication Venue (Co-Authorship) Network. Nodes are sized by Total Degree Centrality and colored by discipline (red for computer science, yellow for psychology, blue for social sciences, black for multidisciplinary, and grey for other)

Figure 2.4 shows the number of papers included in this review, categorized by discipline. Disciplines are determined by the publication venue, and since venues can be affiliated with multiple disciplines, papers may also belong to more than one discipline as well. We observe that the social sciences initiate the literature in this area, with computer science and psychology competing for second. The “other” discipline, which represents all remaining fields, ranging from medicine to environmental science, shows a notable spike near the end of the timeline, indicating a growing interest in researching misinformation interventions across various domains.

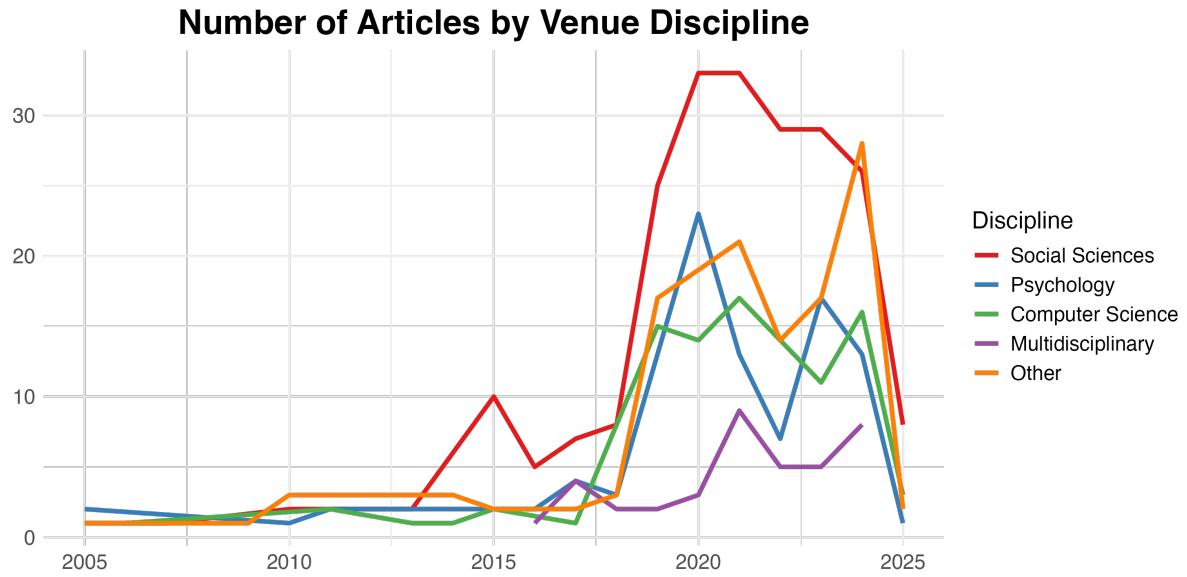


Figure 2.4: Number of articles in this review by venue discipline.

### 2.4.3 Topic Analysis

In this section, we analyze the leading topics, investigating whether certain interventions are over-studied or under-studied, and identifying topics that are often studied together.

First, we calculate the descriptive statistics for the number of papers assigned to each label. Table 2.8 displays the number of papers and unique authors for each intervention label. This table highlights how certain interventions are studied significantly more frequently than others. Fact-checking appears in 127 papers, while 15 of the 35 studied interventions are featured in fewer than 10 papers.

The summary statistics for both paper and author counts are presented in Table 2.9 and further underscore this discrepancy. Fact-checking, debunking, and media literacy are outliers in terms of the number of papers examining those topics based on the calculated interquartile range. Similarly, those three intervention types, along with inoculation, are outliers regarding the total number of authors researching those topics. Furthermore, we found that 404 papers (89.6%) analyzed an intervention’s effectiveness and 40 papers analyzed an intervention’s level of user acceptance (8.9%) (See Table 2.10).

Table 2.8: The number of papers and unique authors who have studied each intervention type.

Intervention	Category	Papers	Authors
fact-checking	Content Moderation	127	335
debunking	Content Moderation	124	359
media literacy	Media Literacy	76	256
inoculation	Media Literacy	58	214
warning labels	Content Labeling	45	170
source credibility labels	Content Labeling	38	155
social corrections	User-Based Measures	35	88
accuracy prompts	Content Distribution	31	116
social norms	User-Based Measures	30	163
fake news games	Media Literacy	26	66
retractions	User-Based Measures	25	45
platform alterations	Content Distribution	23	75
crowdsourcing	Content Labeling	19	67
misinformation detection	Content Moderation	18	72
combining interventions	Other	14	48
institutional measures	Institutional Measures	13	41
friction	Content Distribution	12	63
context labels	Content Labeling	12	67
algorithmic content moderation	Content Moderation	12	41
content labeling	Content Labeling	11	34
other	Other	9	37
government regulation	Institutional Measures	8	14
advertising policy	Content Distribution	7	29
media support	Institutional Measures	6	16
content distribution	Content Distribution	5	13
redirection	Content Distribution	5	12
user-based measures	User-Based Measures	5	14
account removal	Account Moderation	4	12
reporting	User-Based Measures	4	13
content moderation	Content Moderation	3	7
generative AI	Other	3	8
account moderation	Account Moderation	2	8
shadow banning	Account Moderation	1	4
blocking	User-Based Measures	1	4
data sharing	Institutional Measures	1	3

Table 2.9: Statistical summary of the number of papers and unique authors by intervention type. Outliers are identified from the inter-quartile range.

	<b>Mean (SD)</b>	<b>1st Q.</b>	<b>Median</b>	<b>3rd Q.</b>	<b>Outliers</b>
Papers per Intervention	23.2 (30.7)	5	12	28	fact-checking, debunking, media literacy
Authors per Intervention	76.3 (92.4)	13	41	81.5	fact-checking, debunking, media literacy, inoculation

Table 2.10: The number of studies examining the effectiveness or acceptance of one or more interventions.

		<b>Acceptance</b>	
		Yes	No
<b>Effectiveness</b>	Yes	24 (5.3%)	380 (84.3%)
	No	16 (3.5%)	31 (6.9%)

We next analyzed the Co-Topics network, which shows the intervention types that are frequently studied together. A link exists between two topics in the network if a paper discusses both topics and the links are weighted. The network density was 0.40, indicating that countermeasures are frequently studied jointly with other countermeasures. Figure 2.5 shows the Co-Topic network, with nodes sized by total degree centrality and colored based on paper counts. The dark blue nodes represent topics that have low paper counts, and the lighter the blue, the more papers study that topic.

As shown in Figure 2.5, many topics, such as fact-checking and media literacy, are highly central to the network. Not only are they among the most studied interventions, but they are also often studied in conjunction with other interventions. Many countermeasures frequently employed by social media platforms, such as redirection, user-based countermeasures, and intervention combinations, remain relatively understudied. The topics that are overstudied and understudied align with Courchesne et al.’s previous review article, although their metric for categorizing over and understudied topics was based on what the platforms were actually implementing [59].

Next, we analyze whether the types of interventions studied have changed over time. Figure 2.4 shows the number of papers included in this review, categorized by general intervention category studied. Papers could be assigned to more than one intervention label. We see that content moderation interventions were among the first and most prominently studied in the literature. Fact-checking and debunking are interventions that can be examined without access to social media data and would not be affected by a lack of platform transparency. Media literacy has especially taken off in the last few years, perhaps being studied in a wider variety of journal types such as education, medicine, and others. In addition to media literacy, content distribution and institutional measures have more recently reached their peak.

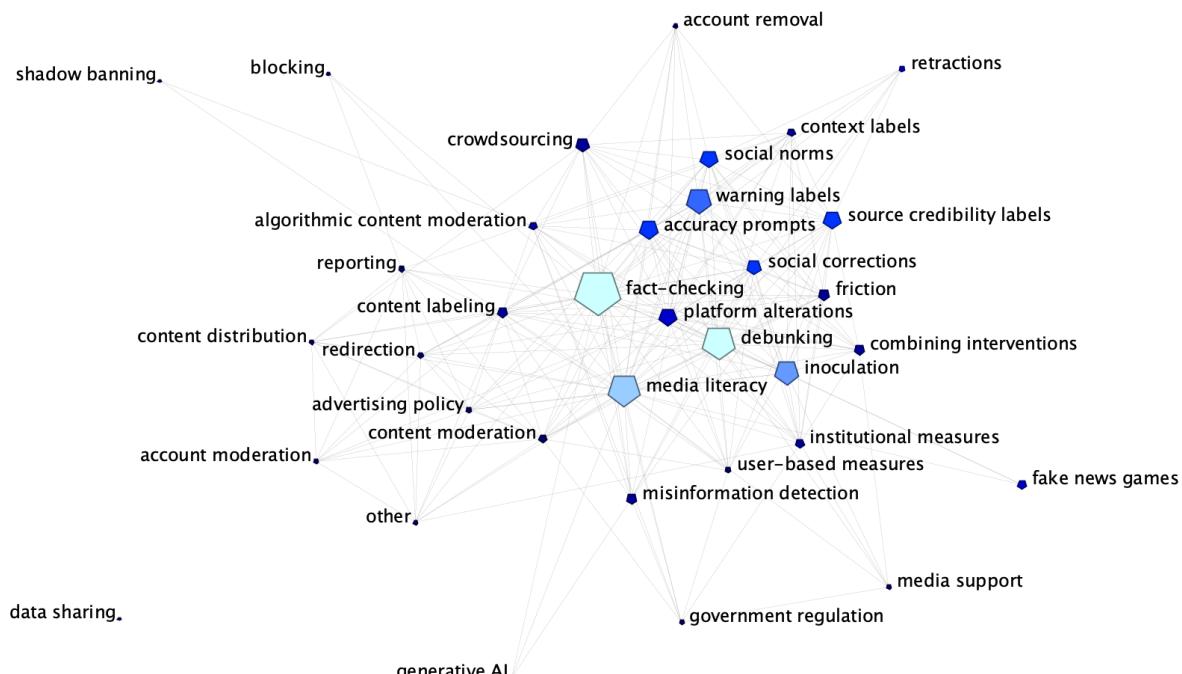


Figure 2.5: The Intervention Topic x Intervention Topic Network. Nodes are sized by total degree centrality and colored by paper count. The lighter the blue, the more papers that study that topic.

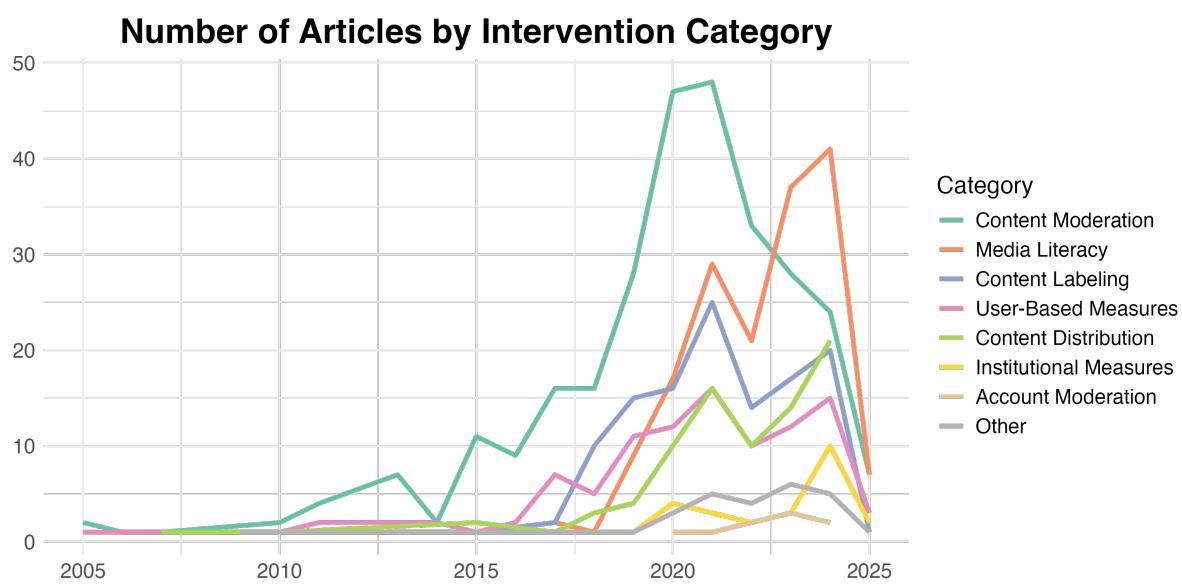


Figure 2.6: Number of articles in this review by general intervention category studied.

## 2.4.4 Author Analysis

In this section, we analyze the top authors in the data and visualize the Co-Authorship network. Table 2.11 displays the top 10 most cited authors, including only those who have published more than once in this dataset. The vast majority of authors (83% or 953 authors) have only one paper in this dataset. These ten authors are also the same authors who have the most number of papers in the dataset.

Table 2.11: The top 10 most cited authors with more than one paper in the data set.

Author Name	Citation Count	Article Count
1. Brendan Nyhan	5549	17
2. Jason Reifler	5305	12
3. David G. Rand	5102	22
4. Gordon Pennycook	5004	21
5. Sander van der Linden	3122	22
6. Ullrich K. H. Ecker	2981	30
7. Stephan Lewandowsky	2735	22
8. Emily K. Vraga	2332	16
9. Leticia Bode	2178	13
10. Jon Roozenbeek	2171	16

Table 2.12 shows the top 10 authors in the *Author x Author (Co-Authorship)* network based on three metrics: total degree centrality, eigenvector centrality, and betweenness centrality. Betweenness centrality is a metric that identifies bridging nodes by quantifying the number of shortest paths that pass through each node. Authors with high betweenness centrality are potentially influential because they can facilitate connections between authors who are otherwise unconnected or who have not previously collaborated.

Table 2.12: The top 10 authors in the Co-Authorship network.

Rank	Total Degree Centrality	Eigenvector Centrality	Betweenness Centrality
1.	Stephan Lewandowsky	Stephan Lewandowsky	Brendan Nyhan
2.	Sander van der Linden	Gordon Pennycook	Jason Reifler
3.	Ullrich K. H. Ecker	David G. Rand	Ullrich K. H. Ecker
4.	Gordon Pennycook	Ullrich K. H. Ecker	Andrew M. Guess
5.	David G. Rand	Sander van der Linden	Sander van der Linden
6.	Brendan Nyhan	Adam J. Berinsky	Gordon Pennycook
7.	Adam J. Berinsky	Briony Swire-Thompson	David G. Rand
8.	Philipp Schmid	Rakoen Maertens	Stephan Lewandowsky
9.	Cornelia Betsch	Philipp Schmid	John Cook
10.	Jason Reifler	Melisa Basol	Adam J. Berinsky

Next, we visualize the *Author x Author (Co-Authorship)* network in Figure 2.7. The density of this network is relatively low at 0.005, which is expected considering the large number of

authors in the dataset (1,147). However, only 326 authors are part of the largest component. In total, there were 218 components, with all other authors in either isolates, dyads, triads, or other relatively small groups. Eight components ranged in size from 10 to 21 authors, likely indicating research groups and authors who have not published outside their own group. This result shows that 821 authors not in the largest component (71.6% of all authors in our dataset) have primarily published on countermeasures within their own research groups and have not collaborated with others.

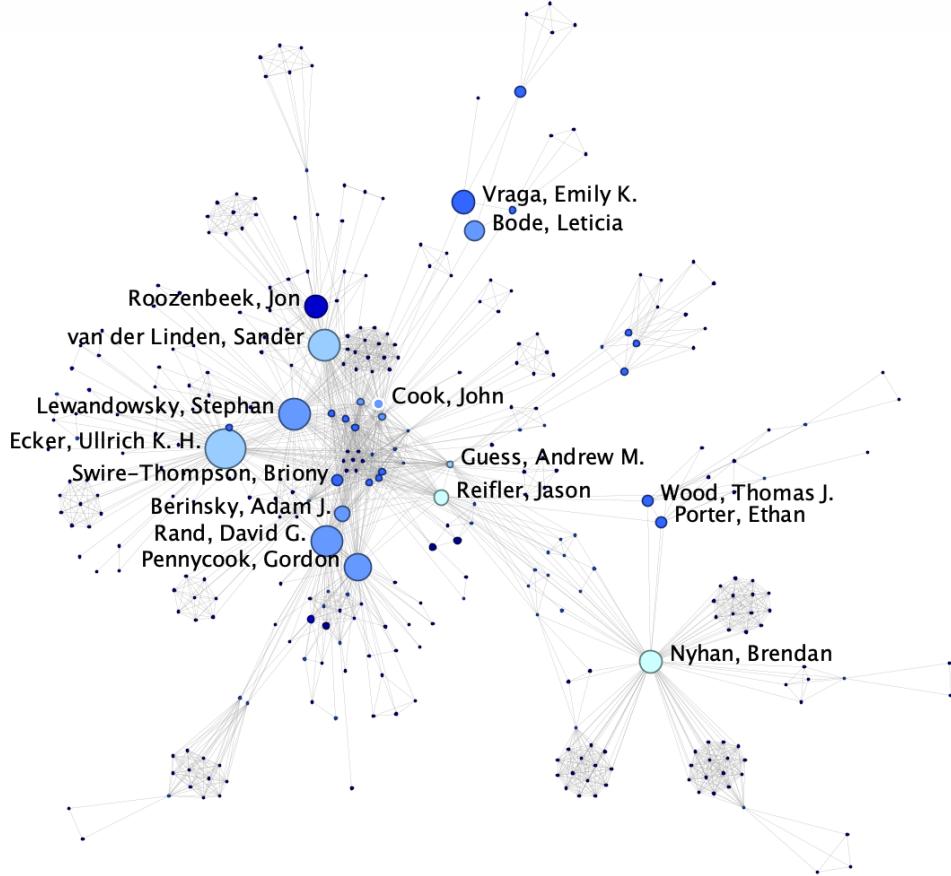


Figure 2.7: The largest component of the Co-Authorship Network. Nodes are sized by total number of papers in the data set and colored by betweenness centrality, with lighter blue colors indicating higher betweenness.

The ORA report “Bibliography and Citation” provides information on the papers with the most authors. Of the 412 papers, 15 have ten or more authors, including one with 30 authors. This article with 30 authors was one of the seed papers [129], indicating it was likely a review article written through the consensus of many leading authors in the field. It features six of the top ten authors by total number of citations (Table 2.11). Removing this single paper from the analysis of the Co-Authorship network causes the largest component to split into two: one of size 230 and another of size 96, as illustrated in Figure 2.8. This finding suggests that many top authors are only connected in the dataset through this one recent review paper, suggesting that

the previous disjointedness in the literature may be decreasing in recent years.

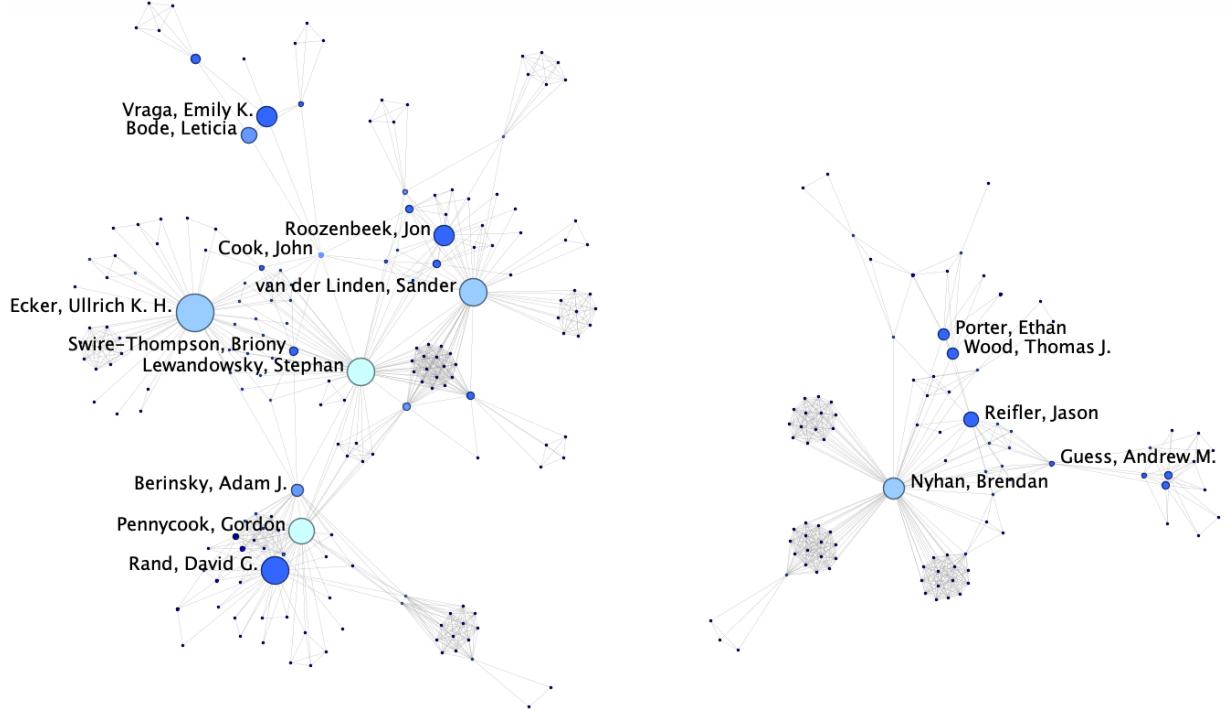


Figure 2.8: The largest component of the Co-Authorship Network excluding one paper with 30 authors. Nodes are sized by total number of papers in the data set and colored by betweenness centrality, with lighter blue colors indicating higher betweenness.

## 2.5 Discussion

We conducted a descriptive and bibliometric analysis of the citation network of prominent papers in the countermeasures literature. The number of articles published in the misinformation intervention space has increased dramatically, indicating that this is a growing field of literature. First, we analyzed the publication venues in this dataset. The Co-Publication Venue network reveals disjointedness in the literature, with most venues clustered near others within the same discipline. While there are several journals, such as the Harvard Misinformation Review and Nature Human Behavior, that bridge the gaps between related fields, nearly half of the venues were not part of the largest component. Multiple fields are conducting research in this area and are visible in their own clusters on the network, including Psychology, Political Science, Journalism and Communication, and Computer Science. This research area is highly interdisciplinary.

Next, we analyzed the top topics studied in the literature. User acceptance is largely overlooked in the literature, with only about 9% of papers exploring this aspect. Acceptance is as important a measure as effectiveness. Without acceptance, platforms may hesitate to implement changes for fear of losing users, and governments might struggle to enact beneficial policies.

Furthermore, similar to a previous review article [59], we find that several critical, frequently used, or highly impactful interventions are underexplored in the literature. These include redirection, user-based countermeasures like user reporting and blocking, institutional measures like media support and data sharing, and emerging interventions involving generative AI. A 2021 review of platform policies indicates that redirection is the most prevalent intervention employed by platforms [243]. However, there were only five papers related to redirection in this list of 412 articles. Additionally, institutional measures overall, including analyses of potential government regulations or actions that civic society can take, are underrepresented in the literature compared to individual or platform-based interventions. Although there is a prominent RAND article that reviews countermeasures based on policy reports [100], it was not used as a seed paper because this article exclusively evaluated policy reports published by think tanks, non-profits, and government entities and did not include any articles from traditional, peer-reviewed journals.

## 2.6 Conclusions

### 2.6.1 Limitations

One of the main limitations of this study is that our method for selecting papers may have missed articles in specific sub-areas. However, the categories of misinformation we identified as over-studied and under-studied are similar to what other related review articles found, mitigating this risk [59]. Additionally, we primarily focused on research conducted in academic peer-reviewed venues. There may be some discrepancies between what is done in academia and what is done in industry or other institutions due to limited data-sharing and access. Lastly, only articles written in English were included in our analysis.

Another limitation is that we did not have a baseline to compare this bibliometric analysis against. For example, is it typical for co-authorship networks to fragment as much as they did in this work when one article with a high number of authors is removed (see Figure 2.7 and 2.8)? Is the clustering of journals by discipline common in other multidisciplinary fields? These questions suggest possible future directions for this work.

### 2.6.2 Contributions

A bibliometric analysis was conducted on the literature surrounding misinformation interventions. This analysis found many under- and over-studied interventions in the literature, including user-based countermeasures, which will be examined further in Chapters 3 and 4. Additionally, we found that the academic literature primarily focuses on the effectiveness of countermeasures without addressing the critical metric of user acceptance. The intervention space should investigate the popularity of these interventions. If user acceptance is low, the implementation of that intervention is unlikely, making its relative effectiveness irrelevant. User acceptance is the main aspect studied in the subsequent chapters of this dissertation. Furthermore, publication venues in this field have largely remained clustered by discipline, although several collaborative and multidisciplinary publications have emerged in recent years.

While analyzing the literature, we encountered several disagreements regarding the effectiveness of various interventions. For example, among the most studied countermeasures, there are several sources of contention; a body of work claims the effectiveness of the “Bad News” game for inoculation [28, 194], while a meta-review finds their results to be insignificant using ROC curves to compare pre and post-treatment classification accuracy [154]. The effectiveness of debunking [51, 157, 165] is also disputed despite a wealth of empirical research. Future research should investigate these disagreements. The lack of consensus on various interventions emphasizes the need for comprehensive evaluation metrics [223] in the field and highlights the importance of meta-reviews. Chapter 6 will review effectiveness and acceptance of interventions in greater detail.



# Chapter 3

## Characterizing User-based Countermeasures

This thesis aims to characterize all types of misinformation interventions, including user-led countermeasures. Examining individual behavior in response to encountering misinformation is critical because previous research has shown that debunking myths is more effective when it comes from a trusted source, such as a friend or family member [144]. In addition, social media corrections have been shown to be highly effective [21, 37, 233]. This suggests that individuals responding directly to misinformation in real-time can help slow or stop its spread.

In this chapter, I investigate the behavior and opinions of social media users when they see or post misinformation. We surveyed 1,010 active social media users residing in the United States. This survey covered the social media platforms where they encounter misinformation, if they have posted misinformation unintentionally, their reactions to seeing or posting misinformation, and their opinions on how they think others should act.

Our primary research question for this chapter is: How are people currently tackling misinformation on social media? More specifically,

1. How do people respond and think others should respond to misinformation posted by others or themselves?
2. Do people respond differently based on who posted the misinformation (**poster**) and where it was posted (**platform**)?
3. What demographic factors, if any, are related to opinions on these topics?

### 3.1 Introduction

As described in Chapter 1, many countermeasures, including user-based countermeasures, have been understudied in the literature to date [59]. Recent research has begun investigating the relationship between seeing misinformation countermeasures online and public perception of those countermeasures [199]. However, the experience of observing or posting misinformation differs from the experience of observing or conducting a countermeasure. Viewers of misinformation on social media can either directly confront the authors of misinformation by offering social cor-

rections or indirectly counter the misinformation by, for example, reporting the misinformation. Studying individual behavior in response to seeing misinformation is critical because previous research has shown that debunking myths is more effective when it comes from a trusted source, like a friend or family member [37, 144]. This suggests that individuals responding directly to misinformation from users in their network can help slow or even stop the spread of misinformation.

This work investigates if these social corrections happen and if they depend on the nature of the relationship between the misinformation poster and the observer or the platform on which it was posted. Given the scale in which social media companies must detect and respond to misinformation [116], users can play a crucial role in limiting the spread in real time. Indeed, social media companies such as X (formerly known as Twitter) have piloted programs like Community Notes where users can add corrections and/or context to tweets they deem misleading [229]. Additionally, previous research has shown that most people do not intend to spread misinformation [13, 27, 237]. Instead of consciously sharing misinformation, cognitive and socio-affective mechanisms (e.g., intuitive thinking, identity motives) facilitate sharing and even belief in misinformation in some cases [52, 74]. If this is the case, nudging social media users to focus on accuracy goals could help limit the unintentional spread of misinformation [177].

In addition, we explore what people believe others should do when they see misinformation or post misinformation themselves. This provides researchers and policymakers a sense of what social media users want the norm response to be, which is essential for public outreach about crowdsourced misinformation mitigation. We also examine how expectations vary from reported actions to understand the extent to which people currently feel empowered to respond to misinformation regardless of (their own) situational constraints. Participants may want others to respond to misinformation with higher effort actions than they do themselves. This act of hypocrisy, failing to practice what one preaches [45], could be leveraged to induce prosocial behavior changes (e.g., directly addressing content they believe contains misinformation) [14, 78, 209]. Making the discrepancy between behavior and advocated norms for behavior salient can activate threats to self-integrity, driving behavior in line with advocated norms to minimize the dissonance [209].

This study surveyed 1,010 United States residents who use social media at least weekly. This survey covered the social media platforms where participants encounter misinformation, if they have posted misinformation (intentionally or unintentionally), their response to seeing or posting misinformation, and their opinions on how they think others should respond to seeing or posting misinformation. Participants report their response and the response they expect others to do when seeing misinformation posted by someone else for three levels of closeness to the sender of misinformation: close (e.g., a close friend or family member), somewhat close (e.g., acquaintances, colleagues, friends, extended family), or not close (e.g., someone you do not know offline). We included seven ways to counter misinformation posted by others, four ways to respond to misinformation posted by oneself, and an option to engage in no action. These responses capture a broad range of actions that involve directly and indirectly interacting with the content containing misinformation.

This chapter will be broken up into two main analysis sections: **closeness analysis** and **platform analysis**. Our research questions, hypotheses, and analysis plan on how responses and beliefs differ based on closeness to the misinformation poster were pre-registered at *Scientific Reports*. The Stage 1 protocol (currently under embargo) can be found here: [\[\[figshare link\]\]](#).

## 3.2 Data and Methods

### 3.2.1 Ethics Information

The Institutional Review Board of Carnegie Mellon University approved this survey, numbered “STUDY2022\_00000143.” They approved this study as exempt from a full review because it is a survey that did not collect personally identifiable information. Informed consent was obtained from all participants. We expected the survey to take 15-18 minutes based on pilot tests. Participants were paid \$3 each, which is equivalent to \$10/hour if they took 18 minutes to complete the survey.

### 3.2.2 Pilot Data

The survey was implemented in Qualtrics and was sent out to a small sample of Cloud Research Mechanical Turk participants to ensure questions were straightforward and the bot and duplicate detection worked. Twenty-two participants attempted the survey: 14 were excluded, most of them automatically by Cloud Research, for either being spam/bots, being a duplicate response, or failing to pass the screening questions (18+, U.S. resident, use social media weekly). Participants excluded for these reasons were removed at the beginning of the survey and were not paid.

Participants were also asked if they had any comments. As a result, a few questions were removed to prevent the survey from being too long or reworded for succinctness and clarity. This document’s hypotheses and research questions are based on the revised survey. While a pilot sample of eight responses is small, it helped improve the research design. It demonstrated that the questions were understandable and that this survey could effectively address the hypotheses and research questions. See the Supplementary Information in the Stage 1 Protocol for more detailed information on the pilot data.

### 3.2.3 Survey Design and Sampling Plan

This section describes our survey design, sample characteristics, sample size determination, data exclusion criteria, and all primary measures. The survey was designed to answer the research questions associated with both Chapters 2 and 4. See [\[figshare link\]](#) to see a copy of the survey.

### Participants

Our survey had 1,010 participants, and the data was collected between July and August 2024. This sample size was deemed appropriate because it provided sufficient power for our proposed hypotheses. Our survey was implemented using Qualtrics and administered through Cloud Research, an online recruiting platform [97], using Mechanical Turk survey participants. Only those respondents who are United States residents, adults, and use social media at least once a week were given the entire survey.

## Procedure

*Qualifying Questions:* We employed several methods to recruit relevant participants and maintain high data quality. Participants were adult U.S. residents who use social media weekly, and they also must have met the following criteria for inclusion:

1. Approved by Cloud Research
2. Had a higher than 95% approval rating on Mechanical Turk
3. Finished the survey
4. Not a bot (both Qualtrics and Cloud Research have bot detection) [1]
5. Not a duplicate response (Qualtrics flags likely duplicate responses) [1]

Data that failed even one of these criteria was excluded. A question at the end of the study asked if participants answered randomly at any point. This measure was for data quality purposes only and was not used to exclude data. 997 participants (99.7%) responded with “no”, while 3 participants responded with “yes” (0.3%), and 10 skipped the question. The median time to complete the survey was 11.0 minutes, while the mean was 14.3 minutes.

Previous research has suggested that Mechanical Turks’ data quality is high and that Turkers are more likely to pass attention-checking questions than other online panels [96]. A previous study also shows that the 95% approval rate cut-off can ensure high-quality data without using attention-checking questions [171]. Therefore, this survey did not have any attention-checking questions.

*Behavioral Questions:* This section asked participants how they respond to seeing misinformation on social media platforms they use and how they react if they realize they have posted misinformation.

First, participants were asked if any of their social media contacts have ever posted something they believe to be misinformation, how often they saw it, and on which platforms they saw it. They were able to select among the top 11 most frequently used platforms in the United States as determined by Pew Research [18]. They also had the option to write in another platform that was not listed. Then, for all platforms they claimed to have seen misinformation on, they were asked how close they were to the people posting misinformation and how they responded. Possible responses are shown in Table 3.1.

Table 3.1: Actions social media users can take when they see misinformation online.

Response	Effort Level
Ignore the post	No Effort
Report the post	Low Effort
Report the user	Low Effort
Block the user	Low Effort
Unfollow or unfriend the user	Low Effort
Privately message the user	High Effort
Comment a correction on the post	High Effort
Create a separate post with the correct information	High Effort

Next, participants were asked if they had ever intentionally or unintentionally posted misinformation. If they have unintentionally posted misinformation, they were asked on which platforms and then asked what they did on each platform once they realized they posted misinformation. Possible actions they could have taken are shown in Table 3.2.

Table 3.2: Actions social media users can take when they realize they have posted misinformation online.

<b>Response</b>	<b>Effort Level</b>
Leave the post as is	No Effort
Delete the post	Low Effort
Comment a correction on the post	High Effort
Update the main post with a correction	High Effort
Create a new post with the correct information	High Effort

*Belief Questions:* This section asked participants how they thought people should respond when seeing misinformation. The questions were broken up by closeness: how should people respond to misinformation posted by a close contact? A somewhat close contact? A not close contact? Again, possible responses they were able to select are described in Table 3.1. Finally, participants were asked what people should do if they realize they have posted misinformation (possible actions are described in Table 3.2).

*Demographic Questions:* This section asked participants for various demographic characteristics. These were age, gender, race, ethnicity, education, household income, religion, political party affiliation, and general political leanings.

### 3.2.4 Measures

We created measures to quantify the amount of effort put into responding to misinformation on social media. For both the closeness and platform analysis measures, anything labeled in Tables 3.1 or 3.2 as no effort received a score of 0, low effort received a score of 1, and high effort a score of 2. Anything without a label was labeled as NA. Participants were given the value of the highest effort level they engaged in per closeness level or per platform.

#### Closeness Analysis Measures

*Measure 1a): Effort Expended to Respond to Misinformation Posted by Others based on Closeness (Behavior)*

*Measure 1b): Effort Expended to Respond to Misinformation Posted by Others based on Closeness (Opinion)*

There were three calculated values for both measures, one per closeness level. Table 3.1 shows a list of possible responses one could have when seeing misinformation on social media, generalized to apply to various social media platforms, and rated as no effort, low effort, or high effort. The only no-effort response is ignoring the post. Respondents could also respond with

“I don’t remember,” in which case their effort level was not recorded, as it is unknown. A low-effort response means an action was taken, but there was no interaction with the content directly. A high-effort response indicates that the user likely took more time to respond and interacted with the content directly. The participants were able to select more than one of these actions. Pilot data values for Measures 1a) and 1b) for somewhat close contacts are in the Supplementary Information file in the Stage 1 Protocol.

*Measure 2a): Effort Expended to Respond to Misinformation Posted by Oneself (Behavior)*

*Measure 2b): Effort Expended to Respond to Misinformation Posted by Oneself (Opinion)*

We created measures 2a) and 2b) to quantify the effort one puts into correcting misinformation they posted online. Table 3.2 shows a list of possible actions someone could take. They are rated in the same way as the efforts described in Table 3.1: no, low, or high effort. Like in Table 3.1, the only no-effort response is leaving the post as is. Respondents could also respond with “I don’t remember,” in which case their effort level was not recorded, as it is unknown. Deleting the post is categorized as low effort. The remaining actions are classified as high effort, as they indicate the user took more time to respond and they placed effort into correcting their mistake. Pilot data values for Measures 2a) and 2b) are in the Supplementary Information file in the Stage 1 Protocol.

### **Platform Analysis Measures**

*Measure 3: Effort Expended to Respond to Misinformation Posted by Others based on Platform (Behavior)*

*Measure 4: Effort Expended to Respond to Misinformation Posted by Oneself based on Platform (Behavior)*

Similarly, for the platform analysis we calculate the effort one puts into correcting others or themselves, but broken up by platform instead of closeness. Opinion questions were not asked based on platform, so measures for those are not included here.

## **3.3 Closeness Analysis**

We analyzed closeness to misinformation poster in our pre-registered Report at *Scientific Reports*. This study involved the following research questions. See the Design Table in Appendix C for more details.

1. How do people respond and think others should respond when they see misinformation? Do response(s) change based on how close the participant is to the poster of misinformation?
2. How do people respond and think others should respond when they realize they have posted misinformation?
3. How do people’s responses and beliefs about how others should respond after seeing misinformation differ from their responses and beliefs when they realize they have posted misinformation?

4. How do beliefs about responses to misinformation differ based on various demographic factors?

### 3.3.1 Related Work

Social media users may refrain from directly responding to (recognized) misinformation due to a myriad of constraints, such as concerns over damaging interpersonal relationships or their own credibility [212]. In addition, verifying and correcting misinformative claims is a time-consuming, effortful process in practice [208]. Users may also feel helpless to counter misinformation given the vast amount available online [212]. We generally expect people will incorporate these constraints more when reporting their own response to misinformation than when describing expectations for others due to cognitive distortions like fundamental attribution error [133, 215]. While users can account for the factors that drive their decision-making about how to respond to misinformation online, it is significantly more challenging to incorporate the hypothetical situational constraints of others. Moreover, asserting that others should respond with high levels of effort can uphold feelings of morality even if participants do not want to engage in the moral behavior (i.e., responding actively to misinformation online) for whatever reason [29, 132]. Therefore, we have the following two hypotheses:

**H1.1:** *People believe individuals should expend more effort to respond to misinformation online than they actually do.*

**H2.1:** *People believe others should expend more effort to respond to misinformation online after realizing they posted misinformation than what they actually do.*

Previous work shows users are more likely to correct a close contact because it is perceived as more worthwhile [212]. If users are going to take the time to engage with misinformative content directly, they want to feel like it will have an impact. If they have a personal relationship with the sender of misinformation, they have more information about the expected effectiveness of their correction. Furthermore, people may be especially concerned about close contacts believing in misinformation due to the potential negative consequences. We expect this will translate to expectations for others as well:

**H1.2:** *People respond with more effort when the sender of misinformation is a close contact than a somewhat close contact and a somewhat close contact than a not close contact.*

**H1.3:** *People believe others should respond with more effort when the sender of misinformation is a close contact than a somewhat close contact and a somewhat close contact than a not close contact.*

When comparing responses to misinformation posted by others and posted by oneself, many individual and social factors may come into play, including wanting to preserve harmony and credibility or avoid embarrassment. Previous research from Singapore has shown that many young people avoid correcting misinformation posted by others to maintain their interpersonal

relationships but do correct themselves to preserve their credibility despite possible embarrassment [163]. Other work suggests that many only correct others if they are close contacts or when it is an issue they care about [212]. Due to the conflicting literature in this area, we have created two non-directional hypotheses where we expect that people respond and expect others to respond with a different level of effort when the sender of misinformation is someone else compared to themselves:

**H3.1:** *People respond with a different level of effort when the sender of misinformation is someone else compared to themselves.*

**H3.2:** *People want others to respond with a different level of effort when the sender of misinformation is someone else compared to themselves.*

Finally, we investigate how behavior and beliefs about responses to misinformation on social media vary by partisanship and other demographic factors (**RQ4**). Extensive previous research has examined differences in misinformation susceptibility across age, gender, education level, income bracket, religious groups [44, 49, 68, 160, 225], and partisan groups [81, 86], as well as the effectiveness of interventions across demographic groups [93]. Increasingly, researchers are studying how individual factors impact support for misinformation interventions, typically at a platform level [70, 123, 127, 199]. In specific contexts, such as highly partisan environments, less susceptibility to misinformation is associated with more support for platform interventions. Left-leaning, Democratic individuals are both more supportive of platform interventions [128, 153, 199], and less likely to observe or spread misinformation online [81, 86], than right-leaning or Republican individuals. In other cases, high susceptibility is linked to more support. Older adults are typically associated with higher susceptibility to sharing and believing misinformation [43, 91]. Yet there is evidence that they support nudge interventions more [123].

Crucially, the type of misinformation seems to substantially affect susceptibility and views of countermeasures (e.g., older adults seem less susceptible to health-related misinformation than younger people [160]). Given the variability in susceptibility and support for platform-level interventions, we conducted exploratory analyses on how demographic attributes and partisanship affect support for and engagement with individual interventions. Understanding how specific populations, particularly those shown to be highly susceptible to misinformation, view and enact individual interventions informs effective public messaging about countering harmful content online.

### 3.3.2 Bayesian Power Analysis

We used a Bayesian approach to test our hypotheses. Unlike frequentist methods and p-values, the Bayes factor can show evidence in favor of either the null or the alternate, not just “reject” or “fail to reject” [131, 203]. Additionally, unlike a traditional confidence interval that gives the range of values that would not be rejected at a specified p-value, the highest density interval includes, say, the 95% high probable values for the estimated parameter [131]. Finally, corrections are typically needed for multiple t-tests due to a concern over the possible detection of false positives, and these corrections can result in reduced power [82]. However, Bayesian tests

typically do not need a correction for multiple tests, as the Bayesian prior places a relatively high probability on null effects [162].

This work used a Bayes Factor fixed-n design, where n is our sample size. Given our budget, our maximum sample size was determined to be approximately 1000 respondents. We used the methods described in the Schönbrodt and Wagenmakers (2018) design analysis paper [203] and the BFDA R package [202] to run a strength of evidence analysis. Like a traditional power analysis in a classical frequentist approach, a strength of evidence analysis can estimate the sample size needed so that a strong Bayes factor is found in a specific percentage of studies. For this analysis, the Bayes factor threshold was set to 10 and 1/10. Researchers consider a Bayes factor of 10 to be in the “strong” evidence range [137], with 10 meaning that the data is 10 times more likely to follow H1 over H0, and 1/10 indicating the reverse [203]. For this design analysis, evidence with a Bayes factor within that range is deemed “inconclusive.” The higher the Bayes factor, the lower the probability of receiving misleading evidence (false positives or false negatives) [203].

We used the JZS Bayes factor, which assumes that the effect of H1 follows a central Cauchy distribution [197, 203]. The JZS Bayes Factor was selected because it is the recommended default when little is known about the expected effect size [197, 203]. The effect size is Cohen’s d, the difference in means divided by the standard deviation. The width parameter of the Cauchy distribution was set to  $\sqrt{2}/2$ , which is the recommended value if expecting smaller effect sizes and is the default used in many software packages, including the R Package BayesFactor [155]. It corresponds to expecting a 50% probability of an effect size between -.707 and .707 for a two-sided test and between 0 and 0.707 for a one-sided test. Pilot data for Hypothesis 1.1 was calculated to have an effect size of 0.80, which we believe is not different enough from the default parameter to warrant changing it, especially with such a small pilot data set.

The strength of evidence analysis was run considering being able to detect a possible effect size of 0.5, which is traditionally interpreted as a “medium” effect size [140]. We chose a medium effect size since we believe a small effect size for our hypotheses would not translate into much practical significance. We used this effect size and a Bayes factor threshold of 10 for the following hypotheses to determine if our sample size was high enough to have a reasonable probability of obtaining strong evidence for an alternate H1. We used the BFDA R package [202], which uses a Monte Carlo method to determine this. It generates 10,000 random samples under H1 with an expected effect size of 0.5 and computes the Bayes factors for those runs given the JZS prior. We ran another 10,000 random samples under H0 and calculated the Bayes factors. Then, we used the distribution of the Bayes factors found to determine the sample size necessary to achieve a high probability (95%) of achieving a Bayes factor of 10.

The following subsections detail the power analysis for each hypothesis. We found that our planned sample size of approximately 1,000 participants was likely sufficient for each hypothesis to detect a medium effect size at a Bayes Factor threshold of 10. The power analysis is additionally summarized in the Design Table in Appendix C.

## Hypothesis 1.1

Our first hypothesis is: “Participants say others should respond with higher effort actions to misinformation compared with the actions they take themselves.” Only participants who have seen

misinformation on social media will have answers to the part of the survey related to Hypothesis 1.1. It is well-established that people are not always successful or consistent at identifying misinformation online [176]. However, according to an Ipsos survey, of the fraction of American participants that use social media (87%), 77% of them have said they have seen misinformation specifically on social media [206]. In the small pre-test of our survey, 87.5% of participants (7/8) reported having seen misinformation on social media,

We expect almost all of the approximately 1000 participants to have seen misinformation on social media since we are explicitly surveying frequent social media users. Using 77% as a conservative baseline for this power analysis would translate to a sample size of approximately 770 participants. However, not all participants will have seen misinformation at each closeness level, and we do not have an estimate from the literature for how many participants will have seen misinformation from a contact in each closeness category. We estimate that for each closeness level, at least 25% of participants who have seen misinformation will have seen misinformation from a contact at that level, indicating we can achieve a sample size of approximately 192 for this hypothesis.

For Hypothesis 1.1, a one-sided paired Bayesian hypothesis test will be run per closeness level. Running a sample size analysis, we find that to achieve a power of 95%, using a Bayes factor threshold of 10 and effect size of 0.5, the sample size must only be of size 81, which we believe is reachable in each closeness category. The results show that at this sample size under H1, 95.1% of studies correctly identify H1, 4.9% get inconclusive results, and 0% are false negatives. At an estimated sample of 192, we find 100% of runs under H1 find evidence for H1. We additionally find that .2% of runs under H0 find false positive evidence for H1, 40% were inconclusive, and 59.7% correctly showed evidence for H0. To find evidence favoring the null in at least 75% of studies, we would need a sample size of approximately 440 at this high Bayes factor threshold of 10.

## Hypothesis 1.2

This one-sided hypothesis states: “People respond with more effort when the sender of misinformation is a close contact than a somewhat close contact and a somewhat close contact than a not close contact.” As mentioned above under Hypothesis 1.1, we expect most participants to have seen misinformation. Two of the seven individuals who had seen misinformation in our pilot data had seen it at two different closeness levels. Since we do not have an estimate for the fraction of individuals who will have seen misinformation at more than one level, we conservatively estimate that 15% of people out of the at least 77% of people who have seen misinformation will have seen it from more than one closeness type. This translates to approximately 115 participants, greater than the necessary sample size of 81, as described under Hypothesis 1.1.

## Hypothesis 1.3

For this hypothesis, all participants will answer the relevant questions, indicating a sample size of approximately 1000. This hypothesis is related to individual opinions on responses to misinformation rather than their actions after seeing it. A one-sided paired Bayesian test will be run for the directional hypothesis. Given such a large sample size of  $n = 1000$ , an effect size of 0.5,

and a Bayes factor threshold of 10, we find that 100% of the Monte Carlo H1 samples correctly show evidence for H1. Like the other paired directional hypotheses, a sample size of just  $n = 81$  is needed to detect an effect size of 0.5 at a Bayes threshold of 10 with 95% power.

### Hypothesis 2.1

This hypothesis is: “People believe others should expend more effort to respond to misinformation online after realizing they posted misinformation than what they actually do.” Only those who said they have unintentionally posted misinformation will answer this part of the survey. A Pew Research survey found that 23% of U.S. adults say they have either posted misinformation intentionally or unintentionally [27]. A Statista survey found that 38.2% of Americans claimed they had accidentally shared misinformation on social media [237]. Additionally, almost half of individuals are afraid they may have shared misinformation unintentionally [204]. In our pre-test of 8 participants, we found that half of them said they had posted misinformation accidentally, and a quarter of them admitted to posting misinformation intentionally. Based on the Statista survey, we conservatively estimate that 25% of our participants will say they have posted misinformation accidentally. This is approximately 250 participants. Similar to Hypothesis 1.1, a one-sided paired Bayesian test will be run, and only a sample size of  $n = 81$  is needed to detect an effect size of 0.5 at a Bayes threshold of 10 with 95% power.

### Hypothesis 3.1

This two-sided hypothesis is: “People respond with a different level of effort when the sender of misinformation is someone else compared to themselves.” Only those who have posted misinformation (we estimate to be at least 77%) and those who have posted misinformation (we estimate 25%) will have data to test this hypothesis. We expect most individuals who have realized they have posted misinformation to have also seen it due to their high awareness of the issue. We conservatively estimate that 15% of all participants will qualify, indicating that the estimated sample size is 150. For a two-sided paired Bayesian test, a sample size of just  $n = 92$  is needed to detect an effect size of 0.5 at a Bayes threshold of 10 with 95% power.

### Hypothesis 3.2

Similar to Hypothesis 1.3, all participants will answer the relevant questions because this hypothesis is related to participants’ opinions on how to respond after posting misinformation. A two-sided paired Bayesian test will be run, as this hypothesis is non-directional. Given such a large sample size of  $n = 1000$ , an effect size of 0.5, and a Bayes factor threshold of 10, we find that 100% of the Monte Carlo H1 samples correctly show evidence for H1. Like Hypothesis 3.1, for a two-sided paired Bayesian test, a sample size of just  $n = 92$  is needed to detect an effect size of 0.5 at a Bayes threshold of 10 with 95% power.

### 3.3.3 Statistical Methods

In this section, we detail all pre-registered analyses per hypothesis. We ran paired hypothesis tests, treating the ordinal data for effort level as interval data. The literature is divided on whether

treating ordinal data as interval is recommended [101, 191, 240]. Given the large sample size, we expect the interpretation will likely not change if we analyze the data as ordinal rather than interval. However, as supplemental work, we conducted robustness checks by analyzing a Chi-square test of independence (treating the data as categorical) to improve the robustness of our results. We performed all statistical tests using the BayesFactor R package [155] and the most recent R and R Studio versions at the time of the analysis (R v4.4.172, RStudio v 2024.04.2+76473).

## Hypothesis 1.1

To test **H1.1**, we ran a one-sided paired Bayesian null hypothesis test comparing the effort level of participants' actions when seeing misinformation (Measure 1a) with the effort level participants say others should do when seeing misinformation (Measure 1b) for each closeness level. To run this analysis, we took each user's maximum effort level, ranging from 0-2, for each closeness level. The effect size is equal to (the mean of Measure 1b - mean of Measure 1a) / standard deviation. The null hypothesis ( $H_0$ ) is that the effect size is  $\leq 0$  (the null interval range is  $-\infty$  to 0). The alternate hypothesis ( $H_1$ ) is that there is a difference in the means with an effect size of greater than 0. The 95% highest density intervals were also calculated.

## Hypothesis 2.1

We tested this hypothesis in a similar manner as Hypothesis 1.1. We ran a one-sided paired Bayesian null hypothesis test comparing the effort level of participants' actions when realizing they have posted misinformation (Measure 2a) with the effort level participants say others should do when they realize they have posted misinformation (Measure 2b) for each closeness level. To run this analysis, we took each user's maximum effort level, ranging from 0-2, on each closeness level. We again calculated the 95% highest density interval.

## Hypothesis 1.2 and 1.3

Similarly, we ran one-sided Bayesian hypothesis tests for these hypotheses, again using a null interval range of  $-\infty$  to 0. The highest density intervals were calculated.

## Hypothesis 3.1 and 3.2

For Hypothesis 3.1, we ran a two-sided paired Bayesian hypothesis test comparing the effort level of participants' actions when seeing others post misinformation (Measure 1a) with the effort level of participants' actions when they realized they posted misinformation themselves (Measure 2a). Similarly, for Hypothesis 3.2, we compared Measure 1b and Measure 2b. For both, the null interval range was -0.2 to 0.2. We added a buffer of 0.2 because we consider an effect size that small to be practically equivalent to no effect size. We additionally calculated the highest density interval and made appropriate visualizations for all the results.

### 3.3.4 Results

We surveyed 1,010 active social media users in the United States. Almost all the participants said they had seen misinformation on at least one social media platform. Table 3.3 shows the number of participants who had seen misinformation or admitted to posting it unintentionally. These numbers indicate how many participants were qualified to answer behavioral questions about what they do after seeing or posting misinformation.

Table 3.3: Misinformation exposure

Survey Question	Yes	No
Have you ever seen misinformation posted or distributed on social media?	93.3% (942)	6.7% (68)
Have you ever posted or linked to something you later realized was misinformation?	25.7% (260)	74.3% (750)

### Registered Analyses

All pre-registered analyses and statistical tests were performed and are described in the following subsections. Tables 3.4-3.6 summarizes the registered analyses for each of the hypothesis tests by research question. For detailed descriptions of the hypotheses and analyses, refer to the Design Table in Appendix C. The “Sample Size” column lists the number of participants whose responses qualified for each paired hypothesis test. The “Interpretation” column follows the standard classification scheme for Bayes Factors [137].

#### RQ1: How do people respond and think others should respond when they see misinformation on social media?

First, we summarize the interventions that participants report using. Figure 3.1 shows the total number of participants who responded at least once with each possible intervention from Table 3.1. We see that ignoring the misinformation was the most common response at every closeness level. Higher-effort actions (like privately messaging the user, commenting, or creating another post) received relatively more traction when responding to misinformation against close contacts compared with somewhat or not close contacts.

Next, we tested our first hypothesis (**H1.1**), which is that participants would be hypocritical when comparing their opinions about what people should do when they see misinformation with their actual actions when seeing misinformation. For **H1.1**, we found overwhelming evidence (Bayes Factor > 100) that participants believe individuals should expend more effort responding to misinformation on social media than those individuals report actually doing when they encounter misinformation. This extremely significant result held no matter how close the participant claimed to be to the poster of the misinformation (close contacts, somewhat close contacts, and not close contacts). The effect size was greater than 0.5 in all three closeness cases, indicating a moderate effect size. The unrestricted 95% highest posterior density interval for the effect sizes were [0.35-0.70], [0.59-0.82], and [0.44-0.58] for close, somewhat, and not close contacts, respectively. This result also held the same strength of evidence when the tests were run using

Table 3.4: Pre-registered Bayesian paired hypothesis test results and interpretation for RQ1. Values in italics are for the hypotheses being run with the effort level summed as a robustness check. The p-value for the generalized McNemar's Chi-square Test is also included.

Hypothesis	Sample Size	Null Interval	Mean Diff (S.D.)	Effect Size	Bayes Factor	95% HDI for Effect Size	Interpretation	McNemar Chi-Sq Test
<b>H1.1:</b> Close contacts	148	( $-\infty$ , 0)	0.47 (0.88) 0.99 (1.93)	0.53  0.51	>100  >100	[0.35,0.70]  [0.33,0.67]	Extreme evidence for H1 <i>Interpretation is the same</i>	p = 1e-7
<b>H1.1:</b> Some-what close	370	( $-\infty$ , 0)	0.65 (0.92) 1.34 (1.94)	0.71  0.58	>100  >100	[0.59,0.82]  [0.47,0.69]	Extreme evidence for H1 <i>Interpretation is the same</i>	p<2e-16
<b>H1.1:</b> Not close	880	( $-\infty$ , 0)	0.45 (0.87) 0.95 (1.87)	0.51  0.51	>100  >100	[0.44,0.58]  [0.44,0.58]	Extreme evidence for H1 <i>Interpretation is the same</i>	p<2e-16
<b>H1.2:</b> Close v. somewhat	122	( $-\infty$ , 0)	0.21 (0.84) 0.16 (1.56)	0.26  0.11	>100  6.94	[0.07,0.43]  [-.07,0.28]	Extreme evidence for H1 <i>Moderate evidence for H1</i>	p = 1e-4
<b>H1.2</b> Some-what v. not	327	( $-\infty$ , 0)	-0.052 (0.88) -0.21 (1.52)	-0.059  -0.14	0.170  <1/100	[-.17,0.05]  [-.24,-.03]	Inconclusive at 1/10 threshold <i>Extreme evidence for H0</i>	p = 2e-4
<b>H1.3:</b> Close v. somewhat	1010	( $-\infty$ , 0)	0.10 (0.58) 0.16 (1.34)	0.18  0.12	>100  >100	[0.12,0.24]  [0.06,0.18]	Extreme evidence for H1 <i>Interpretation is the same</i>	p = 3e-8
<b>H1.3:</b> Some-what vs. not	1010	( $-\infty$ , 0)	0.42 (0.81) 0.34 (1.78)	0.51  0.19	>100  >100	[0.45,0.58]  [0.13,0.25]	Extreme evidence for H1 <i>Interpretation is the same</i>	p<2e-16

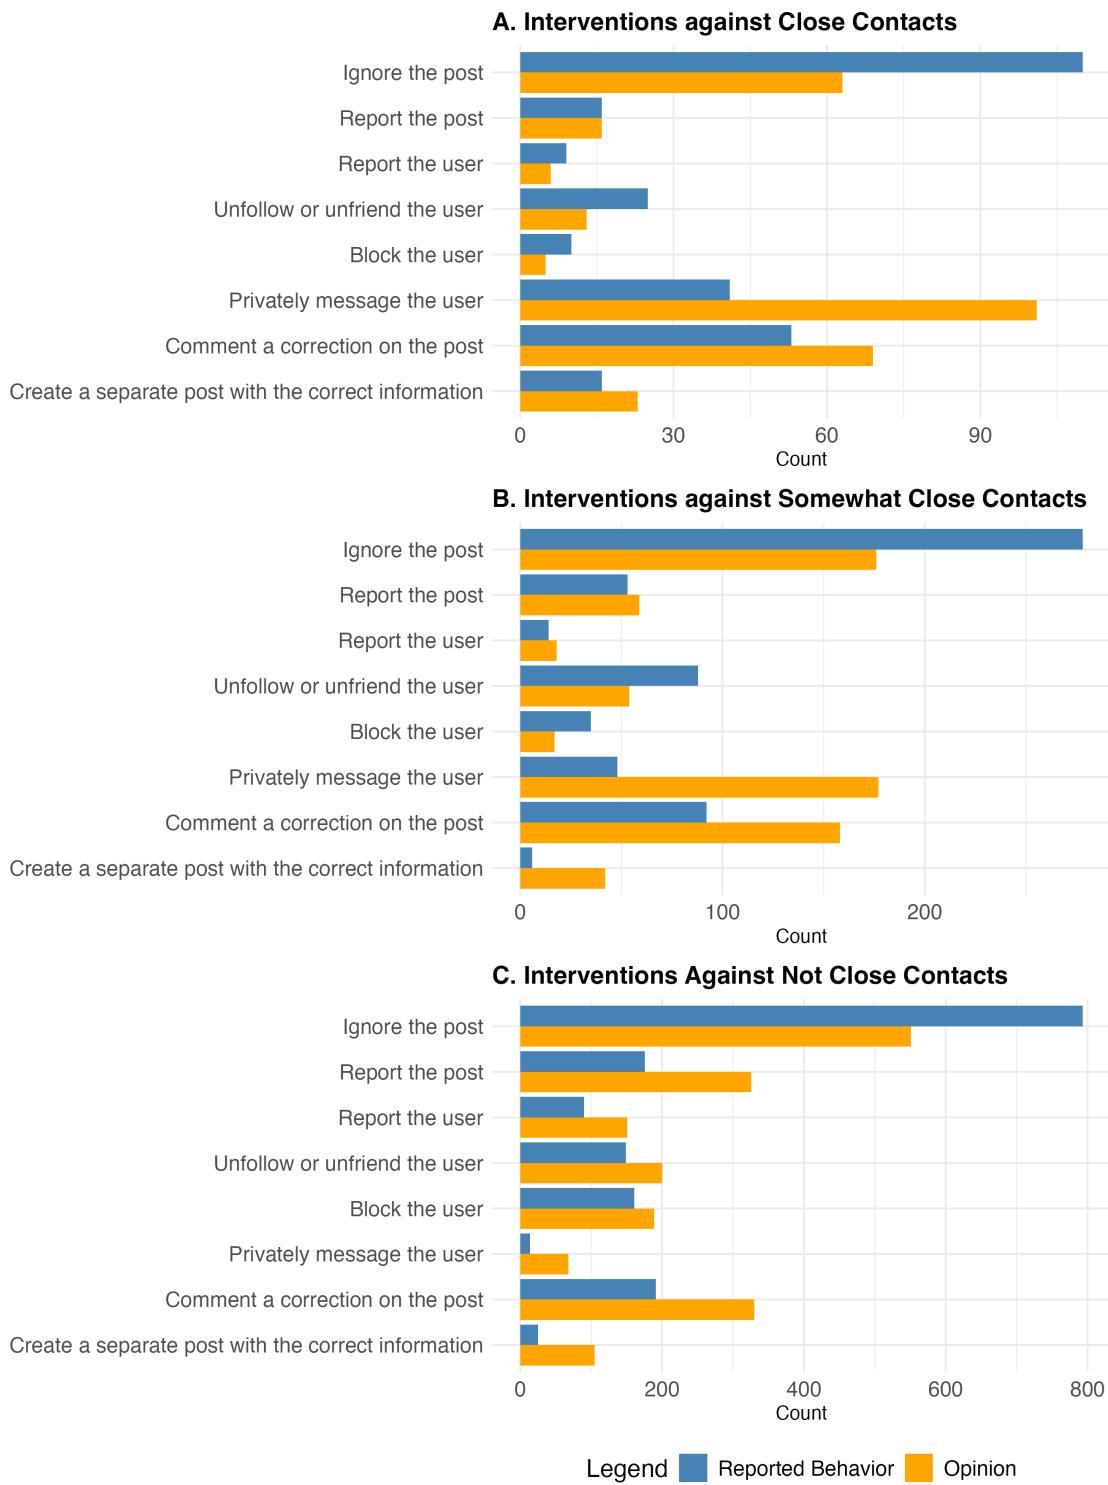


Figure 3.1: Total number of participants who selected each intervention type from Table 3.1 at least once in their behavioral and opinion responses. For comparison purposes, only participants who had seen misinformation at that closeness level have their opinion counts included.

a summed effort level rather than a maximum effort level. Figure 3.2A-C shows the distribution of the maximum effort level reported per closeness level.

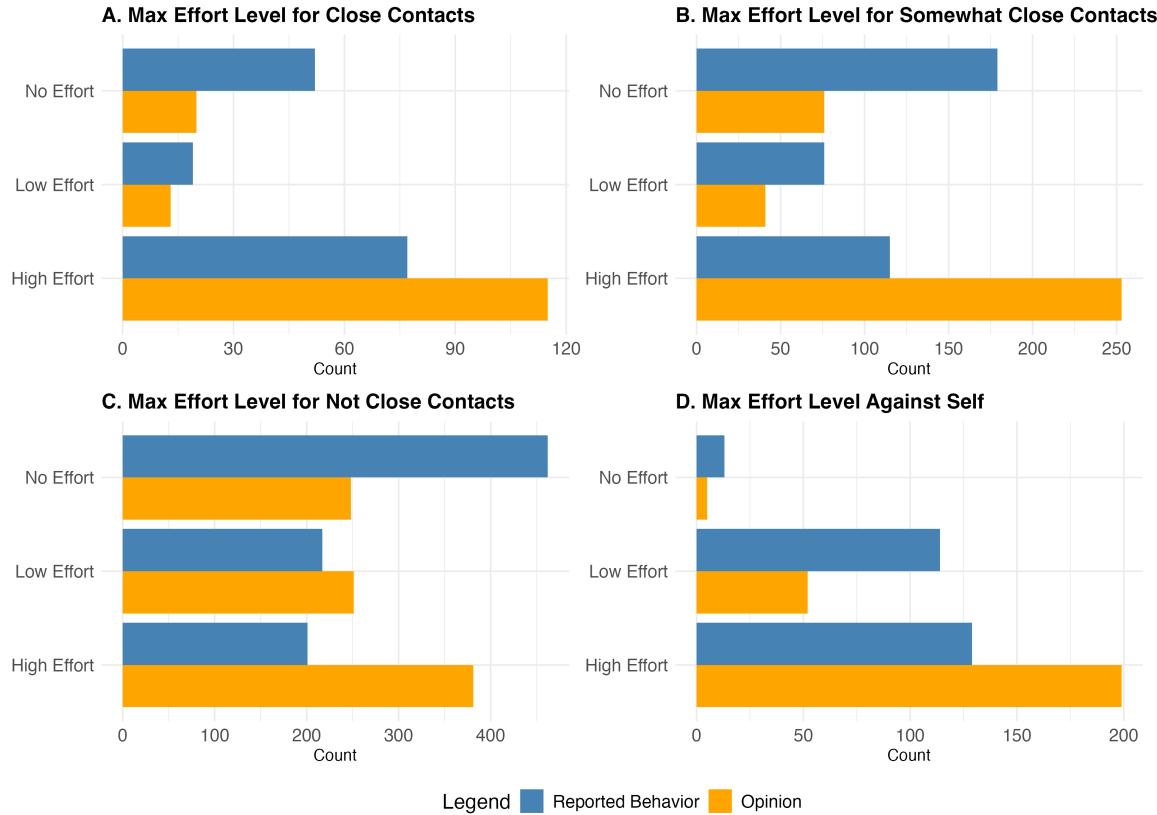


Figure 3.2: Number of participants who reported expending a maximum of no, low, or high effort when seeing misinformation compared with the total number of those same participants who *believe* one should expend no, low, or high effort when seeing misinformation at each of the three closeness levels and against oneself.

For **H1.2**, we found that participants responded with more effort when the misinformation poster was a close contact vs. a somewhat close contact ( $BF > 100$ ) but that there was little difference in responses for somewhat close contacts compared with not close contacts ( $BF$  inconclusive). However, for **H1.3**, we found that participants believe that more effort should be expended on close contacts compared with somewhat close contacts and somewhat close contacts compared with not close contacts ( $BF > 100$ ). Despite the belief that more effort should be put into responding to misinformation posted by a somewhat close contact relative to a not close contact, it seems that participants treated their somewhat close contacts and not close contacts with similar effort levels in practice.

## RQ2: How do people respond and think others should respond when they realize they have posted misinformation?

Over 25% of the participants (see Table 3.3) admitted to accidentally posting misinformation at least once. Figure 3.3 summarizes the interventions people claimed to have taken after real-

izing their mistake. The most frequently reported behavior was deleting the post, followed by updating the post with accurate information.

Table 3.5: Pre-registered Bayesian paired hypothesis test results and interpretation for RQ2. Values in italics are for the hypotheses being run with the effort level summed as a robustness check. The p-value for the generalized McNemar’s Chi-square Test is also included.

Hypothesis	Sample Size	Null Interval	Mean Diff (S.D.)	Effect Size	Bayes Factor	95% HDI for Effect Size	Interpretation	McNemar Chi-Sq Test
<b>H2.1</b>	256	( $-\infty$ , 0)	0.30 (0.55) <i>1.28</i> (1.94)	0.55 0.66	>100 >100	[0.41,0.68] <i>[0.52,0.79]</i>	Extreme evidence for H1 <i>Interpretation is the same</i>	p = 1e-13

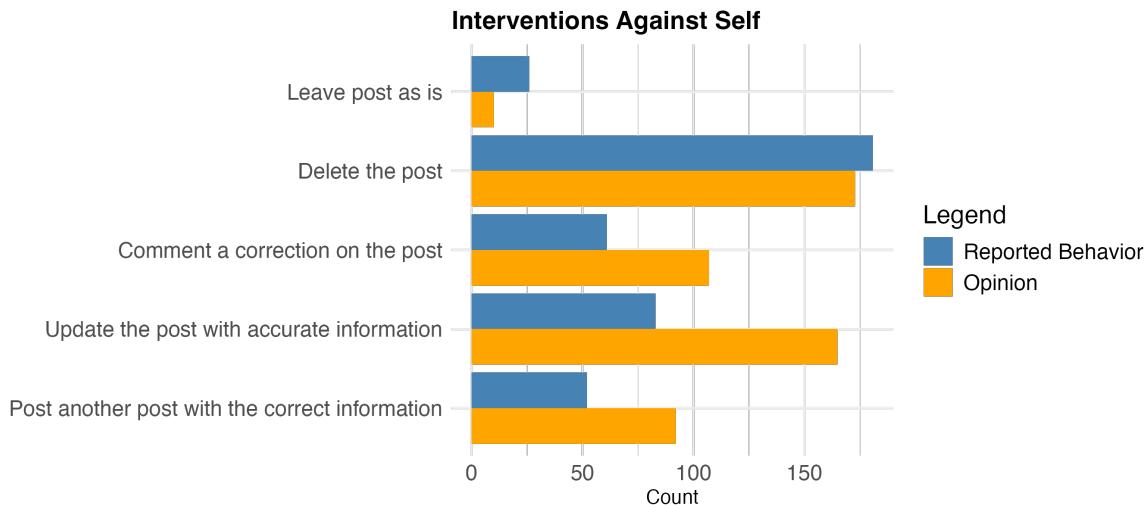


Figure 3.3: Total number of participants who had posted misinformation and selected each intervention type from Table 3.2 at least once in their behavioral and opinion responses. For comparison purposes, only participants who admitted to posting misinformation had their opinion counts included.

We found extreme evidence ( $BF > 100$ ) that people believe that they should expend more effort to respond to the misinformation they posted compared to what they actually do after realizing they have posted misinformation (**H2.1**). The effect size was 0.55, with the 95% HDI of the effect size being [0.41-0.68]. This result held with the same strength of evidence when the test was run using a summed effort level rather than a maximum effort level. See Figure 3.2D for the distribution of maximum effort level after one has posted misinformation.

**RQ3: How do people’s responses and beliefs about how others should respond after seeing misinformation differ from their responses and beliefs when they realize they have posted misinformation?**

Table 3.6: Pre-registered Bayesian paired hypothesis test results and interpretation for RQ3. Values in italics are for the hypotheses being run with the effort level summed as a robustness check. The p-value for the generalized McNemar's Chi-square Test is also included.

Hypothesis	Sample Size	Null Interval	Mean Diff (S.D.)	Effect Size	Bayes Factor	95% HDI for Effect Size	Interpretation	McNemar Chi-Sq Test
<b>H3.1:</b> Self v. close	49	(-0.2,0.2)	0.27 (0.86) <i>0.37</i> (2.26)	0.31 <i>0.16</i>	0.594 0.125	[0.010,0.57] <i>[-.12,0.43]</i>	Inconclusive <i>Interpretation is the same</i>	p = 1e-3
<b>H3.1:</b> Self v. somewhat	133	(-0.2,0.2)	0.41 (0.88) <i>0.56</i> (2.01)	0.46 0.28	79.5 0.805	[0.27,0.63] <i>[0.10,0.44]</i>	Very strong evidence for H1 <i>Inconclusive</i>	p = 9e-9
<b>H3.1:</b> Self v. not close	229	(-0.2,0.2)	0.52 (0.85) <i>0.67</i> (1.80)	0.61 0.37	>100 31.7	[0.46,0.74] <i>[0.23,0.50]</i>	Extreme evidence for H1 <i>Very strong evidence for H1</i>	p<2e-16
<b>H3.1:</b> Self v. max overall	244	(-0.2,0.2)	0.31 (0.85) <i>0.30</i> (1.86)	0.36 0.16	26.8 0.070	[0.23,0.49] <i>[0.031,0.28]</i>	Strong evidence for H1 <i>Strong evidence for H0</i>	p = 3e-13
<b>H3.2:</b> Self v. close	1010	(-0.2,0.2)	0.10 (0.75) <i>0.50</i> (2.15)	0.13 0.23	<1/100 .977	[0.071,0.2] <i>[0.17,0.29]</i>	Extreme evidence for H0 <i>Inconclusive</i>	p<2e-16
<b>H3.2:</b> Self v. somewhat	1010	(-0.2,0.2)	0.20 (0.78) <i>0.76</i> (2.14)	0.26 0.35	6.92 >100	[0.20,0.32] <i>[0.29,0.42]</i>	Inconclusive <i>Extreme evidence for H1</i>	p<2e-16
<b>H3.2:</b> Self v. not close	1010	(-0.2,0.2)	0.62 (0.84) <i>1.71</i> (2.12)	0.74 0.81	>100 >100	[0.67,0.81] <i>[0.74,0.88]</i>	Extreme evidence for H1 <i>Interpretation is the same</i>	p<2e-16
<b>H3.2:</b> Self v. max overall	1010	(-0.2,0.2)	0.031 (0.68) <i>0.20</i> (2.13)	0.045 0.093	<1/100 <1/100	[-.016,.11] <i>[0.031,0.15]</i>	Extreme evidence for H0 <i>Interpretation is the same</i>	p<2e-16

We next investigated how participants' responses and beliefs differ when seeing misinformation versus posting it oneself. To answer this research question, we ran non-directional Bayesian hypothesis tests where we set the null interval to be between [-0.2, 0.2]. For **H3.1**, we found strong evidence that participants respond with more effort when they post misinformation compared with when they see it posted by somewhat close ( $BF = 79.5$ ) and not close contacts ( $BF > 100$ ); however, inconclusive evidence that there was a difference in their responses when compared to close contacts.

For **H3.2**, we found extreme evidence ( $BF > 100$ ) that participants believe that people should respond with more effort when they post misinformation compared to seeing it by not close contacts and inconclusive evidence ( $BF = 6.92$ ) when comparing their misinformation posts to those posted by somewhat close contacts. Finally, we found very strong evidence ( $BF < 1/100$ ) towards the null hypothesis that there is not a difference in the level of effort people believe one should use after posting misinformation oneself vs. seeing a close contact post it. These results indicate that participants believe the most effort should be afforded to counter misinformation posted by close contacts or themselves compared with countering misinformation posted by somewhat and not close contacts.

## Robustness Tests

We ran three robustness checks. First, we ran the registered hypothesis tests using summed effort values instead of maximum effort values. For H3.1 and H3.2, low-effort actions are excluded from this analysis because there are an unequal number of them described in Tables 3.1 and 3.2. In almost all cases, these tests yielded the same or a similar strength of evidence for the hypotheses. In the few instances where the results diverged, the registered test was inconclusive at the Bayes Factor threshold of 1/10 or 10, while the summed version of the test surpassed the threshold in the same direction (**H1.2**: somewhat vs. not close) or vice versa (**H3.1**: oneself vs. somewhat, **H3.2** oneself vs. close). In only one instance did the interpretation completely differ: for **H3.1** (oneself vs. max of all closeness), the registered test showed strong evidence for H1, whereas the summed test showed strong evidence for H0. Notably, the calculated effect size was positive in both cases. However, in the summed version of the hypothesis test, most of the 95% unrestricted highest density interval (HDI) lay below 0.2, placing it in the null interval.

Second, we used the generalized McNemar's Chi-square Test of Independence for categorical paired data to verify that the interpretation is similar if the data are analyzed categorically rather than as interval data. For every hypothesis, the chi-square test produced a p-value of  $<= 0.01$ . This outcome diverged from some of the pre-registered Bayesian analyses, which had found some inconclusive results or evidence pointing towards the null hypothesis for some hypotheses. These discrepancies occurred in tests where the effect size was small, and the HDI overlapped with the null interval used in that test.

Finally, we ran a categorical analysis to analyze the association between the three variables of interest: maximum effort level used when countering, response type (reported behaviors vs. opinions), and closeness level. The three-way interaction term was not significant, indicating that closeness does not moderate the relationship between effort level and response type. Overall, the results were similar to those in our pre-registered analysis. Effort level interacts with both closeness level and response type, with higher counts of high-effort actions when contacts are closer or when people are asked about their opinions rather than their actual behavior. See

Appendix D for additional details.

## Exploratory Analysis

### RQ4. How do beliefs about responses to misinformation differ based on various demographic factors?

Finally, we examine individual differences in beliefs about how individuals should respond to misinformation. This exploratory analysis complements previous work that examines individual differences in support of misinformation countermeasures implemented by governments, social media companies, and other institutions [199]. For example, several previous studies have found that Democratic individuals are more supportive of platform interventions than Republicans [128, 153, 199], but does this translate into increased support for individual-level measures such as social corrections?

Figure 3.4 shows the maximum effort level participants believe one should exert when encountering misinformation posted by close contacts, somewhat close contacts, not close contacts, or oneself for six demographic variables: age, gender, race, education level, income level, and American political party. See Appendix E for details on the percentage of each demographic category that believes one should respond with no, low, or high effort when encountering misinformation posted by others or oneself. Appendix E also shows the detailed Chi-square test results for each demographic category and closeness level.

For political party affiliations, we find differences in belief in response efforts between partisan groups for close and not close contacts. Strong Republicans supported ignoring posts containing misinformation by close contacts more than any other group, although the absolute difference is <10%, which may not have much practical significance. Furthermore, strong Democrats were more likely to support high or low effort responses to not close contacts more than any other group. Notably, we see that, except for not close contacts, at least 70% of respondents in all party affiliations said that one should respond with a high effort action (such as commenting on a correction, updating the post, or messaging the poster).

For age, the chi-squared test shows statistically significant differences in responses among age groups when considering close contacts, somewhat close contacts, and oneself. In general, older participants were more likely to believe one should exert a high level of effort when countering misinformation than younger participants. No significant differences were found between men and women. For racial groups, the only statistically significant difference found was for not close contacts, with Black and Asian Americans more likely to believe in responding with some effort than the other racial groups.

Finally, the percentage of American residents stating that one should use a high level of effort to counter misinformation drops as education or income level increases. The results from the chi-squared test show statistically significant differences in responses among various education groups regarding close contacts and not close contacts, and among various income groups regarding somewhat close contacts and not close contacts.

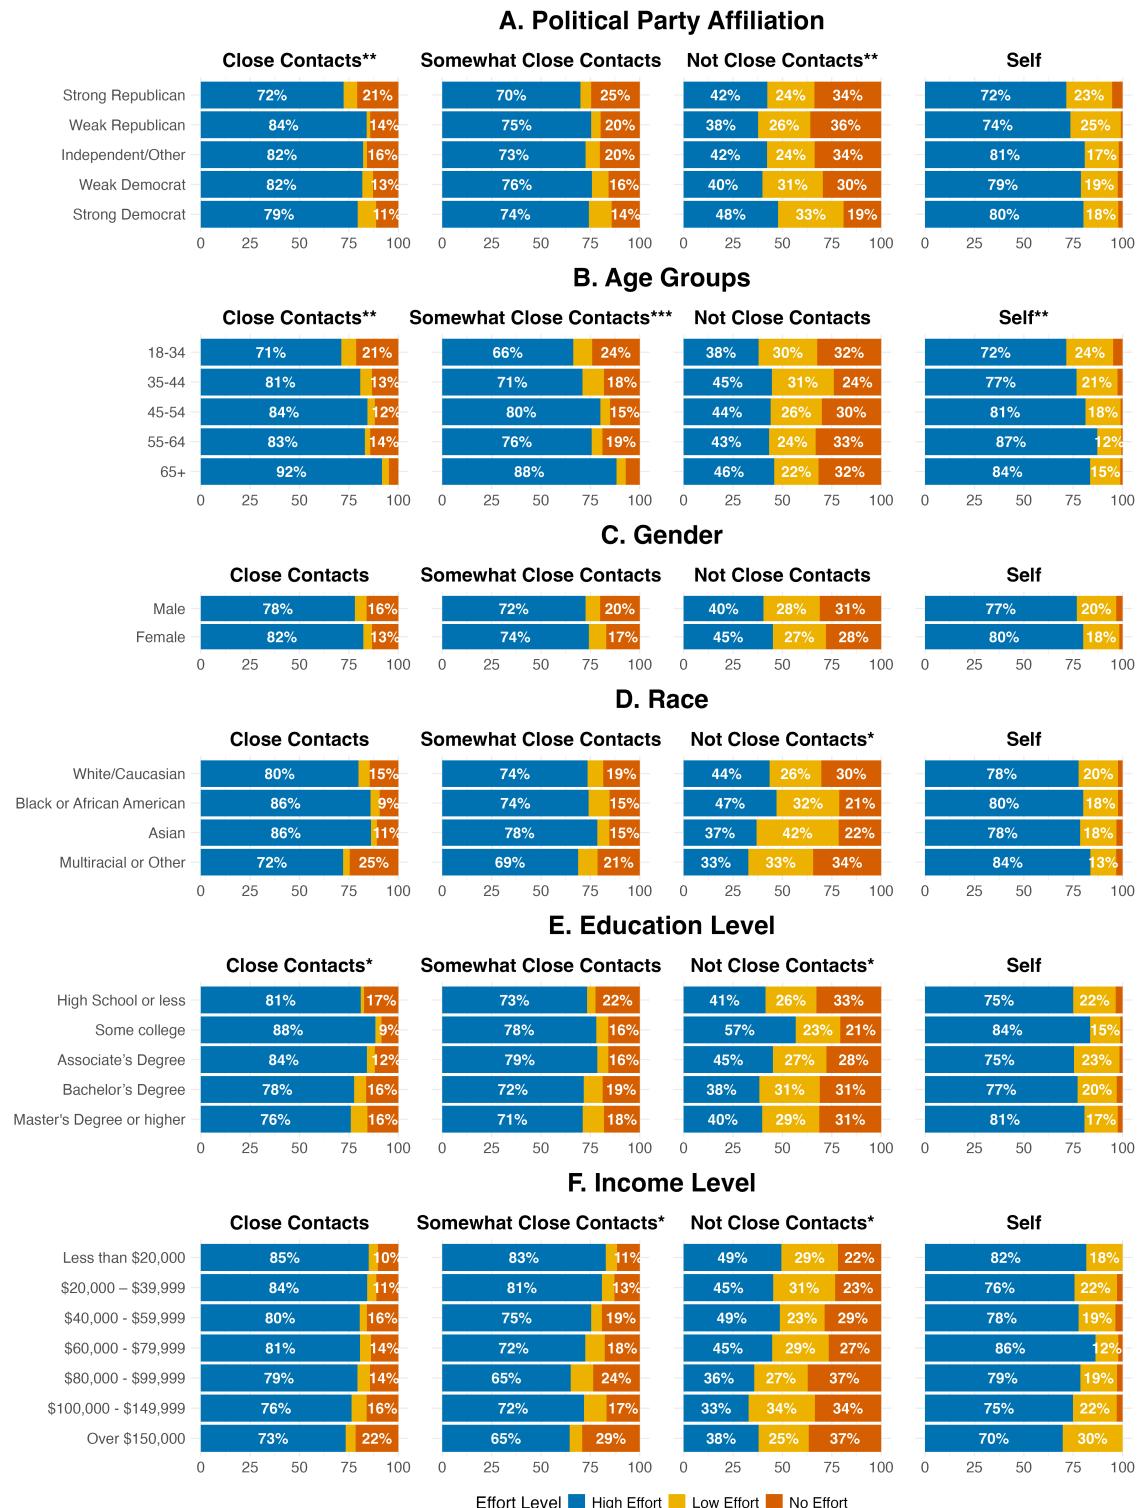


Figure 3.4: Highest effort level participants said one should respond with when seeing misinformation posted by others or oneself broken up by party, age, gender, race, education, and income. Chi-sq tests:  $p < 0.05^*$ ,  $p < 0.01^{**}$ , and  $p < 0.001^{***}$ .

### 3.3.5 Closeness Discussion

In this registered analysis, we compared individuals' beliefs about ideal responses and actual responses to misinformation posted on social media by close contacts, somewhat close contacts, not close contacts, and themselves.

We found overwhelming evidence of hypocrisy in people's responses to misinformation, aligning with our hypotheses (**H1.1, H2.1**). Participants believe others should exert more effort to counter misinformation than they report doing themselves. This pattern holds across all closeness levels, including misinformation posted by oneself, and remains robust to multiple ways of measuring effort. Since there is already a widespread belief that individuals should combat misinformation, efforts to encourage social corrections do not have to convince people to support individual corrections. Instead, they can focus on normalizing these practices and providing strategies to overcome situational constraints (e.g., time and cognitive effort required, social pressures) preventing people from acting.

Furthermore, our results indicate that people not only expect others to exert more effort but also tend to invest more effort themselves when addressing misinformation posted by close contacts compared to those who are somewhat close or not close at all (**H1.2, H1.3**). This increased effort may stem from the impression that correcting a close contact is more likely to be effective due to their relationship, making the effort more worthwhile. Alternatively, people might feel a stronger sense of responsibility to correct a closer contact whose beliefs and behaviors could impact them offline. Additionally, the types of responses differ across closeness as well. For example, people are more likely to privately message a close contact than a less close one. Different approaches may feel more appropriate depending on the source of misinformation. Providing users with a range of options, including private or low-effort methods like reporting, may increase their likelihood of engaging in countering behavior.

When comparing responses to misinformation posted by oneself versus someone else of varying closeness (**H3.1, H3.2**), participants reported putting more effort into responding to misinformation they had posted than to misinformation posted by somewhat or not close contacts. Their beliefs about ideal responses also reflected this pattern. Interestingly, we also found strong evidence that individuals respond with similar levels of effort to misinformation they posted compared with misinformation posted by a close contact. This suggests a similar view of responsibility when the source of misinformation is oneself or a close contact.

Finally, our exploratory analysis revealed demographic differences in beliefs about countering misinformation. Strong Republicans were less likely to believe that high effort should be exerted when countering close contacts, whereas strong Democrats were more inclined to believe some level of effort should be used for not close contacts. This partially aligns with prior research indicating that strong Democrats stood out in their support for institutional countermeasures compared with other partisan groups [199]. We only find this difference holds for not close contacts. Individual-level interventions give people agency to respond to content they believe is misinformation, potentially mitigating distrust in institutional definitions of misinformation. Our findings suggest that this approach to addressing misinformation may be more palatable across the political spectrum.

We found that older Americans were more likely to believe that one should exert high effort to counter misinformation than their younger counterparts. This difference may reflect broader

attitudes towards social media, as older individuals are more likely to perceive adverse effects associated with it [17] and, therefore, may be more motivated to address misinformation. Additionally, higher education and income levels were associated with a decreased belief that high effort should be exerted to counter misinformation. Interestingly, higher education and income are also associated with an increased concern and awareness of the negative impact of misinformation [27]. It may be that this concern does not necessarily translate into a belief in the effectiveness or necessity of individual countermeasures. Rather, these concerns may drive greater support for countermeasures on larger scales (e.g., government or platform), which is beyond the scope of this work but should be examined. Additionally, existing literature suggests that higher-income individuals are less generous overall [60, 178], which may extend to efforts to counter misinformation. This preliminary exploratory work can inform future research and platform policies.

## 3.4 Platform Analysis

Our primary research question here is whether people respond differently to misinformation based on the platform on which it was posted. More specifically:

1. How do people respond when they see misinformation on different platforms?
2. How do people respond when they realize they have posted misinformation on different platforms?

### 3.4.1 Related Work

Social media users frequently engage with multiple platforms for various purposes, as these platforms can fulfill competing needs [227, 249]. Each platform has different levels of content and account moderation [200], which means users may encounter varying amounts of misinformation depending on the platform [158]. Platforms can also prioritize content in diverse ways, from primarily chronological and time-based information to topic- or location-based arrangements [98]. Given the nature of the interactions, the types of connections, and the content on each platform, users might respond differently to misinformation depending on where it was posted. It is important to consider the types of different platforms and their typical uses.

#### Platform Categorization

First, we consider the various types of social media platforms. Previous research has categorized the leading platforms in several ways. Kietzmann et al. (2011) classified platforms according to seven primary building blocks: identity, conversations, sharing, presence, relationships, reputation, and groups [119]:

- **Identity** - This block refers to how users disseminate their identity online. Many platforms require users to build profiles, and on many sites, people predominantly retain their real identities (e.g., Facebook, LinkedIn), while on others, they develop virtual identities or use pseudonyms for their usernames (e.g., Reddit).

- **Conversations** - The Conversations block refers to how users communicate with others on the platform. For example, a brief status update on Twitter, which does not require a response from all followers, or a more thoughtful discussion-based post on Reddit.
- **Sharing** - This refers to the transmission of content on the platform and is often related to the platform's purpose. For example, uploading pictures to Instagram or videos to YouTube and TikTok.
- **Presence** - Presence indicates whether other users are online and available to communicate.
- **Relationships** - This building block refers to how users connect with others. These connections may arise from shared interests or topics (e.g., Reddit), from knowing each other in person (e.g., Facebook), or from being fans of others' work (e.g., Twitter, YouTube, TikTok).
- **Reputation** - Reputation refers to how users assess their social standing with respect to others, which can differ based on the platform. Reputation can sometimes be directly evaluated by likes, shares, view counts, comments, and upvotes.
- **Groups** - Finally, the Groups block illustrates how users form communities with each other. Users can create groups like Twitter lists to categorize their friends and followers. Other groups may be open to everyone or require an invitation.

The idea is that platforms tend to focus on three or four blocks at once rather than just one or attempting to address all of them [119]. For example, on LinkedIn, the main building block is identity, but there is also an emphasis on relationships and reputation since it is primarily a career-oriented platform. On YouTube, the main focus is on sharing content, but attention is also given to reputation, groups, and conversations. In the case of Facebook, the predominant block is relationships, but presence, identity, conversations, and reputations are also a part of it [119].

Another frequently cited typology of platforms was developed by Zhu and Chen (2015), and it categorizes platforms into four general categories based on two dimensions: connection type and message type [249]. Connections can either be profile-based, where the focus is on individual users and their profiles, or content-based, where the focus is instead on the content posted. The messages can either be customized for a specific audience or broadcast to the public. Table 3.7 summarizes their four main categories of social media platforms.

Table 3.7: Social media platform categorization as developed by Zhu and Chen (2015) [249].

	<b>Customized Messages</b>	<b>Broadcast Messages</b>
<b>Profile-Based Connections</b>	<i>Relationship</i> - Platforms based on users connecting, like Facebook, LinkedIn, WhatsApp	<i>Self-Media</i> - Users broadcast to their followers, like Twitter
<b>Content-Based Connections</b>	<i>Collaboration</i> - Discussion-driven platforms to find answers or advice, like Reddit or Quora	<i>Creative Outlets</i> - Users share their interests and creative pursuits, like YouTube, TikTok, Pinterest

While comprehensive, there are aspects that these typologies do not directly incorporate, such as the distinction between private and public content, as well as profiles [227], or the way information is displayed (based on time, location, or topic) [98]. The privacy aspect is likely relevant to countering misinformation, as sharing false information could directly harm one's reputation if associated with one's true identity. I propose a modification to Zhu and Chen's typology that incorporates Kietzmann et al.'s "identity" building block more prominently by dividing each category into two: platforms with predominantly private or anonymous profiles and platforms with predominantly public profiles. Table 3.8 illustrates this updated categorization.

Table 3.8: Modified social media platform categorization.

	Customized Messages		Broadcast Messages	
Profile-Based Connections	<i>Private Relationship:</i> Facebook, WhatsApp, Messenger, Telegram	<i>Public Relationship:</i> LinkedIn, Nextdoor	<i>Private Self-Media:</i> Snapchat, BeReal	<i>Public Self-Media:</i> Twitter, Threads, BlueSky
Content-Based Connections	<i>Personal Collaboration:</i> Reddit	<i>Public Collaboration:</i> Quora	<i>Personal Creative Outlets:</i> Instagram	<i>Public Creative Outlets:</i> YouTube, TikTok, Pinterest, Instagram

In general, users have become more private over time [63], and it is now estimated that nearly half of all social media accounts in the U.S. are set to private [9, 90]. According to a Statista survey in 2018, around 45% of American social media users report that all their social media accounts are private, while another 27% state that some of their accounts are private [65]. Platforms typically do not publicly disclose statistics on the number of private versus public accounts. However, we can make estimates based on various sources, including public opinion polls, the default or typical privacy settings on each platform, and the platform's structure.

When considering the *Relationship* platforms, the sites where most users maintain private or partially private accounts include Facebook, WhatsApp, Messenger, and Telegram. At the same time, LinkedIn and Nextdoor tend to be more public. A 2020 survey conducted for PC Magazine revealed that 57% of Facebook users have accounts that are at least partially private, with nearly 80% saying that they do not share their friends list publicly [90]. Facebook operates as a friendship-based platform rather than a follower-based one, meaning users can usually connect only if both parties agree. For WhatsApp, messages are end-to-end encrypted, and the help center explicitly states that they go to "great lengths to build WhatsApp in a way that helps people communicate privately"<sup>1</sup>. Similarly, Messenger has end-to-end encryption for messages enabled by default<sup>2</sup>. Although Telegram has groups and channels that, in some ways, have similar functionality to a platform like Reddit, it primarily advertises itself as a messaging app that is similar

<sup>1</sup><https://faq.whatsapp.com/595724415641642>

<sup>2</sup><https://messengernews.fb.com/2023/12/06/launching-default-end-to-end-encryption-on-messenger/>

to WhatsApp<sup>3</sup>. Meanwhile, Nextdoor requires users to use their real name and address, as the platform is designed to connect individuals within their local community<sup>4</sup>. By default, users have a public profile on LinkedIn unless they choose to hide it<sup>5</sup>.

For the *Self-Media* platforms, Snapchat and BeReal are primarily private [9], while Twitter, Threads, and BlueSky tend to be more public. By default, Snapchat only permits users to be contacted by their “friends”<sup>6</sup>. BeReal is also a friendship-based platform where posts are shared exclusively with friends by default. This platform does not even have a public option: users can share only with at most “friends of friends”<sup>7</sup>. Twitter, BlueSky, and Threads are follower-based platforms. A 2019 Pew Research survey found that only 13% of U.S. adult Twitter users reported keeping their accounts private [187]. Both Threads [214] and BlueSky are considered Twitter competitors, and as of 2023, BlueSky lacked the functionality to support private accounts<sup>8</sup>.

When considering the *Collaboration* platforms, Reddit has predominantly anonymous users, while Quora accounts are tied to real identities. Although Reddit is a primarily public platform, according to Reddit, most of their users choose to remain anonymous by selecting screen names<sup>9</sup>. On the other hand, Quora encourages users to use their real names and credentials, and the functionality to answer questions anonymously was removed in 2021<sup>10</sup>.

Finally, regarding the *Creative Outlet* platforms, these are generally follower-based instead of friendship-based and are typically public by default unless the user is a minor. YouTube videos are public by default<sup>11</sup>, as are Pinterest accounts<sup>12</sup> and TikTok accounts<sup>13</sup>, although all of them can be set to private. It remains unclear how many adult users maintain private profiles on these platforms. On TikTok, a Pew Research study found that only about half of adults have ever posted a video, and about 40% had posted a public video [32]. This suggests that most users who post videos do so publicly, but it is unclear how many lurkers have private versus public accounts. Instagram users represent a mix of private and public accounts. A 2020 PC Magazine survey found that about half of Instagram users restrict or filter comments on their posts [90]. Additionally, there is a common practice of creating secondary fake accounts, also known as “finstas”, to preserve privacy and share content with a smaller group of people [9].

## Platform Usage

In addition to platform type, it is critical to understand how the different platforms are being used and how platform usage may relate to the issue of countering misinformation. Recent research on

<sup>3</sup><https://telegram.org/faq#q-what-is-telegram-what-do-i-do-here>

<sup>4</sup><https://help.nextdoor.com/s/article/use-your-true-identity>

<sup>5</sup><https://www.linkedin.com/help/linkedin/answer/a528138>

<sup>6</sup><https://help.snapchat.com/hc/en-us/articles/7012343074580-How-do-I-change-my-privacy-settings-on-Snapchat>

<sup>7</sup><https://help.bereal.com/hc/en-us/articles/10444893090205-Audience>

<sup>8</sup><https://bsky.social/about/blog/5-19-2023-user-faq>

<sup>9</sup><https://support.reddithelp.com/hc/en-us/articles/7420342178324-How-does-being-anonymous-work-on-Reddit>

<sup>10</sup><https://productupdates.quora.com/Removing-anonymous-answers>

<sup>11</sup><https://support.google.com/youtube/answer/157177>

<sup>12</sup><https://help.pinterest.com/en/article/make-your-profile-private>

<sup>13</sup><https://support.tiktok.com/en/account-and-privacy/account-privacy-settings/making-your-account-public-or-private>

platform engagement and advertising has assessed the significance of platform type, indicating that users have varying motivations for using these platforms and that engagement can be highly context-specific [172, 227].

Pelletier et al. (2020) [172] conducted an exploratory survey study and identified four main purposes behind platform usage:

- **Social** - Staying connected with friends and family
- **Informational** - Keeping up with the news and trends
- **Entertainment** - Following sports and celebrities, playing games, and sharing memes.
- **Convenience** - Watching random content when bored or scrolling.

These researchers found that among three major platforms (Facebook, Twitter, and Instagram), users preferred Twitter and Instagram for their social needs, Twitter for their informational needs, and Instagram for their entertainment needs [172]. It is unsurprising that Twitter is preferred for informational purposes, given its classification as a *Self-Media* platform where messages are broadcast to followers. It is similarly unsurprising that Instagram is preferred for entertainment, considering that it is a *Creative Outlet* platform. While Facebook had the largest user base, it also had the lowest actual usage, which may explain why it was not identified as the top platform for social purposes, even though it is a *Relationship*-driven platform.

Similarly, Voorveld et al. (2018) considered 11 dimensions of social media engagement across eight platforms: Facebook, YouTube, Twitter, LinkedIn, Instagram, Pinterest, Snapchat, and Google+ [227]. Besides the four dimensions identified by Pelletier et al., they also considered negative emotions related to the content, negative emotions related to the platform, stimulation, identification, practical use, innovation, and empowerment [227]. Table 3.9 shows which platforms scored the highest and lowest in each dimension, with Google+ excluded as it is no longer operational.

Table 3.9: Social media engagement on various platforms per dimension [227]. The innovation and empowerment dimensions are excluded, as all platforms scored low.

	<b>Highest Engagement</b>	<b>Lowest Engagement</b>
<b>Social interaction</b>	Facebook, Snapchat	YouTube, Pinterest
<b>Informational/Topicality</b>	Twitter, LinkedIn	YouTube, Snapchat
<b>Entertainment</b>	Snapchat, YouTube	LinkedIn, Twitter
<b>Convenience/pastime</b>	Instagram, Facebook	LinkedIn
<b>Negative emotion - content</b>	Twitter, Facebook	Pinterest, LinkedIn
<b>Negative emotion - platform</b>	Facebook	Instagram, YouTube
<b>Stimulation</b>	Pinterest	LinkedIn
<b>Identification</b>	Facebook	LinkedIn, YouTube
<b>Practical use</b>	Pinterest	Snapchat

The researchers found that Facebook and Snapchat, both friendship-based platforms that focus on profile-based connections, ranked highest on the social dimension. Twitter and LinkedIn, primarily public platforms, scored highest on the informational dimension. Snapchat and YouTube, both *Creative Outlet* platforms, scored highest on entertainment. Overall, they noted that the

most significant similarity among the platforms was that most of them, except YouTube and Pinterest, were a way to stay informed and up-to-date [227].

## Misinformation Exposure

Despite platforms being used for differing purposes, all platforms have misinformation to some extent, even if it is more common on some platforms than others [158, 210]. For example, platforms that prioritize informational content, such as X (formerly Twitter), may have more news-related misinformation than those focused primarily on entertainment, like Instagram and Pinterest. Given that previous research has found that users often feel powerless to combat misinformation because there is so much of it [212], the extent of exposure might be related to the likelihood of countering it.

Finally, previous work in this chapter has demonstrated that closeness influences the willingness to counter, with individuals being more inclined to counter close contacts. Users are more likely to be connected to close contacts on specific platform types, such as the *Private Relationship* platforms like Facebook and WhatsApp. Additionally, any potential social factors at play, such as the desire to maintain interpersonal relationships and concerns about credibility when countering [163], might be less relevant on platforms where users are primarily anonymous, like Reddit. Understanding how different platform types relate to misinformation exposure and countering is a crucial aspect that could help inform future platform policies and public messaging regarding these issues.

### 3.4.2 Results

#### RQ1: Responses to seeing misinformation by platform

We are interested in examining any differences in exposure to misinformation and responses among the various platforms. Our subquestions are as follows:

1. Which platforms do most people report encountering misinformation? Which platforms do people report seeing the **most** misinformation?
2. Is the frequency of platform usage connected to exposure to misinformation?
3. How do the actions individuals take when encountering misinformation differ across platforms?

Figure 3.5 shows the platforms where people report seeing misinformation compared to their usage of those platforms. The platforms are sorted by the total number of participants who visit each platform at least weekly, from highest to lowest. The most visited platforms among participants were YouTube, Facebook, Reddit, X, and Instagram, with more than half of the participants claiming they visit these platforms at least once a week.

The most frequently used platforms are not fully aligned with those where participants reported encountering misinformation. The total number of participants who claim to have seen misinformation at least once is highest for Facebook, followed by YouTube, X, Reddit, Instagram, and TikTok. In addition, participants were asked to select just one platform on which they believe they have seen the most misinformation, and Facebook (43%) and X (27.8%) dominated.

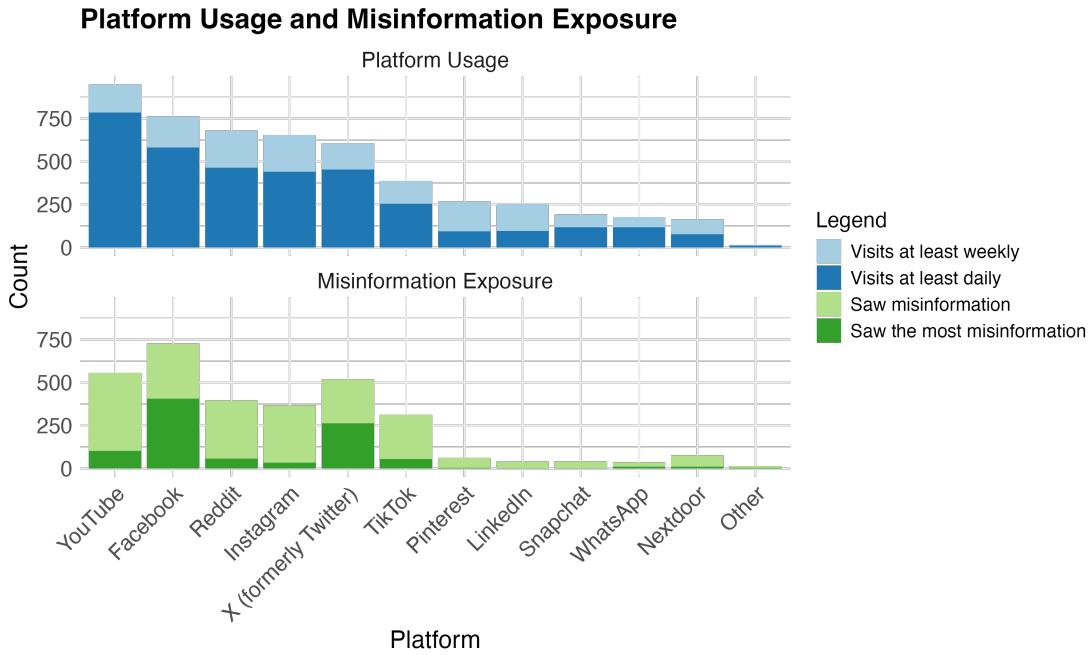


Figure 3.5: The number of participants who use each platform at least weekly or daily compared with the number of participants who have seen misinformation on that platform.

Table 3.10 displays the number of participants who saw misinformation on each platform, categorized by their closeness to the misinformation poster. The percentages indicate the proportion of people who observed misinformation at each closeness level, calculated from the total number of individuals who saw misinformation on that platform. Since individuals can see misinformation from multiple closeness levels, the percentages in each row total more than 100%.

Participants were more likely to encounter misinformation from close or somewhat close contacts on platforms like Facebook, WhatsApp, Snapchat, and Instagram compared to other platforms. On other platforms, such as YouTube, X, Reddit, TikTok, Nextdoor, and Pinterest, the vast majority of misinformation was posted by contacts who were not close. These differences likely relate to the purpose and functionality of the platforms, with people more likely to follow and communicate with closer contacts on more private platforms like Facebook and WhatsApp compared to other platforms.

Next, we examine the maximum effort level employed by participants when encountering misinformation on the top six platforms with the most misinformation: Facebook, YouTube, X, Reddit, Instagram, and TikTok. Figure 3.6 illustrates that on Facebook and Reddit, participants were more likely to report engaging in high-effort actions to counter misinformation. For Facebook, this difference may be due to the greater amounts of misinformation originating from close and somewhat close contacts on that platform compared to others. For Reddit, it could be because it is a *Collaboration* platform designed for discussion, where most users remain anonymous or adopt a virtual identity, reducing the likelihood of in-person conflicts.

However, if we consider the maximum effort level against only not close contacts among these top six platforms, differences remain (Figure 3.7). Reddit had the highest percentage of

Table 3.10: Number and percent of people who saw misinformation from close, somewhat, and not close contacts per platform. The top three platforms by percentage per closeness level are bolded. Platforms are sorted by the number of people who have seen misinformation on the platform.

Platform	Close Contacts	Somewhat Close	Not Close Contacts	Saw Misinformation	Most Misinformation
Facebook	<b>128 (17.6%)</b>	<b>305 (42.0%)</b>	513 (70.7%)	726	<b>405 (43.0%)</b>
YouTube	7 (1.3%)	24 (4.3%)	<b>533 (96.0%)</b>	555	<b>102 (10.8%)</b>
X	10 (1.9%)	57 (11.0%)	484 (93.1%)	520	<b>262 (27.8%)</b>
Reddit	1 (0.3%)	13 (3.3%)	<b>383 (97.0%)</b>	395	58 (6.2%)
Instagram	22 (6.0%)	85 (23.2%)	310 (84.5%)	367	34 (3.6%)
TikTok	6 (1.9%)	18 (5.8%)	<b>296 (95.2%)</b>	311	54 (5.7%)
Nextdoor	1 (1.3%)	6 (7.9%)	69 (90.8%)	76	9 (1.0%)
WhatsApp	<b>8 (22.2%)</b>	<b>13 (36.1%)</b>	17 (47.2%)	36	9 (1.0%)
LinkedIn	1 (2.5%)	12 (30.0%)	31 (77.5%)	40	1 (0.1%)
Snapchat	<b>5 (11.6%)</b>	<b>11 (25.6%)</b>	32 (74.4%)	43	2 (0.2%)
Pinterest	2 (3.2%)	4 (6.5%)	57 (91.9%)	62	3 (0.3%)
Other	1 (6.7%)	3 (20.0%)	12 (80.0%)	15	3 (0.3%)
Total	148	370	880	942	942

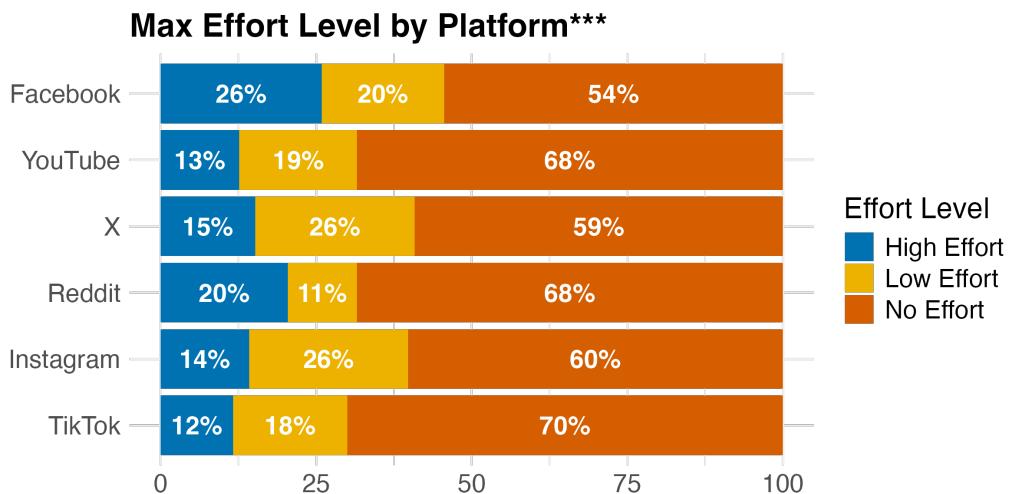


Figure 3.6: Highest effort level participants said they responded with when seeing misinformation posted by others on the top 6 platforms for misinformation. Chi-sq test:  $p < 0.001^{***}$ .

participants who reported engaging in high-effort responses. Participants were more likely to exert effort to counter posts on Facebook, X, and Instagram than on YouTube and TikTok. These differences may be due to the ease of reporting or other platform functionalities and are worth investigating further. Qualitative responses on how closeness and platform may affect countering efforts are further examined in Chapter 3.

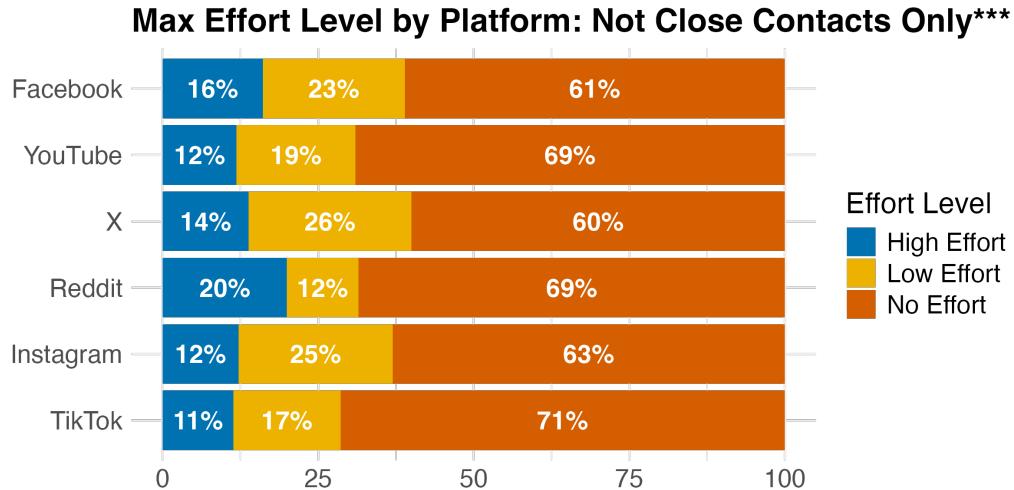


Figure 3.7: Highest effort level participants said they responded with when seeing misinformation posted by **not close contacts only** on the top 6 platforms for misinformation. Chi-sq test:  $p < 0.001^{***}$ .

## RQ2: Responses to posting misinformation by platform

Next, we examine whether there are differences in the posting of misinformation across various platforms. More specifically:

1. Which platforms do people report posting misinformation?
2. How does what people do when they realize they have posted misinformation differ between platforms?

Figure 3.8 shows the platforms where individuals report posting misinformation, whether intentionally or unintentionally. The platforms are sorted by the total number of participants who said that they had unintentionally posted misinformation on that platform, from highest to lowest. Notably, Facebook emerged as the platform with the most posted misinformation, followed by X, and then Reddit and Instagram. Interestingly, the most used platform by participants, YouTube, does not appear in the top three. Table 3.11 shows the top six platforms and details the selected interventions, along with the number of individuals who reported posting misinformation on that platform either intentionally or unintentionally.

Next, we examine the maximum level of effort participants exerted after realizing they posted misinformation on the top four platforms. In this case, we considered only the top four platforms instead of the top six, as no other platform had more than 25 participants admitting to posting

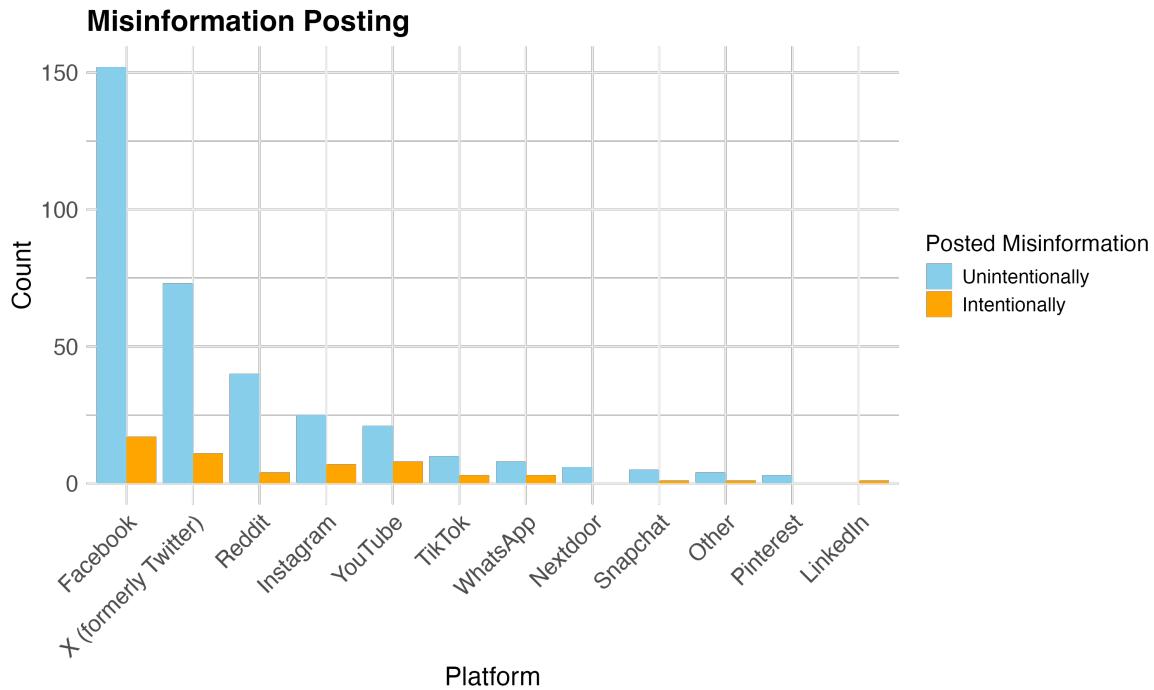


Figure 3.8: Number of participants who admitted to posting misinformation on each platform.

Table 3.11: Summary of Interventions Selected after Posting Misinformation

Platform	Intervention Type					Posted Misinformation	
	Left post	Delete post	Comment correction	Update post	Post correction	On Accident	On Purpose
Facebook	13	111	34	48	27	152	17
X	7	57	12	15	20	73	11
Reddit	7	19	9	19	4	40	4
Instagram	3	18	7	6	3	25	7
YouTube	3	9	7	8	3	21	8
TikTok	3	2	3	1	2	10	3

misinformation. Figure 3.9 illustrates the maximum level of effort exerted after accidentally posting misinformation on Facebook, X, Reddit, and Instagram. The chi-squared test of independence was not statistically significant. While there may be differences among platforms, the sample size for most platforms was small, indicating that more research in this area is needed to gain further insights.

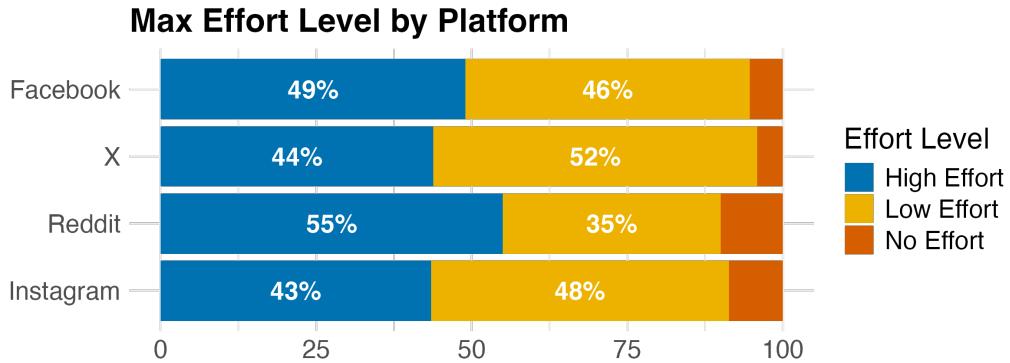


Figure 3.9: Highest effort level participants said they responded with after realizing they had posted misinformation accidentally among the top 4 platforms. Chi-sq test was not significant.

### 3.4.3 Platform Discussion

Overall, we observed differences in misinformation responses depending on the platform where the content was posted. First, we found that the social media platform is related to misinformation exposure. Participants reported encountering more misinformation on Facebook and X (the second and fifth most used platforms) than on other sites. While misinformation exposure is related to platform usage, Facebook and particularly X overperformed in the survey question “Which platform do you see the most misinformation on?” compared to the number of participants using those platforms.

Furthermore, we found that the maximum effort level in countering misinformation varied across platforms and that proximity to the misinformation poster cannot fully explain these differences. While participants reported seeing more misinformation from close and somewhat close contacts on certain platforms (like Facebook and WhatsApp) than on others (such as YouTube, Reddit, or TikTok), when considering responses regarding only not close contacts, we still observed a greater level of effort on some platforms than others. Facebook, X, and Instagram had the highest percentage of participants saying they engaged in any level of effort to counter misinformation. In contrast, Reddit had the highest percentage of participants reporting that they exerted a high level of effort. These differences may be related to platform social norms, the ease of reporting misinformation, or platform functionalities. For example, Reddit is a more discussion-oriented platform that may allow for higher-effort actions like detailed debunking.

Finally, we found that most participants who admitted to accidentally posting misinformation reported that they did so on Facebook. The second-highest platform was X, with only half as many participants claiming to have unintentionally shared misinformation. Again, this may be

related to the functionality and purpose of the platforms. People may be more inclined to post on platforms like Facebook and X while being more passive viewers of content on the *Creative Outlet* platforms like YouTube and TikTok.

### 3.5 Overall Discussion

This work has several practical implications for promoting public participation in countering on-line misinformation in educational and technological contexts. First, our research demonstrates the widespread approval of social corrections online, indicating the social desirability of these behaviors. Prior work shows that highlighting the social desirability of reporting misinformation as an injunctive social norm can motivate reporting [84] and decrease the sharing of misinformation [110]. More broadly, there is strong evidence that social signals (e.g., engagement metrics and comments) influence responses to misinformation [19, 57]. However, they can also encourage harmful behavior, such as sharing misinformation that aligns with one’s pre-existing beliefs [134]. Therefore, platforms and organizations can successfully promote individual interventions by emphasizing their popularity among users, but they must be careful to avoid inadvertently empowering users with questionable motives.

Additionally, the observed disparity between reported behaviors and beliefs could be leveraged to encourage greater public participation. Research on hypocrisy suggests that one of the most effective strategies for driving behavioral changes is to have individuals publicly commit to pro-social actions, such as signing a pledge, and then be privately reminded of times they have failed to follow through [209]. Public call-outs are less effective, as they may prompt people to save face or rationalize their failures by reducing their support for the targeted behavior [209]. Social media platforms could encourage users, such as those who sign up to contribute to Community Notes programs, to publicly support social corrections or similar measures during educational sessions and regularly remind them of their commitment going forward.

Moreover, the result that people believe more should be done to respond to misinformation than they report doing themselves, in conjunction with the differences we see between platforms, indicates that there may be barriers to employing social corrections or reporting features that could be mitigated by platform design, such as improving transparency, usability, and technical support of these features [248]. Platforms should educate users about their reporting systems when joining and provide periodic updates to keep them informed. Users are also more likely to use reporting features if they perceive them as effective. Therefore, sharing information about the outcomes of reports filed or community notes written can incentivize people to use these programs [239, 248]. Furthermore, pop-up windows have been used on several platforms to ask users if they wish to share content they have not reviewed [56], and this method could be used in scenarios where users delete content. Instead of deleting potentially misleading posts, users could be encouraged to edit or update their posts with accurate information.

Recognizing that susceptibility to misinformation varies across demographic groups [160], educational efforts could tailor strategies to effectively reach different populations [42]. For example, older adults are particularly vulnerable to political misinformation, potentially due to lower digital literacy levels [43]. This age group also tends to support higher-effort responses to misinformation encountered online. Therefore, training efforts for older adults might prioritize

digital media literacy over encouraging social corrections. People who are less vulnerable to misinformation, on the other hand, can be promptly educated about operationalizing corrections and leveraging specific platform affordances. Platforms could utilize their internal data to identify these users or implement reputation systems, like X’s Community Notes program, where users earn “Rating Impact” scores based on the helpfulness of their contributions<sup>14</sup>. Overall, educational efforts should be designed to account for individual differences in both vulnerability to misinformation and perceived responsibility to counter it.

Lastly, there are a myriad of individual differences beyond demographics that influence vulnerability to misinformation and the likelihood to correct it that should be examined further in future work. For example, evidence suggests that those with a tendency toward analytical thinking are less susceptible to misinformation [70, 173]. These are likely the same individuals capable of providing accurate and meaningful corrections, as some level of cognitive effort is necessary for higher-effort responses to misinformation (e.g., commenting on a correction). Platforms can encourage users to think critically by using accuracy prompts or similar measures [70, 177]. In addition, platforms and other institutions can target educational resources towards those with a propensity to engage in critical thinking.

## 3.6 Conclusions

### 3.6.1 Limitations

There are several limitations to this work. First, our sample is not demographically representative of all United States residents. We specifically focused on active social media users to better understand current user behavior on social media platforms. While this targeted sample provides relevant information to platforms about how their users act and what they believe, a more demographically representative survey could provide additional information about less active users who can also influence the spread of misinformation. Additionally, while we collected high-level demographic data, we did not investigate the role of more complex individual features, such as analytical reasoning or values, and how they may interact with one’s propensity to intervene against misinformation. We leave this to future work.

Next, participants were asked to recall how they had responded to misinformation they had seen on social media in the past, which they may or may not have encountered recently. This could have led to memory or recall errors. Furthermore, we note the possibility of demand effects or other biases (e.g., social desirability) influencing survey responses. We took care to present the survey to participants as generally about misinformation online without including details that may reveal our expectations. Future work could consider using platform data or conducting a field experiment to observe how people respond to misinformation in real-world contexts. For example, platform data on reporting or social corrections could confirm whether people counter closer contacts more than less close contacts and if these behaviors differ by platform.

Additionally, the results linked to RQ3 in the closeness analysis may have limited generalizability due to the fundamental difference in the potential actions one can employ to correct others compared with correcting oneself. We attempted to enumerate commensurate responses

<sup>14</sup><https://communitynotes.x.com/guide/en/contributing/writing-and-rating-impact>

to misinformation, such as commenting a correction on someone else’s post or one’s own post. However, the low-effort actions are, by nature, not equivalent actions (reporting someone else’s post vs. deleting one’s post). Additionally, there were more listed low-effort actions for responding to others than responding to oneself. We address this in one of our robustness checks, where we compared summed rather than maximum effort levels and excluded all low-effort responses. However, future studies should consider this inherent limitation when conducting this type of analysis.

Finally, we did not investigate possible differences in behavior, not just the beliefs, of various demographic groups. We additionally had limited platform data on the posting of misinformation because the posting of misinformation by participants was strongly concentrated among just a couple of platforms (Facebook and X). Finally, future work could investigate if behavior or beliefs about how to engage on social media are related in any way to support for platform or government measures to counter misinformation. We will explore how misinformation exposure may be related to support of various countermeasures in Chapter 4.

### 3.6.2 Contributions

This study makes an important contribution to the literature on individual-level interventions against misinformation. Our results indicate that facilitating individual responses to misinformation seen or accidentally posted on social media is a viable approach to reducing the spread of misinformation and even preventing belief in it. Using a large sample of active social media users in the US, we demonstrate the widespread belief that individuals should counter misinformation despite a tendency to not always act on this belief themselves. The nature of responses and the willingness to expend effort vary based on the user’s relationship with the misinformation poster and the platform, highlighting opportunities to educate the broader population about different ways to take action depending on their perceived situational constraints. These insights inform efforts to encourage public participation in mitigating the impact of misinformation and suggest ways that platforms can enhance their countering tools to empower users to engage more actively in maintaining the integrity of their online information environment.

# Chapter 4

## Improving User-based Countermeasures

In Chapter 2, I analyze current user opinions and behaviors regarding user-led countermeasures. This chapter will focus on media literacy as a way to empower individuals and combat misinformation at the user level. Media literacy and training games can be developed by the government, platforms, civic society, or even individuals. They are scalable and, therefore, practical.

Our primary research question for this chapter is: How can we improve user-based countermeasures? More specifically,

1. Are media literacy and training games effective at **improving misinformation detection?**
2. Are media literacy and training games effective at **improving the countering of misinformation?**
3. What factors, such as the **poster or platform**, are associated with one's willingness to counter misinformation?

### 4.1 Introduction

As misinformation continues to affect societies around the world, much recent research has focused on developing and deploying effective interventions [35, 59, 100, 129]. One of the most studied intervention types involves media literacy and digital literacy as a preventive measure [59]. Media literacy encompasses many different types of interventions, ranging from short tips [92], in-person training sessions [151], to fake news games [28, 138].

Most media literacy experiments focus on improving participants' skills to better discern truth from falsehoods and improve their critical understanding of the media they encounter [106]. Although many studies also investigate the effectiveness of media literacy in behavioral outcomes, the behaviors studied typically focus on reducing the frequency of harmful, risky, or antisocial behaviors such as engaging with or sharing misinformation or participating in risky sexual encounters. To our knowledge, no studies have yet focused on improving the willingness and ability of participants to counter it [106]. We seek to fill this research gap by running an experiment focused on increasing motivated individuals' willingness, knowledge, and ability to counter misinformation.

This study was conducted with 23 motivated government analysts, where participants were shown a series of social media posts and asked if they believed the posts to be true or false. Par-

ticipants were also shown several explicitly false or misleading social media posts and asked to describe if and how they would respond. A survey was administered before and after an interactive, in-person training session to determine if misinformation detection ability improved and if countering willingness increased. We also included open-ended questions that allowed participants to explain their reasoning so that we could qualitatively analyze the factors associated with willingness to engage in countering behavior.

This chapter will be broken into two main results sections: **improving misinformation detection** and **improving misinformation countering**. The results of this case study will help inform how to encourage and improve user-based countermeasures, such as social corrections, thus contributing to broader efforts to combat misinformation.

## 4.2 Related Work

### 4.2.1 Media Literacy Interventions

Media literacy interventions consist of educational initiatives designed to enhance the public's civic discourse by improving critical thinking ability when reading media content [92, 106]. One type is the development and usage of fake news games. These games include the Bad News Game [28], Go Viral! [154], Troll Spotter [138], and Harmony Square [193]. They are designed to be an interactive and fun way to help players detect misinformation [147, 154].

A related concept to media literacy is the theory of inoculation, sometimes referred to as "prebunking." Inoculation includes interventions such as preemptive warning messages or other anticipatory interventions meant to "inoculate" people, much like a vaccine would, from later believing that misinformation or harmful content [142]. Similarly, media literacy is also intended to improve participant resilience when encountering misleading, harmful, or false messages.

The effectiveness of media literacy as a preventative measure against misinformation has been widely debated in the literature. There is a lack of consensus on whether it is effective, which types are effective, how effective it is, how long the effectiveness lasts, and in which contexts it is most effective [20, 28, 106, 154].

### 4.2.2 User-Based Interventions

In the context of countering misinformation, user-based interventions refer to actions that social media users can take when directly engaging with misinformation [21, 107]. For example, social media platforms typically allow users to report other users or posts [164]. Social media users can also employ social corrections. Social corrections attempt to debunk the misinformation poster by publicly commenting on their post, privately messaging the user, or other related means [21].

User-based measures are an essential type of misinformation countermeasure. Although most platforms employ some automated moderation, they also rely on social media users to report anything those algorithms miss. In addition, users can comment and engage in social corrections to help debunk misinformation. Having a trusted messenger, such as a friend or family member, debunk misinformation has been found to be effective in several studies [21, 37, 144, 213, 233].

User-based interventions are a vital tool in the fight against misinformation, and there is currently little to no work being done utilizing media literacy interventions to improve those measures.

## 4.3 Data and Methods

The participants in this study were government analysts who had signed up for social cybersecurity training. They participated in a two-week training exercise called “OMEN” (Operational Mastery of the Information Environment) [121], which ran from 02/05/24 to 02/16/24 in Orlando, FL. This study was conducted on 02/07/24 during this training exercise.

### 4.3.1 OMEN Overview

OMEN is a project administered by the Office of Naval Research designed to teach analysts how to evaluate their online information environment. The first week of OMEN was reserved for various training sessions during which participants learned about social cybersecurity and how to use various software tools. During the second week, analysts participated in the OMEN game, a training game that teaches analysts and decision-makers how to detect and counter misinformation on social media. It is designed to be a “train-as-you-fight” game, where the storyline is based on real events, and the data is realistic in volume and speed. The game accommodates real tools and workflow, including ORA and NetMapper<sup>1</sup>. It is a multi-day event that generally matches what the analysts would encounter on their day job. See our associated technical report [121] for more details on the design of the OMEN game, the creation of the storyline, the collection of data, the learning objectives, and the lessons learned.

### 4.3.2 Participant Demographics

Twenty-three participants completed the study by filling out the pre-test survey, attending the training sessions, and completing the post-test survey. All participants were members of the defense community from one of the “Five Eyes” nations (US, UK, Canada, Australia, and New Zealand). A high-level summary of participant demographics is provided below.

- **Gender:** 19 men (82.6%), 4 women (17.4%)
- **Race:** 20 white (87.0%), 3 mixed or other (13.0%)
- **Age:** The average age was 35.6 years old, with a standard deviation of 10.5 years. The median age was 34, and the range was 21 to 58. The age distribution:
  - 18-24: 3 (13.0%)
  - 25-34: 9 (39.1%)
  - 35-44: 7 (30.4%)
  - 45-64: 4 (17.4%)
- **Country:** 16 United States, 3 Canada, 2 Australia, 2 New Zealand

<sup>1</sup><https://netanomics.com/netmapper/>

- **Education:** All participants had more than a high school education, with most participants having at least a Bachelor's degree:
  - Some college or Associate's Degree: 8 (34.8%)
  - Bachelor's Degree: 9 (39.1%)
  - Master's Degree: 3 (13.0%)
  - Professional or Doctorate Degree: 3 (13.0%)
- **Political Ideology:** On a Likert scale of 1-5, with “very liberal” as a 1 and “very conservative,” the average ideology was fairly moderate at 2.83, and the median was a 3.
  - Very liberal: 2 (8.7%)
  - Liberal: 6 (26.1%)
  - Moderate: 10 (43.5%)
  - Conservative: 4 (17.4%)
  - Very conservative: 1 (4.3%)
- **Social Media Platforms:** Participants were asked which social media platforms they interact with as part of their job and which platforms they use outside of work. The top platforms used for work were X, Facebook, Instagram, YouTube, and LinkedIn. The top platforms for personal use were Facebook, YouTube, Facebook Messenger, Instagram, and Reddit. These results are summarized in Table 4.1.

### 4.3.3 Survey Design

The assessment survey was implemented in Qualtrics and had two sections: misinformation detection and misinformation countering. The pre and post-tests had identical structures but included different, randomly selected posts.

#### Part 1: Misinformation Detection

In the first section of the survey, participants were presented with a set of 16 social media posts, randomly selected from a pool of posts from four categories. These categories were national news, misinformation (low-credibility and/or misleading news), local news, and pink slime (imposter local news). The participants saw four posts from each category in a random order and were asked a series of questions for each post. They received the following instructions:

In this section, you will be shown a series of 16 generic social media posts. If an organization or group posted the content, those posts will show a username and profile picture. Posts from regular users were modified and anonymized.

Any account that has a blue checkmark indicates that at least one mainstream social media organization considers that account to be a news organization / media company. These accounts may vary in accuracy and bias level.

**NOTE:** You may click on links and image link previews.

The participants were then asked the following questions for each post:

Table 4.1: The number of participants that used various social media platforms for work and for personal reasons.

<b>Platform</b>	<b>Work Count</b>	<b>Personal Count</b>
X (formerly Twitter)	20	10
Facebook	18	20
Instagram	15	14
YouTube	13	17
LinkedIn	12	12
TikTok	8	4
Telegram	8	2
Reddit	7	13
Facebook Messenger	5	15
Snapchat	4	9
Tumblr	4	1
Pinterest	3	4
Discord	3	9
WeChat	2	1
Mastodon	2	2
Twitch	2	4
Nextdoor	1	1
Threads	0	1
None	2	1
Other	3	0

1. What do you believe is the accuracy of the content in this post? [*True, Somewhat true, Neither true nor false, Somewhat false, False*]
2. How trustworthy do you consider the poster of this message to be? [*Trustworthy, Somewhat trustworthy, Neither trustworthy nor untrustworthy, Somewhat untrustworthy, Untrustworthy*]
3. How confident are you in your answers to the previous two questions? *Accuracy: [0-10], Trustworthiness: [0-10]*
4. Do you believe a local reporter wrote this post? [*Yes, No, Unsure*]
5. Would you share this post online (for example, through Facebook or Twitter)? [*Definitely yes, Probably yes, Might or might not, Probably not, Definitely not*]
6. Do you believe the poster of this message is trying to influence you or the audience of this post? [*Definitely yes, Probably yes, Might or might not, Probably not, Definitely not*]
7. Please elaborate on the reasons for your answer to the previous question on influence. [*Free response*]

Several of these questions, particularly those concerning trust and local reporters, are not analyzed in this work and were included in the survey for a concurrent study on “pink slime.” Likewise, local news and pink slime posts are also not analyzed in this work. The results of the pink slime analysis can be found in Christine Lepird’s dissertation [139].

## **Part 2: Countering**

In the second section, participants were shown four explicitly false posts and then asked what, if anything, they would do if they encountered those posts on their social media feeds. They received the following instructions:

In this section, you will be shown a series of four false posts and then asked how you would respond.

When presented with these false posts, the participants could select from the list of possible responses shown in Table 4.2. This list of responses was developed for Chapter 2 [122] but was extended to include an “Other” option. Responses labeled as “Low Effort” refer to indirect or quick actions, while responses labeled as “High Effort” apply to actions that directly engage with the misinformation content and require more time. Participants who selected “Other” were prompted to write in their response, and the level of effort for their response was manually categorized as no, low, or high.

More specifically, the participants were asked the following questions for each post:

1. If you saw this post on your social media feed, what would you do? (select any that apply) [*Options described in Table 4.2*]
2. [*If a “High Effort” action was selected*] What would you write in your comment or post? [*Free response*]
3. Do you think your answer would change depending on how well you knew the person or organization posting it? [*Definitely yes, Probably yes, Might or might not, Probably not, Definitely not*]

Table 4.2: Actions one can take when engaging with misinformation on social media.

Response	Effort Level
Ignore the post	No Effort
Report the post	Low Effort
Report the user	Low Effort
Block the user	Low Effort
Unfollow or unfriend the user	Low Effort
Privately message the user	High Effort
Comment a correction on the post	High Effort
Create a separate post with the correct information	High Effort
Other	-

4. Please elaborate on why or why not your response would change depending on the person or organization posting it. *[Free response]*
5. Do you think your answer would change depending on which platform you saw this post on? *[Definitely yes, Probably yes, Might or might not, Probably not, Definitely not]*
6. Please elaborate on why or why not your response would change depending on which platform you saw this post on. *[Free response]*

### Post-Test Only Questions

The only difference between the pre-test and post-test Qualtrics surveys is that the post-test survey included four summary questions to assess the effectiveness of the training sessions. The participants were given the following instructions and questions:

Instructions: Please rate your level of agreement to each of the following statements using the below rating scale

1. This training has helped me become better at recognizing pink slime news. *[Strongly disagree, Somewhat disagree, Neither agree nor disagree, Somewhat agree, Strongly agree]*
2. This training has helped me become better at recognizing misinformation. *[Strongly disagree, Somewhat disagree, Neither agree nor disagree, Somewhat agree, Strongly agree]*
3. This training has helped me become more knowledgeable about how to counter misinformation. *[Strongly disagree, Somewhat disagree, Neither agree nor disagree, Somewhat agree, Strongly agree]*
4. Which of the following techniques that were taught in training did you utilize when answering the questions in this survey? Select all that apply.
  - Source: Clicking on the link and reading the article
  - Source: Checking a news website's "About" page
  - Author: Looking up the author(s) of the article
  - Other Sources: Reading upstream - clicking links/sources in the article
  - Other Sources: Lateral reading - searching keywords or searching for similar stories

- 3rd Party: Looking up the bias/accuracy ratings of the news organization in question
- Experts: Checking Fact-Checking sites
- Other (please describe) *[Write-in]*

## Post Selection

We designed the survey posts to resemble typical social media posts, complete with profile pictures, user names, time stamps, and the appearance of buttons to like, comment, and share. Figure 4.1 shows an example post from the countering section. The full set of posts used in this study can be found in Appendix F.

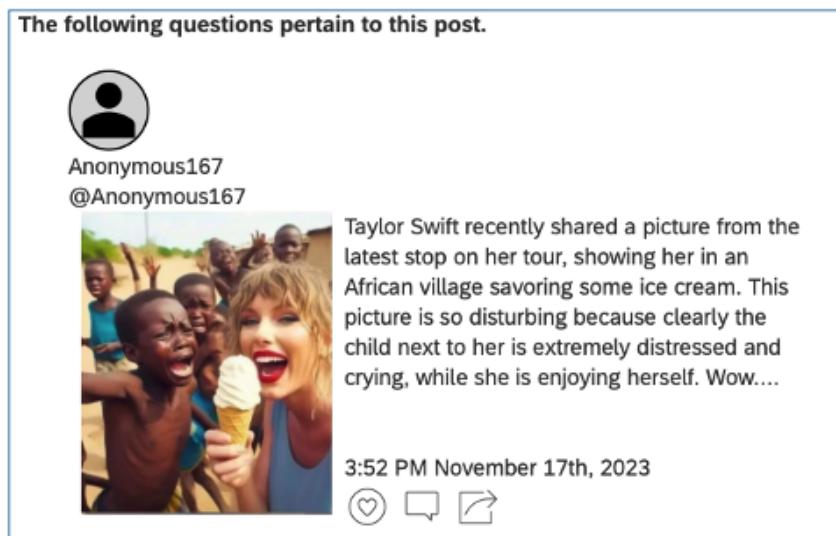


Figure 4.1: An example of a false post shown to participants.

The topics in the posts were selected to be apolitical and timely. They included health (COVID-19, vaccines), science (climate change, flat earth theories), and current entertainment topics (the Barbie movie, Taylor Swift). To account for possible differences in the difficulty of assessing each post, we had a group of experts review the posts for both difficulty and quality. Between four and six CASOS PhD students were tasked with reviewing each post. The reviewers completed the same survey described above but with additional review questions.

For the posts in the misinformation detection section, the reviewers were asked the following questions:

1. How easy or difficult do you think assessing the accuracy and trustworthiness of this post would be for an average American social media user? *[Extremely easy, somewhat easy, Neither easy nor difficult, Somewhat difficult, Extremely difficult]*
2. Should we include this post in the survey? *[Yes, No]*
3. Any comments on this post? *[Free response]*

For the posts in the countering section that only included explicitly false or extremely misleading posts, the reviewers were asked the following questions:

1. Would the average American government employee think this post was true or false? *[True, False, Unsure]*
2. How easy or difficult do you think it is to counter this post for the average American government employee? *[Extremely easy, somewhat easy, Neither easy nor difficult, Somewhat difficult, Extremely difficult]*
3. Should we include this post in the survey? *[Yes, No]*
4. Any comments on this post? *[Free response]*

Based on experts' comments, some posts were removed if they were deemed poor examples or confusing. In the misinformation detection section, two out of the 19 misinformation posts tested and four out of the 16 actual news posts tested were removed. The two misinformation posts were removed because the reviewers noted that it required watching the videos to fact-check, which would not have been conducive in the testing environment. Among the four actual news posts that were removed, one was removed because the headline changed and no longer matched the content of the linked article, one because the study linked in the news article moved and therefore was difficult to verify, one because while the news organization was reputable, it failed to note that the study they quoted had a conflict of interest, and one because most reviewers had trouble verifying the legitimacy of the source. None of the 12 false posts in the countering section were recommended for removal.

The remaining posts were sorted by post type (real/national news, misinformation, local news, pink slime, and false posts for the countering section) and by average difficulty score (rated on a 1-5 Likert scale). They were then randomly divided into the pre- and post-tests. This method was employed to control the item-wise difficulty of the posts and to maintain an equal number of posts per post category in the pre- and post-test. The difficulty scores for the relevant questions in this study are summarized below:

- **Misinformation Posts** - Average difficulty scores ranged from 1.5 to 3, with an overall average of 2.33
- **Real/National News Posts** - Average difficulty scores ranged from 1.6 to 3, with an overall average of 2.31
- **Countering Posts** - Average difficulty scores ranged from 1.2 to 3, with an overall average of 2.09.

The posts in each category were sorted by their average difficulty score and divided into quartiles. Four misinformation posts, four national news posts, and four countering posts, one from each quartile of difficulty, were randomly selected for the pre-test. The same method was applied to randomly select the posts for the post-test. All participants were shown the same posts in the pre-test and post-test because there was a concern that they would discuss among themselves during the breaks throughout the day.

#### **4.3.4 Training Sessions**

After the pre-test survey was administered, participants received training in three main areas: misinformation detection, pink slime detection, and countering misinformation. The pink slime training was an interactive 30-minute session administered by Christine Sowa Lepird [139] and

is not covered further in this work. I led the other two training sessions. Following the training and a break for lunch, participants took the post-test. By the end of the day, participants were briefed on what they did well and the most frequently missed questions.

## Misinformation Detection Training

Participants were given a 45-minute training adapted from my work from previous chapters, various university library trainings, and information from FactCheck.org [118, 179]. First, participants were provided with background information on the topic, including formal definitions of misinformation and disinformation. In addition, participants were given a summary of the various characteristics of misinformation, covering eight main types along with their purpose, audience, and news context (see Section 1.3). Next, they were given ten steps to consider when assessing the accuracy of online information, followed by several examples illustrating each of the ten steps. An example was shown for each major type of misinformation. Figure 4.2 shows the training slide summarizing all ten steps. Participants had access to the slides and could use this information as a reference. The training was interactive, with respondents describing how they would approach the examples presented in the slide deck.



Figure 4.2: Summary training slide on misinformation detection.

## Countering Misinformation Training

Participants were given a 30-minute interactive training session adapted from my literature review work on why and how to counter misinformation effectively. The training was broken down into three main parts:

- 1. Why People Should Counter Misinformation** - We discussed the most common reasons why people do not counter misinformation and how those concerns can be addressed [208, 212]. The main reasons discussed were:

- *Level of Impact* - The first concern discussed was the potential impact one could have when countering misinformation. Research showing the effectiveness of individual-level debunking, especially from trusted messengers, was shared with participants [144].
- *Time Restraints* - Countering misinformation can be time-consuming [208], so participants were encouraged to concentrate on those most likely to change their minds or on individuals closest to them.
- *Conflict Avoidance* - Maintaining harmony in relationships is another primary concern [163]. Participants discussed respectful ways to listen and debunk, and were guided on indirect actions they could take instead, such as reporting, blocking, unfollowing, or writing a community note.
- *Credibility Concerns* - Some individuals raised concerns about countering misinformation on topics where they are not experts. Participants were encouraged to refer to experts and fact-checking organizations, or to engage in indirect actions.

2. **Common Logical Fallacies** - We reviewed logical fallacies and how to spot them. Examples and explanations were adapted from several university research guides<sup>234</sup>. The discussion primarily focused on the following fallacies:

- *Slippery slope* - This fallacy occurs when someone believes that a specific event or action will lead to a series of increasingly severe events without any evidence. It appeals to one's emotions by exploiting the fear of this possible chain of consequences without any real evidence to support this claim.
- *Straw man* - The straw man fallacy occurs when someone misrepresents their opponent's argument by fabricating or exaggerating their position, making it easier to debate them. This tactic is often done to make one appear more reasonable in comparison.
- *Ad hominem* - Ad hominem attacks involve targeting an opponent's personal characteristics to undermine their position. While this argument has legitimate uses if one's characteristics are relevant, it is frequently used to insult and degrade individuals without engaging with their actual arguments.
- *Bandwagon* - This fallacy occurs when individuals argue that something must be true or good simply because it is popular or widely supported by the public. While arguments that have popular support may indeed be true, the fallacy lies in assuming that they must be true solely due to their popularity.
- *False causation* - False causation is often employed by those who intentionally spread misinformation. It refers to the idea that if one event occurs before another, the first event must have caused the second. However, correlation does not always imply causation.

3. **Effective Interventions** - Finally, participants were trained on the types of interventions

<sup>2</sup><https://libguides.princeton.edu/c.php?g=982190&p=7102155>

<sup>3</sup>[https://owl.purdue.edu/owl/general\\_writing/academic\\_writing/](https://owl.purdue.edu/owl/general_writing/academic_writing/)

<sup>4</sup><https://writingcenter.unc.edu/tips-and-tools/fallacies/>

and debunking efforts that are effective [144, 218]. At the individual level, they were encouraged to use trusted community messengers when debunking misinformation. They were also encouraged to ask questions to engage the critical thinking skills of those who believe the misinformation and to demonstrate empathy to emphasize that they share the same overall goals and concerns. Participants were provided with high-level examples of the following types of debunking:

- *Fact or logic-based debunking* - The suggested method was the use of a “fact” sandwich, or starting and ending with the accurate information while debunking the myth or fallacy in the middle to better emphasize the facts over the misinformation [144].
- *Empathy-based debunking* - Empathy-based debunking may sometimes be effective in countering conspiracy theories by pointing out the similarities between the targeted or scapegoated group and one’s own group. [30, 141, 168]
- *Source-based debunking* - This method aims to reduce the credibility of those spreading misinformation or conspiracy theories. Ridicule or humor has been shown to be effective for general audiences, but it may not be advisable when countering the beliefs of an individual [141, 168]

At the organizational level, they were encouraged to focus on careful and transparent research dissemination to prevent the public or news organizations from misinterpreting their results. They were also encouraged to utilize social media campaigns when appropriate to promote counter-narratives against pervasive falsehoods that often go unchecked on social media, such as anti-COVID-19 vaccination narratives [218].

### 4.3.5 Analysis Methods

#### Part 1: Misinformation Detection

Table 4.3 summarizes the measures used in Part 1 of this study. Descriptive statistics will be summarized for all measures. Paired t-tests will be used to compare the average accuracy and confidence scores from the pre-test and post-test.

#### Part 2: Countering

Table 4.4 summarizes the measures used in Part 2 of this study. The chi-squared test will be used to compare the pre- and post-test effort level measures. Paired t-tests will be used to compare the average poster and platform effects per participant in the pre- and post-tests.

Textual analysis will be used to analyze the open-ended questions associated with the poster and platform effect measures. To analyze the free-text responses, we used code mapping to identify common themes throughout the responses. Code mapping, sometimes referred to as affinity diagramming, is frequently used when analyzing open-ended survey data [105, 198]. We first reviewed all the responses for each question and then sorted related comments into groups using Apple’s Freeform program<sup>5</sup>, which is a digital whiteboard. Whenever a participant noted

<sup>5</sup><https://www.apple.com/newsroom/2022/12/apple-launches-freeform-a-powerful-new-app-designed-for-creative-collaboration/>

Table 4.3: The measured variables from Part 1 of the surveys

Variable(s)	Definition	Values
Misinformation Accuracy National News Accuracy	Response to “What do you believe is the accuracy of the content in this post?”	1 = True 2 = Somewhat true 3 = Neither true nor false 4 = Somewhat false 5 = False
Misinformation Confidence National News Confidence	Response to “How confident are you in your answers to the accuracy question?”	0–10 where: 0 = Very unsure 10 = Very confident

something like “same reason as my previous answer,” we placed those in the same group as their previous related response. We iteratively developed the codes by reviewing the responses a second time. Then, we labeled each group of responses with various factors and merged related factors into overarching themes. The primary themes are discussed in the Results section.

#### 4.3.6 Ethics Information

The Carnegie Mellon University Institutional Review Board (IRB) approved this study, numbered “STUDY2023\_00000429.” They determined the study was exempted from full review because it involved a “benign behavioral intervention.” All participants were randomly assigned a user ID to link their pre- and post-training results, but their names were not collected. The survey collected informed consent from all participants. Participants were not paid by the study but instead were paid their typical government salary.

### 4.4 Results: Misinformation Detection

#### 4.4.1 Accuracy

For the four misinformation and four national news posts in both the pre- and post-tests, participants were asked to rate what they believed the accuracy of those posts was on a 1-5 Likert scale. Table 4.5 displays the average and median number of questions participants answered correctly. An accuracy score of 3, 4, or 5 for misinformation posts was labeled as correct, while an accuracy score of 1 or 2 for actual news posts was labeled as correct. The median values are shown to account for potential item effects.

Overall, participant detection of misinformation remained relatively unchanged, as it started high. The average number of misinformation posts correctly identified increased only slightly

Table 4.4: The measured variables from Part 2 of the surveys

<b>Variable(s)</b>	<b>Definition</b>	<b>Values</b>
Effort Level	Response to “If you saw this post on your social media feed, what would you do?”	<b>Set to the max value:</b> Ignore post = 0 Report post = 1 Report user = 1 Block user = 1 Unfollow/unfriend user = 1 Message user = 2 Comment correction = 2 Create post = 2 Other = [write-in]
Poster Effect	Response to “Do you think your answer would change depending on how well you knew the person or organization posting it?”	1 = Definitely yes 2 = Probably yes 3 = Might or might not 4 = Probably not 5 = Definitely not
Platform Effect	Response to “Do you think your answer would change depending on which platform you saw this post on?”	1 = Definitely yes 2 = Probably yes 3 = Might or might not 4 = Probably not 5 = Definitely not

from the pre-test to the post-test, from 87% to 88% correct. Unsurprisingly, the two-sided paired t-test shown in Table 4.6 comparing the average accuracy scores was not significant. Participants were likely more familiar with misinformation detection techniques than we expected, considering the nature of their professions.

Table 4.5: The average and median number of questions answered correctly by participants.

<b>Questions</b>	<b>Pre-Test</b>		<b>Post-Test</b>	
	Average # Correct	Median # Correct	Average # Correct	Median # Correct
Misinfo. Posts	3.48 (87%)	4 (100%)	3.52 (88%)	4 (100%)
Real News Posts	2.43 (61%)	3 (75%)	3.30 (83%)	3 (75%)

Table 4.6: Two-sided paired t-test results ( $df = 22$ ) and the estimated effect size, Cohen's  $d$ , for the average accuracy scores given by each participant as defined in Table 4.3.

<b>Questions</b>	<b>Pre-Test</b>	<b>Post-Test</b>	<b>Paired T-Test</b>		
	Average (SD)	Average (SD)	$t$ -value	p-value	Cohen's $d$
Misinfo. Posts	4.21 (0.63)	4.20 (0.45)	0.065	0.95	0.0135
Real News Posts	2.33 (0.70)	1.77 (0.58)	3.23	0.0038	0.675

However, the detection of real news did increase over time, with only 61% of the questions answered correctly on average in the pre-test compared to 83% in the post-test. The average assigned accuracy scores also improved, rising from an average score of 2.33 in the pre-test to 1.77 in the post-test, indicating that participants were more likely to believe that the actual news posts were true in the post-test. The two-sided paired t-test comparing the average accuracy scores was statistically significant.

#### 4.4.2 Confidence

A concern with media literacy training is that it could lower respondents' confidence in their ability to distinguish true from false news. This possible unintended consequence could make people more skeptical of all news, including accurate information, which tends to be substantially more prevalent than false news. We did not observe this in our data. Table 4.7 shows that the overall average confidence score per question type remained relatively unchanged from the pre-test to the post-test. The two-sided paired t-test comparing the average confidence scores in the pre- and post-tests for both the misinformation posts and for the true news posts were not statistically significant.

Table 4.7: Two-sided paired t-test results ( $df = 22$ ) and the estimated effect size, Cohen's  $d$ , for the average self-reported confidence scores on a scale 1-10 given by each participant as defined in Table 4.3.

Questions	Pre-Test	Post-Test	Paired T-Test		
	Average (SD)	Average (SD)	$t$ -value	p-value	Cohen's $d$
Misinfo. Posts	8.65 (1.21)	8.22 (2.12)	1.14	0.27	0.24
Real News Posts	8.38 (1.30)	8.67 (1.37)	-1.15	0.26	-0.24

## 4.5 Results: Countering Misinformation

### 4.5.1 Selected Interventions

Participants were presented with four clearly marked false posts during both the pre- and post-tests. When asked if and how they would respond to the misinformation shown, the participants could select one or more responses from the options described in Table 4.2. The total number of times participants selected each intervention type is summarized in Figure 4.3.

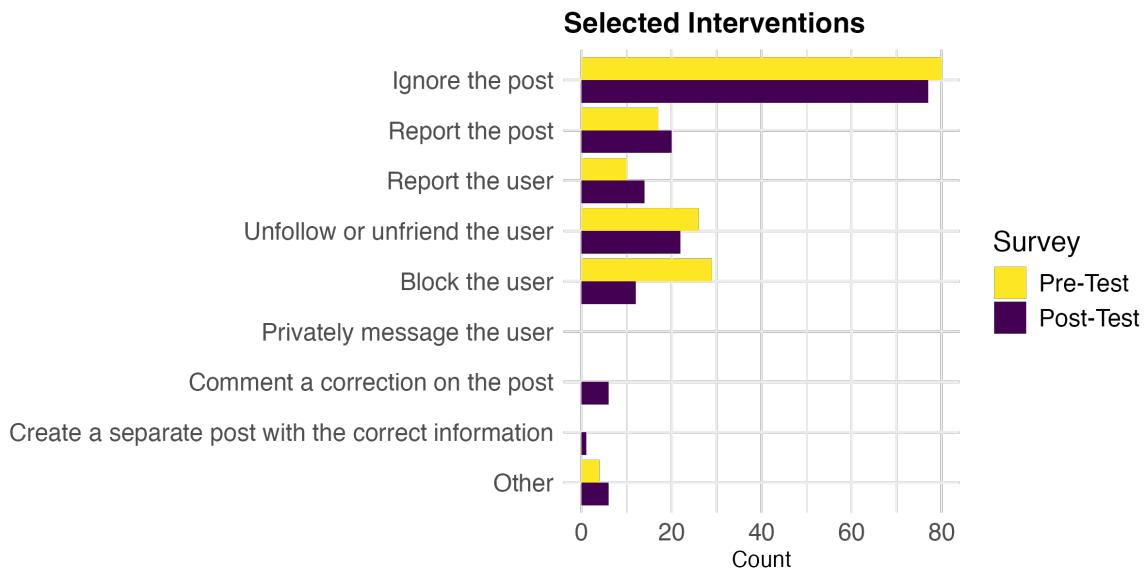


Figure 4.3: The number of times participants said their answer would change depending on either the platform or the poster over both the pre- and post-training surveys.

Unsurprisingly, these results are fairly similar to those observed in Chapter 2 (see Figure 3.1), with ignoring the post being the most frequently selected option. However, many low-effort actions were selected, such as reporting or blocking the user. High-effort actions, such as commenting a correction on the post, were selected more frequently in the post-test.

### 4.5.2 Effort Level

We found that the maximum amount of effort participants selected to counter any misinformation posts increased from the pre-training survey to the post-training survey (see Table 4.8). More people said they would engage in high-effort actions, with this increase primarily coming from those who were already engaging in low-effort actions. The proportion of participants who stated they would not counter any misinformation posts remained unchanged. The chi-squared test comparing the counts in the effort level categories was significant, with a p-value of 0.049. However, due to the small sample size, the chi-squared approximation may be inaccurate as the expected counts per cell are low.

Table 4.8: The percentage and number of participants whose maximum effort level was in each of the three effort level categories described in Table 4.2.

	Pre-Training	Post-Training
No Effort	34.8% (8)	34.8% (8)
Low Effort	65.2% (15)	43.4% (10)
High Effort	0% (0)	21.7% (5)

Table 4.9 shows the contingency table showing the number of participants in each effort level category for both the pre- and post-tests. In the post-test, most participants either maintained the same effort level they selected in the pre-test or increased their maximum effort level. Only three participants (13.0%) opted for a lower effort level. However, the generalized McNemar's chi-squared test, which is applicable for paired data, was not significant with a p-value of 0.16. While the traditional chi-squared test analyzes whether the variables are related, the McNemar's test specifically analyzes whether subjects switched their responses from the pre- to post-tests in a consistent direction. The lack of significance may be related to the small overall sample size.

Table 4.9: The contingency table showing the number of participants that were no, low, and high effort in the pre vs. post-test.

		Post-Test		
		No Effort	Low Effort	High Effort
Pre-Test	No Effort	5 (21.7%)	2 (8.7%)	1 (4.3%)
	Low Effort	3 (13.0%)	8 (34.8%)	4 (17.4%)
	High Effort	0 (0%)	0 (0%)	0 (0%)

### 4.5.3 Platform and Poster Impact

For each post, participants were asked whether their answers would change depending on the person or organization that posted it or the platform on which they saw the misinformation. Figure 4.4 shows the number of times participants selected each possible answer in the pre- and post-test surveys. “Probably Not” and “Definitely Not” were the most frequently chosen responses for both poster and platform. However, across all the posts shown, 15 participants

(65.2%) and seven participants (30.4%) responded with “Probably Yes” or “Definitely Yes” for at least one post when asked if the poster or the platform, respectively, would influence their answer. Participants were more likely to say that the misinformation poster would impact their countering response than the platform.

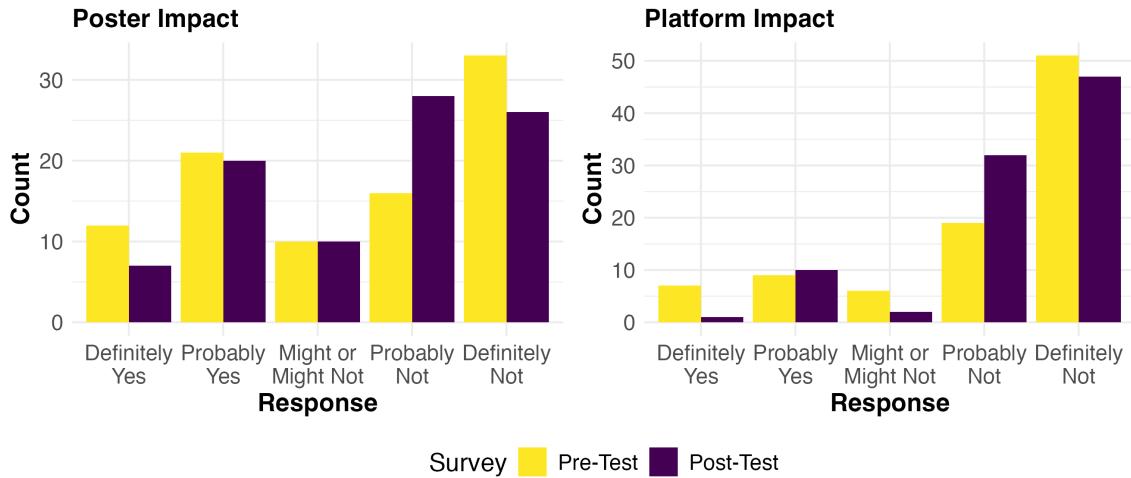


Figure 4.4: The number of times participants said their answer would change depending on either the platform or the poster in both the pre and post-training surveys.

Table 4.10 shows the average scores for the effects of the poster and the platform in both the pre- and post-tests. The two-sided paired t-test, which compared the average scores by participant was not significant for either the poster effect scores or the platform effect scores. Given that the poster and platform effect scores are relatively consistent across the pre- and post-tests, we can combine the total counts for both to better compare the overall poster and platform effects. Figure 4.5 shows the total number of times participants selected each possible answer over both surveys. The chi-squared test comparing the poster with the platform counts was statistically significant, with the p-value < 0.001.

Table 4.10: Two-sided paired t-test results ( $df = 22$ ) and the estimated effect size, Cohen’s  $d$ , for the average poster and platform effect scores given by each participant as defined in Table 4.4.

Questions	Pre-Test	Post-Test	Paired T-Test		
	Average (SD)	Average (SD)	t-value	p-value	Cohen’s $d$
Poster Impact	3.40 (1.11)	3.49 (0.99)	-0.30	0.77	-0.063
Platform Impact	4.07 (1.07)	4.24 (0.80)	-0.95	0.35	-0.20

#### 4.5.4 Factors Affecting Countering Actions

The survey asked participants to explain their reasoning about whether and how the misinformation poster or the platform would affect their countering behavior. When elaborating on their

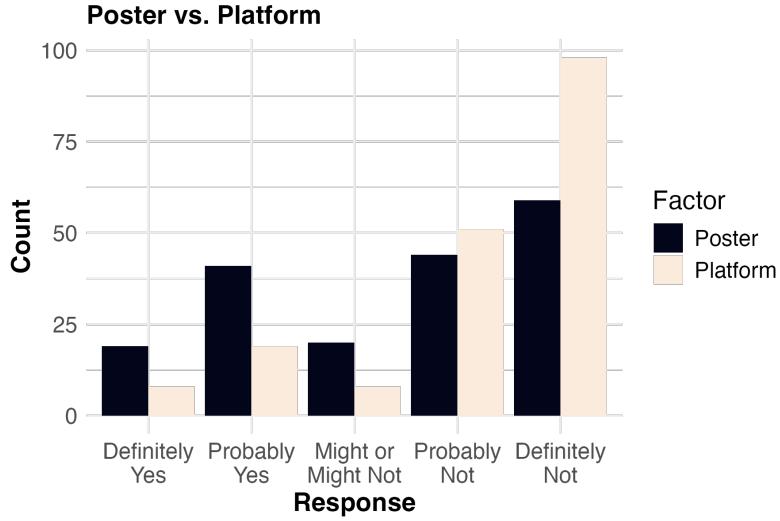


Figure 4.5: The number of times participants said their answer would change depending on either the poster or the platform combined over both the pre- and post-training surveys.

responses regarding how these factors would impact their countering efforts, we identified four main recurring themes:

1. *Platform / Account Preferences*: This theme addresses the participant's preference, or lack thereof, regarding which types of posters or platforms they are more likely to engage in countering actions.
2. *Content*: Content refers to how the content of the misinformation post affects an individual's likelihood of engaging in countering efforts.
3. *Platform / Account Features*: This theme focuses on the features of the misinformation accounts and social media platforms and their potential influence on whether someone chooses to take action.
4. *Impact*: Finally, many participants discussed the long-term consequences of their potential countering actions or lack of actions.

All relevant comments fall into these four main themes. Other comments were either irrelevant, supplementary to the participants' main comments, or expressed that they did not care about the post. These other comments are classified under a theme called "Other." Table 4.11 summarizes the major themes and how frequently they were mentioned when considering the misinformation poster or platform, sorted in descending order of total overall mentions. Table 4.12 displays the specific sub-themes mentioned by the participants when considering the misinformation poster. Similarly, Table 4.13 shows the sub-themes when considering platform.

Figure 4.6 illustrates the total number of times each sub-theme was referenced when considering either the misinformation poster or the platform. It emphasizes how prominent the *Platform / Account Preferences* theme was: closeness to the poster is a huge factor when considering the likelihood of countering. At the same time, most participants indicated that they had no platform preferences.

Table 4.11: The number of unique participants and total mentions of each major theme when considering the misinformation poster or the platform.

<b>Theme</b>	<b>Poster</b>		<b>Platform</b>		<b>Overall</b>	
	# Unique	# Total	# Unique	# Total	# Unique	# Total
Platform/Account Preferences	13 (56.5%)	40	13 (56.5%)	55	17 (73.9%)	95
Content	18 (78.3%)	44	12 (52.2%)	20	20 (87.0%)	64
Platform/Account Features	4 (17.4%)	7	3 (13.0%)	10	6 (26.1%)	17
Impact	5 (21.7%)	10	3 (13.0%)	5	7 (30.4%)	15
Other	8 (34.8%)	12	8 (34.8%)	16	12 (52.2%)	28

Table 4.12: The number of unique participants that mentioned each sub-theme and the total number of times each sub-theme was mentioned when participants were asked to consider the poster of the misinformation.

<b>Theme</b>	<b>Sub-Theme</b>	<b># Unique</b>	<b># Total</b>
Account Preferences	More effort if closer to poster	13	40
Content	Content is too far gone	13	21
	Content is clearly false / easy to debunk	9	12
	Possible offline harms / very serious	7	9
	Familiarity with or care about the content	3	3
Account Features	More effort for verified / news accounts	3	5
	Less direct effort / more indirect for organizations or less close contacts	2	3
	More effort if recognizable source	2	2
	More effort for accounts with large followings	1	1
Impact	Avoid conflict	1	4
	Perceived lack of countering impact	3	4
	Time-intensive to counter	2	3
Other	Unsure of poster's motive	1	1
	Unrelated comment	7	11

Table 4.13: The number of unique participants that mentioned each sub-theme and the total number of times each sub-theme was mentioned when participants were asked to consider the misinformation platform.

Theme	Sub-Theme	# Unique	# Total
Platform Preferences	Treat all platforms equally	11	47
	Anonymity	1	5
	Platform preference	2	3
Content	Content is too far gone	4	6
	Content is clearly false / easy to debunk	4	9
	Possible offline harms / very serious	5	5
Platform Features	Reporting functionality	1	6
	Platform type	2	2
	Ability to contact poster directly	1	2
Impact	Audience	2	4
	Platform is serious about misinformation	1	1
Other	Don't care	3	5
	Unrelated comment	6	11

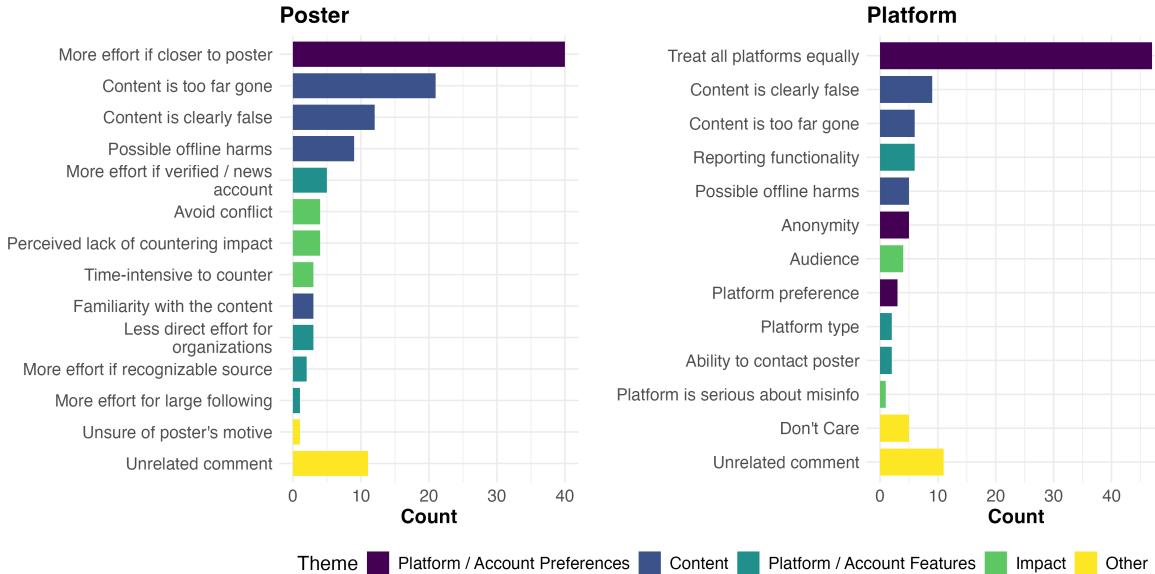


Figure 4.6: This figure displays the total number of times over both surveys that participants mentioned a specific sub-theme when asked if and how their answer would change depending on the misinformation poster and the platform.

## **Platform and Account Preferences**

Most participants (73.9%) mentioned preferences in some way across both surveys and factors. When considering the misinformation poster, many participants (13 or 56.5%) repeatedly stated that they prefer engaging with close contacts over less close contacts. These participants suggested that if they knew the misinformation poster, they would be more inclined to participate in direct debunking efforts, whether on social media, via private messages, or in person. This sentiment was by far the most frequently mentioned comment overall when participants were asked if their responses would change depending on the poster. Two examples of quotes from participants are shown below.

“If I knew the person, I would be more inclined to messaging them about it first.”

“If I knew the poster personally i [sic] would privately message them to try and keep them from posting fake news.”

On the other hand, while most participants (13 or 56.5%) mentioned the theme of platform preferences, most of them mentioned it because they believe they treat all platforms equally (11 or 47.8%). Only two respondents (8.7%) mentioned having a platform preference. Two examples of quotes are provided below.

“I am not likely to engage with users on sites where my identity is directly tied to the account. Accounts like Reddit, where I am more anonymous, makes discussion easier to partake in and exit from.”

“The platform would not determine my decision.”

## **Content**

The content of the post was one of the most frequently cited reasons for engaging or not engaging in countering efforts. Twenty respondents (87%) mentioned content at least once when explaining their reasoning. Specifically, if participants perceived the content as extreme or incredulous, many expressed that they thought any effort would be wasted. For example:

“I just don’t think i [sic] could change the mind of someone who believed the earth was flat.”

“If they believe that they would be too far gone for reasoning.”

Conversely, if something was easy to debunk, such as a straightforward or factual error, a topic they were knowledgeable about, or if the post had the potential for severe offline consequences, participants indicated they would be more likely to engage.

“i [sic] would report this one and try to get it to stop promulgating because i think it could be harmful.”

## **Platform and Account Features**

Four participants (17.4%) mentioned the significance of account characteristics and said they would be more inclined to intervene if the poster was a verified account, a news agency, a recognizable source, or had a large following.

“It was a verified channel posting it so the misinformation will spread faster. If it was a random person posting it, I might not repost it and just ignore it instead”

“I would be more likely to report this if it was from a more recognizable source.”

While many participants mentioned that they treat all their social media platforms equally, three participants (13%) expressed that they believed they would be more likely to respond on certain platforms than others due to the features of those platforms. For example, one participant stated:

“If I were in a channel specific to that topic, I may look to validate/invalidated the content. But on a platform like facebook [sic] I would be less likely to give it a second look.”

The ease with which one could report a post or contact the misinformation poster was also mentioned as a possible factor. For example:

“Misinformation is everywhere, trying to fight it is exhausting and not every platform has the same method to report it.”

“On different platforms it might be easier to directly contact the poster.”

## **Impact**

Seven participants (30.4%) mentioned the potential impact of countering, and most cited it as a reason for their reluctance to take action. They expressed concerns that debunking would be too time-consuming and have little to no effect, that they lacked sufficient knowledge to counter the post, or that they wanted to avoid conflict.

“Where do I even begin? An anon account posting disinfo is hardly worth my time, but there are so many misguided people out there, it’s easier to report and block rather than engage.”

“I feel that reporting would do little to change anything.”

One participant mentioned impact positively and said that taking action on some platforms that take moderation seriously may be more impactful than those with less moderation. However, another participant felt differently:

“I would be more likely to report it on an application or site with worse media literacy.”

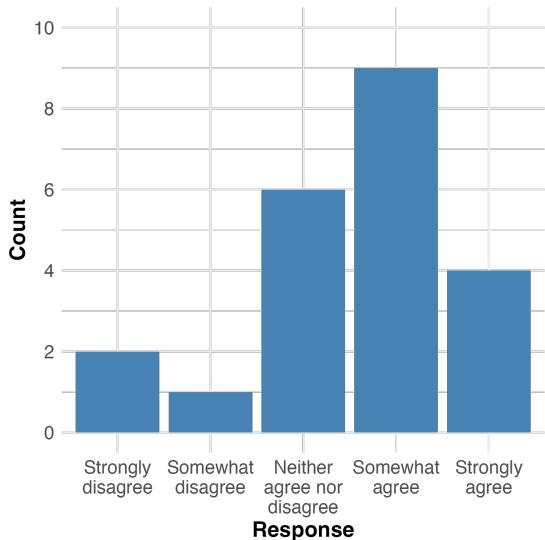
## **4.6 Participant Feedback**

Feedback from participants after the session was generally positive. At the end of the post-test, participants anonymously answered feedback questions, allowing us to evaluate the effectiveness of the training sessions. First, participants were asked if they felt the training improved their ability to recognize misinformation. They indicated their beliefs on a Likert scale of 1 to 5, with 1 meaning they “strongly disagreed” that the training was helpful and 5 meaning they “strongly agreed.” Similarly, they were asked whether they believed the training helped them become more knowledgeable about countering misinformation. Table 4.14 shows the high-level summary statistics, while Figure 4.7 shows the distribution of the feedback scores. In both cases,

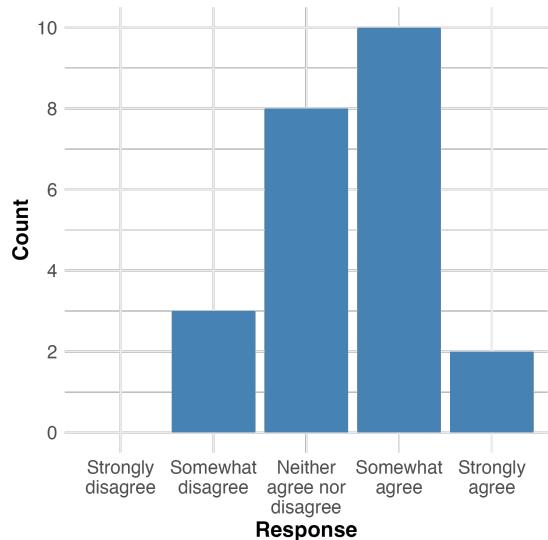
the median response was a 4 (“somewhat agree”). Considering that many of these analysts were already highly skilled at detecting misinformation, it is not surprising that some did not find the training helpful. However, the majority of participants did. Only three of the 23 participants indicated they disagreed with either feedback statement.

Table 4.14: The average, standard deviation, and median feedback score on a 1-5 Likert scale, where 1 represents “strongly disagree” and 5 represents “strongly agree.”

<b>Feedback</b>	<b>Average</b>	<b>Std. Dev.</b>	<b>Median</b>
This training helped me become better at recognizing misinformation	3.55	1.14	4
This training helped me become more knowledgeable about how to counter misinformation	3.48	0.85	4



(a) Part 1: Misinformation Detection Feedback



(b) Part 2: Countering Ability Feedback

Figure 4.7: This figure summarizes the responses to the feedback questions at the end of the study, indicating whether participants agreed or disagreed that the training was helpful.

Next, we asked participants to select which techniques covered in the training sessions, if any, they used when answering the questions in the post-test survey. Figure 4.8 summarizes the results. The most frequently selected answer was to click on the link and read the article. Most participants employed all the listed techniques except for fact-checking websites. This result demonstrates that the techniques we taught were helpful, although we may want to dedicate more time to expert fact-checking websites in future training sessions.

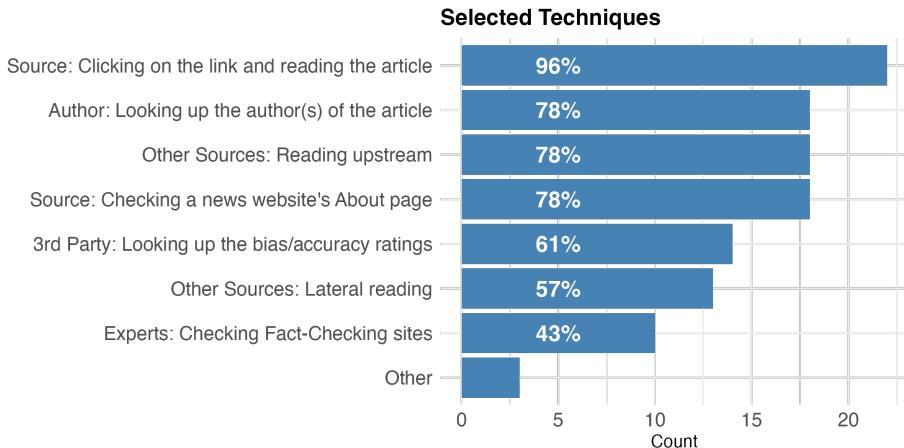


Figure 4.8: The number of participants who said they used each technique in the post-training survey.

## 4.7 Discussion

We conducted an experiment to test the effectiveness of media literacy training on improving misinformation detection and countering misinformation. The evidence for improved misinformation detection was limited partly because participants already excelled at detecting misinformation in the pre-test, resulting in minimal improvement in the post-test that was not statistically significant. However, actual news detection increased in the post-test, with more participants accurately classifying real news stories as true (see Table 4.5). Participants' confidence in their ability to discern real news posts from misinformation posts remained steady in the post-test (see Table 4.7), indicating that the training did not have the unintended side effect of making participants more skeptical of all news posts.

Next, we analyzed the effectiveness of the training on participant willingness and likelihood of utilizing countermeasures, finding that more respondents claimed that they would intervene with more effort and more directly in the post-test compared to the pre-test (see Table 4.8). However, the number of participants who stated they would exert no effort remained the same. This increase in high-effort actions came primarily from individuals who were already engaging in low-effort countering efforts, such as reporting or blocking.

We also qualitatively analyzed the participants' explanations when asked if and how their likelihood of countering would change depending on the account posting the misinformation and the platform on which it was posted. Overall, we found that even though the **Content** theme was mentioned by the highest number of unique participants (20), **Platform and Account Preferences** were the most frequently cited factors in all posts. As shown in Figure 4.6, closeness to the poster and platform neutrality dominate all other factors. These results suggest that individuals holding these beliefs feel strongly about them and felt the need to mention them multiple times across posts. Other frequently discussed themes were **Platform and Account Features** and potential **Impact** (or lack thereof) when countering.

Finally, the feedback from participants was overall positive. They rated the usefulness of both training sessions with a median Likert score of 4 ("somewhat agree"). Additionally, most

participants used the techniques we covered during the training. These results are encouraging, considering the nature of the participants' profession and the fact that many were already proficient at identifying most of the misinformation stories. Each training session lasted at least 15 minutes longer than originally scheduled due to high audience participation. Similarly, the pre- and post-tests also ran longer than expected (45 minutes instead of 30), largely because of the participants' high motivation levels. These analysts demonstrated a strong desire to get the answers correct, even though they knew their responses would remain anonymous. They also provided incredibly detailed and thoughtful responses to all qualitative sections.

## 4.8 Conclusions

### 4.8.1 Limitations

There are multiple limitations to this work. First, the sample size is relatively small. While this allowed us to gather detailed qualitative feedback and analysis, further research is needed to validate these results and generalize them to other populations. Second, the participants were government analysts who were generally more educated than the average American. All participants had at least some college education, with most having a bachelor's degree or higher. However, participants were motivated by the fact that they would continue earning their full-time salary if they took this training and received training credits. The types of analysts who enrolled in the training were highly motivated, resulting in very detailed and thoughtful qualitative responses. Finally, the two interactive, in-person training sessions took over an hour to administer, which is not practical for scaling to the general population. Nonetheless, this work demonstrates the potential usability of similar training sessions among motivated individuals who are already skilled in detecting misinformation.

### 4.8.2 Contributions

In a study conducted with motivated government analysts, we analyzed the effectiveness of media literacy training on misinformation detection and participants' willingness and ability to counter it. Previous research on media literacy has shown conflicting results regarding its effectiveness, and little to no research has been done on its usefulness in training individuals to combat misinformation. Overall, we found some, though limited, evidence that participant truth discernment increased in the post-test without lowering participant confidence or increasing skepticism toward all news. We found more substantial evidence that the training effectively increased participants' willingness and likelihood to counter misinformation on social media.

Additionally, we conducted valuable qualitative work on why some individuals choose to counter misinformation while others do not. Understanding why people do or do not counter and showing that training can boost individuals' willingness to intervene is crucial in determining how to improve social corrections and other user-based countermeasures. For example, closeness was a significant factor. People may feel more comfortable addressing misinformation with those close to them and feel like they can make a more significant impact or should at least attempt to do so among loved ones. Verified accounts or those with large followings were con-

sidered more important to counter. The simplicity of reporting a post or user was also mentioned. Knowing these factors, social media companies can better design their platforms. For example, platforms can encourage more reporting by improving reporting functions and promote more social corrections by highlighting posts from closer contacts.



# Chapter 5

## Characterizing Platform and Government Countermeasures

In addition to research on user-based countermeasures, increasing attention has been given to misinformation mitigation efforts at both the platform [243] and government levels [192, 242]. The literature consistently shows that public opinion plays a critical role in shaping policy implementation and effectiveness [48]. A review of recent work in the misinformation space demonstrates that public support is a key factor to consider when developing successful interventions [66, 125, 126]. Understanding why people support or oppose specific countermeasures is therefore essential.

This chapter investigates several factors that have been previously identified as relevant to public policy support across various domains [89], including climate change initiatives [31, 103] and public health interventions [40, 64]. These features are fairness, intrusiveness, and effectiveness, and they are particularly relevant to misinformation interventions, as concerns over censorship and equity frequently arise among social media users [66, 226].

The primary research question for this chapter is: What types of countermeasures are supported by the general public and why? More specifically,

1. To what extent does a misinformation intervention's perceived **fairness**, **intrusiveness**, and **effectiveness** predict **support**?
2. How do the attributes people consider when forming preferences change due to the **implementer** (social media companies vs. governments) of the intervention?
3. What demographic factors, if any, are related to opinions on these topics?

### 5.1 Introduction

While the extent of misinformation's impact remains a subject of debate [6, 47, 71], its documented consequences are far-reaching. Research has linked misinformation to the erosion of democratic norms and institutions [71, 220], as well as the proliferation of violent and extreme conspiracy theories [61, 76, 196]. In response, a growing body of research has focused on developing effective and practical misinformation countermeasures, including social corrections [21],

warning labels [152], accuracy prompts [175], and various government regulations [184].

Several review articles have evaluated the extensive literature on misinformation interventions [59, 100, 242]. Interventions are typically evaluated by their effectiveness in reducing the creation, spread, or belief in misinformation by evaluating methods to improve truth discernment and corresponding behavior. However, to practically counter misinformation at scale, the public needs to trust and engage with the interventions.

Indeed, there have been calls for public participation in misinformation countermeasures [66, 126]. Research consistently shows the significant impact public opinion can have on the development and efficacy of public policy [48]. Furthermore, social media platforms are unlikely to implement unpopular countermeasures, as they are responsive to the desires of their users and any potential revenue implications [145]. Public acceptance of various countermeasures is a necessary yet understudied component of misinformation interventions. Therefore, we are motivated to address *why* people support or do not support certain inventions with the assumption that support is required for engagement.

For this chapter, we surveyed 1,010 residents of the United States who use social media at least once a week. This survey represents the second half of the one described in Chapter 2. Participants were asked to rate their support for various potential interventions on a Likert scale ranging from 1 to 5. Half of the participants were asked as if the government were implementing these policies, while the other half were asked as if the social media companies were doing so. Additionally, they were asked to rate how effective, fair, and intrusive they believed each intervention to be. We included ten interventions designed to span most of the categories outlined in Chapter 2.

## 5.2 Related Work

Recent research has examined the relationship between various personal attributes, including partisanship, trust in institutions, and previous exposure to misinformation or interventions, with support for interventions [149, 199]. However, the qualities of the countermeasures themselves that predict preferences remain unaddressed. Previous studies on the features influencing support for climate change policies identified three primary attributes: fairness, intrusiveness, and effectiveness [103]. These attributes are directly relevant to the misinformation context due to concerns about the potential infringement on free speech rights and the possible disproportionate impact of countermeasures on specific groups, such as Republican social media users [55, 184, 189].

Free speech is a fundamental right and value in the U.S., essential for a functioning democracy. However, Americans are divided on whether free speech or the restriction of false content should be prioritized [153]. We generally expect Americans to want countermeasures that protect free speech as much as possible while limiting the distribution and impact of false or misleading information. In other words, people aim to maximize effectiveness and fairness while minimizing the intrusiveness of these measures. Countermeasures involving user participation tend to limit the need for content removal and algorithmic manipulation, which arguably threaten free speech rights the most and provide the least transparency. Therefore, we believe that these attributes are timely and relevant for the implementation and communication strategies regarding

misinformation mitigation.

Our first research question focuses on how perceived fairness, intrusiveness, and effectiveness of an intervention are related to the support of an intervention. We also explore whether user preferences differ based on whether the intervention is implemented by social media platforms or the government. With respect to countering misinformation, American citizens have expressed more concern about the government infringing on their free speech rights than private companies [153]. Therefore, perceptions of fairness, intrusiveness, and effectiveness may vary depending on which entity is responsible for these mitigation measures.

**RQ1.1: To what extent does a misinformation intervention’s perceived fairness, intrusiveness, and effectiveness predict support?**

**RQ1.2: How do the attributes people consider when forming preferences change due to the implementer of the intervention?**

Next, we compare the general support, perceived fairness, perceived effectiveness, and perceived intrusiveness for each intervention.

**RQ2: What is the average and variance in support, perceived fairness, perceived intrusiveness, and perceived effectiveness for each intervention?**

Furthermore, certain segments of the U.S. population on social media may be more or less accepting of misinformation interventions due to different attributes. Understanding demographic and partisan differences in intervention support informs public messaging and intervention design, as well as larger trends in values involved in policy support judgments.

**RQ3.1: How strongly do demographic differences predict support for misinformation interventions?**

**RQ3.2: Does support depend on different attributes for different demographic groups?**

This work will have implications for the public communication strategies that governments and social media companies will use to foster support and engagement with misinformation mitigation efforts. Furthermore, to the best of our knowledge, this study is the first attempt to directly measure public support for a wide range of possible intervention strategies implemented by government entities and social media companies.

### 5.2.1 Intervention Selection

We examined ten interventions that could be implemented by either a social media platform or a government entity in this study. These interventions were selected to represent a broad range of possible countermeasures. Existing review articles in the misinformation intervention area have categorized interventions similarly; however, there is no common typology [59, 93, 100, 242]. After reviewing previous categorizations and drawing from my intervention categorization in

Chapter 2, we included the six general categories of countermeasures that could apply to both platforms or governments: content distribution, content moderation, account moderation, content labeling, media literacy, and institutional measures. User-based measures were excluded. We identified 1-2 representative interventions per category to present to study participants. Participants were told in advance whether the implementer of the intervention was social media platforms or government entities, explicitly mentioning that the determiner of what misinformation is would fall on the intervention implementer (e.g., platforms could fact-check internally or use an external, independent organization). The interventions are described in Table 5.1.

Table 5.1: Selected misinformation interventions. Text in [brackets] was included when the specified implementer was government.

<b>Category</b>	<b>Intervention</b>	<b>Ref(s)</b>
1. Content distribution/ friction	[Require social media companies to] temporarily delay users posting content the user did not open or spent less than a certain amount of time viewing.	[115, 177, 182]
2. Content distribution/ advertising policy	[Require social media companies to] put all advertising through a fact-checking process.	[26, 100]
3. Content moderation/ alg. downranking	[Require social media companies to] de-emphasize posts that are verified to contain misinformation.	[26, 83]
4. Content moderation/ content removal	[Require social media companies to] remove posts verified to contain misinformation.	[26, 108]
5. Account moderation/ account removal	[Require social media companies to] permanently ban users who post misinformation a certain number of times.	[23, 186]
6. Content labeling/ misinfo. disclosure	[Require social media companies to] notify users if they posted content verified to contain misinformation.	[59, 123]
7. Content labeling / fact-check labels	[Require social media companies to] publicly label posts verified to contain misinformation with information about and from verified sources.	[170, 243]
8. Media literacy	Invest in digital media literacy and promote educational content about detecting misinformation on and offline.	[92, 124, 195]
9. Institutional measures/ media support	Promote and invest in local media, which is thought to be most in tune with local norms, culture, and context.	[25, 41]
10. Institutional measures/ data sharing	[Require social media companies to] regularly release data and/or internal research reports about misinformation prevalence, spread, and mitigation to the public and outside researchers.	[15, 26, 41]

## 5.3 Data and Methods

The survey questions associated with this Chapter were included in the same survey referenced in Chapter 3. The same ethics information and sampling plan apply. See Section 3.2 for more details.

### 5.3.1 Platform and Government Survey Questions

After the informed consent, qualifying questions, and individual behavior and opinion questions, participants progressed to the second half of the survey. Each participant was randomly assigned to see interventions implemented by either the government or social media companies. Each participant saw a random subset of 8 (of 10) interventions. They received the following instructions:

This next section concerns policies that could be implemented by [the government / social media companies] to limit the spread and influence of misinformation.

This means that, when applicable, **the [government / social media companies] would determine what misinformation is** when implementing the intervention(s).

Next, participants were instructed to evaluate one potential policy at a time, selected from those described in Table 5.1. They were first asked how much they support or oppose the proposed policy on a five-point Likert scale, which ranged from “Strongly support” to “Strongly oppose.” After that, participants were asked about perceived effectiveness, fairness, and intrusiveness. Each participant was asked about these three factors in a random order to mitigate potential question-order biases. This random order was maintained for all policies presented to each participant. The three factor questions are displayed below:

In your opinion, would the proposed policy be **effective** or **ineffective** in reducing misinformation on social media? [*Very effective, Somewhat effective, Neither effective nor ineffective, Somewhat ineffective, Very ineffective*]

In your opinion, would the proposed policy be **fair** or **unfair** when applied to different types of social media users? [*Very fair, Somewhat fair, Neither fair nor unfair, Somewhat unfair, Very unfair*]

In your opinion, would the proposed policy be **intrusive** or **unintrusive** on the experience of social media users? [*Very intrusive, Somewhat intrusive, Neither intrusive nor unintrusive, Somewhat unintrusive, Very unintrusive*]

The survey design and the phrasing of the factor questions were modeled off of prior work [103]. In particular, the phrasing of the fairness question was designed in order to frame it in terms of perceived distributional fairness (effects on various groups) rather than personal fairness (effects on oneself) [31, 89, 112]. This choice was deliberate, as not every participant may inherently consider the same definition of fairness when evaluating the fairness of different policies. Furthermore, a meta-analysis found that perceived distributional fairness had a significantly stronger effect on support levels for climate change policies compared to perceptions of personal fairness [31].

### 5.3.2 Measures

*Support for intervention(s)* We asked participants to rate a subset of interventions as {strongly support, somewhat support, neither support nor oppose, somewhat oppose, strongly oppose}. These responses are coded from 1 to 5 (least to most support).

*Perceived effectiveness of intervention(s)* We asked participants to rate a subset of interventions as {very effective, somewhat effective, neither effective nor ineffective, somewhat ineffective, very ineffective}. These responses are coded from 1 to 5 (least to most effective).

*Perceived fairness of intervention(s)* We asked participants to rate a subset of interventions as {very fair, somewhat fair, neither fair nor unfair, somewhat unfair, very unfair}. These responses are coded from 1 to 5 (least to most fair).

*Perceived intrusiveness of intervention(s)* We asked participants to rate a subset of interventions as {very intrusive, somewhat intrusive, neither intrusive nor unintrusive, somewhat unintrusive, very unintrusive}. These responses are coded from 1 to 5 (least to most intrusive).

### 5.3.3 Pre-Registered Analysis Plan

In our pre-registration (<https://osf.io/b2yjt/>), we included the primary research questions and analysis plan for this study.

#### RQ1: Predicting Support

To model support levels as a function of factor perceptions and implementer, we initially planned to run a multilevel model to account for random effects (i.e., random slope and intercept) of interventions and participants, as each participant saw 8 of the 10 selected interventions, drawn randomly, and multiple participants rated each intervention. We also pre-registered that if this model did not converge, we would fit an OLS regression model with robust standard errors clustered on participants and interventions. Since the multilevel model did not converge, the model output reported in this chapter is an OLS regression with robust standard errors clustered on participants and interventions. We ran the planned robustness checks by including participants who responded to a part of the survey but did not complete it. This does not change the direction or significance of the effects found in the primary model.

We calculated adjusted fractional Bayes factors with Gaussian approximations for the primary models using the BFPack R package [159]. We report BF10 for each estimate where the alternative hypothesis is directional based on the sign of the estimate (i.e.,  $b < 0$ ,  $b > 0$ ) and the null hypothesis is  $b = 0$ . Thus, if  $BF > 1$ , the evidence is more consistent with the alternative hypothesis; if  $BF < 1$ , the evidence is more consistent with the null hypothesis.

#### RQ2: Descriptive Analyses

We described how we would also report descriptive analyses of the average and spread of support for each intervention, as well as for the perceptions of effectiveness, fairness, and intrusiveness.

### RQ3: Individual Differences

We pre-registered our planned regression modeling analysis to determine how ratings for support, effectiveness, fairness, and intrusiveness vary based on individual differences of the participant. We included standard demographic variables (gender, age, income, education), partisanship (party id and political ideology), and self-reported exposure to misinformation. The question of misinformation exposure was asked in part 1 of the survey on individual behaviors and opinions (see Chapter 3).

## 5.4 Results

### 5.4.1 Factors that Influence Support for Interventions

**RQ1: To what extent does a misinformation intervention's perceived fairness, intrusiveness, and effectiveness predict support? How do the attributes people consider when forming preferences change due to the implementer of the intervention?**

Figure 5.1 and Table 5.2 show the regression results for RQs 1.1 and 1.2. Perceived fairness is most strongly associated with support ( $\beta = 0.624, SE = 0.016, p < 0.001$ ), followed by perceived effectiveness ( $\beta = 0.302, SE = 0.015, p < 0.001$ ) and intrusiveness ( $\beta = -0.065, SE = 0.010, p < 0.001$ ). These effects are moderated by the implementer, though implementer on its own is not significant in the model ( $\beta = 0.123, SE = 0.089, p = 0.167$ ). Fairness is less associated with support when the implementer is government than social media platforms ( $\beta = -0.080, SE = 0.024, p < 0.001$ ), while intrusiveness ( $\beta = -0.036, SE = 0.015, p = 0.015$ ) and effectiveness ( $\beta = 0.077, SE = 0.022, p < 0.001$ ) are more strongly associated.

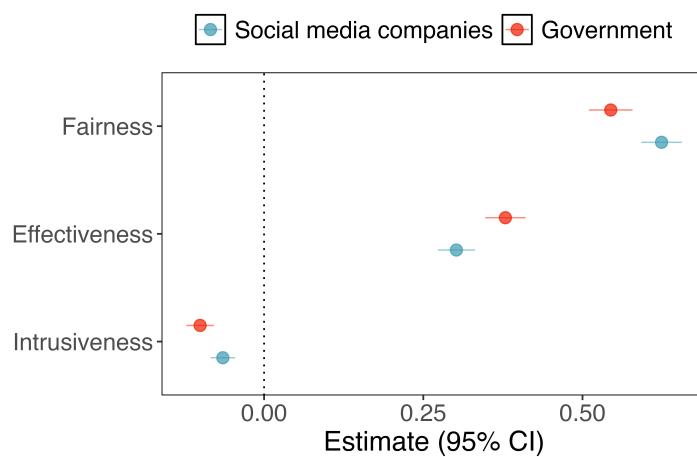


Figure 5.1: Estimate and 95% CI of the effect of perceived fairness, effectiveness and intrusiveness on support depending on implementer.

Table 5.2: OLS regression predicting support for a misinformation intervention as a function of perceived fairness, perceived effectiveness, perceived intrusiveness, and implementer (reference level: social media company) with robust standard errors clustered on participant and intervention. Model 1 is the main model. Model 2 is the robustness check that includes responses from all participants who responded to at least part of the survey.

	Model 1	Model 2
	<i>Estimate (Standard error)</i>	
(Intercept)	0.586*** (0.060)	0.582*** (0.060)
	BF > 100	BF > 100
Implementer	0.123 (0.089)	0.123 (0.089)
	BF = 0.053	BF = 0.054
Perceived fairness	0.624*** (0.016)	0.626*** (0.016)
	BF > 100	BF > 100
Perceived effectiveness	0.302*** (0.015)	0.300*** (0.015)
	BF > 100	BF > 100
Perceived intrusiveness	-0.065*** (0.010)	-0.064*** (0.010)
	BF > 100	BF > 100
Implementer × perceived fairness	-0.080*** (0.024)	-0.083*** (0.024)
	BF = 5.91	BF = 9.34
Implementer × perceived effectiveness	0.077*** (0.022)	0.080*** (0.022)
	BF = 10.31	BF = 17.93
Implementer × perceived intrusiveness	-0.036* (0.015)	-0.036* (0.015)
	BF = 0.42	BF = 0.43
Observations	8,071	8,102
R <sup>2</sup>	0.760	0.761
Adjusted R <sup>2</sup>	0.760	0.760

Note: \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

## 5.4.2 Overall Support and Perceptions of Interventions

**RQ2: What is the average and variance in support, perceived fairness, perceived intrusiveness, and perceived effectiveness for each intervention?**

Figure 5.2 contains the estimate and 95% CI for support, perceived fairness, perceived effectiveness, and perceived intrusiveness for each intervention by each implementer (government and social media company). Participants were more supportive of interventions in the content labeling category and less supportive of those in the content distribution or moderation categories.

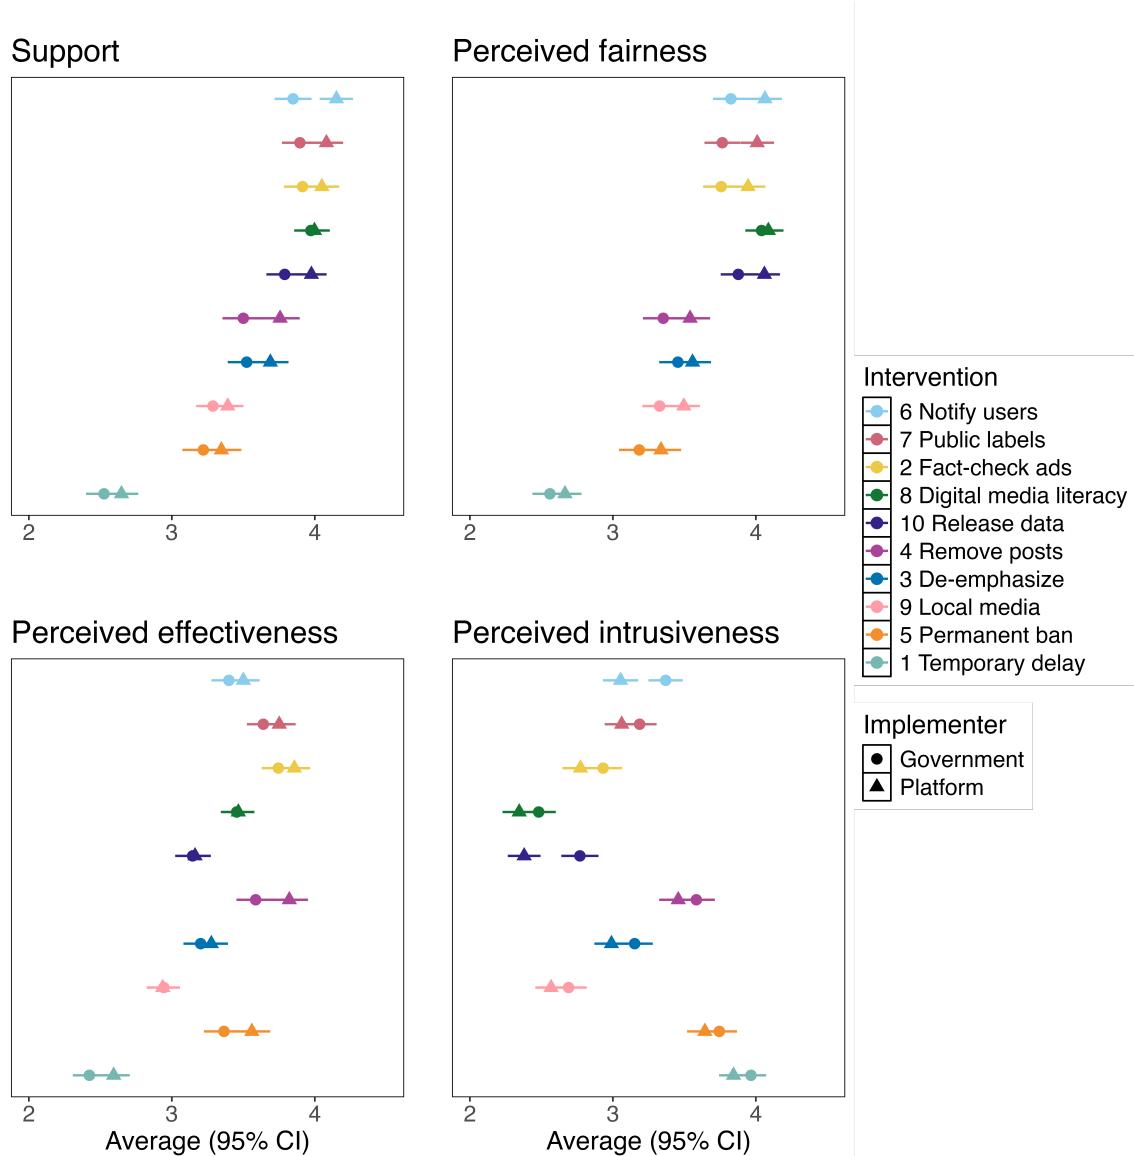


Figure 5.2: Average and 95% CI support, perceived fairness, perceived effectiveness, and perceived intrusiveness for each intervention (1-10) and implementer (government and platform) on a 1-5 Likert scale.

### 5.4.3 Individual Differences in Support and Perceptions of Interventions

Next, we investigate individual differences in support and perceived attributes of misinformation interventions (full regression output in Table 5.3). We do not find significant differences in perceptions of interventions across age groups, education level, or frequency of previous exposure to misinformation. However, we do for partisan and gender groups. Figure 5.3 shows the average support level, perceived fairness, perceived effectiveness, and perceived intrusiveness broken down by various demographic groups.

Republicans and Independents support misinformation interventions 23.8% and 15.3% less than Democrats on average, respectively. They also perceive interventions as less fair (22% and 14.1%) and less effective (17.8% and 12.1%) than Democrats on average ( $p < 0.05$  for all). These results are robust to separating the “Independent” and “Other/unaffiliated” categories and mapping partisan categories to corresponding numeric values (see SI tables S4 and S5). Similarly, being more liberal is associated with more support for interventions, higher fairness ratings, higher effectiveness ratings, and lower intrusiveness ratings ( $p < 0.001$  for all). Finally, for gender, we find that men support misinformation countermeasures 6.7% less than women on average ( $p < 0.001$ ). Furthermore, men rate interventions as 5.1% less fair ( $p = 0.012$ ), 6.9% less effective ( $p < 0.001$ ), and 3.8% more intrusive ( $p = 0.034$ ) than women on average.

Next, we performed ad hoc analyses to examine how the variables that varied in support (gender and partisanship) interact with the implementer of the intervention and perceived attributes to predict support (i.e., add demographic variables to the model used in RQ1). We find that partisanship interacts with implementer and fairness, with Republicans caring more about fairness than Democrats ( $\beta = 0.061, SE = 0.030, p = 0.040$ ). There is also a larger difference in support for interventions implemented by government versus social media companies for Republicans and Independents than for Democrats ( $\beta = -0.094, SE = 0.039, p = 0.017$ ;  $\beta = -0.108, SE = 0.036, p = 0.003$ ), where interventions implemented by governments are less supported. The implementer is now significant in the model. If perceptions of the interventions were equivalent, then government-led interventions would receive greater support ( $\beta = 0.213, SE = 0.094, p = 0.024$ ). However, as shown in Figure 5.2, interventions implemented by companies are generally supported more and perceived as more effective, fairer, and less intrusive than those implemented by the government. See Table 5.4 for the full results.

## 5.5 Discussion

In this work, we surveyed American social media users to examine public acceptance of interventions against misinformation implemented by the government and social media companies. We found that belief in fairness was most strongly associated with support for an intervention, followed by effectiveness, and finally intrusiveness. Fairness was more of a concern when the implementer was social media companies than the government, while effectiveness and intrusiveness were more salient when the government was the implementer (Figure 5.1). However, in general, the same intervention implemented by social media companies received more support, was perceived as fairer and more effective, and was viewed as less intrusive than when implemented by the government (Figure 5.2). These findings may reflect public attitudes towards

Table 5.3: OLS regressions predicting average support, perceived fairness, perceived effectiveness, and perceived intrusiveness of interventions as a function of age (18-24 to 65+ age brackets mapped to numeric values 0 to 5), gender (reference level: Female), partisanship (reference level: Democrat, combined “Other” party with “Independent”), political ideology (very liberal to very conservative categories mapped to numeric values 0 to 4), highest education level (less than high school diploma to Doctorate or Professional Degree categories mapped to numeric values 0 to 6), income (less than \$20,000 to over \$200,000 income brackets mapped to numeric values 0 to 7), frequency seeing misinformation (never to very often mapped to numeric values 0 to 4).

	Dependent variable:			
	Avg. support rating	Avg. fairness rating	Avg. effectiveness rating	Avg. intrusiveness rating
Age	0.020 (0.021)	0.044* (0.021)	-0.023 (0.019)	0.038* (0.018)
Male : Female	-0.199*** (0.054)	-0.137* (0.055)	-0.214*** (0.050)	0.102* (0.048)
Other (e.g., non-binary) : Female	-0.114 (0.173)	-0.296 (0.176)	-0.325* (0.161)	0.139 (0.154)
Education	0.013 (0.022)	0.003 (0.022)	0.006 (0.020)	0.016 (0.019)
Income	0.031* (0.016)	0.045** (0.016)	0.012 (0.015)	-0.005 (0.014)
Independent / Other : Democrat	-0.284*** (0.074)	-0.270*** (0.075)	-0.218** (0.068)	0.212** (0.066)
Republican : Democrat	-0.285** (0.107)	-0.294** (0.109)	-0.208* (0.099)	0.161 (0.095)
Political Ideology	-0.295*** (0.037)	-0.263*** (0.037)	-0.184*** (0.034)	0.130*** (0.033)
Misinformation Exposure	-0.014 (0.027)	-0.007 (0.027)	-0.047 (0.025)	0.038 (0.024)
Constant	4.236*** (0.122)	4.038*** (0.124)	3.997*** (0.114)	2.496*** (0.109)
Observations	1,010	1,010	1,010	1,010
R <sup>2</sup>	0.244	0.206	0.151	0.095
Adjusted R <sup>2</sup>	0.237	0.199	0.144	0.087
Residual Std. Error (df = 1000)	0.834	0.847	0.775	0.744
F Statistic (df = 9; 1000)	35.809***	28.792***	19.798***	11.711***

Note:

\*p<0.05; \*\*p<0.01; \*\*\*p<0.001

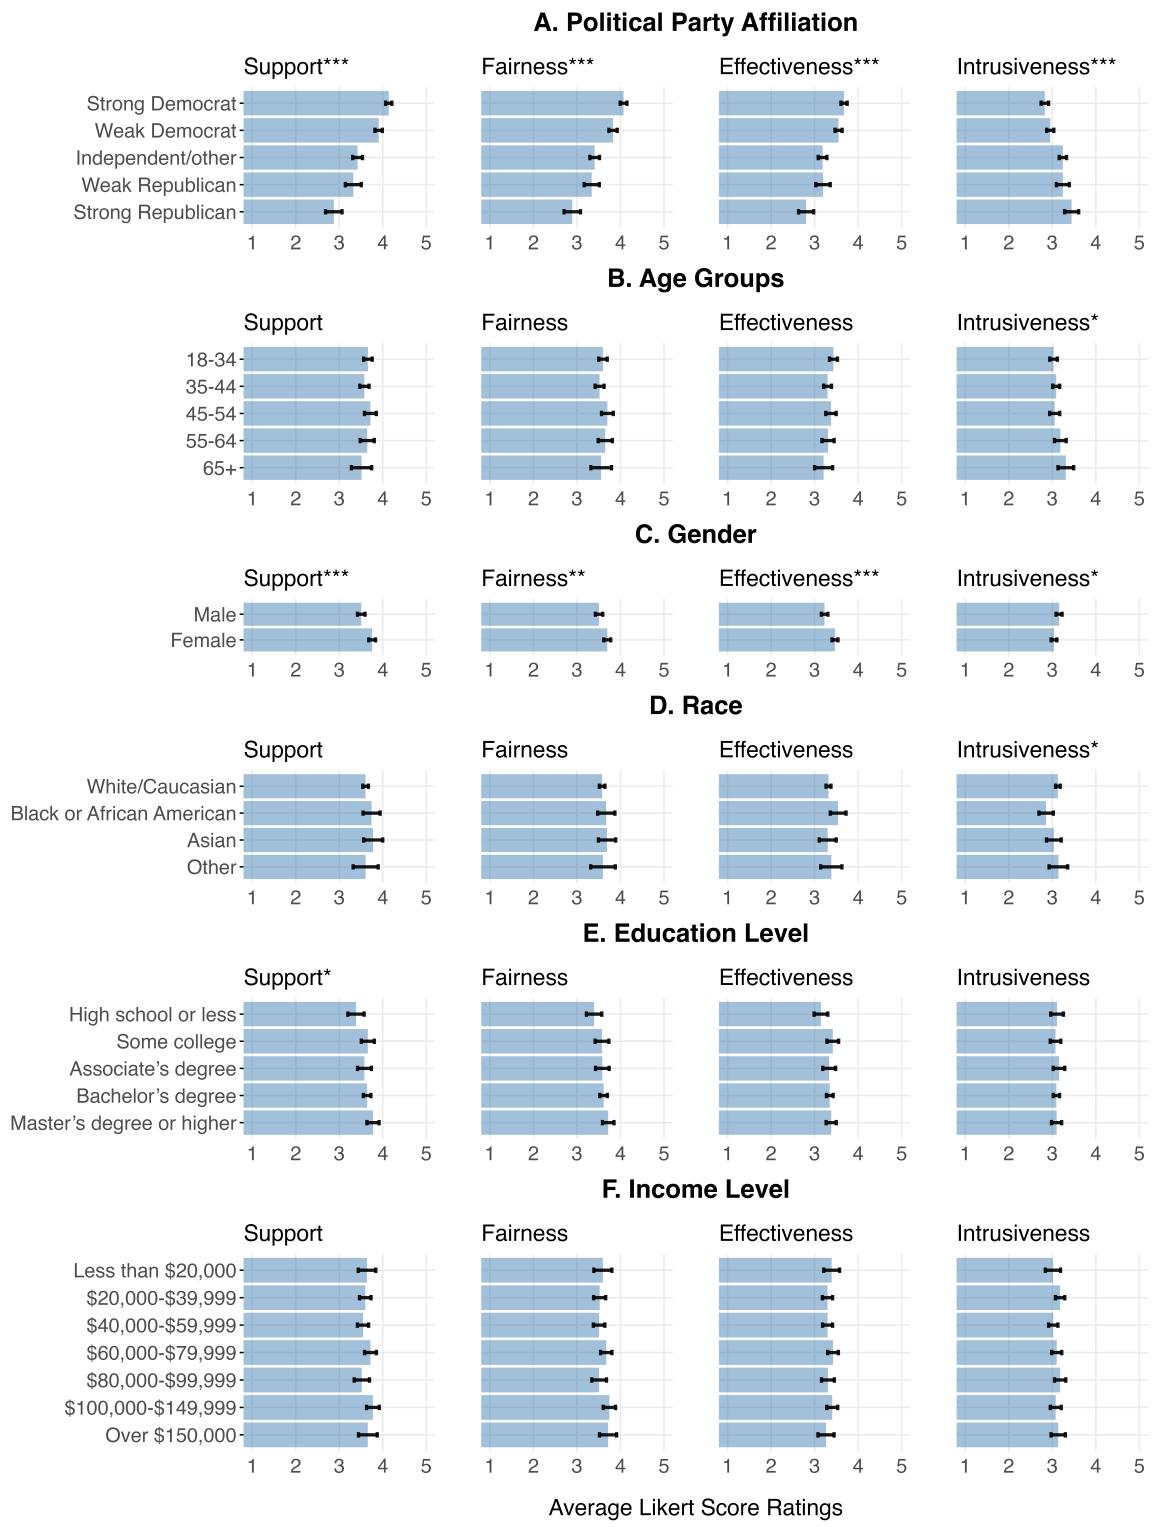


Figure 5.3: Average ratings by political party, age, gender, race, education, and income. One-way ANOVA tests were run on each grouping, with stars indicating the level of significance:  $p < 0.05^*$ ,  $p < 0.01^{**}$ , and  $p < 0.001^{***}$

Table 5.4: Adhoc analysis: OLS regression predicting support for a misinformation intervention as a function of perceived fairness, perceived effectiveness, perceived intrusiveness, implementer (reference level: social media company), gender (reference level: Female), and partisanship (reference level: Democrat, with Independents combined with “Other”) with robust standard errors clustered on participant and intervention.

	<i>Dependent variable: Support</i>	
	Estimate	Std. Err.
Implementer	0.213*	(0.094)
Perceived fairness	0.564***	(0.025)
Perceived effectiveness	0.303***	(0.022)
Perceived intrusiveness	-0.063***	(0.015)
Male : Female	-0.231**	(0.089)
Other gender : Female	0.706*	(0.296)
Independent/Other : Democrat	-0.251*	(0.111)
Republican : Democrat	-0.242*	(0.113)
Implementer x Perceived fairness	-0.086***	(0.024)
Implementer x Perceived effectiveness	0.066**	(0.022)
Implementer x Perceived intrusiveness	-0.034*	(0.014)
Implementer x Male	0.010	(0.030)
Implementer x Other gender	0.188	(0.107)
<b>Implementer x Independent/Other</b>	<b>-0.108**</b>	(0.036)
<b>Implementer x Republican</b>	<b>-0.094*</b>	(0.039)
Perceived fairness x Male	0.047	(0.024)
Perceived fairness x Other gender	-0.087	(0.060)
Perceived fairness x Independent/Other	0.038	(0.029)
<b>Perceived fairness x Republican</b>	<b>0.061*</b>	(0.030)
Perceived effectiveness x Male	-0.014	(0.022)
Perceived effectiveness x Other gender	-0.037	(0.068)
Perceived effectiveness x Independent/Other	0.018	(0.026)
Perceived effectiveness x Republican	-0.021	(0.027)
Perceived intrusiveness x Male	0.020	(0.015)
Perceived intrusiveness x Other gender	-0.059	(0.048)
Perceived intrusiveness x Independent/Other	-0.009	(0.018)
Perceived intrusiveness x Republican	-0.021	(0.019)
Constant	0.889***	(0.096)
Observations	8071	
R <sup>2</sup>	0.766	
Adjusted R <sup>2</sup>	0.766	

*Note:*

\*p<0.05; \*\*p<0.01; \*\*\*p<0.001

businesses and government, where companies are more trusted to address misinformation in a timely manner at scale. Alternatively, people may believe that social media companies have a greater responsibility to address misinformation than the government.

Our results further indicate that people desire agency and transparency in misinformation interventions, echoing findings from Saltz et al. [199] and research from other policy contexts [62, 89]. They support interventions that provide information to users that they can use when deciding how to interact with certain content, such as notifying them if they have posted misinformation, adding public labels to content containing misinformation, implementing digital media literacy programs, and requiring platforms to release of data or internal research reports related to misinformation. There was also strong support for holding advertising accountable through fact-checking. People were generally less supportive of interventions that involve removing or de-emphasizing posts identified as containing misinformation. While banning users who repeatedly post misinformation is also largely not supported, many believed that it would be relatively effective. Belief in effectiveness is simply not enough to support certain interventions that are considered unfair or intrusive. These results are consistent with the literature in other policy areas, which finds that the public generally prefers informational interventions over more restrictive measures even though they are often less effective [64, 94].

Interestingly, two of the least supported interventions do not directly threaten free speech or involve any censoring activity. Promoting and investing in local media was perceived to be largely ineffective. It may be that this intervention was too vague for participants to envision how it could help mitigate misinformation. Finally, the least popular intervention by a large margin was temporarily delaying users when attempting to post content they did not or barely viewed. It was rated as the least fair and effective and the most intrusive. This result is fairly unexpected considering that some platforms currently implement this intervention, including Facebook and X [56, 182]. Social media affords instantaneous communication and content, which could drive impatience for even the slightest inconvenience or delay (e.g., commercial breaks [46]). Therefore, when employing nudge-based approaches like accuracy prompts [175] or temporary delays in posting, it is imperative to minimize intrusiveness in the user experience, such as through design choices.

The analysis of individual differences in support for and perceptions of interventions revealed that men tend to support misinformation interventions less than women, and that Republicans and Independents support them less than Democrats. In addition to supporting interventions less, men and Republicans viewed them as less fair, less effective and more intrusive than women and Democrats, respectively. The gender gap reflects broader trends of women supporting more (government) regulations than men across policy domains [2]. The gap in support is, in part, explained by differences in perceptions of the proposed interventions. Future work should examine how perceived features of policies interact with other factors like emotional reactions and issue awareness to predict differences in support between genders [201].

Furthermore, the partisan differences align with findings from Saltz et al. [199]. Previous studies have shown that Republicans are more likely to perceive interventions as biased against them [199, 226]. A 2022 Pew Research poll found that approximately 70% of Republicans believe that major technology companies favor the views of liberals over conservatives, while only 22% of Democrats say they believe companies favor conservatives over liberals [226]. It is likely that because Republicans believe that they are more likely to be censored for their viewpoints,

they perceive interventions as less fair and are, therefore, less supportive of all interventions potentially employed by companies or the government. Whether this perception is accurate or not, social media companies must work on rebuilding trust among all their users.

Fairness emerged as the most significant predictor of support for interventions, and it varied across demographic groups, with fairness being especially important for Republicans. This emphasis on fairness is not unique to misinformation and social media policies. Across a variety of policy contexts, including health, environment, and transportation, perceived fairness and effectiveness are often among the most predictive factors associated with support [24, 31, 40, 89]. In fact, a systematic review found that enhancing the communication of a potential policy’s effectiveness to participants can boost support levels by 4%, a small but meaningful increase [188]. Fairness also plays a significant role and has been found to be the most important factor in a variety of public health interventions, such as those promoting healthy food [40], as well as support for climate change policies [31].

Intrusiveness is also a related factor, but it is not surprising that it holds less importance in the social media space than in other policy contexts. Previous studies have shown that when interventions are perceived to more directly impact an individual’s personal choices or daily life, like increased costs associated with owning a car [112], they tend to be unpopular [64, 94, 103]. It may be that the potential intrusiveness of social media policies on the user experience does not result in as severe of consequences for the typical individual as those in other policy areas, which could impact access to food, health care, transportation, and more. Acceptance levels may also be affected by the public’s awareness of the problem [64, 89], which was not investigated in this study.

## 5.6 Conclusion

### 5.6.1 Limitations

While this study is among the first to analyze the factors associated with support for misinformation interventions, several limitations could be addressed in future work. First, we focused our survey on active social media users, as those individuals would be the most affected by any potential policies and the most familiar with current interventions. We also only included ten interventions to limit the length of the survey, allowing us to focus more on the factors behind support. However, several new and emerging interventions were not included. For example, X’s (formerly Twitter) Community Notes program has been relatively successful at increasing the volume of fact-checks and boosting trust in misinformation flags by using crowd-sourcing misinformation detection and labeling [54, 67]. However, reports of its overall effectiveness in reducing engagement with misleading content are mixed [54, 113]. More factors (e.g., transparency) should be considered in future surveys about public acceptance of policies. Problem awareness may also be an important aspect to consider in future research [64, 89]. Finally, in the survey, we indicate that the implementer of the intervention is responsible for misinformation detection, which is an oversimplification. Future work should assess how people think misinformation should be detected.

## **5.6.2 Contributions**

Collectively, our findings suggest that fairness is valued above intrusiveness and effectiveness when determining support for misinformation interventions, and it is especially critical for specific demographic groups like Republicans. When designing and implementing misinformation interventions, mitigating any possible disproportionate impacts on certain groups or individuals is critical. In addition, public messaging should emphasize why each intervention is needed and how they are being implemented fairly, in addition to providing recourse for users when necessary. Furthermore, there is more support for and positive perceptions of interventions deployed by social media companies than the government, which may reflect broader trends in institutional trust. Most likely, effective misinformation interventions require collaboration across institutions. However, broader support for company-implemented interventions can be leveraged in public communications and education.

Our analysis of support levels and perceived features of interventions highlights the importance of promoting user agency to garner widespread support. For example, platforms can allow users to engage with misinformation warnings behind interstitials rather than strictly and opaquely removing violating content. At the same time, interventions should be carefully designed and implemented to minimize disruption to the user experience. Overall, this work has important implications for designing misinformation interventions and messaging that will be positively received by social media users.

# Chapter 6

## Recommendations for Effective and Practical Countermeasures

While earlier chapters have characterized a wide range of countermeasures and sought to evaluate and improve those interventions, this chapter consolidates that prior work to establish a framework for developing effective and practical countermeasures. The objective of this chapter is to offer a comprehensive set of recommendations across the intervention landscape that researchers, companies, and policymakers can use.

To develop this framework, I integrate the research from previous chapters and evaluate interventions based on several critical features: **user acceptance, effectiveness, effort level, cost, and political feasibility**. Additionally, an expert survey was conducted to gather professional opinions on these factors and incorporate them into this analysis.

The main research question for this chapter is: How can we identify which interventions are both practical and effective, and under what circumstances? More specifically:

1. What are the characteristics of different interventions?
2. What features do effective and practical countermeasures have in common?

### 6.1 Introduction

Interventions are unlikely to be implemented unless they are both effective and have the support of the affected users or the public, at least to some extent. However, while effectiveness and acceptance may be seen as necessary conditions for implementation, they are not sufficient on their own. The literature on comparative policy analysis indicates that other critical factors are associated with the likelihood of a policy's implementation, whether at the platform or government level: the level of implementation effort, intervention cost, and overall political feasibility [130, 192].

In this chapter, the operationalized interventions outlined in Section 1.6 will be compared and ranked based on five criteria: user acceptance, effectiveness, effort level, cost, and political feasibility. These factors will be evaluated using a combination of sources, including the existing literature, results from previous chapters, and findings from an expert survey conducted for this chapter.

As we observed in Chapter 2, there are several over- and under-studied interventions in the literature, making it challenging to compare interventions if only reviewing the existing research. Additionally, user acceptance is a relatively under-studied aspect of interventions in the current academic literature, although it was supplemented by work in some of the previous chapters. Values for effort level, cost, and feasibility are limited in the literature and will primarily be drawn from the expert survey results.

For this chapter, we surveyed X misinformation researchers who had been invited to or attended a recent *Disinformation and AI* summit held at Carnegie Mellon University from January 22-24, 2025. These researchers were asked for their professional opinions on the effectiveness and acceptability of the general intervention categories. They were then asked to rate the operationalized interventions on the five primary factors.

## 6.2 Related Work

So far, the misinformation countermeasures in this dissertation have primarily been categorized by general types, such as content distribution and account moderation. These categories were created because they share several *descriptive characteristics*. Additionally, as discussed in Chapter 4, the features of the interventions and how they are implemented can influence *user perceptions* of them. We found that an intervention's perceived effectiveness, fairness, and intrusiveness were all strongly related to user acceptance of the interventions. *Descriptive characteristics* and *user perceptions* of interventions both affect the five *main criteria* analyzed in this chapter: user acceptance, effectiveness, effort level, cost, and political feasibility.

### 6.2.1 Descriptive Characteristics

Descriptive characteristics of interventions refer to the relatively static features that outline how the intervention is implemented and by whom. These features include:

- **Implementer:** Which organization or organizations are implementing the intervention? For example, a social media company or the federal government.
- **Policy Changes:** Does the intervention require a new organizational policy, law, or regulation? This is most often dictated by the implementer.
- **Targeted Aspect:** Does the intervention primarily target the creation, spread, or belief in misinformation? If targeting beliefs, is it primarily preventative (increasing competence) or reactive (correcting mis-beliefs)?
- **Targeted Content:** What type of content is targeted by the intervention? This includes the news context, such as harmful political or health misinformation.
- **Information Changes:** Does the intervention require informational changes, such as adding, altering, removing, or tagging the misinformation?

For example, the intervention *redirection* is typically implemented by platforms and focuses on the spread of misinformation rather than its creation or belief. Platforms target specific news contexts according to their policies, and redirection policies likely require an updated platform policy that specifies what is targeted, how it is targeted, and why. For example, platform policies

could clarify that authoritative public health sources, like the CDC or WHO, will be prominently featured in search results while unverified sources will not be. Additionally, implementing redirection requires platforms to modify the information shown to users or the order in which it is shown, which requires informational changes.

These descriptive characteristics likely affect users' perceptions of the interventions, which in turn affect public policy support. For example, transparency is an important factor in public policy support [89, 108]. A descriptive and transparent redirection policy or clearly marked informational changes may influence an individual's perception of an intervention's level of fairness or intrusiveness.

### 6.2.2 User Perceptions

User perception refers to how individuals view the different characteristics of interventions and includes factors such as perceived effectiveness, fairness, intrusiveness, transparency, and problem awareness [89]. Specifically, it addresses how people believe the intervention will be implemented and their expectations about its effects. For example, do they think the intervention will effectively address the issue at hand? Do they feel the problem is something that even requires a solution? Even if they believe the strategy would be effective, do they trust the implementing organization to carry out the intervention in a fair, transparent, and non-intrusive way? The descriptive characteristics of the interventions influence user perceptions:

- **Implementer:** Some individuals may trust certain implementers more than others. For example, various demographic characteristics, such as partisanship, are associated with lower levels of support for government regulation [226].
- **Policy changes:** Transparent policy changes can enhance trust and, consequently, support for misinformation interventions [199]
- **Targeted Aspect:** Chapter 4 and prior literature shows that individuals prefer informational interventions over intrusive ones [64, 94]. This could mean greater support for interventions aimed at addressing the belief in misinformation rather than trying to control its creation or restrict its spread.
- **Targeted Content:** Research indicates that problem awareness is linked to support for policy interventions [89]. Prior work in the public health domain shows a higher level of support for anti-smoking initiatives compared to other similar interventions aimed at addressing alcohol consumption or diet. This difference in support can be partially explained by the public's awareness of the negative consequences associated with smoking [64]. Therefore, the severity of the content and its potential risks are likely connected to how seriously individuals perceive the issue. People may be more willing to accept interventions for content they consider to be particularly severe. In Chapter 3, we observed individuals clearly articulate their greater willingness to engage in social corrections for serious content.
- **Informational changes:** Informational changes associated with interventions can also be implemented in explainable and transparent ways, increasing support levels [199]. Additionally, improved tool usability and platform reporting support have been shown to both

increase users' willingness to engage with reporting systems as well as their perceived effectiveness [248].

In addition to the descriptive characteristics that affect user perceptions, increased educational and messaging efforts could also improve perceptions. For example, a meta-analysis of public policy interventions showed that effectively communicating a potential policy's effectiveness noticeably increased support by approximately 4% [188].

### 6.2.3 Main Criteria

Each defined intervention will be compared and rated across five primary characteristics: acceptance, effectiveness, effort level, cost, and political feasibility. These criteria were selected by reviewing previous systematic review articles or policy analyses. In Blair et al., the authors note that in addition to effectiveness as a standard measure of an intervention to consider, that policymakers and practitioners designing interventions should also consider feasibility, scalability, and durability [35]. Feasibility refers to how easily an intervention could be implemented and scalability refers to how easy it is to scale that intervention to reach more individuals. Durability is an aspect of effectiveness and refers to how long the desired effect lasts. It is often studied in the academic literature, at least in the most commonly studied intervention types such as media literacy [92, 102], inoculation [147], and fact-checking [180]. In a comparative policy analysis of various types of social media regulation, Rochefort uses effectiveness, administrative difficulty, cost, and political acceptability as criteria [192]. Administrative burden, in the form of cost or implementation effort, is often used in policy analysis [69, 130].

The definitions of these main factors are as follows:

- **Acceptance:** Interventions may be acceptable and supported by certain types of users, in specific situations, and on some or all platforms. Acceptance is impacted by numerous factors including perceived transparency, intrusiveness, privacy, fairness, as well as descriptive factors like the implementer of the intervention.
- **Effectiveness:** Interventions may be effective at reducing the creation, spread, or belief in misinformation for some users, in some circumstances, or on certain platforms.
- **Effort level:** The effort level to implement the intervention or engage with it for the user, platform, or the government entity may vary substantially for different intervention types.
- **Cost:** The cost for the user, platform, or government entity may vary substantially for different intervention types.
- **Political feasibility:** Political feasibility refers to the likelihood that an intervention would be implemented and incorporates all previous factors as well as external factors, such as stakeholder perspectives and support from relevant officials.

The descriptive characteristics of each intervention and how individuals perceive these interventions influence the five key evaluative criteria used throughout this chapter. For example, user perceptions impact both **acceptance** levels and the degree of policy buy-in, which can increase policy **effectiveness**. Similarly, interventions that require large amounts of policy or informational changes may increase administrative burden, which in turn increases implementation **effort** and **cost**. **Political feasibility** incorporates these factors as well as several external elements,

such as governmental or platform structures, legal considerations, the influence of other stakeholders, and the current political climate. For example, if the administrative burden of a policy (high cost and implementation effort) outweighs the potential benefits (effectiveness of the intervention), such policies may not be implemented even if they are both effective and acceptable to the public [130].

An example of the impact of political feasibility is Meta’s recent removal of its fact-checking program [114]. This program has been found to be fairly effective [87, 135] and accepted by most users across the political spectrum [185]. While fact-checking is popular, including among a majority of Republicans, support is higher among Democrats and Independents than among Republicans [85, 185]. However, recent work also finds that not only do Republicans perceive professional fact-checkers as a more legitimate means of doing content moderation compared with regular social media users [148], but relying solely on community notes is actually extremely unpopular across the political spectrum [185]. External factors, such as public pressure, fear of potential government regulation, or perceived anti-conservative bias, may have influenced their decision to remove the program [246].

## 6.3 Data and Methods

A survey was designed to gather expert opinions on intervention acceptance, effectiveness, effort level, cost, and political feasibility. The survey was promoted to participants and invitees for the *Disinformation and AI Summit* at Carnegie Mellon University, which took place from January 22 to 24, 2025. Data was collected between January 22 and March 9, 2025.

### 6.3.1 Survey Design

The survey was conducted using Qualtrics and consisted of four main sections: Introduction and Consent, General Category Questions, Intervention Questions, and Demographics and Survey Feedback.

#### Introduction and Consent

Participants were informed that the survey focused on interventions to counter misinformation. They were also told that the survey would take approximately 10 minutes, that there were no risks or compensation associated with the survey, that they could exit the survey at any time, and that all responses would be anonymized in any analysis of the results. They were asked to give their informed consent to participate before proceeding to the survey.

#### General Categories

In this section, participants were asked for their professional opinions on the eight general categories of misinformation interventions defined in this dissertation (Table 1.5). The only difference is that the **Other** category was redefined as a **Generative AI** category. More specifically, participants were asked two primary questions.

1. To what extent do you agree that the following categories of misinformation interventions **are effective** at combatting misinformation if adopted widely by social media companies or institutions? [*Strongly disagree, Disagree, Tend to disagree, Neither agree nor disagree, Tend to agree, Agree, Strongly agree*]
2. To what extent do you agree that the following categories of misinformation interventions **would be acceptable** to social media users or citizens if adopted widely by social media companies or institutions? [*Strongly disagree, Disagree, Tend to disagree, Neither agree nor disagree, Tend to agree, Agree, Strongly agree*]

For each question, participants were instructed to consider, compare, and rate all eight categories simultaneously using a 1-7 Likert Scale. They were provided with the following definitions for each general category, summarized from Appendix A:

- **Content Distribution:** Refers to how content is distributed on social media. This includes interventions such as redirection, accuracy prompts, friction, platform alterations, and advertising policy
- **Content Moderation:** Refers to how content is shown or not shown on social media. This includes interventions such as fact-checking, debunking, misinformation detection, and algorithmic content moderation
- **Account Moderation:** Refers to how accounts are moderated on social media. This includes suspending, banning, shadow banning, or demonetizing user accounts.
- **Content Labeling:** Refers to misinformation disclosure involving labels of any kind. This includes fact-checking labels, crowdsourcing labels like community notes, warning labels, source credibility labels, and context labels
- **User-based Measures:** Refers to measures that involve other users seeing misinformation and how they respond to it. This includes users reporting or blocking other users or their posts, using social corrections, or updating social norms.
- **Media Literacy and Education:** Refers to any educational or training effort meant to increase the public's civic reasoning and critical thinking skills. Includes fake news games and other inoculation methods
- **Institutional Measures:** Refers to how governments and other public or civic institutions can help manage the spread of misinformation. This includes interventions such as investing in or promoting local news, allowing researchers access to data, and government regulation.
- **Generative AI:** The usage of generative AI to detect misinformation, generate rebuttals, create educational resources, or dialoguing with those who believe misinformation.

## Specific Interventions

In this section, participants were asked to provide their professional opinions on the specific, operationalized interventions (see Tables 1.7 and 1.8). Each participant randomly viewed 12 of the 40 interventions described in those tables. They were informed that they would be asked about each potential intervention's effectiveness, user acceptance, effort level, cost, and political

feasibility. They rated each intervention on these five factors using a 1-5 Likert scale, similar to the one implemented in Chapter 4.

- Would the proposed policies be **effective or ineffective** in combatting misinformation if adopted widely by social media companies or institutions? [*Very ineffective, Somewhat ineffective, Neither effective nor ineffective, Somewhat effective, Very effective*]
- Would the proposed policies be **acceptable or unacceptable** to users or citizens if adopted widely by social media companies or institutions? [*Very unacceptable, Somewhat unacceptable, Neither acceptable nor unacceptable, Somewhat acceptable, Very acceptable*]
- Would the proposed policies require a **high or low amount of time and effort** for companies or institutions to implement? [*Very low effort, Somewhat low effort, Neither high nor low effort, Somewhat high effort, Very high effort*]
- Would the proposed policies require a **high or low amount of financial investment** for companies or institutions to implement? [*Very low cost, Somewhat low cost, Neither high nor low cost, Somewhat high cost, Very high cost*]
- Would the proposed policies be **politically feasible or infeasible** for companies or institutions to implement? [*Very infeasible, Somewhat infeasible, Neither feasible nor infeasible, Somewhat feasible, Very feasible*]

## Demographics and Survey Feedback

Participants were given optional demographic questions to answer, including academic rank (PhD student, Postdoc, Faculty, Industry, or Other), age category, gender, and primary area of expertise (Psychology, Communication and Media Studies, Political Science, Journalism, Computational Social Sciences, Sociology, Computer Science, or Other). Finally, participants were given the opportunity to share their thoughts and feedback.

### 6.3.2 Participant Demographics

Thirty-eight participants completed the survey. A high-level summary of participant demographics is provided below:

- **Gender:** 15 women (39.5%), 22 men (57.9%), 1 other/rrefer not to answer
- **Age:** 18-34 (14 people, 36.8%), 35-44 (12 people, 31.6%), 45+ (10 people, 26.3%), 2 Prefer not to answer
- **Academic Rank:** 21 faculty (55.3%), 11 PhD students (28.9%), 2 Industry (5.3%), 1 Postdoc (2.6%), 3 Other (7.9%).
- **Primary Discipline:**
  - *Computational social sciences:* 17 (44.7%)
  - *Computer science:* 9 (23.7%)
  - *Sociology:* 4 (10.5%)
  - *Communications and media studies:* 4 (10.5%)

- *Political science:* 2 (5.3%)
- *Other:* 2 (5.3%)

### **6.3.3 Analysis Methods**

This chapter will mainly analyze the values from the expert survey, but it will also discuss the results from the previous chapters and the literature review if applicable. Meta-analyses and systematic review articles on the effectiveness of specific interventions, if available, will be incorporated in the discussion. Any available sources for the other four metrics, such as public opinion polls or policy analyses, will also be addressed.

In general, an ideal intervention would have high acceptance, effectiveness, and political feasibility while requiring low effort and cost. Therefore, to better compare interventions, effort level and cost scores were inverted so that 1 indicates very high effort or cost and 5 indicates very low effort and cost. This was done so that high Likert scores for all five factors represented the ideal outcome. The average Likert scores across the five metrics will be calculated per intervention and will be used to compare and rank interventions.

### **6.3.4 Measures**

*Effectiveness for interventions by general category:* These responses are coded from 1 to 7 (strongly disagree to strongly agree).

*Acceptance for interventions by general category:* These responses are coded from 1 to 7 (strongly disagree to strongly agree).

*Effectiveness for specific intervention(s):* These responses are coded from 1 to 5 (least to most effective)

*Acceptance for specific intervention(s):* These responses are coded from 1 to 5 (least to most acceptable).

*Effort level for specific intervention(s):* These responses are coded from 1 to 5 (very high to very low effort).

*Cost for specific intervention(s):* These responses are coded from 1 to 5 (very high to very low cost).

*Political feasibility for specific intervention(s):* These responses are coded from 1 to 5 (very infeasible to very feasible).

## **6.4 Results**

First, we compare how effective and practical experts consider each general category of misinformation intervention to be. Then, each category of interventions will be analyzed in more detail to

compare the operationalized interventions within each category. The descriptive characteristics of the interventions are outlined. The results of the expert survey analysis on the five main factors are presented and compared with the work from the previous chapters or the existing literature. The results section concludes with a summary of the top interventions for each of the five main metrics.

#### 6.4.1 Comparison of General Categories

Figure 6.1 and Table 6.1 show the fraction of participants who believe each category to be effective or acceptable to the public, sorted by the fraction who believe it is effective. Content Moderation, Account Moderation, and Media Literacy and Education are regarded as the most effective, with over 70% of respondents agreeing they would be effective. However, moderation, especially Account Moderation, is viewed by experts as one of the least acceptable to the public.

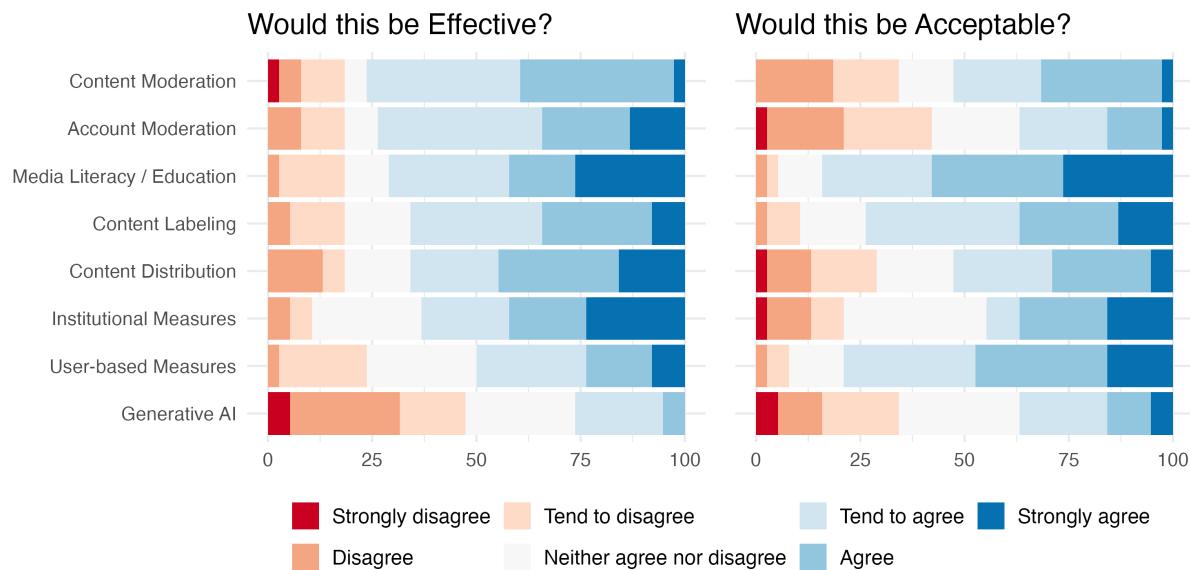


Figure 6.1: Stacked bar plot representing the percentage of responses on the effectiveness and acceptance of the general categories of interventions.

Figure 6.2 illustrates the average Likert scores and 95% confidence intervals for the effectiveness and acceptance ratings, sorted by the highest average effectiveness score. In this figure, we observe that Media Literacy / Education and Institutional Measures rise in the effectiveness rankings, bolstered by the proportion of participants who “strongly agree” regarding their effectiveness rather than more weakly agreeing. Content and Account Moderation remain at high levels. This figure more clearly highlights the differences in effectiveness and acceptance across various intervention categories. Account Moderation and User-based Measures exhibit the largest gap between effectiveness and acceptance, with effectiveness exceeding acceptance in Account Moderation, and acceptance exceeding effectiveness in User-based Measures.

These results align with a similar survey of experts, which found that they believed media

Category	Effectiveness			Acceptance		
	Agree	Disagree	Neither	Agree	Disagree	Neither
Content Moderation	76.3	18.4	5.26	52.6	34.2	13.2
Account Moderation	73.7	18.4	7.89	36.8	42.1	21.1
Media Literacy / Education	71.1	18.4	10.5	84.2	5.26	10.5
Content Labeling	65.8	18.4	15.8	73.7	10.5	15.8
Content Distribution	65.8	18.4	15.8	52.6	28.9	18.4
Institutional Measures	63.2	10.5	26.3	44.7	21.1	34.2
User-based Measures	50.0	23.7	26.3	78.9	7.89	13.2
Generative AI	26.3	47.4	26.3	36.8	34.2	28.9

Table 6.1: The percentage of respondents who agree or disagree that each general category of interventions is effective or acceptable to the public.

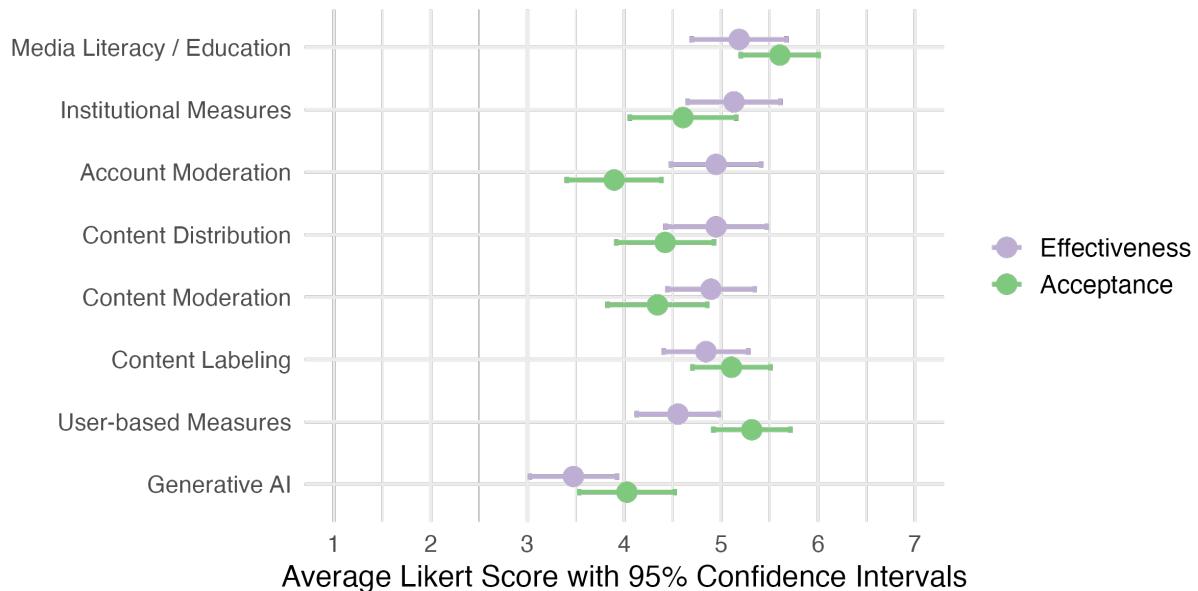


Figure 6.2: Average effectiveness and acceptance Likert scores per general intervention category.

literacy, labeling of false content, and fact-checking to be the most effective interventions, each receiving around 70% approval. Meanwhile, accuracy prompts (a form of content distribution), source labeling, and inoculation ranked lower on the list. These were the only interventions included in this survey in a similarly structured question [10]. In this survey conducted by Altay et al., experts were also asked if certain types of actions *should* be taken, with platform design changes, algorithmic changes, and general content moderation as the most popular (over 75% agree), and penalizing misinformation sharing and shadow banning (a form of account moderation) as the least popular. In fact, shadow banning was the only suggested intervention where more experts disagreed with its implementation than agreed [10]. Similarly, Blair et al. found in a survey of experts that they recommended implementing educational and institutional interventions (such as media literacy, platform alterations, and journalist training) over other types of interventions [35].

## 6.4.2 Content Distribution

This section presents a comparative analysis of the operationalized content distribution interventions (refer to Table 1.7 or Appendix A for detailed definitions).

### Descriptive Characteristics

Table 6.2 shows the descriptive characteristics of the eight operationalized content distribution interventions, listed in the order they were initially presented in Table 1.7.

Intervention	Implementer	Targeted Aspect	Information Changes
Redirection	Platforms	Spread, Belief	Alters, Removes
Accuracy Prompts	Platforms	Spread	Adds
Friction	Platforms	Spread	Adds
Platform Alterations	Platforms	Spread	Alters
Ban Political Ads	Platforms, Govts	Spread	Removes
Fact-check Ads	Platforms, Govts, Institutions	Belief	Tags, Removes
Limit Forwarding	Platforms	Spread	Removes
Limit Resharing	Platforms	Spread	Removes

Table 6.2: Descriptive characteristics of Content Distribution interventions.

The implementer column outlines the typical implementer of each intervention; sometimes, these interventions could be carried out by one or more implementers. For example, banning political advertising on social media platforms may originate from the platforms themselves or be mandated by the government, which would still require the platforms to implement it. Fact-checking advertisements might involve third-party fact-checking organizations. Policy changes depend on the implementer. When platforms adopt many of these content distribution interventions, such as redirection, friction, advertising policies, and restrictions on resharing or forwarding, they are usually detailed in their platform policies (see Section X), though to varying degrees of detail.

Most content distribution interventions target the distribution or spread of misinformation. However, redirecting users to official content or fact-checking advertisements also addresses the belief in misinformation. When considering the informational changes associated with implementing the interventions, both accuracy prompts and friction require adding interstitials or other content or UI changes. Redirection and platform alterations alter the way information is displayed to users. Banning political ads removes content, while limiting forwarding or resharing also removes functionality.

## Expert Survey Results

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
Limit Forwarding	13	3.69 (0.63)	3.54 (1.05)	3.92 (1.04)	4.31 (0.85)	4.08 (0.86)	3.91
Limit Resharing	8	3.75 (0.46)	2.88 (1.25)	4.12 (0.64)	4.50 (0.53)	3.88 (1.25)	3.83
Friction	12	3.50 (1.00)	3.33 (1.07)	3.33 (1.15)	3.83 (1.47)	3.75 (0.97)	3.55
Accuracy Prompts	8	3.62 (0.92)	3.50 (1.07)	3.38 (1.06)	2.88 (1.13)	3.38 (1.51)	3.35
Platform Alterations	10	3.80 (1.32)	3.00 (0.94)	2.50 (0.97)	3.10 (1.37)	3.60 (1.07)	3.20
Redirection	10	3.10 (0.99)	3.10 (1.10)	3.10 (1.10)	3.90 (0.88)	2.60 (1.35)	3.16
Ban Political Ads	15	3.00 (0.93)	3.21 (1.37)	3.60 (1.30)	3.27 (1.49)	2.20 (1.32)	3.06
Fact-Check Ads	9	3.44 (0.73)	3.75 (0.71)	2.00 (0.93)	2.00 (0.93)	3.33 (1.00)	2.91

Table 6.3: Expert scores for content distribution interventions. Mean values are shown with their corresponding standard deviations in parentheses.

## Discussion

**Redirection / Platform Alterations** - Redirection and platform alterations are generally effective, have moderate acceptance levels, and are relatively easy to implement and scale, although more studies are needed [34, 36, 123]. Redirection through related articles and similar methods has generally been found to be effective in the literature [34, 36, 123]. Regarding user acceptance, one prior study found that over half of participants agreed with the use of related articles combined with a reduction in the size of misinformation, while less than a quarter disagreed.

**Accuracy Prompts** - Accuracy prompts are generally effective, although effect sizes vary, and they are relatively easy to implement and scale. However, more studies are needed on user acceptance [35, 174]. A meta-analysis found that accuracy prompts can reduce the sharing of false information by approximately 10% [174]. User acceptance of accuracy prompts has not been studied as extensively, though it may be similar to related interventions such as friction.

**Friction** – While friction may be effective, more work is needed [22, 34, 205]. It may also be relatively unpopular, though additional research should be conducted. Friction was analyzed in Chapter 4 and was the least popular intervention, with an average Likert score of 2.6 (or “low”), despite the frequent usage of this intervention by platforms. Furthermore, it received low perceived effectiveness and fairness ratings and was viewed as highly intrusive.

**Ban Political Ads / Fact-Check Ads** - Banning political ads is generally supported by the public, but more work is needed to understand its effectiveness [16]. Similarly, the public generally supports fact-checking ads, and fact-checking has generally been found to be effective

[5, 232]. Fact-checking ads ranked among the most popular interventions tested in Chapter 4, with an average Likert score of 4 (indicating “high” support). It additionally received high scores for perceived effectiveness and fairness and was rated at a medium level of intrusiveness.

**Limit Forwarding/Limit Resharing** - Limiting the rampant forwarding or sharing of messages is likely effective and relatively easy to implement and scale [3, 181]. However, more studies are needed on both its effectiveness and user acceptance

### 6.4.3 Content Moderation

This section presents a comparative analysis of the operationalized content moderation interventions (refer to Table 1.7 or Appendix A for detailed definitions).

#### Descriptive Characteristics

Table 6.4 shows the descriptive characteristics of the five operationalized content moderation interventions, in the order they were initially presented in Table 1.7.

Intervention	Implementer	Targeted Aspect	Information Changes
Content Removal	Platforms, Govts	Spread, Belief	Removes
Downranking	Platforms	Spread	Alters, Removes
Algorithmic Changes	Platforms	Spread	Adds, Alters, Removes
Virality Circuit Breakers	Platforms	Spread	Removes
User Control	Platforms, Users	Spread	Adds, Removes

Table 6.4: Descriptive characteristics of Content Moderation interventions.

#### Expert Survey Results

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
User Control	14	3.00 (1.18)	4.36 (0.84)	2.31 (0.95)	3.38 (1.39)	4.21 (1.12)	3.45
Algorithmic Changes	11	4.09 (0.54)	3.55 (1.04)	2.55 (0.93)	3.00 (1.26)	3.64 (1.03)	3.36
Virality Circuit Breakers	12	3.58 (1.38)	3.08 (0.90)	2.92 (1.00)	3.00 (1.28)	3.58 (1.08)	3.23
Downranking	12	3.92 (0.79)	3.17 (1.19)	2.92 (0.90)	3.00 (1.21)	3.00 (1.41)	3.20
Content Removal	8	3.75 (0.89)	3.25 (0.89)	2.25 (1.16)	3.62 (0.92)	3.00 (1.20)	3.17

Table 6.5: Expert scores for content moderation interventions. Mean values are shown with their corresponding standard deviations in parentheses.

<b>Intervention</b>	<b>Implementer</b>	<b>Targeted Aspect</b>	<b>Information Changes</b>
Account Suspensions	Platforms	Creation	Removes
Deplatforming	Platforms	Creation	Removes
Shadow Banning	Platforms	Spread	Removes
Demonetization	Platforms	Creation	Removes

Table 6.6: Descriptive characteristics of Account Moderation interventions.

## Discussion

### 6.4.4 Account Moderation

#### Descriptive Characteristics

#### Expert Survey Results

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
Demonetization	12	4.00 (0.6)	3.75 (1.06)	2.92 (1.24)	2.92 (1.51)	3.58 (1.31)	3.43
Shadow Banning	9	3.33 (1.32)	2.89 (1.05)	3.38 (1.51)	3.50 (1.69)	3.44 (1.24)	3.31
Account Suspensions	15	3.53 (0.99)	2.86 (1.29)	2.40 (0.83)	3.47 (1.13)	2.80 (1.52)	3.01
Deplatforming	11	3.73 (0.79)	2.82 (0.98)	2.36 (1.12)	3.27 (1.42)	2.64 (1.36)	2.96

Table 6.7: Expert scores for Account Moderation interventions. Mean values are shown with their corresponding standard deviations in parentheses.

## Discussion

### 6.4.5 Content Labeling

#### Descriptive Characteristics

<b>Intervention</b>	<b>Implementer</b>	<b>Targeted Aspect</b>	<b>Information Changes</b>
Crowdsourcing	Platforms, Users	Belief	Adds, Tags
Click-through Warnings	Platforms	Belief	Adds, Alters
Source Labeling	Platforms	Belief	Adds, Tags
Government Labels	Platforms	Belief	Adds, Tags
Fact-Check Labels	Platforms, Institutions	Belief	Adds, Tags

Table 6.8: Descriptive characteristics of Content Labeling interventions.

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
Government Labels	5	3.80 (1.30)	4.00 (0.00)	4.00 (0.00)	4.40 (0.55)	4.00 (0.00)	4.04
Crowdsourcing	13	3.31 (0.75)	4.17 (0.58)	2.92 (1.19)	4.08 (1.19)	4.23 (0.60)	3.74
Source Labeling	9	3.67 (0.87)	3.67 (0.87)	2.33 (1.00)	3.56 (1.24)	3.11 (1.17)	3.27
Fact-check Labels	15	3.87 (0.74)	3.80 (0.94)	2.21 (1.12)	2.79 (1.12)	3.53 (1.06)	3.24
Click-through Warnings	12	3.75 (0.97)	3.17 (1.03)	2.42 (1.08)	3.17 (1.40)	3.42 (1.08)	3.18

Table 6.9: Expert scores for Content Labeling interventions. Mean values are shown with their corresponding standard deviations in parentheses.

## Expert Survey Results

### Discussion

#### 6.4.6 User-based Measures

##### Descriptive Characteristics

Intervention	Implementer	Targeted Aspect	Information Changes
Reporting	Platforms, Users	Spread	Adds, Alters
Social Norms	Platforms, Users	Spread	Adds, Alters
Alter Platform Metrics	Platforms	Creation	Alters

Table 6.10: Descriptive characteristics of User-based interventions.

## Expert Survey Results

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
Alter Platform Metrics	12	3.92 (1.00)	3.75 (1.36)	2.50 (1.09)	3.00 (1.41)	4.08 (0.79)	3.45
Reporting	9	3.78 (0.83)	3.89 (0.93)	2.22 (0.97)	3.00 (1.58)	4.33 (0.50)	3.44
Social Norms	13	3.62 (0.96)	3.83 (0.94)	2.67 (1.37)	3.33 (1.23)	3.62 (0.77)	3.41

Table 6.11: Expert scores for User-based Measures. Mean values are shown with their corresponding standard deviations in parentheses.

<b>Intervention</b>	<b>Implementer</b>	<b>Targeted Aspect</b>	<b>Information Changes</b>
Digital Media Literacy	Platforms, Govts, Institutions	Belief	Adds
Inoculation	Platforms, Govts, Institutions	Belief	Adds

Table 6.12: Descriptive characteristics of Media Literacy interventions.

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
Digital Media Literacy	9	3.78 (1.20)	4.11 (0.93)	2.00 (1.12)	2.00 (1.12)	3.78 (0.83)	3.13
Inoculation	10	3.40 (0.84)	3.90 (0.88)	2.10 (0.99)	2.50 (0.97)	3.60 (0.84)	3.10

Table 6.13: Expert scores for Media Literacy and Education interventions. Mean values are shown with their corresponding standard deviations in parentheses.

## Discussion

### 6.4.7 Media Literacy and Education

#### Descriptive Characteristics

#### Expert Survey Results

#### Discussion

### 6.4.8 Institutional Measures

#### Descriptive Characteristics

<b>Intervention</b>	<b>Implementer</b>	<b>Targeted Aspect</b>	<b>Information Changes</b>
Media Support	Platforms, Govts, Institutions	Prevention	-
Journalism Support	Platforms, Govts, Institutions	Prevention	-
Data Sharing	Platforms, Govts, Institutions	-	Adds
Government Regulation	Govts	Creation, Spread	Alters, Removes
Privacy Legislation	Govts	-	Alters, Removes
Anti-Trust Action	Govts	-	-
Taxes / Fines	Govts	-	Alters, Removes
Targeted Advertising	Platforms, Govts	-	Alters, Removes

Table 6.14: Descriptive characteristics of Institutional Measures interventions.

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
Targeted Ads	12	3.42 (0.67)	3.83 (0.83)	3.18 (1.25)	2.91 (1.30)	3.42 (1.24)	3.35
Data Sharing	10	4.10 (0.88)	4.10 (0.74)	2.10 (1.45)	2.80 (1.03)	3.10 (1.45)	3.24
Journalism Support	13	4.00 (0.71)	3.77 (1.01)	1.85 (0.80)	1.38 (0.51)	3.46 (1.20)	2.89
Media Support	10	3.40 (1.07)	3.78 (0.67)	2.00 (0.94)	1.80 (0.92)	3.30 (1.34)	2.86
Taxes / Fines	12	3.00 (1.28)	3.83 (1.34)	2.00 (1.26)	2.64 (1.57)	2.17 (1.03)	2.73
Government Regulation	11	3.64 (1.12)	3.55 (0.82)	1.55 (1.04)	1.82 (0.60)	2.91 (1.14)	2.69
Anti-Trust Action	9	3.67 (1.12)	4.11 (0.60)	1.78 (1.09)	1.22 (0.44)	2.33 (1.41)	2.62
Privacy Legislation	12	3.25 (1.36)	3.80 (0.92)	1.09 (0.30)	1.36 (0.50)	2.75 (1.48)	2.45

Table 6.15: Expert scores for Institutional Measures. Mean values are shown with their corresponding standard deviations in parentheses.

## Expert Survey Results

### Discussion

#### 6.4.9 Generative AI

##### Descriptive Characteristics

Intervention	Implementer	Targeted Aspect	Information Changes
Gen AI Chatbots	Platforms, Institutions	Belief	-
Gen AI Content	Platforms, Institutions	Prevention, Belief	-
Deepfakes	Platforms, Govts	Creation	Removes
AI in Ads	Platforms, Govts	Creation	Removes
AI Disclosure	Platforms, Govts	Belief	Tags

Table 6.16: Descriptive characteristics of Generative AI interventions.

## Expert Survey Results

Intervention	N	Effectiveness	Acceptance	Effort	Cost	Feasibility	Average
AI Disclosure	14	3.71 (0.99)	4.07 (0.83)	3.31 (1.18)	3.31 (1.25)	3.71 (1.59)	3.62
Deepfakes	11	2.55 (1.29)	4.10 (0.99)	2.40 (1.17)	3.00 (1.25)	4.00 (0.89)	3.21
AI in Ads	17	3.00 (1.00)	3.88 (0.99)	2.88 (1.17)	2.71 (1.26)	3.53 (1.33)	3.20
Gen AI Content	12	3.00 (0.85)	3.09 (0.83)	2.73 (0.79)	2.64 (0.81)	3.33 (0.65)	2.96
Gen AI Chatbots	16	3.19 (0.75)	2.75 (0.93)	2.75 (1.00)	2.69 (1.20)	3.35 (0.93)	2.95

Table 6.17: Expert scores for Generative AI interventions. Mean values are shown with their corresponding standard deviations in parentheses.

## Discussion

### 6.4.10 Summary

Top interventions by overall rank

Overall	Intervention	Effectiveness	Acceptance	Effort	Cost	Feasibility
1	Government Labels	8	8	2	2	6
2	Limit Forwarding	17	25	3	3	4
3	Limit Resharing	12	37	1	1	8
4	Crowdsourcing	31	2	11	4	2
5	AI Disclosure	16	7	8	13	11
6	Friction	25	27	7	6	10
7	User Control	35	1	27	11	3
8	Alter Platform Metrics	5	19	20	18	5
9	Reporting	10	10	29	19	1
10	Demonetization	3	20	12	24	16

Table 6.18: Ranking of the top 10 overall interventions by the five metrics.

# **Chapter 7**

## **Conclusion**

### **7.1 Summary**

### **7.2 Contributions**

### **7.3 Limitations**

### **7.4 Future Work**



# Appendix A

## Countermeasures Categorization

This appendix describes the categorization of specific misinformation interventions developed in Chapter 1, which was used to label the literature in the citation network analysis project described in Chapter 2.

### A.1 Content Distribution

Content distribution refers to a broad category of interventions concerning how content is distributed on social media. This includes methods such as redirection, accuracy prompts, friction, platform alterations, and advertising policies.

- **Redirection** - Redirection is a form of content distribution where users are directed to alternative content (such as official resources) or no content at all when searching for something that could be problematic or harmful [36, 207]. This can be implemented by presenting users with related articles or official information from authoritative sources on the topic, such as the CDC or WHO, when looking for COVID-19 information [59, 243].
- **Accuracy Prompts** - Sometimes referred to as “nudging”, accuracy prompts are a type of content distribution designed to encourage individuals to consider accuracy before posting or sharing content [177]. These prompts can be implemented in various ways, such as by asking users to evaluate the accuracy of one or more headlines or to reflect on the importance of sharing accurate news before continuing on the platform [77].
- **Friction** - Friction encourages individuals to pause and reflect before engaging with content [22, 115]. This can be operationalized by temporarily preventing users from resharing content they have not opened, prompting them to consider reading the article through interstitials or pop-up windows [182].
- **Platform Alterations** - Any modifications to the design or architecture of social media platforms that influence how content is distributed or displayed to users or how they are encouraged to engage with the platform [123]. Examples include reducing the size or visibility of a post [83, 123, 205] or altering platform architecture to guide users to move posts to specific rather than general groups [120].
- **Advertising Policy** - Advertising policy refers to how platforms (or governments) regulate,

correct, or display advertisements to the public [26, 59, 100]. Examples include banning political ads, requiring ads to undergo a fact-checking process before posting, prohibiting certain fake news websites from advertising, or labeling ads as “paid for” [5, 53, 59].

- **Other Content Distribution** - Other types of content distribution involve limiting users’ forwarding capabilities, which caps the number of recipients to whom a given message can be forwarded [181], and limiting resharing, which restricts rampant resharing by removing share buttons on posts after several levels of sharing [3].

## A.2 Content Moderation

Content moderation refers to a broad range of interventions regarding how content is displayed or not displayed on social media. This encompasses fact-checking, narrative counterspeech like debunking, and the use of algorithms to assist in moderation and misinformation detection [108, 200].

- **Fact-checking** - The process of verifying information, typically performed by experts. This verification can be done by experts, journalists, and platforms, and includes multi-modal fact-checking, such as fact-checking videos [39, 230].
- **Debunking** - Debunking is a stronger form of fact-checking, where context and coherence are typically provided in addition to verifying or correcting content. It can also be described as a “narrative intervention” [51, 144].
- **Algorithmic Content Moderation** - This refers to automated content moderation, including automated fact-checking, labeling, or removing posts containing misinformation or other violating content. It can also involve downranking or de-emphasizing low-quality content, as well as upranking or promoting high-quality or authoritative news sources [26, 37, 83, 88]. Additionally, it can be employed as a virality circuit breaker, which automatically flags certain fast-spreading and unverified content, temporarily halting algorithmic amplification until the information is verified [3].
- **Misinformation Detection** - The algorithmic detection of misinformation, typically for content moderation purposes [117, 146].
- **Other Content Moderation** - Other forms of content moderation could include user control, which would involve transferring some moderation responsibilities currently done by platforms to users. This approach would give users greater control over the content displayed in their own news feeds [107].

## A.3 Account Moderation

Account moderation involves moderating user accounts by implementing actions such as account suspensions, removals, shadowbanning users, or demonetizing accounts [59, 238].

- **Account Removal** - Account removal refers to the permanent or temporary banning of users who share misinformation or violate other platform policies a certain number of times. A specific type of account removal, where platforms coordinate their efforts to

remove particularly problematic or dangerous user accounts, is typically referred to as deplatforming [186, 216].

- **Shadow Banning** - The practice of limiting the reach of posts from certain policy-violating accounts without explicitly banning or suspending them, typically conducted in a concealed or opaque manner [109, 238].
- **Other Account Moderation** - Another form of account moderation is the demonetization of user accounts [150]. This process refers to removing or restricting monetization features for a user account that is found to violate a platform's policies repeatedly.

## A.4 Content Labeling

Content labeling includes all general types of misinformation disclosure through labeling. Labels are commonly used to present fact-checks, source information or credibility, or to provide additional context on a post [156, 243].

- **Crowdsourcing** - Crowdsourcing generally involves utilizing regular individuals to verify information and label content instead of relying on journalists or expert fact-checkers [7]. The most recognized operational version of a crowdsourcing intervention is X's Community Notes program [8].
- **Warning Labels** - Warning labels refer to general warnings about misinformation and can address the source, content, or context [170]. One way to implement this type of intervention is by using click-through warning labels or interstitials, which both warn users and encourage them to reflect before viewing the content [205].
- **Source Credibility Labels** - This type of intervention involves disclosing or labeling the credibility of a post's source. This intervention can be implemented by labeling the reliability of a news source[79] or by labeling the accounts of government officials or state-run media sources [161].
- **Context Labels** - Context labels are labels that specifically add context or additional information to a post, such as Community Notes programs [8].
- **Other Content Labeling** - A related type of content labeling involves specifically notifying users when they have posted or interacted with content verified to contain misinformation or originating from a state-run media source, through the use of misinformation disclosure [59].

## A.5 Media Literacy and Education

This general intervention category involves any educational or training effort aimed at enhancing the public's civic reasoning, digital literacy, and critical thinking skills when interacting with media messages [92, 106].

- **Fake News Games** - Games designed to help players detect misinformation and improve their critical thinking skills [147, 154, 194].

- **Inoculation** - Commonly referred to as “pre-bunking,” inoculation consists of warning messages or information about misleading rhetorical techniques meant to prevent people from later believing misinformation [142].
- **Other Media Literacy Efforts** - Other types of educational initiatives and relevant research studies, such as providing individuals with tips on recognizing fake news [92], evaluating people’s information, digital, or news literacy [111], or developing strategies to improve media literacy messaging [222].

## A.6 User-based Measures

User-based measures are interventions that focus on individuals’ responses to encountering misinformation [212].

- **Reporting** - Users can report other users or their posts [164, 248].
- **Blocking** - Users have the ability to block other users or specific topics [212].
- **Social Corrections** - Users who fact-check or debunk other users directly [21, 37]. This intervention may involve publicly commenting on a post or privately messaging the misinformation poster.
- **Social Norms** - The use of social or community influence to change behavior and promote social and self-corrections [75, 84]. This could be operationalized by adjusting platform metrics to reward accuracy instead of engagement, aiming to discourage the sharing of misinformation [19].
- **Retractions** - This category includes instances when users or organizations retract misinformation they have posted and how that impacts individuals who have already encountered the misinformation [12, 72].
- **Other User-based Measures** - Other forms of user-based measures might include encouraging users to deactivate their social media accounts [224], along with other user-driven behaviors.

## A.7 Institutional Measures

Institutional measures are those implemented by civic society, governments, the media, or other organizations [41].

- **Media Support** - Investing in local news or promoting trustworthy local news on social media platforms [217]. Supporting and training the next generation of journalists to engage in high-quality, independent reporting [25, 41].
- **Data Sharing** - Social media companies should regularly release data and internal research reports on the prevalence, spread, and mitigation of misinformation to the public and outside researchers [15, 26, 33].
- **Government Regulation** - This category includes any laws, rules, or regulations at local, state, or federal levels [164, 192, 242]. Regulation may involve holding companies

accountable for the content shared on their platforms, such as by modifying Section 230 or regulating them like utility or media companies. Legislation could involve developing comprehensive privacy laws, similar to Europe's GDPR, or taking anti-trust action by breaking up monopolistic technology companies [192]. Furthermore, companies could be taxed or fined for using personal user data, and micro-targeted advertising could be banned or restricted [183].

- **Other Institutional Measures** - Additional institutional measures may involve researching and developing tools to support civil society with these issues, as well as enhancing collaboration among various types of institutions [41, 235].

## A.8 Other

This category includes any interventions that do not fit into the previous categories or are newly introduced.

- **Generative AI** - Using generative AI to combat or detect misinformation. Examples include employing AI chatbots to reduce belief in conspiracy theories or misinformation [58] and using AI-generated content to create rebuttals against misinformation or to develop educational initiatives. Policy responses may involve prohibiting or labeling the use of AI or manipulated content to produce deepfakes, banning the use of AI in advertising, or requiring clear disclosure on any ads that incorporate AI-generated images, videos, or audio [99].
- **Combining Interventions** - Studies that explicitly compare the effects of using multiple interventions simultaneously with using one intervention [23].
- **Other** - Any other intervention not previously described.



# Appendix B

## Chapter 2: ChatGPT Prompts

This appendix includes the final ChatGPT prompts that were created to assist with the labeling of papers in Chapter 2.

### B.1 Effectiveness Labeling Prompt

#### **Effectiveness Labeling - Task Overview:**

You are an academic researcher tasked with labeling research papers (PDFs) focused on interventions aimed at countering misinformation. Your goal is to determine whether each paper studies the effectiveness of intervention(s) or action(s) in reducing misinformation.

#### **Instructions for Labeling:**

##### **Label as “Yes” (Studies Effectiveness):**

Mark papers as “Yes” if they empirically measure or analyze, either directly or indirectly, the effectiveness of an intervention in reducing the creation, belief, sharing, or spread of misinformation. This includes studies that assess effectiveness through:

- Direct measurement, such as experimental data or controlled trials, or
- Indirect analysis, such as observational studies or secondary analyses that provide insights into intervention impact
- Papers that find an intervention is not effective should still be labeled “Yes” for studying effectiveness

##### **Label as “No” (Does Not Study Effectiveness):**

Mark papers as “No” if they do not empirically evaluate the effectiveness of an intervention, directly or indirectly.

- This includes theory-driven papers focused on conceptual frameworks or hypotheses without testing outcomes, as well as discussion-based papers that explore ideas, potential strategies, or frameworks without empirical testing.
- Literature reviews that summarize existing research without conducting empirical analysis related to effectiveness also fall under this category.

## B.2 Acceptance Labeling Prompt

### Acceptance Labeling - Task Overview:

You are an academic researcher tasked with labeling research papers (PDFs) focused on interventions aimed at countering misinformation. Your goal is to determine whether each paper studies the public's or users' acceptance or support of intervention(s) or action(s) to counter misinformation.

### Instructions for Labeling:

#### Label as “Yes” (Studies Acceptance):

Mark papers as “Yes” if they explicitly explore users’ or the public’s opinions, perceptions, or acceptance of an intervention. Papers that qualify for this label often:

- Include surveys, interviews, or other methods of feedback collection on user preferences.
- Explore potential improvements to the intervention based on user feedback.

#### Label as “No” (Does Not Study Acceptance):

Mark papers as “No” if they do not examine user acceptance or support. Papers that should be labeled “No” often:

- Focus exclusively on the effectiveness, impact, technical performance, or societal impact of the intervention without assessing user feedback
- Some papers may include surveys or interviews but still do not qualify if these methods are used solely to analyze intervention performance, not user acceptance.

## B.3 Interventions Labeling Prompt

### Interventions Labeling - Task Overview:

You are an academic researcher tasked with labeling research papers (PDFs) focused on interventions aimed at countering misinformation. Your goal is to determine the primary intervention(s) or action(s) studied in each paper.

### Instructions for Labeling:

- Assign only the most relevant label(s) that best capture the central intervention(s) investigated in each paper.
- Do not label interventions that are only briefly mentioned. Only label interventions that are a key focus of the study.
- If a paper discusses multiple interventions, prioritize the most rigorously studied or most emphasized intervention(s).
- Use only the labels provided below and format them in a comma-separated list.
- If no label fits, use “other” instead of forcing an inappropriate label.

**Label Definitions:** Interventions are broken up by category for clarity.

- Content Distribution interventions relate to the methods by which content is distributed on social media.
  - redirection: directing users to alternative or official content, or showing related content.
  - accuracy prompts: reminding users to assess content accuracy before sharing.
  - friction: adding delays, prompts, or other mechanisms to slow down engagement and encourage thoughtful interaction.
  - platform alterations: structural changes affecting how content is distributed or displayed to users.
  - advertising policy: rules governing advertisements and their effects.
  - content distribution: unspecified content distribution intervention that is not described by one of the other labels.
- Content Moderation interventions determine what content is displayed or suppressed.
  - fact-checking: verifying information accuracy, often by expert fact-checkers or journalists.
  - debunking: providing context or narrative coherence to correct or refute misinformation.
  - misinformation detection: using algorithmic techniques to identify misinformation.
  - algorithmic content moderation: the use of automated moderation techniques.
  - content moderation: unspecified content moderation intervention that is not described by one of the other labels.
- Account Moderation interventions are policies and actions to moderate user accounts.
  - account removal: banning users or coordinated deplatforming.
  - shadow banning: restricting account or post visibility without an official ban.
  - demonetization: restricting or removing monetization options from a user or content.
  - account moderation: unspecified account moderation intervention that is not described by one of the other labels.
- Content Labeling interventions are labels that disclose or provide context for misinformation.
  - crowdsourcing: user-generated contributions that help label or assess misinformation.
  - warning labels: explicit labels warning users about misinformation.
  - source credibility labels: labels that indicate an information source's reliability.
  - context labels: adding background or contextual information to a post.
  - content labeling: unspecified content labeling intervention that is not described by one of the other labels.
- User-based interventions are actions users can take to respond to misinformation.
  - reporting: enabling users to flag misinformation for review.

- blocking: user action that prevents engagement with specific accounts.
- social norms: encouraging platform-specific social or community-based norms that discourage misinformation.
- social corrections: direct corrections provided by other users in response to misinformation.
- retractions: users voluntarily retracting or correcting their own misinformation.
- user-based measures: unspecified user-based intervention that is not described by one of the other labels.
- Media Literacy interventions are educational efforts for critical media engagement.
  - fake news games: interactive games designed to train users to recognize misinformation
  - inoculation: pre-bunking techniques that prepare users to recognize and resist misinformation.
  - media literacy: unspecified media literacy intervention that is not described by one of the other labels.
- Institutional interventions are interventions by organizations or authorities.
  - media support: supporting or promoting reliable, high-quality journalism and local news sources.
  - data sharing: enabling ethical data access for misinformation research.
  - government regulation: laws or policies designed to regulate misinformation or social media platforms.
  - institutional measures: unspecified institutional intervention that is not described by one of the other labels.
- Other interventions are miscellaneous or multiple combined interventions.
  - generative AI: AI usage to counter misinformation.
  - combining interventions: analyzing the use of multiple intervention strategies at once instead of one at a time.
  - other: interventions that do not fit any of the above categories.

# Appendix C

## Chapter 3: Pre-Registered Hypotheses and Analyses

The design table summarizes the research questions, hypotheses, power calculations, planned analyses, and possible interpretations for the closeness analysis in Chapter 2.

Question	Details
<b>RQ1: How do people respond and think others should respond when they see misinformation?</b>	<p><b>H1.1:</b> People believe individuals should expend more effort to respond to misinformation online than they actually do.</p> <p><b>Sampling Plan:</b> The necessary sample size for a one-sided paired Bayesian test to achieve 95% power to detect a medium effect size of 0.5 at a Bayes threshold of 10 is estimated to be 81. The estimated number of participants who have seen misinformation per closeness level is approximately 192.</p> <p><b>Analysis Plan:</b></p> <ul style="list-style-type: none"><li>• Calculate Measures 1a and 1b at each of the three closeness levels.</li><li>• Run three one-sided Bayesian paired hypothesis tests (one at each closeness level).</li><li>• Calculate the 95% highest density interval (HDI) for the effect size.</li><li>• Visualize results with bar charts or similar visualizations comparing possible actions.</li></ul> <p><b>Interpretation:</b> The resulting Bayes factor will inform the probability of the alternate or null hypothesis being true. For example, a Bayes factor of 10 implies the data is 10 times more likely under H1 than H0.</p>

Question	Details
<b>RQ1: How do people respond and think others should respond when they see misinformation?</b>	<p><b>H1.2:</b> People respond with more effort when the sender of misinformation is a close contact than a somewhat close contact and a somewhat close contact than a not close contact</p> <p><b>Sampling Plan:</b> The necessary sample size for a one-sided paired Bayesian test to achieve 95% power to detect a medium effect size of 0.5 at a Bayes threshold of 10 is estimated to be 81. The estimated sample size is approximately 115.</p> <p><b>Analysis Plan:</b></p> <ul style="list-style-type: none"> <li>• Calculate Measures 1a at each of the three closeness levels.</li> <li>• Run a one-sided Bayesian paired hypothesis test comparing the effort level (Measure 1a) for close contacts and somewhat close contacts. Similarly, run a one-sided Bayesian paired hypothesis test comparing the effort level (Measure 1a) for somewhat close contacts and not close contacts.</li> <li>• For both tests, calculate the 95% highest density interval (HDI) for the effect size. Visualize the results.</li> </ul> <p><b>Interpretation:</b> The resulting Bayes factor will inform the probability of the alternate or null hypothesis being true. Visualizations will help the reader interpret the results.</p>
<b>RQ1: How do people respond and think others should respond when they see misinformation?</b>	<p><b>H1.3:</b> People <i>believe</i> others should respond with more effort when the sender of misinformation is a close contact than a somewhat close contact and a somewhat close contact than a not close contact.</p> <p><b>Sampling Plan:</b> The necessary sample size for a one-sided paired Bayesian test to achieve 95% power to detect a medium effect size of 0.5 at a Bayes threshold of 10 is estimated to be 81. All participants will be included in this test, indicating a sample size of approx. n = 1000.</p> <p><b>Analysis Plan:</b></p> <ul style="list-style-type: none"> <li>• Calculate Measures 1b at each of the three closeness levels.</li> <li>• Run a one-sided Bayesian paired hypothesis test comparing the effort level (Measure 1b) for close contacts and somewhat close contacts. Similarly, run a one-sided Bayesian paired hypothesis test comparing the effort level (Measure 1b) for somewhat close contacts and not close contacts.</li> <li>• For both tests, calculate the 95% highest density interval (HDI) for the effect size. Visualize the results.</li> </ul> <p><b>Interpretation:</b> The resulting Bayes factor will inform the probability of the alternate or null hypothesis being true. Visualizations will help the reader interpret the results.</p>

Question	Details
<b>RQ2: How do people behave after realizing they have posted misinformation?</b>	<p><b>H2.1:</b> People believe others should expend more effort to respond to misinformation online after realizing they posted misinformation than they actually do.</p> <p><b>Sampling Plan:</b> The necessary sample size for a one-sided paired Bayesian test to achieve 95% power to detect a medium effect size of 0.5 at a Bayes threshold of 10 is estimated to be 81. The estimated number of participants who have posted misinformation accidentally is <math>n = 320</math>.</p> <p><b>Analysis Plan:</b></p> <ul style="list-style-type: none"> <li>• Calculate Measures 2a and 2b to determine response effort.</li> <li>• Run a one-sided Bayesian paired test with these two measures.</li> <li>• Calculate the 95% HDI for the effect size and visualize the results.</li> </ul> <p><b>Interpretation:</b> The resulting Bayes factor will inform the probability of the alternate or null hypothesis being true. The visualizations will help readers understand the effort differences.</p>
<b>RQ3: How do people's responses and beliefs about how others should respond after seeing misinformation differ from their responses and beliefs when they realize they have posted misinformation?</b>	<p><b>H3.1:</b> People respond with a different level of effort when the sender of misinformation is someone else compared to themselves.</p> <p><b>Sampling Plan:</b> The necessary sample size for a two-sided paired Bayesian test to achieve 95% power to detect a medium effect size of 0.5 at a Bayes threshold of 10 is estimated to be 92. The estimated number of participants is 150.</p> <p><b>Analysis Plan:</b></p> <ul style="list-style-type: none"> <li>• Calculate Measures 1a (for each closeness level) and 2a.</li> <li>• Run a two-sided paired Bayesian test comparing these two measures.</li> <li>• The null hypothesis is these two measures are equal. The null interval is -0.2 to 0.2 (effect size is effectively 0). The alternate hypothesis is that the absolute value of the effect size is at least 0.2.</li> <li>• Calculate the 95% HDI and visualize the results.</li> </ul> <p><b>Interpretation:</b> The resulting Bayes factor will inform the probability of the alternate or null hypothesis being true. If the 95% highest density interval falls entirely within the “region of practical equivalence” (ROPE) range, of -0.2 to 0.2, then we can accept the null. If it falls entirely outside the range, then we can accept the alternate. Otherwise, using this method we will state inconclusive results. The visualizations will help readers understand the effort differences.</p>

Question	Details
<b>RQ3: How do people's responses and beliefs about how others should respond after seeing misinformation differ from their responses and beliefs when they realize they have posted misinformation?</b>	<p><b>H3.2:</b> People want others to respond with a different level of effort when the sender of misinformation is someone else compared to themselves.</p> <p><b>Sampling Plan:</b> The necessary sample size for a two-sided paired Bayesian test to achieve 95% power to detect a medium effect size of 0.5 at a Bayes threshold of 10 is estimated to be 92. All participants will be included in this test, indicating a sample size of approx. <math>n = 1000</math>.</p> <p><b>Analysis Plan:</b></p> <ul style="list-style-type: none"> <li>• Calculate Measures 1b (for each closeness level) and 2b.</li> <li>• Run a two-sided paired Bayesian test comparing these two measures.</li> <li>• The null hypothesis is these two measures are equal. The null interval is -0.2 to 0.2 (effect size is effectively 0). The alternate hypothesis is that the absolute value of the effect size is at least 0.2.</li> <li>• Calculate the 95% HDI and visualize the results.</li> </ul> <p><b>Interpretation:</b> The resulting Bayes factor will inform the probability of the alternate or null hypothesis being true. If the 95% highest density interval falls entirely within the “region of practical equivalence” (ROPE) range, of -0.2 to 0.2, then we can accept the null. If it falls entirely outside the range, then we can accept the alternate. Otherwise, using this method we will state inconclusive results. The visualizations will help readers understand the effort differences.</p>

# Appendix D

## Chapter 3: Supplemental Categorical Analysis

Categorical analysis was conducted as a robustness check to better understand the association between the three categorical variables in the model: **effort level** when countering, **response type** (reported behavior vs. opinion), and **closeness level**. Only the beliefs of participants who saw misinformation at each closeness level are included to ensure this analysis is analogous to those of our pre-registered hypothesis tests. Considering the different actions available for countering oneself (Table 3.2) versus others (Table 3.1), we first only consider responses from others. Table D.1 shows the three-way contingency table, which underpins Figure 3.2 in Chapter 2. The generalized McNemar's chi-squared tests for categorical paired data were run to assess the relationship between effort level and response type at each closeness level, and all yielded statistically significant results. These findings are consistent with our results from H1.1.

Closeness Level	Effort Level	Response Type		McNemar's Chi-sq. Test
		Behavior	Opinion	
Close contacts (n = 148)	No Effort	52	20	$\chi^2 = 34.9^{***}$ $p = 1e-7$
	Low Effort	19	13	
	High Effort	77	115	
Somewhat close contacts (n = 370)	No Effort	179	76	$\chi^2 = 126.3^{****}$ $p < 2e-16$
	Low Effort	76	41	
	High Effort	115	253	
Not close contacts (n = 880)	No Effort	462	248	$\chi^2 = 201.5^{***}$ $p < 2e-16$
	Low Effort	217	251	
	High Effort	201	381	

Table D.1: Contingency table of effort level, response type, and closeness level. Each cell contains the frequency of occurrences for each combination of the categorical variables. The chi-sq test statistic is shown:  $p < 0.05^*$ ,  $p < 0.01^{**}$ , and  $p < 0.001^{***}$ .

Since the contingency table includes some repeated responses from the same participants, we analyzed the data using generalized estimating equations (GEE) with the “geepack” R package

[104] to estimate the average population-level effects of the categorical variables. Table D.2 shows the two best model results from the analysis based on QIC values, which are modified AIC values applicable to GEE models [169]. We compared three models to predict the count frequencies in the contingency table: the model with main effects only, the model with main effects and two-way interaction terms, and the saturated model with the three-way interaction term. According to the QIC values, the saturated model with the three-way interaction term was the best, even though the three-way interaction term was not statistically significant. Several two-way interaction terms are significant. There is a higher frequency of “high effort” responses when contacts are closer or when people are asked about their opinions instead of their actual behavior, echoing our results from RQ1.

We conducted a similar analysis of the responses against oneself. Table D.3 shows the two-way contingency table and the generalized McNemar’s chi-squared test result, which was statistically significant. This finding is consistent with our result from H2.1.

Table D.4 shows the model to predict the frequency of counts in the “self” contingency table. Again, effort level interacts with response type, with higher effort actions seeing higher counts when people are asked about their opinion rather than their actual behavior, echoing our results from H2.1. We also see higher counts for low and high-effort actions compared to no-effort actions. The estimates in this model for low and high effort counts are also higher than the comparable estimates in Table D.2, indicating that individuals may be exerting more effort to correct themselves than others.

Table D.2: The generalized estimating equation (GEE) model predicting the average frequency counts as a function of effort level (reference level: no effort), response type (reference level: behavior), and closeness (reference level: close contacts) when considering responses against others.

	<i>Dependent variable: Frequency counts</i>			
	Model with Three-Way Terms		Model with Two-Way Terms	
	Estimate	Std. Err.	Estimate	Std. Err.
Low Effort	-1.007***	(0.264)	-1.038***	(0.213)
High Effort	0.393*	(0.173)	0.380**	(0.143)
Opinion	-0.956***	(0.259)	-0.994***	(0.132)
Not Close	2.184***	(0.139)	2.179***	(0.123)
Somewhat Close	1.236***	(0.151)	1.210***	(0.134)
Low Effort : Opinion	0.576	(0.441)	0.664***	(0.091)
High Effort : Opinion	1.357***	(0.294)	1.398***	(0.082)
Low Effort : Not Close	0.251	(0.273)	0.339	(0.216)
High Effort : Not Close	-1.225***	(0.187)	-1.299***	(0.149)
Low Effort : Somewhat	0.150	(0.294)	0.003	(0.237)
High Effort : Somewhat	-0.835***	(0.205)	-0.676***	(0.160)
Opinion : Not Close	0.333	(0.267)	0.358**	(0.128)
Opinion : Somewhat	0.099	(0.289)	0.188	(0.137)
Low Effort : Opinion : Not Close	0.192	(0.453)		
High Effort : Opinion : Not Close	-0.095	(0.310)		
Low Effort : Opinion : Somewhat	-0.337	(0.496)		
High Effort : Opinion : Somewhat	0.288	(0.337)		
Constant	-2.882***	(0.135)	-2.871***	(0.120)

Note:

\*p<0.05; \*\*p<0.01; \*\*\*p<0.001

	<b>Effort Level</b>	<b>Response Type</b>		<b>McNemar's Chi-sq. Test</b>
		Behavior	Opinion	
Self (n = 256)	No Effort	13	5	$\chi^2 = 63.2^{***}$
	Low Effort	114	52	$p = 1e-13$
	High Effort	129	199	

Table D.3: Contingency table of effort level and response type for actions against oneself. Each cell contains the frequency of occurrences for each combination of the categorical variables. The chi-sq test statistic is shown:  $p < 0.05^*$ ,  $p < 0.01^{**}$ , and  $p < 0.001^{***}$ .

Table D.4: The generalized estimating equation (GEE) model predicting the average frequency counts as a function of effort level (reference level: no effort) and response type (reference level: behavior) when considering responses to oneself.

<i>Dependent variable: Frequency counts</i>		
Model with Two-Way Terms		
	Estimate	Std. Err.
Low Effort	2.171***	(0.279)
High Effort	2.295*	(0.277)
Opinion	-0.956	(0.519)
Low Effort : Opinion	0.171	(0.538)
High Effort : Opinion	1.389**	(0.524)
Constant	-2.980***	(0.270)

*Note:*

\* $p < 0.05$ ; \*\* $p < 0.01$ ; \*\*\* $p < 0.001$

## Appendix E

### Chapter 3: Demographics Summary Table

Table E.1 shows the percentage of each demographic category that believes one should respond with a maximum of no, low, or high effort when encountering misinformation posted by others or oneself. Any category with a sample size of fewer than 50 participants was combined with the next closest demographic category, except the “Other” category for gender. The chi-square test of independence was conducted for each demographic category and each closeness level. This test is used to determine whether the beliefs of the demographic groups are independent of one another.

Table E.1: Demographics summary table.

Category	Sample Size (n)	% Effort (No / Low / High)				Chi-sq Test of Independence Results
		Close contacts	Somewhat close	Not close	Self	
Overall	n = 1010	14.5 / 5.2 / 80.3	18.2 / 8.1 / 73.7	29.2 / 27.0 / 42.8	2.5 / 19.2 / 78.3	p < 0.05*, p < 0.01**, p < 0.001***
Age	18-34 (253)	21.3 / 7.5 / 71.1	24.1 / 9.5 / 66.4	32.4 / 29.6 / 37.9	4.7 / 23.7 / 71.5	<b>Test statistic:</b> Close: 25.5** Somewhat: 26.4*** Not Close: 10.9 Self: 20.7**
	35-44 (338)	13.3 / 5.9 / 80.8	18.0 / 10.9 / 71.0	24.0 / 31.4 / 44.7	2.7 / 20.7 / 76.6	
	45-54 (186)	11.8 / 3.8 / 84.4	15.1 / 4.8 / 80.1	30.1 / 25.8 / 44.1	1.1 / 17.7 / 81.2	
	55-64 (148)	14.2 / 2.7 / 83.1	18.9 / 5.4 / 75.7	33.1 / 23.6 / 43.2	1.0 / 12.2 / 87.2	
	65+ (85)	4.7 / 3.5 / 91.8	7.1 / 4.7 / 88.2	31.8 / 22.4 / 45.9	1.2 / 15.3 / 83.5	

Category	Sample Size (n)	% Effort (No / Low / High)				Chi-sq Test of Independence Results
		Close contacts	Somewhat close	Not close	Self	
Gender	Female (465)	13.3 / 4.3 / 82.4	17.0 / 8.82 / 74.2	28.0 / 26.9 / 45.2	1.7 / 18.3 / 80	<b>Test statistic:</b> Close: 3.02 Somewhat: 1.82 Not Close: 2.38, Self: 3.03
	Male (520)	16.0 / 6.0 / 78.1	20.0 / 7.5 / 72.5	31.2 / 28.5 / 40.4	3.3 / 20 / 76.7	
	Other (25) (excl. from stat. tests)	4.0 / 8.0 / 88.0	4.0 / 8.0 / 88.0	12.0 / 40.0 / 48.0	0 / 20 / 80	
Race	White/ Caucasian (799)	14.5 / 5.6 / 79.8	18.5 / 7.9 / 73.6	30.3 / 26.2 / 43.6	2.4 / 19.9 / 77.7	<b>Test statistic:</b> Close: 8.88 Somewhat: 2.67 Not Close: 12.64* Self: 2.08
	Black or African American (85)	9.4 / 4.7 / 85.9	15.3 / 10.6 / 74.1	21.2 / 31.8 / 47.1	2.4 / 17.6 / 80.0	
	Asian (65)	10.8 / 3.1 / 86.2	15.4 / 6.2 / 78.5	21.5 / 41.5 / 36.9	3.1 / 18.5 / 78.5	
	Multiracial/ Other (61)	24.6 / 3.3 / 72.1	21.3 / 9.8 / 68.9	34.4 / 32.8 / 32.8	3.3 / 13.1 / 83.6	
Education	High school or less (116)	17.2 / 1.7 / 81.0	22.4 / 4.3 / 73.3	32.8 / 25.9 / 41.4	3.5 / 21.6 / 75.0	<b>Test statistic:</b> Close: 17.2* Somewhat: 9.9 Not Close: 18.5* Self: 6.6
	Some college (164)	8.5 / 3.1 / 88.4	15.9 / 6.1 / 78.0	20.7 / 22.6 / 56.7	1.2 / 15.2 / 83.5	
	Associate's degree (126)	11.9 / 4.0 / 84.1	15.9 / 5.6 / 78.6	27.8 / 27.0 / 45.2	1.6 / 23.0 / 75.4	
	Bachelor's degree (438)	16.2 / 6.2 / 77.6	18.7 / 9.6 / 71.7	31.1 / 30.6 / 38.4	3.0 / 19.9 / 77.2	
	Master's degree or higher (166)	15.7 / 8.4 / 75.9	18.1 / 10.8 / 71.1	31.3 / 28.9 / 39.8	2.4 / 16.9 / 80.7	
Income	Less than \$20,000 (87)	10.3 / 4.6 / 85.1	11.5 / 5.8 / 82.8	21.8 / 28.7 / 49.4	0 / 18.4 / 81.6	<b>Test statistic:</b> Close: 10.2 Somewhat: 25.1* Not Close: 22.4* Self: 20.2
	\$20,000 - \$39,999 (172)	11.0 / 4.7 / 84.3	12.8 / 6.4 / 80.8	23.3 / 31.4 / 45.3	2.9 / 21.5 / 75.6	
	\$40,000 - \$59,999 (220)	15.9 / 3.6 / 80.5	19.1 / 5.5 / 75.5	28.6 / 22.7 / 48.6	3.6 / 18.6 / 77.7	
	\$60,000 - \$79,999 (181)	13.8 / 5.5 / 80.7	17.7 / 9.9 / 72.4	26.5 / 28.7 / 44.8	2.2 / 11.6 / 86.2	
	\$80,000 - \$99,999 (140)	14.3 / 6.4 / 79.3	23.6 / 11.4 / 65.0	37.1 / 27.1 / 35.7	2.9 / 18.6 / 78.6	
	\$100,000 - \$149,999 (131)	16.0 / 7.6 / 76.3	16.8 / 11.5 / 71.8	33.6 / 33.6 / 32.8	3.1 / 22.1 / 74.8	
	Over \$150,000 (79)	21.5 / 5.1 / 73.4	29.1 / 6.3 / 64.6	36.7 / 25.3 / 38.0	0 / 30.4 / 69.6	

Category	Sample Size (n)	% Effort (No / Low / High)				Chi-sq Test of Independence Results
		Close contacts	Somewhat close	Not close	Self	
Party	Strong Republican (130)	20.8 / 6.9 / 72.3	24.6 / 5.4 / 70.0	33.8 / 23.8 / 42.3	5.4 / 23.1 / 71.5	<b>Test statistic:</b> Close: 25.4** Somewhat: 14.1 Not Close: 22.3** Self: 11.5
	Weak Republican (106)	14.2 / 1.9 / 84.0	19.8 / 4.7 / 75.5	35.8 / 26.4 / 37.7	1.0 / 25.5 / 73.6	
	Independent/ Other (303)	15.8 / 2.0 / 82.2	20.1 / 7.3 / 72.6	33.7 / 24.1 / 42.2	2.0 / 17.2 / 80.9	
	Weak Democrat (203)	12.8 / 5.4 / 81.8	15.8 / 8.4 / 75.9	29.6 / 30.5 / 39.9	2.5 / 18.7 / 78.8	
	Strong Democrat (268)	11.2 / 9.3 / 79.5	14.2 / 11.6 / 74.3	19.0 / 33.2 / 47.8	2.2 / 17.5 / 80.2	



# Appendix F

## Chapter 4: OMEN Survey Posts

This appendix includes all the survey posts used in the survey experiment in Chapter 3.

### F.1 Misinformation Detection Posts

These posts were used in the misinformation detection section of the survey experiment. Posts labeled as “misinformation” indicate that they are considered false or partially false and misleading by mainstream fact-checking organizations. Posts labeled as “real news” are from reputable sources or accounts posting reputable links and are considered true.



Anonymous836  
@Anonymous836

Just saw on Channel 4 that Facebook will start charging all users starting in April 2024. I can't believe there doing this. You can opt-out if you hold your fingers over and copy this – don't share it they will know:

I don't give permission to Facebook to charge me \$4.99 a month.

11:33AM January 29th, 2024

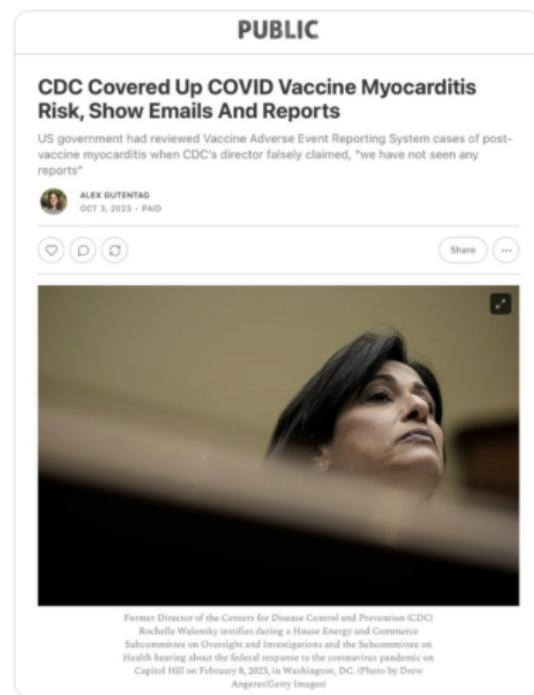


Figure F.1: Pre-Test Misinformation Post #1.



Anonymous5512  
@Anonymous5512

The CDC believes that giving the latest COVID19 booster to kids and young people is safer than not getting it, but it is not. Now we know that the CDC lied about the dangers of myocarditis.  
Groundbreaking new report from @Anonymous247.



1:25PM October 3rd, 2023



Figure F.2: Pre-Test Misinformation Post #2.



Amazon Wholesale Liquidation Pallet  
@AmazonPallet

Every day, many packages get lost in warehouses. Normally, Amazon throws them away, but right now they have a special offer. You can get one of these lost packages for just \$1! The packages often include items such as iPhones, appliances, and more. Just message me and I'll send you the form to fill in to get a package. But remember, there are only a limited number of packages available, so act fast!



10:12AM August 17th, 2023

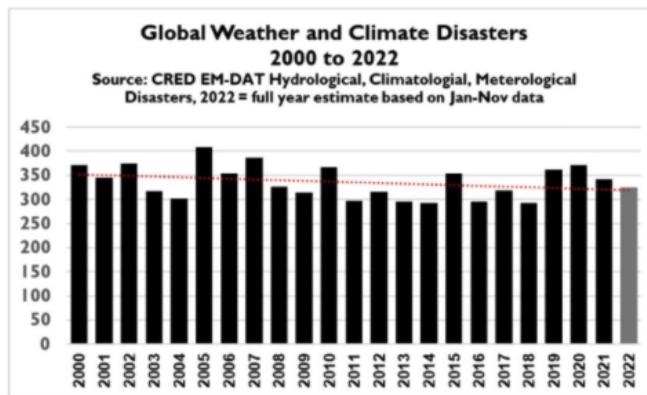


Figure F.3: Pre-Test Misinformation Post #3.



Anonymous179  
@Anonymous179

Remember, weather and climate are not the same thing. Not every storm, flood, wildfire, or blizzard means climate change is happening. There is a [#ClimateScam](#). There isn't a [#ClimateEmergency](#).



2:55PM February 9th, 2023



Figure F.4: Pre-Test Misinformation Post #4.



NPR ✅  
@NPR

The EPA is stepping in to take control of the cleanup of the train derailment in Ohio earlier this month.

Under the legal order, Norfolk Southern is responsible for the costs and cleanup of contaminated soil and water.



2:00PM November 17th, 2022

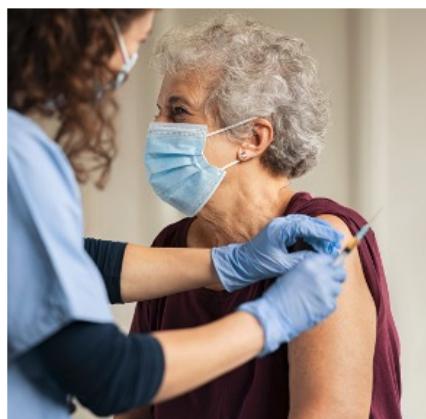


Figure F.5: Pre-Test Real News Post #1.



Anonymous3131  
@Anonymous3131

I got my **#COVID19** booster today! **#Covid** is still here and it's dangerous for many people. Please get vaccinated to help save lives.



5:39PM October 29th, 2023



Figure F.6: Pre-Test Real News Post #2.



Anonymous99  
@Anonymous99

Encouraging developments in the battle against climate change: experts indicate that the swift expansion of global wind and solar energy aligns closely with the necessary pace to meet the 2030 target within the 1.5°C pathway. However, here in Alberta, the UCP appears to be heading in the opposite direction. 😞



6:03 PM October 22nd, 2023

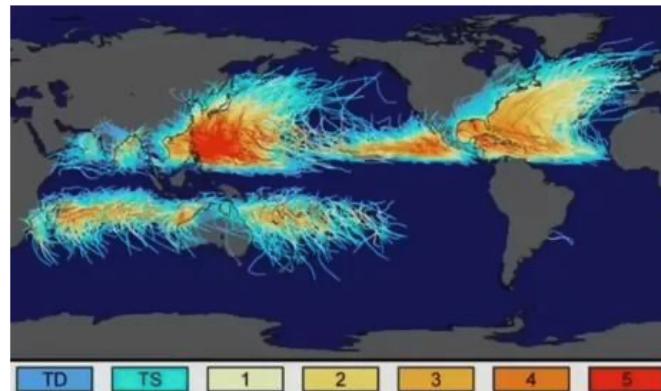


Figure F.7: Pre-Test Real News Post #3.



Anonymous3366  
@Anonymous3366

Really interesting fact I just learned - there's never been a hurricane that has ever crossed the equator (see picture below). They don't generate near the equator because there is no Coriolis effect there.



11:17 AM January 19th, 2024



Figure F.8: Pre-Test Real News Post #4.

**The following questions pertain to this post.**



Anonymous1989  
@Anonymous1989

Serious questions:

- In 2019, what did Gates offer to DARPA/UN/CIA to make him the one public figure calling the 'shots' in the '5G securitization' context?
- Why did Gates rekindle his passion with 'vaccines' in 2008?
- What has the 'nerd-in-chief' worked on since the late 90s?

11:22AM October 3rd, 2023



Figure F.9: Post-Test Misinformation Post #1.



Anonymous1027  
@Anonymous1027

Ohio: There are reports of dead fish and cattle found up to 100 miles away from the trainderailment.

Some journalists trying to cover this story have been arrested.

What in the world is happening?



11:54PM February 16th, 2023

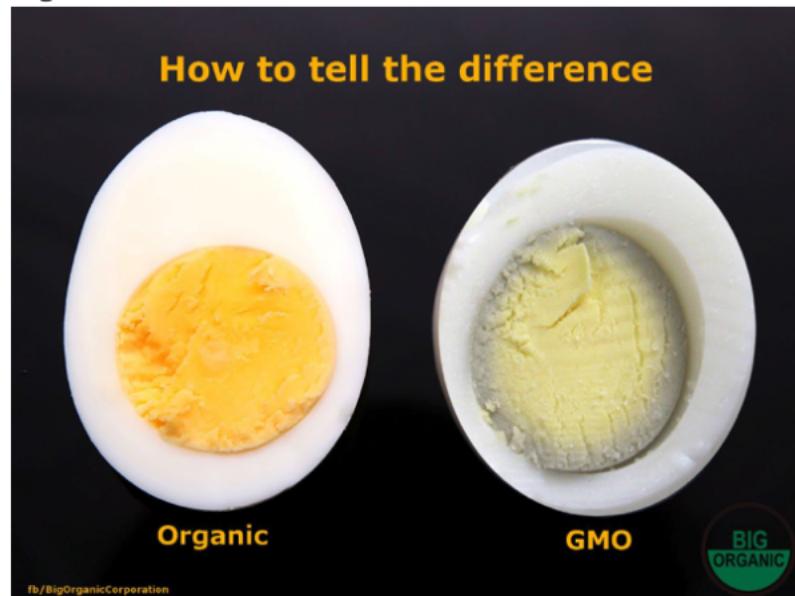


Figure F.10: Post-Test Misinformation Post #2.



Support Organic Food  
@supportorganicfood

Here's an important message from the kind people at "Support Organic Food"



7:07PM September 29th, 2023



Figure F.11: Post-Test Misinformation Post #3.



Daily Loud @DailyLoud

Flowers are now starting to bloom in Antarctica and experts say this is not good news.

This would be the first evidence of accelerated ecosystem response in Antarctica that is directly associated as a consequence of global warming, according to Nicoletta Cannone, a professor of ecology at the University of Insubria.



5:17PM September 21st, 2023



Figure F.12: Post-Test Misinformation Post #4.



The Washington Post   
@washingtonpost

Malnourished children in war-torn Yemen struggle to survive as temperatures rise.

In the city of Hodeida, climate change and hunger are converging in devastating ways. More families than usual were turned away from bursting hospital wards this summer.



WASHINGTONPOST.COM

**Where heat worsens hunger**

Malnourished children in war-torn Yemen struggle to survive as temperatures rise. In the cit...

2:04PM October 13th, 2023



Figure F.13: Post-Test Real News Post #1.



The New York Times   
@nytimes

"Barbie" will finish this weekend with more than \$1 billion in ticket sales at the global box office, according to Warner Bros., making Greta Gerwig the first woman to have a sole directing credit on a billion-dollar movie.



12:12 PM August 6th, 2023



Figure F.14: Post-Test Real News Post #2.

WSJ

The Wall Street Journal   
@WSJ

Breaking: The Food and Drug Administration said the overdose reversal medication Narcan could be sold over-the-counter for the first time since the opioid crisis began began.



9:07 AM March 29th, 2023



Figure F.15: Post-Test Real News Post #3.



Anonymous797  
@Anonymous797

Today I learned that honeybees are actually not endangered in the US.



10:11 PM November 29th, 2023



Figure F.16: Post-Test Real News Post #4.

## F.2 Countering Posts

These posts were used in the countering section of the survey experiment. All posts were described as being false to participants and were selected so as to be uncontroversially and obviously considered false or “misinformation” by all participants.



AnonymousXYZ  
@AnonymousXYZ

Don't forget to turn off your cell phones on October 4th.  
Unfortunately, the Emergency Broadcast System will be "testing"  
5G, and this might trigger the Marburg virus in vaccinated  
individuals, sadly transforming some into zombies.

10:33 PM September 29th, 2023



Figure F.17: Pre-Test Countering Post #1.



Anonymous412  
@Anonymous412

Isn't the moon supposed to be on the other side of the Earth where it is dark? But as you can see in this picture, it's over here with the sun.

Just more proof that the Earth is flat, and all they do is lie to us.



5:43 PM August 2nd, 2023



Figure F.18: Pre-Test Countering Post #2.



The People's Voice @realtpv

"Bill Gates mRNA 'Air Vaccine' Approved for Use Against Non-Consenting Humans"



12:17 PM October 5th, 2023



Figure F.19: Pre-Test Countering Post #3.



Anonymous03  
@Anonymous03

What do they do differently? Maybe it is because they don't take vaccines...



SLAYNEWS.COM

**Zero Amish Children Diagnosed with Cancer, Diabetes or Autism**

A comprehensive study has found that no Amish children have been diagnosed with chroni...

10:27 AM August 6th, 2023

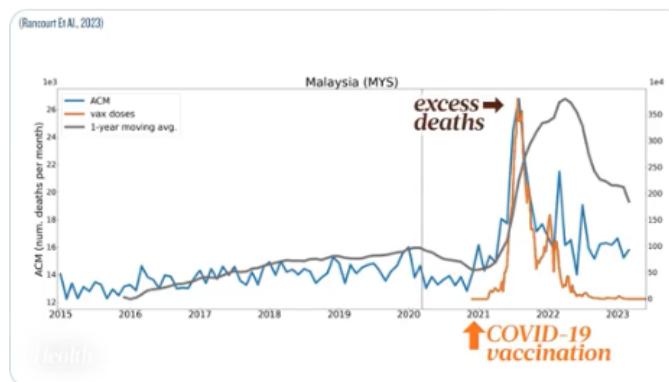


Figure F.20: Pre-Test Countering Post #4.



The Epoch Times   
@EpochTimes

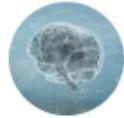
'No Lives Were Saved' by COVID-19 Vaccines, Scientists Estimate  
As the number of deaths clearly increased, upon closer examination, they noticed that the excess deaths coincided with the timing of the #COVID19 vaccine program rollout



3:50 PM October 8th, 2023



Figure F.21: Post-Test Countering Post #1.



Educated Brains  
@EducatedBrains

Scientists find  
sniffing rosemary can  
increase memory by  
**75%**



1:39 PM July 13th, 2021



Figure F.22: Post-Test Countering Post #2.



Anonymous72  
@Anonymous72

Whoa...

The Washington Post [Sign In](#)

## Henry Kissinger turned down OpenAI board seat days before death

The diplomat had an offer to join the OpenAI board but turned it down due to poor health sources reveal

By Thomas W. Lippman  
November 29, 2023 at 9:01 p.m. EST



8:07 PM November 30th, 2023



Figure F.23: Post-Test Countering Post #3.



Anonymous167  
@Anonymous167

Taylor Swift recently shared a picture from the latest stop on her tour, showing her in an African village savoring some ice cream. This picture is so disturbing because clearly the child next to her is extremely distressed and crying, while she is enjoying herself. Wow....



3:52 PM November 17th, 2023



Figure F.24: Post-Test Countering Post #4.

# Appendix G

## Chapter 5: Robustness Analysis

In the main analysis of individual differences in support and perceived attributes of misinformation interventions (see Chapter 5 Table 5.3), we ran OLS regressions to predict average support, perceived fairness, perceived effectiveness, and perceived intrusiveness of interventions as a function of various demographic variables. To supplement this analysis, we varied the partisanship groups to show that the results were robust to different partisan category mappings.

Table G.1 separates out the “Independent” and “Other/unaffiliated” categories rather than combining them. Table G.2 maps partisan categories to corresponding numeric values (strong Democrat to strong Republican categories mapped to numeric values 0 to 4). The results are substantially similar to the primary analysis.

Table G.1: OLS regressions predicting average support, perceived fairness, perceived effectiveness, and perceived intrusiveness of interventions as a function of demographic variables. Same analysis as Table 5.3, except partisanship separated out the “Other” party from “Independent”.

	<i>Dependent variable:</i>			
	Avg. support rating	Avg. fairness rating	Avg. effectiveness rating	Avg. intrusiveness rating
Age	0.020 (0.021)	0.044* (0.021)	-0.023 (0.019)	0.038* (0.018)
Male : Female	-0.201*** (0.054)	-0.139* (0.054)	-0.217*** (0.050)	0.103* (0.048)
Other (e.g., non-binary) : Female	-0.074 (0.175)	-0.252 (0.178)	-0.271 (0.162)	0.105 (0.156)
Education	0.013 (0.022)	0.003 (0.022)	0.007 (0.020)	0.016 (0.019)
Income	0.032* (0.016)	0.046** (0.016)	0.012 (0.015)	-0.005 (0.014)
Independent : Democrat	-0.260*** (0.075)	-0.243** (0.076)	-0.185** (0.070)	0.191** (0.067)
Other party : Democrat	-0.576** (0.199)	-0.596** (0.202)	-0.620*** (0.185)	0.461** (0.178)
Republican : Democrat	-0.275* (0.107)	-0.281** (0.109)	-0.193 (0.099)	0.152 (0.095)
Political Ideology	-0.300*** (0.037)	-0.268*** (0.037)	-0.190*** (0.034)	0.134*** (0.033)
Misinformation Exposure	-0.015 (0.027)	-0.008 (0.027)	-0.048 (0.025)	0.039 (0.024)
Constant	4.240*** (0.122)	4.042*** (0.124)	4.002*** (0.114)	2.493*** (0.109)
Observations	1,010	1,010	1,010	1,010
R <sup>2</sup>	0.246	0.208	0.156	0.097
Adjusted R <sup>2</sup>	0.238	0.200	0.147	0.088
Residual Std. Error (df = 999)	0.833	0.846	0.773	0.743
F Statistic (df = 10; 999)	32.525***	26.264***	18.443***	10.782***

Note:

\*p<0.05; \*\*p<0.01; \*\*\*p<0.001

Table G.2: OLS regressions predicting average support, perceived fairness, perceived effectiveness, and perceived intrusiveness of interventions as a function of demographic variables. Same analysis as Table 5.3, except partisanship was mapped to numeric values 0 to 4, with strong Democrats mapped to 0 and strong Republicans mapped to 4.

	<i>Dependent variable:</i>			
	Avg. support rating	Avg. fairness rating	Avg. effectiveness rating	Avg. intrusiveness rating
Age	0.018 (0.021)	0.043* (0.021)	-0.024 (0.019)	0.039* (0.019)
Male : Female	-0.199*** (0.054)	-0.136* (0.054)	-0.214*** (0.050)	0.103* (0.048)
Other (e.g., non-binary) : Female	-0.130 (0.173)	-0.306 (0.175)	-0.338* (0.161)	0.158 (0.154)
Education	0.014 (0.022)	0.003 (0.022)	0.007 (0.020)	0.015 (0.019)
Income	0.031* (0.016)	0.045** (0.016)	0.012 (0.015)	-0.005 (0.014)
Partisanship	-0.128*** (0.033)	-0.135*** (0.034)	-0.098** (0.031)	0.080** (0.030)
Political Ideology	-0.263*** (0.038)	-0.227*** (0.038)	-0.157*** (0.035)	0.108** (0.034)
Misinformation Exposure	-0.015 (0.027)	-0.008 (0.027)	-0.048 (0.025)	0.040 (0.024)
Constant	4.246*** (0.120)	4.054*** (0.122)	4.001*** (0.112)	2.503*** (0.108)
Observations	1,010	1,010	1,010	1,010
R <sup>2</sup>	0.244	0.208	0.151	0.092
Adjusted R <sup>2</sup>	0.238	0.202	0.144	0.085
Residual Std. Error (df = 1001)	0.833	0.846	0.775	0.745
F Statistic (df = 8; 1001)	40.316***	32.834***	22.280***	12.742***

Note:

\*p<0.05; \*\*p<0.01; \*\*\*p<0.001



# Bibliography

- [1] Fraud Detection. URL <https://www.qualtrics.com/support/survey-platform/survey-module/survey-checker/fraud-detection/>. 4, 5
- [2] Auto Bailout Now Backed, Stimulus Divisive. Technical report, Pew Research Center, February 2012. URL <https://www.pewresearch.org/politics/2012/02/23/auto-bailout-now-backed-stimulus-divisive/>. 5.5
- [3] Democracy by Design: Social Media’s Policy Scores. Technical report, Accountable Tech, 2024. URL <https://accountabletech.org/research/democracy-by-design-social-medias-policy-scores/>. 1.6, 6.4.2, A.1, A.2
- [4] Zhila Aghajari, Eric P. S. Baumer, and Dominic DiFranzo. Reviewing Interventions to Address Misinformation: The Need to Expand Our Vision Beyond an Individualistic Focus. In *Proceedings of the ACM on Human-Computer Interaction*, volume 7 of *CSCW*, pages 87:1–87:34. Association for Computing Machinery, April 2023. doi: 10.1145/3579520. URL <https://dl.acm.org/doi/10.1145/3579520>. 1.4.1, 1.4, 1.4.1, 1.4.2, 2.1, 2.2
- [5] Kathryn J. Aikin, Brian G. Southwell, Ryan S. Paquin, Douglas J. Rupert, Amie C. O’Donoghue, Kevin R. Betts, and Philip K. Lee. Correction of misleading information in prescription drug television advertising: The roles of advertisement similarity and time delay. *Research in Social and Administrative Pharmacy*, 13(2):378–388, March 2017. ISSN 1551-7411. doi: 10.1016/j.sapharm.2016.04.004. URL <https://www.sciencedirect.com/science/article/pii/S1551741116300146>. 2.2, 6.4.2, A.1
- [6] Jennifer Allen and David Rand. Combating Misinformation Runs Deeper Than Swatting Away ‘Fake News’, September 2024. URL <https://www.scientificamerican.com/article/combating-misinformation-runs-deeper-than-swatting-away-fake-news/>. 5.1
- [7] Jennifer Allen, Antonio A. Arechar, Gordon Pennycook, and David G. Rand. Scaling up fact-checking using the wisdom of crowds. *SCIENCE ADVANCES*, 2021. A.4
- [8] Jennifer Allen, Cameron Martel, and David G. Rand. Birds of a feather don’t fact-check each other: Partisanship and the evaluation of news in Twitter’s Birdwatch crowdsourced fact-checking program. In *CHI Conference on Human Factors in Computing Systems*, pages 1–19, New Orleans LA USA, April 2022. ACM. ISBN 978-1-4503-9157-3. doi: 10.1145/3491102.3502040. A.4

- [9] Bobby Allyn. The Rise of 'Grid Zero': Why more Instagram users are hiding their profile. *NPR*, April 2024. URL <https://www.npr.org/2024/04/19/1245316400/grid-zero-instagram-users-hiding-profile-privacy>. 3.4.1
- [10] Sacha Altay, Manon Berriche, Hendrik Heuer, Johan Farkas, and Steven Rathje. A survey of expert views on misinformation: Definitions, determinants, solutions, and future of the field. *Harvard Kennedy School Misinformation Review*, July 2023. doi: 10.37016/mr-2020-119. URL <https://misinforeview.hks.harvard.edu/article/a-survey-of-expert-views-on-misinformation-definitions-determinants-solutions-and-future-of-the-field/>. 6.4.1
- [11] Antonio A. Arechar, Jennifer Allen, Adam J. Berinsky, Rocky Cole, Ziv Epstein, Kiran Garimella, Andrew Gully, Jackson G. Lu, Robert M. Ross, Michael N. Stagnaro, Yunhao Zhang, Gordon Pennycook, and David G. Rand. Understanding and combatting misinformation across 16 countries on six continents. *Nature Human Behaviour*, 7(9):1502–1513, September 2023. ISSN 2397-3374. doi: 10.1038/s41562-023-01641-6. URL <https://www.nature.com/articles/s41562-023-01641-6>. Publisher: Nature Publishing Group. 2.2
- [12] Ahmer Arif, John J. Robinson, Stephanie A. Stanek, Elodie S. Fichet, Paul Townsend, Zena Worku, and Kate Starbird. A Closer Look at the Self-Correcting Crowd: Examining Corrections in Online Rumors. In *Proceedings of the 2017 ACM Conference on Computer Supported Cooperative Work and Social Computing*, CSCW, pages 155–168, New York, NY, USA, February 2017. Association for Computing Machinery. ISBN 978-1-4503-4335-0. doi: 10.1145/2998181.2998294. URL <https://dl.acm.org/doi/10.1145/2998181.2998294>. A.6
- [13] K. Peren Arin, Deni Mazrekaj, and Marcel Thum. Ability of detecting and willingness to share fake news. *Scientific Reports*, 13(1):7298, May 2023. ISSN 2045-2322. doi: 10.1038/s41598-023-34402-6. URL <https://www.nature.com/articles/s41598-023-34402-6>. Publisher: Nature Publishing Group. 3.1
- [14] E Aronson, C Fried, and J Stone. Overcoming denial and increasing the intention to use condoms through the induction of hypocrisy. *American Journal of Public Health*, 81(12):1636–1638, December 1991. ISSN 0090-0036, 1541-0048. doi: 10.2105/AJPH.81.12.1636. URL <https://ajph.aphapublications.org/doi/full/10.2105/AJPH.81.12.1636>. 3.1
- [15] Dennis Assemacher, Derek Weber, Mike Preuss, André Calero Valdez, Alison Bradshaw, Björn Ross, Stefano Cresci, Heike Trautmann, Frank Neumann, and Christian Grimme. Benchmarking Crisis in Social Media Analytics: A Solution for the Data-Sharing Problem. *Social Science Computer Review*, 40(6):1496–1522, December 2022. ISSN 0894-4393. doi: 10.1177/08944393211012268. URL <https://doi.org/10.1177/08944393211012268>. Publisher: SAGE Publications Inc. 5.1, A.7
- [16] Brooke Auxier. 54% of Americans say social media companies shouldn't allow any political ads, September 2020. URL <https://www.pewresearch.org/fact-tank/2020/09/24/54-of-americans-say-social-media-companies-shouldnt-allow-any-political-ads/>. 6.4.2

- [17] Brooke Auxier. 64% of Americans say social media have a mostly negative effect on the way things are going in the U.S. today, October 2020. URL <https://www.pewresearch.org/short-reads/2020/10/15/64-of-americans-say-social-media-have-a-mostly-negative-effect-on-the-way-things-are-going-in-the-u-s-today/>. 3.3.5
- [18] Brooke Auxier and Monica Anderson. Social Media Use in 2021, April 2021. URL <https://www.pewresearch.org/internet/2021/04/07/social-media-use-in-2021/>. 3.2.3
- [19] Mihai Avram, Nicholas Micallef, Sameer Patil, and Filippo Menczer. Exposure to social engagement metrics increases vulnerability to misinformation. *Harvard Kennedy School Misinformation Review*, July 2020. doi: 10.37016/mr-2020-033. URL <https://misinforeview.hks.harvard.edu/article/exposure-to-social-engagement-metrics-increases-vulnerability-to-misinformation/>. 3.5, A.6
- [20] Sumitra Badrinathan. Educative Interventions to Combat Misinformation: Evidence from a Field Experiment in India. *American Political Science Review*, 115(4):1325–1341, November 2021. ISSN 0003-0554, 1537-5943. doi: 10.1017/S0003055421000459. 4.2.1
- [21] Sumitra Badrinathan and Simon Chauchard. “I Don’t Think That’s True, Bro!” Social Corrections of Misinformation in India. *The International Journal of Press/Politics*, page 19401612231158770, February 2023. ISSN 1940-1612. doi: 10.1177/19401612231158770. 1.4.1, 2.2, 3, 4.2.2, 5.1, A.6
- [22] Bence Bago, David G. Rand, and Gordon Pennycook. Fake news, fast and slow: Deliberation reduces belief in false (but not true) news headlines. *Journal of Experimental Psychology: General*, 149(8):1608–1613, 2020. ISSN 1939-2222. doi: 10.1037/xge0000729. 2.2, 6.4.2, A.1
- [23] Joseph B. Bak-Coleman, Ian Kennedy, Morgan Wack, Andrew Beers, Joseph S. Schafer, Emma S. Spiro, Kate Starbird, and Jevin D. West. Combining interventions to reduce the spread of viral misinformation. *Nature Human Behaviour*, pages 1–9, June 2022. ISSN 2397-3374. doi: 10.1038/s41562-022-01388-6. URL <https://www.nature.com/articles/s41562-022-01388-6>. Publisher: Nature Publishing Group. 5.1, A.8
- [24] Sebastian Bamberg and Daniel Rölle. Determinants of People’s Acceptability of Pricing Measures – Replication and Extension of a Causal Model. In Jens Schade and Bernhard Schlag, editors, *Acceptability of Transport Pricing Strategies*, pages 235–248. Emerald Group Publishing Limited, January 2003. ISBN 978-0-08-044199-3 978-1-78635-950-6. doi: 10.1108/9781786359506-015. URL <https://doi.org/10.1108/9781786359506-015>. 5.5
- [25] Anton Barashin. Improving the Western Strategy to Combat Kremlin Propaganda and Disinformation. Technical report, Atlantic Council: Eurasia Center, May 2018. URL [https://www.atlanticcouncil.org/wp-content/uploads/2018/06/Improving\\_the\\_Western\\_Strategy.pdf](https://www.atlanticcouncil.org/wp-content/uploads/2018/06/Improving_the_Western_Strategy.pdf). 5.1, A.7
- [26] Paul Barrett. Tackling Domestic Disinformation: What the Social Media Companies Need

- to Do. Technical report, NYU Stern Center for Business and Human Rights, March 2019. URL [https://issuu.com/nyusterncenterforbusinessandhumanri/docs/nyu Domestic\\_Disinformation\\_Digital](https://issuu.com/nyusterncenterforbusinessandhumanri/docs/nyu Domestic_Disinformation_Digital). 5.1, A.1, A.2, A.7
- [27] Michael Barthel, Amy Mitchell, and Jesse Holcomb. Many Americans Believe Fake News Is Sowing Confusion, December 2016. URL <https://www.pewresearch.org/journalism/2016/12/15/many-americans-believe-fake-news-is-sowing-confusion/>. 3.1, 3.3.2, 3.3.5
- [28] Melisa Basol, Jon Roozenbeek, and Sander van der Linden. Good News about Bad News: Gamified Inoculation Boosts Confidence and Cognitive Immunity Against Fake News. *Journal of Cognition*, 3(1):1–9, 2020. ISSN 2514-4820. doi: 10.5334/joc.91. 2.6.2, 4.1, 4.2.1
- [29] C. Daniel Batson, Diane Kobrynowicz, Jessica L. Dinnerstein, Hannah C. Kampf, and Angela D. Wilson. In a very different voice: Unmasking moral hypocrisy. *Journal of Personality and Social Psychology*, 72(6):1335–1348, 1997. ISSN 1939-1315. doi: 10.1037/0022-3514.72.6.1335. Place: US Publisher: American Psychological Association. 3.3.1
- [30] Rony Berger, Joy Benatov, Hisham Abu-Raiya, and Carmit T. Tadmor. Reducing prejudice and promoting positive intergroup attitudes among elementary-school children in the context of the Israeli–Palestinian conflict. *Journal of School Psychology*, 57:53–72, August 2016. ISSN 0022-4405. doi: 10.1016/j.jsp.2016.04.003. URL <https://www.sciencedirect.com/science/article/pii/S0022440516300097>. 3
- [31] Magnus Bergquist, Andreas Nilsson, Niklas Harring, and Sverker C. Jagers. Meta-analyses of fifteen determinants of public opinion about climate change taxes and laws. *Nature Climate Change*, 12(3):235–240, March 2022. ISSN 1758-6798. doi: 10.1038/s41558-022-01297-6. URL <https://www.nature.com/articles/s41558-022-01297-6>. Publisher: Nature Publishing Group. 5, 5.3.1, 5.5
- [32] Samuel Bestvater. How U.S. Adults Use TikTok, February 2024. URL <https://www.pewresearch.org/internet/2024/02/22/how-u-s-adults-use-tiktok/>. 3.4.1
- [33] Libby Bishop and Daniel Gray. Ethical Challenges of Publishing and Sharing Social Media Research Data. In Kandy Woodfield, editor, *The Ethics of Online Research*, volume 2 of *Advances in Research Ethics and Integrity*, pages 159–187. Emerald Publishing Limited, January 2017. ISBN 978-1-78714-486-6 978-1-78714-485-9. <https://doi.org/10.1108/S2398-601820180000002007>. 1.1, A.7
- [34] Robert A. Blair, Jessica Gottlieb, Brendan Nyhan, Laura Paler, Pablo Argote, and Charlene J. Stainfield. Interventions to counter misinformation: Lessons from the Global North and applications to the Global South. Technical report, USAID Development Experience Clearinghouse, July 2023. URL <https://democratic-erosion.org/wp-content/uploads/2024/09/INTERVENTIONS-TO-COUNTER-MISINFORMATION-LESSONS-FROM-THE-GLOBAL-NORTH-AND-APPLICATIONS-TO-THE-GLOBAL-SOUTH.pdf>. 6.4.2

- [35] Robert A. Blair, Jessica Gottlieb, Brendan Nyhan, Laura Paler, Pablo Argote, and Charlene J. Stainfield. Interventions to counter misinformation: Lessons from the Global North and applications to the Global South. *Current Opinion in Psychology*, 55:101732, February 2024. ISSN 2352-250X. doi: 10.1016/j.copsyc.2023.101732. URL <https://www.sciencedirect.com/science/article/abs/pii/S2352250X2300177X>. 1.4.1, 1.4, 1.4.1, 1.4.2, 2.1, 4.1, 6.2.3, 6.4.1, 6.4.2
- [36] Leticia Bode and Emily K. Vraga. In Related News, That was Wrong: The Correction of Misinformation Through Related Stories Functionality in Social Media. *Journal of Communication*, 65(4):619–638, August 2015. ISSN 0021-9916. doi: 10.1111/jcom.12166. URL <https://doi.org/10.1111/jcom.12166>. 2.2, 6.4.2, A.1
- [37] Leticia Bode and Emily K. Vraga. See Something, Say Something: Correction of Global Health Misinformation on Social Media. *Health Communication*, 33(9):1131–1140, September 2018. ISSN 1041-0236. doi: 10.1080/10410236.2017.1331312. 1.4.1, 3, 3.1, 4.2.2, A.2, A.6
- [38] Andrew Booth, Diana Papaioannou, and Anthea Sutton. *Systematic Approaches to a Successful Literature Review*. Sage Publications, 1st edition, January 2012. ISBN 978-1-4739-1245-8. URL [https://www.researchgate.net/profile/Andrew-Booth-2/publication/235930866\\_Systematic\\_Approaches\\_to\\_a\\_Successful\\_Literature\\_Review/links/5da06c7f45851553ff8705fa/Systematic-Approaches-to-a-Successful-Literature-Review.pdf](https://www.researchgate.net/profile/Andrew-Booth-2/publication/235930866_Systematic_Approaches_to_a_Successful_Literature_Review/links/5da06c7f45851553ff8705fa/Systematic-Approaches-to-a-Successful-Literature-Review.pdf). 2.2.4
- [39] Alexander Bor, Mathias Osmundsen, Stig Hebbelstrup Rye Rasmussen, Anja Bechmann, and Michael Bang Petersen. "Fact-checking" videos reduce belief in misinformation and improve the quality of news shared on Twitter, September 2020. URL <https://osf.io/a7huq>. A.2
- [40] Colin Bos, Ivo Van Der Lans, Frank Van Rijnsoever, and Hans Van Trijp. Consumer Acceptance of Population-Level Intervention Strategies for Healthy Food Choices: The Role of Perceived Effectiveness and Perceived Fairness. *Nutrients*, 7(9):7842–7862, September 2015. ISSN 2072-6643. doi: 10.3390/nu7095370. URL <https://www.mdpi.com/2072-6643/7/9/5370>. Number: 9 Publisher: Multidisciplinary Digital Publishing Institute. 5, 5.5
- [41] Samantha Bradshaw and Lisa-Maria Neudert. The Road Ahead: Mapping Civil Society Responses to Disinformation. Technical report, National Endowment for Democracy, January 2021. URL <https://www.ned.org/wp-content/uploads/2021/01/The-Road-Ahead-Mapping-Civil-Society-Responses-to-Disinformation-Bradshaw-Neudert-Jan-2021-2.pdf>. 5.1, A.7
- [42] Nadia M. Brashier. Fighting misinformation among the most vulnerable users. *Current Opinion in Psychology*, 57:101813, June 2024. ISSN 2352-250X. doi: 10.1016/j.copsyc.2024.101813. URL <https://www.sciencedirect.com/science/article/abs/pii/S2352250X24000265>. 3.5

- [43] Nadia M. Brashier and Daniel L. Schacter. Aging in an Era of Fake News. *Current Directions in Psychological Science*, 29(3):316–323, June 2020. ISSN 0963-7214. doi: 10.1177/0963721420915872. URL <https://doi.org/10.1177/0963721420915872>. Publisher: SAGE Publications Inc. 3.3.1, 3.5
- [44] Aengus Bridgman, Eric Merkley, Peter John Loewen, Taylor Owen, Derek Ruths, Lisa Teichmann, and Oleg Zhilin. The causes and consequences of COVID-19 misperceptions: Understanding the role of news and social media. *Harvard Kennedy School Misinformation Review*, June 2020. doi: 10.37016/mr-2020-028. URL <https://misinforeview.hks.harvard.edu/?p=1832>. 3.3.1
- [45] Nils Brunsson. *The Organization of Hypocrisy: Talk, Decisions and Actions in Organizations*. Wiley, New York, NY, 1989. 3.1
- [46] Hanne Bruun. The delay economy of “continuity” and the emerging impatience culture of the digital era. *Nordic Journal of Media Studies*, 1(1):85–101, June 2019. doi: 10.2478/njms-2019-0006. URL <https://sciendo.com/article/10.2478/njms-2019-0006>. 5.5
- [47] Ceren Budak, Brendan Nyhan, David M. Rothschild, Emily Thorson, and Duncan J. Watts. Misunderstanding the harms of online misinformation. *Nature*, 630(8015):45–53, June 2024. ISSN 0028-0836, 1476-4687. doi: 10.1038/s41586-024-07417-w. URL <https://www.nature.com/articles/s41586-024-07417-w>. 5.1
- [48] Paul Burstein. The Impact of Public Opinion on Public Policy: A Review and an Agenda. *Political Research Quarterly*, 56(1):29–40, 2003. ISSN 1065-9129. doi: 10.2307/3219881. URL <https://www.jstor.org/stable/3219881>. Publisher: [University of Utah, Sage Publications, Inc.]. 5, 5.1
- [49] John M. Carey, Victoria Chi, D. J. Flynn, Brendan Nyhan, and Thomas Zeitzoff. The effects of corrective information about disease epidemics and outbreaks: Evidence from Zika and yellow fever in Brazil. *Science Advances*, 6(5):eaaw7449, January 2020. doi: 10.1126/sciadv.aaw7449. URL <https://www.science.org/doi/10.1126/sciadv.aaw7449>. Publisher: American Association for the Advancement of Science. 3.3.1
- [50] Kathleen M. Carley. ORA: A Toolkit for Dynamic Network Analysis and Visualization. In Reda Alhajj and Jon Rokne, editors, *Encyclopedia of Social Network Analysis and Mining*, pages 1–10. Springer, New York, NY, 2017. ISBN 978-1-4614-7163-9. doi: 10.1007/978-1-4614-7163-9\_309-1. URL [https://doi.org/10.1007/978-1-4614-7163-9\\_309-1](https://doi.org/10.1007/978-1-4614-7163-9_309-1). 2.3
- [51] Man-pui Sally Chan, Christopher R. Jones, Kathleen Hall Jamieson, and Dolores Albaracín. Debunking: A Meta-Analysis of the Psychological Efficacy of Messages Counteracting Misinformation. *Psychological Science*, 28(11):1531–1546, November 2017. ISSN 0956-7976. doi: 10.1177/0956797617714579. 1.1, 2.6.2, A.2
- [52] Xinran Chen, Sei-Ching Joanna Sin, Yin-Leng Theng, and Chei Sian Lee. Why Students Share Misinformation on Social Media: Motivation, Gender, and Study-level Differences. *The Journal of Academic Librarianship*, 41(5):583–592, September

2015. ISSN 0099-1333. doi: 10.1016/j.acalib.2015.07.003. URL <https://www.sciencedirect.com/science/article/pii/S0099133315001494>. 3.1

- [53] Lesley Chiou and Catherine Tucker. Fake News and Advertising on Social Media: A Study of the Anti-Vaccination Movement, November 2018. URL <https://www.nber.org/papers/w25223>. A.1
- [54] Yuwei Chuai, Haoye Tian, Nicolas Pröllochs, and Gabriele Lenzini. Did the Roll-Out of Community Notes Reduce Engagement With Misinformation on X/Twitter? *Proc. ACM Hum.-Comput. Interact.*, 8(CSCW2):428:1–428:52, November 2024. doi: 10.1145/3686967. URL <https://dl.acm.org/doi/10.1145/3686967>. 5.6.1
- [55] Giovanni Luca Ciampaglia, Alexios Mantzarlis, Gregory Maus, and Filippo Menczer. Research Challenges of Digital Misinformation: Toward a Trustworthy Web. *AI Magazine*, 39(1):65–74, March 2018. ISSN 0738-4602, 2371-9621. doi: 10.1609/aimag.v39i1.2783. URL <https://onlinelibrary.wiley.com/doi/10.1609/aimag.v39i1.2783>. 5.2
- [56] Mitchell Clark. Facebook wants to make sure you've read the article you're about to share, May 2021. URL <https://www.theverge.com/2021/5/10/22429174/facebook-article-popup-read-misinformation>. 3.5, 5.5
- [57] Jonas Colliander. “This is fake news”: Investigating the role of conformity to other users’ views when commenting on and spreading disinformation in social media. *Computers in Human Behavior*, 97:202–215, August 2019. ISSN 07475632. doi: 10.1016/j.chb.2019.03.032. URL <https://linkinghub.elsevier.com/retrieve/pii/S074756321930130X>. 3.5
- [58] Thomas H. Costello, Gordon Pennycook, and David G. Rand. Durably reducing conspiracy beliefs through dialogues with AI. *Science*, 385(6714):eadq1814, September 2024. doi: 10.1126/science.adq1814. URL <https://www.science.org/doi/10.1126/science.adq1814>. Publisher: American Association for the Advancement of Science. A.8
- [59] Laura Courchesne, Julia Ilhardt, and Jacob N. Shapiro. Review of social science research on the impact of countermeasures against influence operations. *Harvard Kennedy School Misinformation Review*, 2, September 2021. doi: 10.37016/mr-2020-79. 1.1, 1.4.1, 1.4, 1.4.1, 1.4.2, 2.1, 2.4.1, 2.4.3, 2.5, 2.6.1, 3.1, 4.1, 5.1, 5.2.1, 5.1, A.1, A.3, A.4
- [60] Stéphane Côté, Julian House, and Robb Willer. High economic inequality leads higher-income individuals to be less generous. *Proceedings of the National Academy of Sciences*, 112(52):15838–15843, December 2015. doi: 10.1073/pnas.1511536112. URL <https://www.pnas.org/doi/10.1073/pnas.1511536112>. Publisher: Proceedings of the National Academy of Sciences. 3.3.5
- [61] David De Coninck, Thomas Frissen, Koen Matthijs, Leen d’Haenens, Grégoire Lits, Olivier Champagne-Poirier, Marie-Eve Carignan, Marc D. David, Nathalie Pignard-Cheynel, Sébastien Salerno, and Melissa Généreux. Beliefs in Conspiracy Theories and Misinformation About COVID-19: Comparative Perspectives

- on the Role of Anxiety, Depression and Exposure to and Trust in Information Sources. *Frontiers in Psychology*, 12:646394, April 2021. ISSN 1664-1078. doi: 10.3389/fpsyg.2021.646394. URL <https://www.frontiersin.org/articles/10.3389/fpsyg.2021.646394/full>. 5.1
- [62] Jenny de Fine Licht. Policy Area as a Potential Moderator of Transparency Effects: An Experiment. *Public Administration Review*, 74(3):361–371, 2014. ISSN 1540-6210. doi: 10.1111/puar.12194. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/puar.12194>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1111/puar.12194>. 5.5
- [63] Ratan Dey, Zubin Jelveh, and Keith Ross. Facebook users have become much more private: A large-scale study. In *2012 IEEE International Conference on Pervasive Computing and Communications Workshops*, pages 346–352, March 2012. doi: 10.1109/PerComW.2012.6197508. URL <https://ieeexplore.ieee.org/abstract/document/6197508>. 3.4.1
- [64] Stephanie Diepeveen, Tom Ling, Marc Suhrcke, Martin Roland, and Theresa M. Marteau. Public acceptability of government intervention to change health-related behaviours: a systematic review and narrative synthesis. *BMC Public Health*, 13(1):756, August 2013. ISSN 1471-2458. doi: 10.1186/1471-2458-13-756. URL <https://doi.org/10.1186/1471-2458-13-756>. 5, 5.5, 5.6.1, 6.2.2
- [65] Stacy Jo Dixon. U.S. social media user account privacy 2018, 2018. URL <https://www.statista.com/statistics/934874/users-have-private-social-media-account-usa/>. 3.4.1
- [66] Joan Donovan. Why social media can't keep moderating content in the shadows, November 2020. URL <https://www.technologyreview.com/2020/11/06/1011769/social-media-moderation-transparency-censorship/>. 5, 5.1
- [67] Chiara Patricia Drolsbach, Kirill Solovev, and Nicolas Pröllochs. Community notes increase trust in fact-checking on social media. *PNAS Nexus*, 3(7):pgae217, July 2024. ISSN 2752-6542. doi: 10.1093/pnasnexus/pgae217. URL <https://doi.org/10.1093/pnasnexus/pgae217>. 5.6.1
- [68] James N. Druckman, Katherine Ognyanova, Matthew A. Baum, David Lazer, Roy H. Perlis, John Della Volpe, Mauricio Santillana, Hanyu Chwe, Alexi Quintana, and Matthew Simonson. The role of race, religion, and partisanship in misperceptions about COVID-19. *Group Processes & Intergroup Relations*, 24(4):638–657, June 2021. ISSN 1368-4302. doi: 10.1177/1368430220985912. URL <https://doi.org/10.1177/1368430220985912>. Publisher: SAGE Publications Ltd. 3.3.1
- [69] Susan E. Dudley and Jerry Brito, editors. *Regulation: a primer*. Mercatus Center at George Mason University, Arlington, Va, 2nd ed edition, 2012. ISBN 978-0-9836077-3-1 978-0-9836077-4-8. 6.2.3
- [70] Francesca D'Errico, Paolo Giovanni Cicirelli, Giuseppe Corbelli, and Marinella Paciello. Addressing racial misinformation at school: a psycho-social intervention aimed at re-

- ducing ethnic moral disengagement in adolescents. *Social Psychology of Education*, 27(3):611–630, June 2024. ISSN 1573-1928. doi: 10.1007/s11218-023-09777-z. URL <https://doi.org/10.1007/s11218-023-09777-z>. 3.3.1, 3.5
- [71] Ullrich Ecker, Jon Roozenbeek, Sander van der Linden, Li Qian Tay, John Cook, Naomi Oreskes, and Stephan Lewandowsky. Misinformation remains a threat to democracy. *Nature*, 630:29–32, June 2024. doi: 10.1038/d41586-024-01587-3. URL <https://www.nature.com/articles/d41586-024-01587-3>. 5.1
  - [72] Ullrich K. H. Ecker and Luke M. Antonio. Can you believe it? An investigation into the impact of retraction source credibility on the continued influence effect. *Memory & Cognition*, 49(4):631–644, May 2021. ISSN 1532-5946. doi: 10.3758/s13421-020-01129-y. URL <https://doi.org/10.3758/s13421-020-01129-y>. 2.2, A.6
  - [73] Ullrich K. H. Ecker, Joshua L. Hogan, and Stephan Lewandowsky. Reminders and repetition of misinformation: Helping or hindering its retraction? *Journal of Applied Research in Memory and Cognition*, 6(2):185–192, June 2017. ISSN 2211-369X, 2211-3681. doi: 10.1037/h0101809. 1.1
  - [74] Ullrich K. H. Ecker, Stephan Lewandowsky, John Cook, Philipp Schmid, Lisa K. Fazio, Nadia Brashier, Panayiota Kendeou, Emily K. Vraga, and Michelle A. Amazeen. The psychological drivers of misinformation belief and its resistance to correction. *Nature Reviews Psychology*, 1(1):13–29, January 2022. ISSN 2731-0574. doi: 10.1038/s44159-021-00006-y. URL <https://www.nature.com/articles/s44159-021-00006-y>. Number: 1 Publisher: Nature Publishing Group. 3.1
  - [75] Ullrich K. H. Ecker, Jasmyne A. Sanderson, Paul McIlhiney, Jessica J. Rowsell, Hayley L. Quekett, Gordon DA Brown, and Stephan Lewandowsky. Combining refutations and social norms increases belief change. *Quarterly Journal of Experimental Psychology (2006)*, 76(6):1275–1297, June 2023. ISSN 1747-0226. doi: 10.1177/17470218221111750. URL <https://journals.sagepub.com/doi/10.1177/17470218221111750>. 2.2, A.6
  - [76] Adam M. Enders, Joseph E. Uscinski, Michelle I. Seelig, Casey A. Klofstad, Stefan Wuchty, John R. Funchion, Manohar N. Murthi, Kamal Premaratne, and Justin Stoler. The Relationship Between Social Media Use and Beliefs in Conspiracy Theories and Misinformation. *Political Behavior*, 45(2):781–804, June 2023. ISSN 0190-9320, 1573-6687. doi: 10.1007/s11109-021-09734-6. URL <https://link.springer.com/10.1007/s11109-021-09734-6>. 5.1
  - [77] Ziv Epstein, Adam J. Berinsky, Rocky Cole, Andrew Gully, Gordon Pennycook, and David G. Rand. Developing an accuracy-prompt toolkit to reduce COVID-19 misinformation online. *Harvard Kennedy School Misinformation Review*, May 2021. doi: 10.37016/mr-2020-71. URL <https://misinforeview.hks.harvard.edu/article/developing-an-accuracy-prompt-toolkit-to-reduce-covid-19-misinformation-online/>. 2.2, A.1
  - [78] Valerie Fointiat. "I know what I have to do, but..." When hypocrisy leads to behavioral change. *Social Behavior and Personality: An International Journal*, 32(8):741–

- 746, January 2004. ISSN 0301-2212. doi: 10.2224/sbp.2004.32.8.741. URL <https://www.ingentaconnect.com/content/10.2224/sbp.2004.32.8.741>. 3.1
- [79] Mingkun Gao, Ziang Xiao, Karrie Karahalios, and Wai-Tat Fu. To Label or Not to Label: The Effect of Stance and Credibility Labels on Readers' Selection and Perception of News Articles. In *Proceedings of the ACM on Human-Computer Interaction*, volume 2 of *CSCW*, pages 55:1–55:16, November 2018. doi: 10.1145/3274324. URL <https://dl.acm.org/doi/10.1145/3274324.2.2>, A.4
- [80] Naman Garg and Monika Yadav. Learning to Resist Misinformation: A Field Experiment in India, 2022. URL <https://www.aeaweb.org/doi/10.1257/rct.7923.2.2>
- [81] R. Kelly Garrett and Robert M. Bond. Conservatives' susceptibility to political misperceptions. *Science Advances*, 7(23):eabf1234, June 2021. doi: 10.1126/sciadv.abf1234. URL <https://www.science.org/doi/full/10.1126/sciadv.abf1234>. Publisher: American Association for the Advancement of Science. 3.3.1
- [82] Andrew Gelman, Jennifer Hill, and Masanao Yajima. Why We (Usually) Don't Have to Worry About Multiple Comparisons. *Journal of Research on Educational Effectiveness*, 5(2):189–211, April 2012. ISSN 1934-5747. doi: 10.1080/19345747.2011.618213. URL <https://doi.org/10.1080/19345747.2011.618213>. Publisher: Routledge eprint: <https://doi.org/10.1080/19345747.2011.618213>. 3.3.2
- [83] Tarleton Gillespie. Do Not Recommend? Reduction as a Form of Content Moderation. *Social Media + Society*, 8(3):20563051221117552, July 2022. ISSN 2056-3051. doi: 10.1177/20563051221117552. URL <https://doi.org/10.1177/20563051221117552>. Publisher: SAGE Publications Ltd. 5.1, A.1, A.2
- [84] Henner Gimpel, Sebastian Heger, Christian Olenberger, and Lena Utz. The Effectiveness of Social Norms in Fighting Fake News on Social Media. *Journal of Management Information Systems*, 38(1):196–221, January 2021. ISSN 0742-1222. URL <https://www.tandfonline.com/doi/abs/10.1080/07421222.2021.1870389>. Publisher: Routledge. 3.5, A.6
- [85] Burt Glass. Survey: Across parties, Americans accept removal of false health info by social media companies, survey says, February 2025. URL <https://sites.bu.edu/crc/2025/02/17/survey-across-parties-americans-accept-removal-of-false-health-info-by-social-media-companies-survey-says/>. 6.2.3
- [86] Sandra González-Bailón, David Lazer, Pablo Barberá, Meiqing Zhang, Hunt Allcott, Taylor Brown, Adriana Crespo-Tenorio, Deen Freelon, Matthew Gentzkow, Andrew M. Guess, Shanto Iyengar, Young Mie Kim, Neil Malhotra, Devra Moehler, Brendan Nyhan, Jennifer Pan, Carlos Velasco Rivera, Jaime Settle, Emily Thorson, Rebekah Tromble, Arjun Wilkins, Magdalena Wojcieszak, Chad Kiewiet de Jonge, Annie Franco, Winter Mason, Natalie Jomini Stroud, and Joshua A. Tucker. Asymmetric ideological segregation in exposure to political news on Facebook. *Science*, 381(6656):392–398, July 2023. doi: 10.1126/science.adc7138. URL <https://www.science.org/doi/full/10.1126/science.adc7138>. Publisher: American Association for the Ad-

vancement of Science. 3.3.1

- [87] Sandra González-Bailón, David Lazer, Pablo Barberá, William Godel, Hunt Allcott, Taylor Brown, Adriana Crespo-Tenorio, Deen Freelon, Matthew Gentzkow, Andrew M. Guess, Shanto Iyengar, Young Mie Kim, Neil Malhotra, Devra Moehler, Brendan Nyhan, Jennifer Pan, Carlos Velasco Rivera, Jaime Settle, Emily Thorson, Rebekah Tromble, Arjun Wilkins, Magdalena Wojcieszak, Chad Kiewiet de Jonge, Annie Franco, Winter Mason, Natalie Jomini Stroud, and Joshua A. Tucker. The Diffusion and Reach of (Mis)Information on Facebook During the U.S. 2020 Election. *Sociological Science*, 11:1124–1146, December 2024. ISSN 2330-6696. doi: 10.15195/v11.a41. URL <https://sociologicalscience.com/articles-v11-41-1124/>. 6.2.3
- [88] Robert Gorwa, Reuben Binns, and Christian Katzenbach. Algorithmic content moderation: Technical and political challenges in the automation of platform governance. *Big Data & Society*, 7(1):2053951719897945, January 2020. ISSN 2053-9517. doi: 10.1177/2053951719897945. URL <https://doi.org/10.1177/2053951719897945>. Publisher: SAGE Publications Ltd. A.2
- [89] Sonja Grelle and Wilhelm Hofmann. When and Why Do People Accept Public-Policy Interventions? An Integrative Public-Policy-Acceptance Framework. *Perspectives on Psychological Science*, 19(1):258–279, January 2024. ISSN 1745-6916. doi: 10.1177/17456916231180580. URL <https://doi.org/10.1177/17456916231180580>. Publisher: SAGE Publications Inc. 5, 5.3.1, 5.5, 5.6.1, 6.2.1, 6.2.2
- [90] Eric Griffith. Social Privacy Is on the Rise: Almost Half of Social Media Accounts Are Kept Private, June 2020. URL <https://www.pcmag.com/news/social-privacy-is-on-the-rise-almost-half-of-social-media-accounts-are>. 3.4.1
- [91] Andrew Guess, Brendan Nyhan, and Jason Reifler. Selective Exposure to Misinformation: Evidence from the consumption of fake news during the 2016 U.S. presidential campaign. Technical report, 2018. URL <https://about.fb.com/wp-content/uploads/2018/01/fake-news-2016.pdf>. 3.3.1
- [92] Andrew M. Guess, Michael Lerner, Benjamin Lyons, Jacob M. Montgomery, Brendan Nyhan, Jason Reifler, and Neelanjan Sircar. A digital media literacy intervention increases discernment between mainstream and false news in the United States and India. *Proceedings of the National Academy of Sciences*, 117(27):15536–15545, July 2020. doi: 10.1073/pnas.1920498117. 4.1, 4.2.1, 5.1, 6.2.3, A.5
- [93] Paweł Gwiaździński, Aleksander B. Gundersen, Michał Piąka, Izabela Krysińska, Jonas R. Kunst, Karolina Noworyta, Agata Olejniuk, Mikołaj Morzy, Rafal Rygula, Tomi Wójtowicz, and Jan Piasecki. Psychological interventions countering misinformation in social media: A scoping review. *Frontiers in Psychiatry*, 13, 2023. ISSN 1664-0640. doi: 10.3389/fpsyg.2022.974782. 3.3.1, 5.2.1
- [94] Désirée Hagmann, Michael Siegrist, and Christina Hartmann. Taxes, labels, or nudges? Public acceptance of various interventions designed to reduce sugar intake. *Food Policy*, 79:156–165, August 2018. ISSN 0306-9192. doi: 10.1016/

- j.foodpol.2018.06.008. URL <https://www.sciencedirect.com/science/article/pii/S0306919217310096>. 5.5, 6.2.2
- [95] Stuart Hall. Encoding and decoding in the television discourse. Technical report, Birmingham: Centre for Contemporary Cultural Studies, 1973. 1.3.2
- [96] David J. Hauser and Norbert Schwarz. Attentive Turkers: MTurk participants perform better on online attention checks than do subject pool participants. *Behavior Research Methods*, 48(1):400–407, March 2016. ISSN 1554-3528. doi: 10.3758/s13428-015-0578-z. URL <https://doi.org/10.3758/s13428-015-0578-z>. 3.2.3
- [97] David J. Hauser, Aaron J. Moss, Cheskie Rosenzweig, Shalom N. Jaffe, Jonathan Robinson, and Leib Litman. Evaluating CloudResearch’s Approved Group as a solution for problematic data quality on MTurk. *Behavior Research Methods*, 55(8):3953–3964, December 2023. ISSN 1554-3528. doi: 10.3758/s13428-022-01999-x. URL <https://doi.org/10.3758/s13428-022-01999-x>. 3.2.3
- [98] Moreen Heine, Susanna Kuper, and Thomas Neururer. Which Platform to Use?: Social Media Platform Types and their Suitability for Sound Decision Making by Voluntary Helpers. In *Proceedings of the 11th International Conference on Theory and Practice of Electronic Governance*, pages 395–402, Galway Ireland, April 2018. ACM. ISBN 978-1-4503-5421-9. doi: 10.1145/3209415.3209485. URL <https://dl.acm.org/doi/10.1145/3209415.3209485>. 3.4.1, 3.4.1
- [99] Todd C. Helmus and Bilva Chandra. Generative Artificial Intelligence Threats to Information Integrity and Potential Policy Responses. Technical report, RAND Corporation, April 2024. URL <https://www.rand.org/pubs/perspectives/PEA3089-1.html>. A.8
- [100] Todd C. Helmus and Marta Kepe. A Compendium of Recommendations for Counteracting Russian and Other State-Sponsored Propaganda. Technical report, RAND Corporation, June 2021. URL [https://www.rand.org/pubs/research\\_reports/RRA894-1.html](https://www.rand.org/pubs/research_reports/RRA894-1.html). 1.4.1, 2.5, 4.1, 5.1, 5.2.1, 5.1, A.1
- [101] David M. Howcroft and Verena Rieser. What happens if you treat ordinal ratings as interval data? Human evaluations in NLP are even more under-powered than you think. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 8932–8939, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-main.703. URL <https://aclanthology.org/2021.emnlp-main.703>. 3.3.3
- [102] Guanxiong Huang, Wufan Jia, and Wenting Yu. Media Literacy Interventions Improve Resilience to Misinformation: A Meta-Analytic Investigation of Overall Effect and Moderating Factors. *Communication Research*, page 00936502241288103, October 2024. ISSN 0093-6502. doi: 10.1177/00936502241288103. URL <https://doi.org/10.1177/00936502241288103>. Publisher: SAGE Publications Inc. 6.2.3
- [103] Robert A. Huber, Michael L. Wicki, and Thomas Bernauer. Public support for environmental policy depends on beliefs concerning effectiveness, intrusiveness, and fair-

- ness. *Environmental Politics*, 29(4):649–673, June 2020. ISSN 0964-4016. doi: 10.1080/09644016.2019.1629171. URL <https://www.tandfonline.com/doi/10.1080/09644016.2019.1629171>. Publisher: Routledge. 5, 5.2, 5.3.1, 5.5
- [104] Søren Højsgaard, Ulrich Halekoh, and Jun Yan. The R Package geepack for Generalized Estimating Equations. *Journal of Statistical Software*, 15:1–11, 2006. ISSN 1548-7660. doi: 10.18637/jss.v015.i02. URL <https://doi.org/10.18637/jss.v015.i02>. D
  - [105] Kristin M. Jackson and William M. K. Trochim. Concept Mapping as an Alternative Approach for the Analysis of Open-Ended Survey Responses. *Organizational Research Methods*, 5(4):307–336, October 2002. ISSN 1094-4281. doi: 10.1177/109442802237114. URL <https://doi.org/10.1177/109442802237114>. Publisher: SAGE Publications Inc. 4.3.5
  - [106] Se-Hoon Jeong, Hyunyi Cho, and Yoori Hwang. Media Literacy Interventions: A Meta-Analytic Review. *The Journal of Communication*, 62(3):454–472, June 2012. ISSN 0021-9916. doi: 10.1111/j.1460-2466.2012.01643.x. 1.4.1, 4.1, 4.2.1, A.5
  - [107] Shagun Jhaver and Amy X. Zhang. Do users want platform moderation or individual control? Examining the role of third-person effects and free speech support in shaping moderation preferences. *New Media & Society*, page 14614448231217993, December 2023. ISSN 1461-4448. doi: 10.1177/14614448231217993. URL <https://doi.org/10.1177/14614448231217993>. Publisher: SAGE Publications. 4.2.2, A.2
  - [108] Jialun Aaron Jiang, Peipei Nie, Jed R. Brubaker, and Casey Fiesler. A Trade-off-centered Framework of Content Moderation. *ACM Transactions on Computer-Human Interaction*, 30(1):3:1–3:34, March 2023. ISSN 1073-0516. doi: 10.1145/3534929. URL <https://dl.acm.org/doi/10.1145/3534929>. 1.4.1, 5.1, 6.2.1, A.2
  - [109] Amelia Johns, Francesco Bailo, Emily Booth, and Marian-Andrei Rizoiu. Labelling, shadow bans and community resistance: did Meta’s strategy to suppress rather than remove COVID misinformation and conspiracy theory on Facebook slow the spread? *Media International Australia*, page 1329878X241236984, March 2024. ISSN 1329-878X, 2200-467X. doi: 10.1177/1329878X241236984. URL <https://journals.sagepub.com/doi/10.1177/1329878X241236984>. A.3
  - [110] Christopher M. Jones, Daniel Diethei, Johannes Schöning, Rehana Shrestha, Tina Jahnel, and Benjamin Schüz. Impact of Social Reference Cues on Misinformation Sharing on Social Media: Series of Experimental Studies. *Journal of Medical Internet Research*, 25(1):e45583, August 2023. doi: 10.2196/45583. URL <https://www.jmir.org/2023/1/e45583>. Company: Journal of Medical Internet Research Distributor: Journal of Medical Internet Research Institution: Journal of Medical Internet Research Label: Journal of Medical Internet Research Publisher: JMIR Publications Inc., Toronto, Canada. 3.5
  - [111] S. Mo Jones-Jang, Tara Mortensen, and Jingjing Liu. Does Media Literacy Help Identification of Fake News? Information Literacy Helps, but Other Literacies Don’t. *American Behavioral Scientist*, 65(2):371–388, February 2021. ISSN 0002-7642. doi: 10.1177/0002764219869406. URL <https://doi.org/10.1177/0002764219869406>.

Publisher: SAGE Publications Inc. 2.2, A.5

- [112] Steffen Kallbekken, Jorge H. Garcia, and Kristine Korneliussen. Determinants of public support for transport taxes. *Transportation Research Part A: Policy and Practice*, 58: 67–78, December 2013. ISSN 09658564. doi: 10.1016/j.tra.2013.10.004. URL <https://linkinghub.elsevier.com/retrieve/pii/S0965856413001845>. 5.3.1, 5.5
- [113] Sarawut Kankham and Jian-Ren Hou. Community Notes vs. Related Articles: Assessing Real-World Integrated Counter-Rumor Features in Response to Different Rumor Types on Social Media. *International Journal of Human–Computer Interaction*, pages 1–15, September 2024. ISSN 1044-7318. doi: 10.1080/10447318.2024.2400389. URL <https://doi.org/10.1080/10447318.2024.2400389>. Publisher: Taylor & Francis \_eprint: <https://doi.org/10.1080/10447318.2024.2400389>. 5.6.1
- [114] Joel Kaplan. More Speech and Fewer Mistakes, January 2025. URL <https://about.fb.com/news/2025/01/meta-more-speech-fewer-mistakes/>. 1.5.2, 6.2.3
- [115] Matthew Katsaros, Kathy Yang, and Lauren Fratamico. Reconsidering Tweets: Intervening during Tweet Creation Decreases Offensive Content. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 16 of *ICWSM*, pages 477–487. Association for the Advancement of Artificial Intelligence, May 2022. doi: 10.1609/icwsm.v16i1.19308. URL <https://ojs.aaai.org/index.php/ICWSM/article/view/19308>. 5.1, A.1
- [116] Kulvinder Kaur and Samrat Gupta. Towards dissemination, detection and combating misinformation on social media: a literature review. *Journal of Business & Industrial Marketing*, 38(8):1656–1674, January 2022. ISSN 0885-8624. doi: 10.1108/JBIM-02-2022-0066. URL <https://doi.org/10.1108/JBIM-02-2022-0066>. Publisher: Emerald Publishing Limited. 3.1
- [117] Tanveer Khan, Antonis Michalas, and Adnan Akhunzada. Fake news outbreak 2021: Can we stop the viral spread? *Journal of Network and Computer Applications*, 190:103112, September 2021. ISSN 1084-8045. doi: 10.1016/j.jnca.2021.103112. URL <https://www.sciencedirect.com/science/article/pii/S1084804521001326>. A.2
- [118] Eugene Kiely and Lori Robertson. How to Spot Fake News, November 2016. URL <https://www.factcheck.org/2016/11/how-to-spot-fake-news/>. 4.3.4
- [119] Jan H. Kietzmann, Kristopher Hermkens, Ian P. McCarthy, and Bruno S. Silvestre. Social media? Get serious! Understanding the functional building blocks of social media. *Business Horizons*, 54(3):241–251, May 2011. ISSN 0007-6813. doi: 10.1016/j.bushor.2011.01.005. URL <https://www.sciencedirect.com/science/article/pii/S0007681311000061>. 3.4.1
- [120] Jisu Kim, Curtis McDonald, Paul Meosky, Matthew Katsaros, and Tom Tyler. Promoting Online Civility Through Platform Architecture. *Journal of Online Trust and Safety*, 1 (4), September 2022. ISSN 2770-3142. doi: 10.54501/jots.v1i4.54. URL <https://doi.org/10.54501/jots.v1i4.54>.

tsjournal.org/index.php/jots/article/view/54. Number: 4. 2.2, A.1

- [121] Catherine King, Christine Sowa Lepird, and Dr Kathleen M Carley. Project OMEN: Designing a Training Game to Fight Misinformation on Social Media. Technical Report CMU-ISR-21-110, Carnegie Mellon University, August 2021. URL <http://reports-archive.adm.cs.cmu.edu/anon/isr2021/abstracts/21-110.html>. 4.3, 4.3.1
- [122] Catherine King, Samantha Phillips, and Kathleen Carley. Registered Report: A path forward on online misinformation mitigation based on current user behavior [Registered Report Stage 1 protocol]. *Scientific Reports*, 2024. URL <https://figshare.com/s/683b1e7c2f2bad96f604>. Publisher: figshare. 4.3.3
- [123] Jan Kirchner and Christian Reuter. Countering Fake News: A Comparison of Possible Solutions Regarding User Acceptance and Effectiveness. *Proceedings of the ACM on Human-Computer Interaction*, 4(CSCW2):1–27, October 2020. ISSN 2573-0142. doi: 10.1145/3415211. URL <https://dl.acm.org/doi/10.1145/3415211>. 3.3.1, 5.1, 6.4.2, A.1
- [124] Marcus Kolga. Stemming the VIRUS: Understanding and responding to the threat of Russian disinformation. Technical report, Macdonald-Laurier Institute Publication, January 2019. 5.1
- [125] Nadejda Komendantova, Love Ekenberg, Mattias Svahn, Aron Larsson, Syed Iftikhar Hussain Shah, Myrsini Glinos, Vasilis Koulolias, and Mats Danielson. A value-driven approach to addressing misinformation in social media. *Humanities and Social Sciences Communications*, 8(1):1–12, January 2021. ISSN 2662-9992. doi: 10.1057/s41599-020-00702-9. URL <https://www.nature.com/articles/s41599-020-00702-9>. Number: 1 Publisher: Palgrave. 5
- [126] Vasilis Koulolias, Gideon Mekonnen Jonathan, Miriam Fernandez, and Dimitris Sotirchos. Combating Misinformation: An ecosystem in co-creation. Technical report, OECD Publishing, January 2018. 5, 5.1
- [127] Anastasia Kozyreva, Stephan Lewandowsky, and Ralph Hertwig. Citizens Versus the Internet: Confronting Digital Challenges With Cognitive Tools. *Psychological Science in the Public Interest*, 21(3):103–156, 2020. doi: <https://doi.org/10.1177/1529100620946707>. 3.3.1
- [128] Anastasia Kozyreva, Stefan M. Herzog, Stephan Lewandowsky, Ralph Hertwig, Philipp Lorenz-Spreen, Mark Leiser, and Jason Reifler. Resolving content moderation dilemmas between free speech and harmful misinformation. *Proceedings of the National Academy of Sciences*, 120(7):e2210666120, February 2023. doi: 10.1073/pnas.2210666120. URL <https://www.pnas.org/doi/10.1073/pnas.2210666120>. Publisher: Proceedings of the National Academy of Sciences. 3.3.1, 3.3.4
- [129] Anastasia Kozyreva, Philipp Lorenz-Spreen, Stefan M. Herzog, Ullrich K. H. Ecker, Stephan Lewandowsky, Ralph Hertwig, Ayesha Ali, Joe Bak-Coleman, Sarit Barzilai, Melisa Basol, Adam J. Berinsky, Cornelia Betsch, John Cook, Lisa K. Fazio,

Michael Geers, Andrew M. Guess, Haifeng Huang, Horacio Larreguy, Rakoen Maertens, Folco Panizza, Gordon Pennycook, David G. Rand, Steve Rathje, Jason Reifler, Philipp Schmid, Mark D. Smith, Briony Swire-Thompson, Paula Szewach, Sander van der Linden, and Sam Wineburg. Toolbox of individual-level interventions against online misinformation. *Nature Human Behaviour*, pages 1–9, May 2024. ISSN 2397-3374. doi: 10.1038/s41562-024-01881-0. URL <https://www.nature.com/articles/s41562-024-01881-0>. Publisher: Nature Publishing Group. 1.4.1, 1.4, 1.4.1, 1.4.2, 2.1, 2.4.4, 4.1

- [130] Michael E. Kraft and Scott R. Furlong. *Public Policy: Politics, Analysis, and Alternatives*. CQ Press, 6th edition, April 2017. ISBN 978-1-5063-5814-7. Google-Books-ID: VWRCDgAAQBAJ. 6.1, 6.2.3
- [131] John K. Kruschke and Torrin M. Liddell. The Bayesian New Statistics: Hypothesis testing, estimation, meta-analysis, and power analysis from a Bayesian perspective. *Psychonomic Bulletin & Review*, 25(1):178–206, February 2018. ISSN 1531-5320. doi: 10.3758/s13423-016-1221-4. URL <https://doi.org/10.3758/s13423-016-1221-4>. 3.3.2
- [132] Joris Lammers, Diederik A. Stapel, and Adam D. Galinsky. Power Increases Hypocrisy: Moralizing in Reasoning, Immorality in Behavior. *Psychological Science*, 21(5):737–744, May 2010. ISSN 0956-7976. doi: 10.1177/0956797610368810. URL <https://doi.org/10.1177/0956797610368810>. Publisher: SAGE Publications Inc. 3.3.1
- [133] Darren Langridge and Trevor Butt. The fundamental attribution error: A phenomenological critique. *British Journal of Social Psychology*, 43(3):357–369, 2004. ISSN 2044-8309. doi: 10.1348/0144666042037962. URL <https://onlinelibrary.wiley.com/doi/abs/10.1348/0144666042037962>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1348/0144666042037962>. 3.3.1
- [134] M. Asher Lawson, Shikhar Anand, and Hemant Kakkar. Tribalism and tribulations: The social costs of not sharing fake news. *Journal of Experimental Psychology: General*, 152(3):611–631, March 2023. ISSN 1939-2222, 0096-3445. doi: 10.1037/xge0001374. URL <https://doi.apa.org/doi/10.1037/xge0001374>. 3.5
- [135] David Lazer and Sandra González-Bailón. Mark Zuckerberg’s Immoderate Proposal | TechPolicy.Press, January 2025. URL <https://techpolicy.press/mark-zuckerbergs-immoderate-proposal>. 6.2.3
- [136] David M. J. Lazer, Matthew A. Baum, Yochai Benkler, Adam J. Berinsky, Kelly M. Greenhill, Filippo Menczer, Miriam J. Metzger, Brendan Nyhan, Gordon Pennycook, David Rothschild, Michael Schudson, Steven A. Sloman, Cass R. Sunstein, Emily A. Thorson, Duncan J. Watts, and Jonathan L. Zittrain. The science of fake news. *Science*, 359(6380):1094–1096, March 2018. doi: 10.1126/science.aaq2998. URL <https://www.science.org/doi/full/10.1126/science.aaq2998>. Publisher: American Association for the Advancement of Science. 1.1, 1.2
- [137] Michael D Lee and Eric-Jan Wagenmakers. *Bayesian Cognitive Modeling: A Practical Course*. Cambridge University Press, 2013. 3.3.2, 3.3.4

- [138] Jeffrey Lees, John A Banas, Darren Linvill, Patrick C Meirick, and Patrick Warren. The Spot the Troll Quiz game increases accuracy in discerning between real and inauthentic social media accounts. *PNAS Nexus*, 2(4):pgad094, April 2023. ISSN 2752-6542. doi: 10.1093/pnasnexus/pgad094. 4.1, 4.2.1
- [139] Christine Sowa Lepird. *Digital Pink Slime: Measuring, Finding, and Countering Online Threats to Local News*. PhD thesis, Carnegie Mellon University, November 2024. URL [https://christine-lepird.github.io/Lepird\\_Thesis\\_Oct2024.pdf](https://christine-lepird.github.io/Lepird_Thesis_Oct2024.pdf). 4.3.3, 4.3.4
- [140] Jimmie Leppink, Patricia O’Sullivan, and Kal Winston. Effect size – large, medium, and small. *Perspectives on Medical Education*, 5(6):347–349, December 2016. ISSN 2212-277X. doi: 10.1007/s40037-016-0308-y. URL <https://doi.org/10.1007/s40037-016-0308-y>. 3.3.2
- [141] Stephan Lewandowsky and John Cook. The Conspiracy Theory Handbook. Technical report, March 2020. URL <http://sks.to/conspiracy>. 3
- [142] Stephan Lewandowsky and Sander van der Linden. Countering Misinformation and Fake News Through Inoculation and Prebunking. *European Review of Social Psychology*, 0 (0):1–38, February 2021. ISSN 1046-3283. doi: 10.1080/10463283.2021.1876983. 1.1, 4.2.1, A.5
- [143] Stephan Lewandowsky, Ullrich K. H. Ecker, Colleen M. Seifert, Norbert Schwarz, and John Cook. Misinformation and Its Correction: Continued Influence and Successful Debiasing. *Psychological Science in the Public Interest*, 13(3):106–131, December 2012. ISSN 1529-1006. doi: 10.1177/1529100612451018. 1.3.2, 1.3.2
- [144] Stephan Lewandowsky, John Cook, Ullrich Ecker, Dolores Albarracín, Michelle A. Amazeen, Panayiota Kendeou, Doug Lombardi, Eryn J. Newman, Gordon Pennycook, Ethan Porter, David G. Rand, David N. Rapp, Jason Reifler, Jon Roozenbeek, Philipp Schmid, Colleen M. Seifert, Gale M. Sinatra, Briony Swire-Thompson, Sander van der Linden, Emily K. Vraga, Thomas J. Wood, and Maria S. Zaragoza. Debunking Handbook 2020, 2020. URL <https://sks.to/db2020>. 1.4.1, 3, 3.1, 4.2.2, 1, 3, A.2
- [145] Yi Liu, T Pinar Yildirim, and Z John Zhang. Implications of Revenue Models and Technology for Content Moderation Strategies. *SSRN Electronic Journal*, December 2021. doi: 10.2139/ssrn.3969938. URL [https://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=3969938](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3969938). 5.1
- [146] Zhuoran Lu, Patrick Li, Weilong Wang, and Ming Yin. The Effects of AI-based Credibility Indicators on the Detection and Spread of Misinformation under Social Influence. In *Proceedings of the ACM on Human-Computer Interaction*, volume 6 of *CSCW*, pages 461:1–461:27. Association for Computing Machinery, November 2022. doi: 10.1145/3555562. URL <https://dl.acm.org/doi/10.1145/3555562>. 2.2, A.2
- [147] Rakoen Maertens, Jon Roozenbeek, Melisa Basol, and Sander van der Linden. Long-term effectiveness of inoculation against misinformation: Three longitudinal experiments. *Journal of Experimental Psychology: Applied*, 27(1):1–16, 2021. ISSN 1939-2192. doi: 10.1037/xap0000315. 2.2, 4.2.1, 6.2.3, A.5

- [148] Cameron Martel, Adam J. Berinsky, David Rand, Amy Xian Zhang, and Paul Resnick. Perceived legitimacy of layperson and expert content moderators, June 2024. URL [https://osf.io/5n9tp\\_v1](https://osf.io/5n9tp_v1). 6.2.3
- [149] Justin D Martin and Fouad Hassan. Testing Classical Predictors of Public Willingness to Censor on the Desire to Block Fake News Online. *Convergence: The International Journal of Research into New Media Technologies*, 28(3):867–887, June 2022. ISSN 1354-8565, 1748-7382. doi: 10.1177/13548565211012552. URL <https://journals.sagepub.com/doi/10.1177/13548565211012552>. 5.2
- [150] Francisco J. Martínez-López, Yangchun Li, and Susan M. Young. Social Media Monetization and Demonetization: Risks, Challenges, and Potential Solutions. In Francisco J. Martínez-López, Yangchun Li, and Susan M. Young, editors, *Social Media Monetization: Platforms, Strategic Models and Critical Success Factors*, pages 185–214. Springer International Publishing, Cham, 2022. ISBN 978-3-031-14575-9. doi: 10.1007/978-3-031-14575-9\_13. URL [https://doi.org/10.1007/978-3-031-14575-9\\_13](https://doi.org/10.1007/978-3-031-14575-9_13). A.3
- [151] Sarah McGrew. Learning to evaluate: An intervention in civic online reasoning. *Computers & Education*, 145:103711, February 2020. ISSN 0360-1315. doi: 10.1016/j.compedu.2019.103711. 4.1
- [152] Paul Mena. Cleaning Up Social Media: The Effect of Warning Labels on Likelihood of Sharing False News on Facebook. *Policy & Internet*, 12(2): 165–183, 2020. ISSN 1944-2866. doi: 10.1002/poi3.214. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/poi3.214>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/poi3.214>. 5.1
- [153] Amy Mitchell and Mason Walker. More Americans now say government should take steps to restrict false information online than in 2018, August 2021. URL <https://www.pewresearch.org/fact-tank/2021/08/18/more-americans-now-say-government-should-take-steps-to-restrict-false-information-online-than-in-2018/>. 3.3.1, 3.3.4, 5.2
- [154] Ariana Modirrousta-Galian and Philip A. Higham. Gamified inoculation interventions do not improve discrimination between true and fake news: Reanalyzing existing research with receiver operating characteristic analysis. *Journal of Experimental Psychology: General*, 152(9):2411–2437, 2023. ISSN 1939-2222. doi: 10.1037/xge0001395. Place: US Publisher: American Psychological Association. 2.6.2, 4.2.1, A.5
- [155] Richard D. Morey, Jeffrey N. Rouder, Tahira Jamil, Urbanek, Karl Forner, and Alexander Ly. Package ‘BayesFactor’, 2022. URL <https://richarddmorey.github.io/BayesFactor/>. 3.3.2, 3.3.3
- [156] Garrett Morrow, Briony Swire-Thompson, Jessica Montgomery Polny, Matthew Kopec, and John P. Wihbey. The emerging science of content labeling: Contextualizing social media content moderation. *Journal of the Association for Information Science and Technology*, 73(10):1365–1386, 2022. ISSN 2330-1643. doi: 10.1002/asi.24637. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/asi.24637>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/asi.24637>. A.4

- [157] Mohsen Mosleh, Cameron Martel, Dean Eckles, and David G. Rand. Perverse Downstream Consequences of Debunking: Being Corrected by Another User for Posting False Political News Increases Subsequent Sharing of Low Quality, Partisan, and Toxic Content in a Twitter Field Experiment. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, CHI, pages 1–13, New York, NY, USA, May 2021. Association for Computing Machinery. ISBN 978-1-4503-8096-6. doi: 10.1145/3411764.3445642. URL <https://dl.acm.org/doi/10.1145/3411764.3445642>. 2.6.2
- [158] Mohsen Mosleh, Jennifer Nancy Lee Allen, and David Rand. Divergent patterns of engagement with partisan and low-quality news across seven social media platforms, December 2024. URL <https://osf.io/9csy3>. 3.4.1, 3.4.1
- [159] Joris Mulder, Donald R. Williams, Xin Gu, Andrew Tomarken, Florian Böing-Messing, Anton Olsson-Collentine, Marlyne Meijerink, Janosch Menke, Robbie Van Aert, Jean-Paul Fox, Herbert Hoijtink, Yves Rosseel, Eric-Jan Wagenmakers, and Caspar Van Lissa. **BFpack** : Flexible Bayes Factor Testing of Scientific Theories in R. *Journal of Statistical Software*, 100(18), 2021. ISSN 1548-7660. doi: 10.18637/jss.v100.i18. URL <https://www.jstatsoft.org/v100/i18/>. 5.3.3
- [160] Xiaoli Nan, Yuan Wang, and Kathryn Thier. Why do people believe health misinformation and who is at risk? A systematic review of individual differences in susceptibility to health misinformation. *Social Science & Medicine*, 314:115398, December 2022. ISSN 0277-9536. doi: 10.1016/j.socscimed.2022.115398. URL <https://www.sciencedirect.com/science/article/pii/S0277953622007043>. 3.3.1, 3.5
- [161] Jack Nassetta and Kimberly Gross. State media warning labels can counteract the effects of foreign disinformation. *Harvard Kennedy School Misinformation Review*, 1(Special Issue on US Elections and Disinformation), October 2020. doi: 10.37016/mr-2020-45. URL <https://misinforeview.hks.harvard.edu/article/state-media-warning-labels-can-counteract-the-effects-of-foreign-misinformation/>. 2.2, A.4
- [162] Andrew A. Neath, Javier E. Flores, and Joseph E. Cavanaugh. Bayesian multiple comparisons and model selection. *WIREs Computational Statistics*, 10(2):e1420, 2018. ISSN 1939-0068. doi: 10.1002/wics.1420. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/wics.1420>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/wics.1420>. 3.3.2
- [163] Sheryl Wei Ting Ng. Self- and Social Corrections on Instant Messaging Platforms. *International Journal Of Communication*, 17:426–446, 2023. URL <https://ijoc.org/index.php/ijoc/article/view/19142/4005>. 3.3.1, 3.4.1, 1
- [164] Konrad Niklewicz. Weeding Out Fake News: An Approach to Social Media Regulation. *European View*, 16(2):335–335, December 2017. ISSN 1781-6858, 1865-5831. doi: 10.1007/s12290-017-0468-0. 4.2.2, A.6, A.7
- [165] Brendan Nyhan, Jason Reifler, Sean Richey, and Gary L. Freed. Effective Messages in

- Vaccine Promotion: A Randomized Trial. *Pediatrics*, 133(4):e835–e842, April 2014. ISSN 0031-4005. doi: 10.1542/peds.2013-2365. URL <https://doi.org/10.1542/peds.2013-2365>. 2.6.2
- [166] Chitu Okoli. A Guide to Conducting a Standalone Systematic Literature Review. *Communications of the Association for Information Systems*, 37, 2015. ISSN 15293181. doi: 10.17705/1CAIS.03743. URL <https://aisel.aisnet.org/cais/vol37/iss1/43/>. 2.2, 2.2.4
- [167] Tomasz Oleksy, Anna Wnuk, Dominika Maison, and Agnieszka Łyś. Content matters. Different predictors and social consequences of general and government-related conspiracy theories on COVID-19. *Personality and Individual Differences*, 168:110289, January 2021. ISSN 01918869. doi: 10.1016/j.paid.2020.110289. URL <https://linkinghub.elsevier.com/retrieve/pii/S0191886920304797>. 1.1
- [168] Gábor Orosz, Péter Krekó, Benedek Paskuj, István Tóth-Király, Beáta Bőthe, and Christine Roland-Lévy. Changing Conspiracy Beliefs through Rationality and Ridiculing. *Frontiers in Psychology*, 7, October 2016. ISSN 1664-1078. doi: 10.3389/fpsyg.2016.01525. URL <https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2016.01525/full>. Publisher: Frontiers. 3
- [169] Wei Pan. Akaike's Information Criterion in Generalized Estimating Equations. *Biometrics*, 57(1):120–125, March 2001. ISSN 0006-341X. doi: 10.1111/j.0006-341X.2001.00120.x. URL <https://doi.org/10.1111/j.0006-341X.2001.00120.x>. D
- [170] Orestis Papakyriakopoulos and Ellen Goodman. The Impact of Twitter Labels on Misinformation Spread and User Engagement: Lessons from Trump's Election Tweets. In *Proceedings of the ACM Web Conference 2022*, WWW '22, pages 2541–2551, New York, NY, USA, April 2022. Association for Computing Machinery. ISBN 978-1-4503-9096-5. doi: 10.1145/3485447.3512126. URL <https://doi.org/10.1145/3485447.3512126>. 5.1, A.4
- [171] Eyal Peer, Joachim Vosgerau, and Alessandro Acquisti. Reputation as a sufficient condition for data quality on Amazon Mechanical Turk. *Behavior Research Methods*, 46(4):1023–1031, December 2014. ISSN 1554-3528. doi: 10.3758/s13428-013-0434-y. URL <https://doi.org/10.3758/s13428-013-0434-y>. 3.2.3
- [172] Mark J. Pelletier, Alexandra Krallman, Frank G. Adams, and Tyler Hancock. One size doesn't fit all: a uses and gratifications analysis of social media platforms. *Journal of Research in Interactive Marketing*, 14(2):269–284, June 2020. ISSN 2040-7122. doi: 10.1108/JRIM-10-2019-0159. URL <https://www.emerald.com/insight/content/doi/10.1108/jrim-10-2019-0159/full/html>. Publisher: Emerald Publishing Limited. 3.4.1
- [173] Gordon Pennycook and David G. Rand. Lazy, not biased: Susceptibility to partisan fake news is better explained by lack of reasoning than by motivated reasoning. *Cognition*, 188:39–50, July 2019. ISSN 0010-0277. doi: 10.1016/j.cognition.2018.06.011. URL <http://www.sciencedirect.com/science>

article/pii/S001002771830163X. 3.5

- [174] Gordon Pennycook and David G. Rand. Accuracy prompts are a replicable and generalizable approach for reducing the spread of misinformation. *Nature Communications*, 13(1):2333, April 2022. ISSN 2041-1723. doi: 10.1038/s41467-022-30073-5. URL <https://www.nature.com/articles/s41467-022-30073-5>. Publisher: Nature Publishing Group. 6.4.2
- [175] Gordon Pennycook and David G. Rand. Nudging Social Media toward Accuracy. *The ANNALS of the American Academy of Political and Social Science*, 700(1):152–164, March 2022. ISSN 0002-7162, 1552-3349. doi: 10.1177/00027162221092342. URL <https://journals.sagepub.com/doi/10.1177/00027162221092342>. 5.1, 5.5
- [176] Gordon Pennycook, Jonathon McPhetres, Yunhao Zhang, Jackson G. Lu, and David G. Rand. Fighting COVID-19 Misinformation on Social Media: Experimental Evidence for a Scalable Accuracy-Nudge Intervention. *Psychological Science*, 31(7):770–780, July 2020. ISSN 0956-7976. doi: 10.1177/0956797620939054. URL <https://doi.org/10.1177/0956797620939054>. Publisher: SAGE Publications Inc. 3.3.2
- [177] Gordon Pennycook, Ziv Epstein, Mohsen Mosleh, Antonio A. Arechar, Dean Eckles, and David G. Rand. Shifting attention to accuracy can reduce misinformation online. *Nature*, pages 1–6, March 2021. ISSN 1476-4687. doi: 10.1038/s41586-021-03344-2. URL <https://www.nature.com/articles/s41586-021-03344-2>. Publisher: Nature Publishing Group. 3.1, 3.5, 5.1, A.1
- [178] Paul K. Piff, Michael W. Kraus, Stéphane Côté, Bonnie Hayden Cheng, and Dacher Keltner. Having less, giving more: The influence of social class on prosocial behavior. *Journal of Personality and Social Psychology*, 99(5):771–784, 2010. ISSN 1939-1315, 0022-3514. doi: 10.1037/a0020092. URL <https://doi.apa.org/doi/10.1037/a0020092>. 3.3.5
- [179] Mark Aaron Polger. Misinformation and Disinformation: Thinking Critically about Information Sources, October 2024. URL <https://library.csi.cuny.edu/misinformation/spotfakenews>. 4.3.4
- [180] Ethan Porter and Thomas J. Wood. The global effectiveness of fact-checking: Evidence from simultaneous experiments in Argentina, Nigeria, South Africa, and the United Kingdom. *Proceedings of the National Academy of Sciences*, 118(37):e2104235118, September 2021. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.2104235118. URL <https://pnas.org/doi/full/10.1073/pnas.2104235118>. 1.1, 6.2.3
- [181] Jon Porter. WhatsApp says its forwarding limits have cut the spread of viral messages by 70 percent, April 2020. URL <https://www.theverge.com/2020/4/27/21238082/whatsapp-forward-message-limits-viral-misinformation-decline>. 6.4.2, A.1
- [182] Carlie Porterfield. Twitter Begins Asking Users To Actually Read Articles Before Sharing Them, June 2020. URL <https://www.forbes.com/sites/carlieporterfield/2020/06/10/twitter-begins-asking-users-to-actually-read-articles-before-sharing-them/>. 5.1, 5.5, A.1

- [183] Elaine S. Povich. Internet Ads Are a Popular Tax Target for Both Parties, June 2021. URL <https://pew.org/3psT1Rq>. A.7
- [184] Roxana Radu. Fighting the ‘Infodemic’: Legal Responses to COVID-19 Disinformation. *Social Media + Society*, 6(3):2056305120948190, July 2020. ISSN 2056-3051. doi: 10.1177/2056305120948190. URL <https://doi.org/10.1177/2056305120948190>. Publisher: SAGE Publications Ltd. 5.1, 5.2
- [185] David Rand and Cameron Martel. Americans actually do want expert content moderation, January 2025. URL <https://thehill.com/opinion/technology/5109667-sorry-zuckerberg-americans-actually-do-want-expert-content-moderation/>. 6.2.3
- [186] Adrian Rauchfleisch and Jonas Kaiser. The impact of deplatforming the far right: an analysis of YouTube and BitChute. *Information, Communication & Society*, 27(7):1478–1496, May 2024. ISSN 1369-118X, 1468-4462. doi: 10.1080/1369118X.2024.2346524. URL <https://www.tandfonline.com/doi/full/10.1080/1369118X.2024.2346524>. 5.1, A.3
- [187] Emma Remy. How public and private Twitter users in the U.S. compare — and why it might matter for your research, July 2019. URL <https://www.pewresearch.org/decoded/2019/07/how-public-and-private-twitter-users-in-the-u-s-compare-and-why-it-might-matter-for-your-research/>. 3.4.1
- [188] J. P. Reynolds, K. Stautz, M. Pilling, S. van der Linden, and T. M. Marteau. Communicating the effectiveness and ineffectiveness of government policies and their impact on public support: a systematic review with meta-analysis. *Royal Society Open Science*, 7(1):190522, January 2020. doi: 10.1098/rsos.190522. URL <https://royalsocietypublishing.org/doi/10.1098/rsos.190522>. Publisher: Royal Society. 5.5, 6.2.2
- [189] Timothy S. Rich, Ian Milden, and Mallory Treece Wagner. Research note: Does the public support fact-checking social media? It depends who and how you ask. *Harvard Kennedy School Misinformation Review*, November 2020. doi: 10.37016/mr-2020-46. URL <https://misinforeview.hks.harvard.edu/?p=3861>. 5.2
- [190] Diana Ridley. *The Literature Review: A Step-by-Step Guide for Students*. SAGE Publications Ltd, Los Angeles London New Delhi, second edition edition, July 2012. ISBN 978-1-4462-0142-8. 2.2.4
- [191] Alexander Robitzsch. Why Ordinal Variables Can (Almost) Always Be Treated as Continuous Variables: Clarifying Assumptions of Robust Continuous and Ordinal Factor Analysis Estimation Methods. *Frontiers in Education*, 5, 2020. ISSN 2504-284X. URL <https://www.frontiersin.org/articles/10.3389/feduc.2020.589965>. 3.3.3
- [192] Alex Rochefort. Regulating Social Media Platforms: A Comparative Policy Analysis. *Communication Law and Policy*, 25(2):225–260, April 2020. ISSN 1081-1680, 1532-6926. doi: 10.1080/10811680.2020.1735194. URL <https://doi.org/10.1080/10811680.2020.1735194>

[www.tandfonline.com/doi/full/10.1080/10811680.2020.1735194](http://www.tandfonline.com/doi/full/10.1080/10811680.2020.1735194). 1.6,  
5, 6.1, 6.2.3, A.7

- [193] Jon Roozenbeek and Sander van der Linden. Breaking Harmony Square: A game that “inoculates” against political misinformation. *Harvard Kennedy School Misinformation Review*, November 2020. doi: 10.37016/mr-2020-47. 4.2.1
- [194] Jon Roozenbeek, Sander van der Linden, and Thomas Nygren. Prebunking interventions based on “inoculation” theory can reduce susceptibility to misinformation across cultures. *Harvard Kennedy School Misinformation Review*, 1(2), February 2020. doi: 10.37016/mr-2020-008. 1.1, 2.6.2, A.5
- [195] Jon Roozenbeek, Sander van der Linden, Beth Goldberg, Steve Rathje, and Stephan Lewandowsky. Psychological inoculation improves resilience against misinformation on social media. *Science Advances*, 8(34):eabo6254, August 2022. doi: 10.1126/sciadv.abo6254. URL <https://www.science.org/doi/10.1126/sciadv.abo6254>. Publisher: American Association for the Advancement of Science. 5.1
- [196] Bettina Rottweiler and Paul Gill. Conspiracy Beliefs and Violent Extremist Intentions: The Contingent Effects of Self-efficacy, Self-control and Law-related Morality. *Terrorism and Political Violence*, 34(7):1485–1504, October 2020. ISSN 0954-6553. doi: 10.1080/09546553.2020.1803288. URL <https://doi.org/10.1080/09546553.2020.1803288>. Publisher: Routledge eprint: <https://doi.org/10.1080/09546553.2020.1803288>. 5.1
- [197] Jeffrey N. Rouder, Paul L. Speckman, Dongchu Sun, Richard D. Morey, and Geoffrey Iverson. Bayesian t tests for accepting and rejecting the null hypothesis. *Psychonomic Bulletin & Review*, 16(2):225–237, April 2009. ISSN 1069-9384, 1531-5320. doi: 10.3758/PBR.16.2.225. URL <http://link.springer.com/10.3758/PBR.16.2.225>. 3.3.2
- [198] Johnny Saldaña. *The Coding Manual for Qualitative Researchers*. Sage, 2nd edition, 2013. ISBN 978-1-4462-4737-2. 4.3.5
- [199] Emily Saltz, Soubhik Barari, Claire Leibowicz, and Claire Wardle. Misinformation interventions are common, divisive, and poorly understood. *Harvard Kennedy School Misinformation Review*, October 2021. doi: 10.37016/mr-2020-81. URL <https://misinforeview.hks.harvard.edu/article/misinformation-interventions-are-common-divisive-and-poorly-understood/>. 3.1, 3.3.1, 3.3.4, 3.3.5, 5.2, 5.5, 6.2.2
- [200] Brennan Schaffner, Arjun Nitin Bhagoji, Siyuan Cheng, Jacqueline Mei, Jay L. Shen, Grace Wang, Marshini Chetty, Nick Feamster, Genevieve Lakier, and Chenhao Tan. “Community Guidelines Make this the Best Party on the Internet”: An In-Depth Study of Online Platforms’ Content Moderation Policies. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, pages 1–16, May 2024. doi: 10.1145/3613904.3642333. URL <http://arxiv.org/abs/2405.05225>. arXiv:2405.05225 [cs]. 3.4.1, A.2

- [201] Mark Schlesinger and Caroline Heldman. Gender Gap or Gender Gaps? New Perspectives on Support for Government Action and Policies. *Journal of Politics*, 63(1):59–92, 2001. ISSN 1468-2508. doi: 10.1111/0022-3816.00059. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/0022-3816.00059>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1111/0022-3816.00059>. 5.5
- [202] F. D. Schönbrodt and A. M. Stefan. BFDA: An R package for Bayes factor design analysis (version 0.5.0), 2019. URL <https://github.com/nicebread/BFDA>. 3.3.2
- [203] Felix D. Schönbrodt and Eric-Jan Wagenmakers. Bayes factor design analysis: Planning for compelling evidence. *Psychonomic Bulletin & Review*, 25(1):128–142, February 2018. ISSN 1531-5320. doi: 10.3758/s13423-017-1230-y. URL <https://doi.org/10.3758/s13423-017-1230-y>. 3.3.2
- [204] Amanda Seitz and Hannah Fingerhut. Americans agree misinformation is a problem, poll shows, October 2021. URL <https://apnews.com/article/coronavirus-pandemic-technology-business-health-misinformation-fbe9d09024d7b92e1600e411d5f931dd>. Section: Coronavirus pandemic. 3.3.2
- [205] Filipo Sharevski, Raniem Alsaadi, Peter Jachim, and Emma Pieroni. Misinformation warnings: Twitter’s soft moderation effects on COVID-19 vaccine belief echoes. *Computers & Security*, 114:102577, March 2022. ISSN 0167-4048. doi: 10.1016/j.cose.2021.102577. URL <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC8675217/>. 2.2, 6.4.2, A.1, A.4
- [206] Sean Simpson. Fake News: A Global Epidemic Vast Majority (86%) of Online Global Citizens Have Been Exposed to it. Technical report, Ipsos, Toronto, Canada, June 2019. URL <https://www.ipsos.com/en-us/news-polls/cigi-fake-news-global-epidemic>. 3.3.2
- [207] Ciarra N. Smith and Holli H. Seitz. Correcting Misinformation About Neuroscience via Social Media. *Science Communication*, 41(6):790–819, December 2019. ISSN 1075-5470. doi: 10.1177/1075547019890073. URL <https://doi.org/10.1177/1075547019890073>. Publisher: SAGE Publications Inc. A.1
- [208] Brian G. Southwell, Jeff Niederdeppe, Joseph N. Cappella, Anna Gaysinsky, Danielle E. Kelley, April Oh, Emily B. Peterson, and Wen-Ying Sylvia Chou. Misinformation as a Misunderstood Challenge to Public Health. *American Journal of Preventive Medicine*, 57(2):282–285, August 2019. ISSN 07493797. doi: 10.1016/j.amepre.2019.03.009. URL <https://linkinghub.elsevier.com/retrieve/pii/S074937971930159X>. 3.3.1, 1
- [209] Jeff Stone and Nicholas C. Fernandez. To Practice What We Preach: The Use of Hypocrisy and Cognitive Dissonance to Motivate Behavior Change. *Social and Personality Psychology Compass*, 2(2):1024–1051, 2008. ISSN 1751-9004. doi: 10.1111/j.1751-9004.2008.00088.x. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1751-9004.2008.00088.x>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1111/j.1751-9004.2008.00088.x>.

- [210] Victor Suarez-Lledo and Javier Alvarez-Galvez. Prevalence of Health Misinformation on Social Media: Systematic Review. *Journal of Medical Internet Research*, 23(1): e17187, January 2021. doi: 10.2196/17187. URL <https://www.jmir.org/2021/1/e17187>. Company: Journal of Medical Internet Research Distributor: Journal of Medical Internet Research Institution: Journal of Medical Internet Research Label: Journal of Medical Internet Research Publisher: JMIR Publications Inc., Toronto, Canada. 3.4.1
- [211] Edson Tandoc, Zheng Lim, and Rich Ling. Defining “Fake News”: A typology of scholarly definitions. *Digital Journalism*, 6:1–17, August 2017. doi: 10.1080/21670811.2017.1360143. 1.2, 1.1, 1.3.1
- [212] Edson C Tandoc, Darren Lim, and Rich Ling. Diffusion of disinformation: How social media users respond to fake news and why. *Journalism*, 21(3):381–398, March 2020. ISSN 1464-8849. doi: 10.1177/1464884919868325. 3.3.1, 3.4.1, 1, A.6
- [213] Li Qian Tay, Mark J. Hurlstone, Tim Kurz, and Ullrich K. H. Ecker. A comparison of prebunking and debunking interventions for implied versus explicit misinformation. *British Journal of Psychology*, 113(3):591–607, 2022. ISSN 2044-8295. doi: 10.1111/bjop.12551. 1.4.1, 4.2.2
- [214] Josh Taylor. Threads: how do I sign up and is it any different to Twitter? *The Guardian*, July 2023. ISSN 0261-3077. URL <https://www.theguardian.com/technology/2023/jul/06/thread-how-do-i-sign-up-and-is-it-any-different-to-twitter>. 3.4.1
- [215] Philip E. Tetlock. Accountability: A Social Check on the Fundamental Attribution Error. *Social Psychology Quarterly*, 48(3):227–236, 1985. ISSN 0190-2725. doi: 10.2307/3033683. URL <https://www.jstor.org/stable/3033683>. Publisher: [Sage Publications, Inc., American Sociological Association]. 3.3.1
- [216] Daniel Robert Thomas and Laila A. Wahedi. Disrupting hate: The effect of deplatforming hate organizations on their online audience. *Proceedings of the National Academy of Sciences*, 120(24):e2214080120, June 2023. doi: 10.1073/pnas.2214080120. URL <https://www.pnas.org/doi/10.1073/pnas.2214080120>. Publisher: Proceedings of the National Academy of Sciences. 2.2, A.3
- [217] Benjamin Toff and Nick Mathews. Is Social Media Killing Local News? An Examination of Engagement and Ownership Patterns in U.S. Community News on Facebook. *Digital Journalism*, 0(0):1–20, October 2021. ISSN 2167-0811. doi: 10.1080/21670811.2021.1977668. URL <https://doi.org/10.1080/21670811.2021.1977668>. Publisher: Routledge eprint: <https://doi.org/10.1080/21670811.2021.1977668>. A.7
- [218] Samuel P. Trethewey. Strategies to combat medical misinformation on social media. *Postgraduate Medical Journal*, 96(1131):4–6, January 2020. ISSN 0032-5473, 1469-0756. doi: 10.1136/postgradmedj-2019-137201. URL <https://pmj.bmjjournals.com/content/96/1131/4>. Publisher: The Fellowship of Postgraduate Medicine Section: Editorial. 3

- [219] Andrea C. Tricco, Erin Lillie, Wasifa Zarin, Kelly K. O'Brien, Heather Colquhoun, Danielle Levac, David Moher, Micah D. J. Peters, Tanya Horsley, Laura Weeks, Susanne Hempel, Elie A. Akl, Christine Chang, Jessie McGowan, Lesley Stewart, Lisa Hartling, Adrian Aldcroft, Michael G. Wilson, Chantelle Garrity, Simon Lewin, Christina M. Godfrey, Marilyn T. Macdonald, Etienne V. Langlois, Karla Soares-Weiser, Jo Moriarty, Tammy Clifford, Özge Tunçalp, and Sharon E. Straus. PRISMA Extension for Scoping Reviews (PRISMA-ScR): Checklist and Explanation. *Annals of Internal Medicine*, 169(7):467–473, October 2018. ISSN 1539-3704. doi: 10.7326/M18-0850. 2.2
- [220] Joshua Tucker, Andrew Guess, Pablo Barbera, Cristian Vaccari, Alexandra Siegel, Sergey Sanovich, Denis Stukal, and Brendan Nyhan. Social Media, Political Polarization, and Political Disinformation: A Review of the Scientific Literature. *SSRN Electronic Journal*, 2018. ISSN 1556-5068. doi: 10.2139/ssrn.3144139. URL <https://www.ssrn.com/abstract=3144139>. 5.1
- [221] Joshua A. Tucker, Yannis Theocharis, Margaret E. Roberts, and Pablo Barberá. From Liberation to Turmoil: Social Media and Democracy. *Journal of Democracy*, 28(4):46–59, October 2017. doi: 10.1353/jod.2017.0064. URL <https://www.journalofdemocracy.org/articles/from-liberation-to-turmoil-social-media-and-democracy/>. 1.1
- [222] Melissa Tully, Emily K. Vraga, and Leticia Bode. Designing and Testing News Literacy Messages for Social Media. *Mass Communication and Society*, 23(1):22–46, January 2020. ISSN 1520-5436. doi: 10.1080/15205436.2019.1604970. URL <https://doi.org/10.1080/15205436.2019.1604970>. Publisher: Routledge eprint: <https://doi.org/10.1080/15205436.2019.1604970>. 2.2, A.5
- [223] Sander van der Linden, Jon Roozenbeek, Rakoen Maertens, Melisa Basol, Ondřej Kácha, Steve Rathje, and Cecilie Steenbuch Traberg. How Can Psychological Science Help Counter the Spread of Fake News? *The Spanish Journal of Psychology*, 24:e25, 2021. ISSN 1138-7416, 1988-2904. doi: 10.1017/SJP.2021.23. 2.6.2
- [224] Tiago Ventura, Rajeshwari Majumdar, Jonathan Nagler, and Joshua A. Tucker. Misinformation Exposure Beyond Traditional Feeds: Evidence from a WhatsApp Deactivation Experiment in Brazil, May 2023. URL <https://papers.ssrn.com/abstract=4457400>. A.6
- [225] Santosh Vijaykumar, Yan Jin, Daniel Rogerson, Xuerong Lu, Swati Sharma, Anna Maughan, Bianca Fadel, Mariella Silva de Oliveira Costa, Claudia Pagliari, and Daniel Morris. How shades of truth and age affect responses to COVID-19 (Mis)information: randomized survey experiment among WhatsApp users in UK and Brazil. *Humanities and Social Sciences Communications*, 8(1):1–12, March 2021. ISSN 2662-9992. doi: 10.1057/s41599-021-00752-7. URL <https://www.nature.com/articles/s41599-021-00752-7>. Number: 1 Publisher: Palgrave. 3.3.1
- [226] Emily A. Vogels. Support for more regulation of tech companies has declined in U.S., especially among Republicans, May 2022. URL <https://www.pewresearch.org/short-reads/2022/05/13/support-for-more-regulation-of-tech-companies-has-declined-in-u-s->

especially-among-republicans/. 5, 5.5, 6.2.2

- [227] Hilde A. M. Voorveld, Guda van Noort, Daniël G. Muntinga, and Fred Bronner. Engagement with Social Media and Social Media Advertising: The Differentiating Role of Platform Type. *Journal of Advertising*, 47(1):38–54, January 2018. ISSN 0091-3367. doi: 10.1080/00913367.2017.1405754. URL <https://doi.org/10.1080/00913367.2017.1405754>. Publisher: Routledge \_eprint: <https://doi.org/10.1080/00913367.2017.1405754>. (document), 3.4.1, 3.4.1, 3.4.1, 3.9, 3.4.1
- [228] Soroush Vosoughi, Deb Roy, and Sinan Aral. The spread of true and false news online. *Science*, 359(6380):1146–1151, March 2018. doi: 10.1126/science.aap9559. URL <https://www.science.org/doi/10.1126/science.aap9559>. Publisher: American Association for the Advancement of Science. 1.1, 1.2
- [229] Kurt Wagner. Inside Twitter’s Plan to Fact-Check Tweets. *Bloomberg.com*, March 2021. URL <https://www.bloomberg.com/news/newsletters/2021-03-04/birdwatch-inside-twitter-s-plan-to-fact-check-tweets>. 3.1
- [230] Nathan Walter and Sheila T. Murphy. How to unring the bell: A meta-analytic approach to correction of misinformation. *Communication Monographs*, 85(3):423–441, July 2018. ISSN 0363-7751. doi: 10.1080/03637751.2018.1467564. 1.1, 1.3.2, A.2
- [231] Nathan Walter and Nikita A. Salovich. Unchecked vs. Uncheckable: How Opinion-Based Claims Can Impede Corrections of Misinformation. *Mass Communication and Society*, 24(4):500–526, July 2021. ISSN 1520-5436. doi: 10.1080/15205436.2020.1864406. URL <https://doi.org/10.1080/15205436.2020.1864406>. Publisher: Routledge \_eprint: <https://doi.org/10.1080/15205436.2020.1864406>. 2.2
- [232] Nathan Walter, Jonathan Cohen, R. Lance Holbert, and Yasmin Morag. Fact-Checking: A Meta-Analysis of What Works and for Whom. *Political Communication*, 37(3): 350–375, May 2020. ISSN 1058-4609. doi: 10.1080/10584609.2019.1668894. URL <https://doi.org/10.1080/10584609.2019.1668894>. Publisher: Routledge \_eprint: <https://doi.org/10.1080/10584609.2019.1668894>. 6.4.2
- [233] Nathan Walter, John J. Brooks, Camille J. Saucier, and Sapna Suresh. Evaluating the Impact of Attempts to Correct Health Misinformation on Social Media: A Meta-Analysis. *Health Communication*, 36(13):1776–1784, November 2021. ISSN 1532-7027. doi: 10.1080/10410236.2020.1794553. 1.4.1, 3, 4.2.2
- [234] Min Wang, Mingke Rao, and Zhipeng Sun. Typology, Etiology, and Fact-Checking: A Pathological Study of Top Fake News in China. *Journalism Practice*, 16:1–19, August 2020. doi: 10.1080/17512786.2020.1806723. 1.1, 1.3.1
- [235] Claire Wardle and Hossein Derakhshan. Information Disorder: Toward an interdisciplinary framework for research and policy making. Technical report, Council of Europe, September 2017. URL <https://rm.coe.int/information-disorder-toward-an-interdisciplinary-framework-for-research/168076277c>. 1.2, 1.3.1, 1.1, 1.3.1, 1.3.2, 1.3.2, 1.3.2, 1.3.2, A.7
- [236] Benjamin R. Warner and Ryan Neville-Shepard. Echoes of a Conspiracy: Birthers,

- Truthers, and the Cultivation of Extremism. *Communication Quarterly*, 62(1):1–17, January 2014. ISSN 0146-3373. doi: 10.1080/01463373.2013.822407. URL <https://doi.org/10.1080/01463373.2013.822407>. Publisher: Routledge eprint: <https://doi.org/10.1080/01463373.2013.822407>. 1.1
- [237] Amy Watson. Sharing of made-up news on social networks in the U.S. 2020, June 2022. URL <https://www.statista.com/statistics/657111/fake-news-sharing-online/>. 3.1, 3.3.2
- [238] Sarah Myers West. Censored, suspended, shadowbanned: User interpretations of content moderation on social media platforms. *New Media & Society*, 20(11):4366–4383, November 2018. ISSN 1461-4448. doi: 10.1177/1461444818773059. URL <https://doi.org/10.1177/1461444818773059>. Publisher: SAGE Publications. A.3
- [239] Randy Yee Man Wong, Christy M. K. Cheung, Bo Xiao, and Jason Bennett Thatcher. Standing Up or Standing By: Understanding Bystanders' Proactive Reporting Responses to Social Media Harassment. *Information Systems Research*, 32(2):561–581, June 2021. ISSN 1047-7047. doi: 10.1287/isre.2020.0983. URL <https://pubsonline.informs.org/doi/abs/10.1287/isre.2020.0983>. Publisher: INFORMS. 3.5
- [240] Huiiping Wu and Shing-On Leung. Can Likert Scales be Treated as Interval Scales?—A Simulation Study. *Journal of Social Service Research*, 43(4):527–532, August 2017. ISSN 0148-8376. doi: 10.1080/01488376.2017.1329775. URL <https://doi.org/10.1080/01488376.2017.1329775>. Publisher: Routledge eprint: <https://doi.org/10.1080/01488376.2017.1329775>. 3.3.3
- [241] Liang Wu, Fred Morstatter, Kathleen M. Carley, and Huan Liu. Misinformation in Social Media: Definition, Manipulation, and Detection. *ACM SIGKDD Explorations Newsletter*, 21(2):80–90, November 2019. ISSN 1931-0145, 1931-0153. doi: 10.1145/3373464.3373475. URL <https://dl.acm.org/doi/10.1145/3373464.3373475>. 1.2
- [242] Kamya Yadav. Countering Influence Operations: A Review of Policy Proposals Since 2016. Technical report, Carnegie Endowment for International Peace, November 2020. URL <https://carnegieendowment.org/2020/11/30/countering-influence-operations-review-of-policy-proposals-since-2016-pub-83333>. 5, 5.1, 5.2.1, A.7
- [243] Kamya Yadav. Platform Interventions: How Social Media Counters Influence Operations. Technical report, Carnegie Endowment for International Peace, January 2021. URL <https://carnegieendowment.org/2021/01/25/platform-interventions-how-social-media-counters-influence-operations-pub-83698>. 2.5, 5, 5.1, A.1, A.4
- [244] Juanita Zainudin, Nazlena Mohamad Ali, Alan F. Smeaton, and Mohamad Taha Ijab. Intervention Strategies for Misinformation Sharing on Social Media: A Bibliometric Analysis. *IEEE Access*, 12:140359–140379, 2024. ISSN 2169-3536. doi: 10.1109/ACCESS.2024.3469248. URL <https://ieeexplore.ieee.org/document/>

- [245] Savvas Zannettou, Michael Sirivianos, Jeremy Blackburn, and Nicolas Kourtellis. The Web of False Information: Rumors, Fake News, Hoaxes, Clickbait, and Various Other Shenanigans. *Journal of Data and Information Quality*, 11(3):1–37, September 2019. ISSN 1936-1955, 1936-1963. doi: 10.1145/3309699. URL <https://dl.acm.org/doi/10.1145/3309699>. 1.1, 1.3.1, 1.3.2
- [246] Max Zhan. Here's why Meta ended fact-checking, according to experts, January 2025. URL <https://abcnews.go.com/US/why-did-meta-remove-fact-checkers-experts-explain/story?id=117417445>. 6.2.3
- [247] Xichen Zhang and Ali A. Ghorbani. An overview of online fake news: Characterization, detection, and discussion. *Information Processing & Management*, 57(2):102025, March 2020. ISSN 03064573. doi: 10.1016/j.ipm.2019.03.004. URL <https://linkinghub.elsevier.com/retrieve/pii/S0306457318306794>. 1.1, 1.2, 1.3.2, 1.3.2, 1.3.2
- [248] Hong Zhou, Yaobin Lu, Ling Zhao, Bin Wang, and Ting Li. Effective reporting system to encourage users' reporting behavior in social media platforms: an empirical study based on structural empowerment theory. *Behaviour & Information Technology*, 43(14): 3490–3509, October 2024. ISSN 0144-929X. doi: 10.1080/0144929X.2023.2281491. URL <https://doi.org/10.1080/0144929X.2023.2281491>. Publisher: Taylor & Francis \_eprint: <https://doi.org/10.1080/0144929X.2023.2281491>. 3.5, 6.2.2, A.6
- [249] Yu-Qian Zhu and Houn-Gee Chen. Social media and human need satisfaction: Implications for social media marketing. *Business Horizons*, 58(3):335–345, May 2015. ISSN 0007-6813. doi: 10.1016/j.bushor.2015.01.006. URL <https://www.sciencedirect.com/science/article/pii/S0007681315000075>. (document), 3.4.1, 3.4.1, 3.7